



What's the Scoop with ES-Hadoop?

James Baiera and Anoop Sunke

Elastic

2017-03-09

@JimmyThaHat & @apsunke





James Baiera

- Software Engineer @ Elastic
- ES-Hadoop Maintainer
- [@jbaiera](#) on Github
- [@JimmyThaHat](#) on Twitter

Agenda

45 minute block

1

Introduction to ES-Hadoop

2

Anatomy of a Job Execution

3

Connector Feature Tour and What's to Come

4

User Success Stories and Use Cases

5

Q&A

Hadoop Overview

In two minutes or less

Hadoop Ecosystem



Hadoop Ecosystem



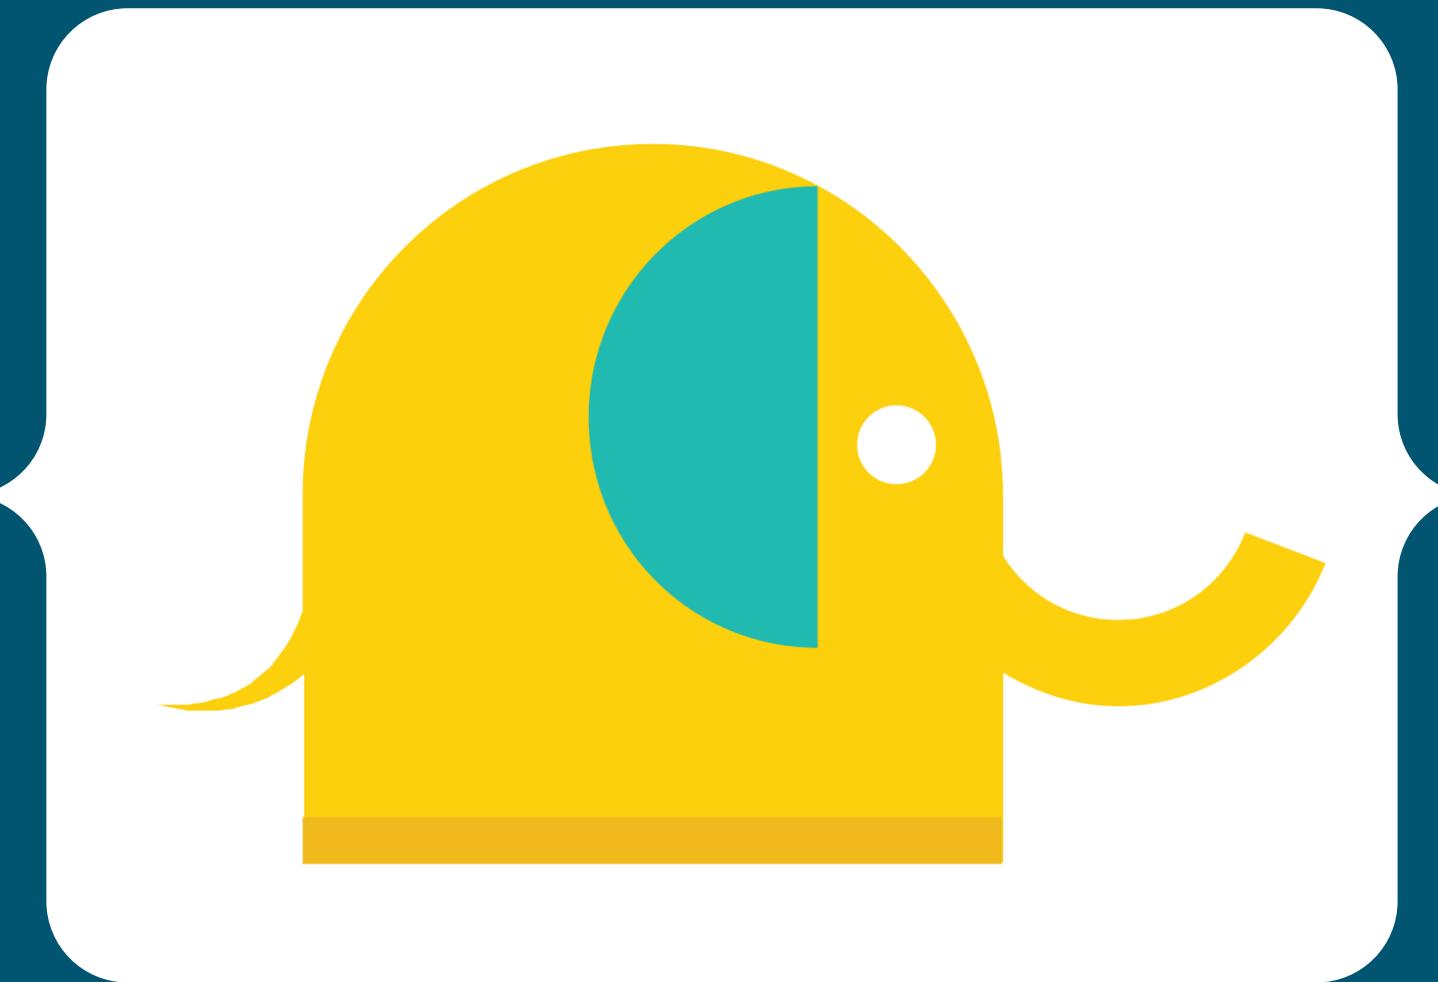
Map / Reduce

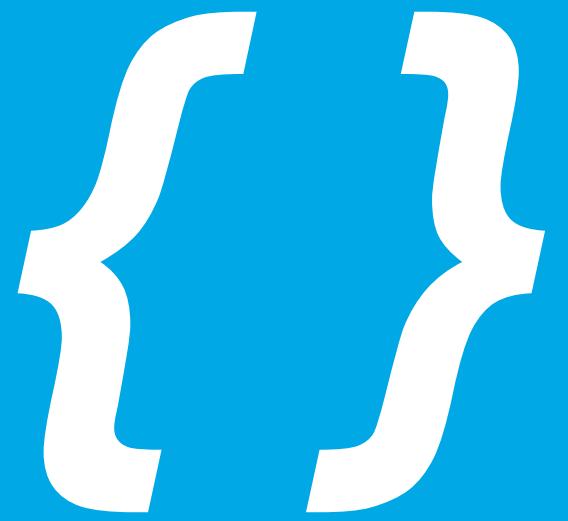
Other

YARN

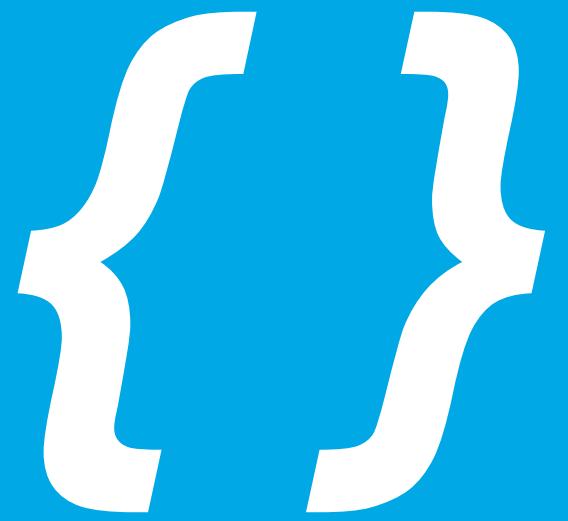
Hadoop Distributed File System (HDFS)

Introduction to ES-Hadoop

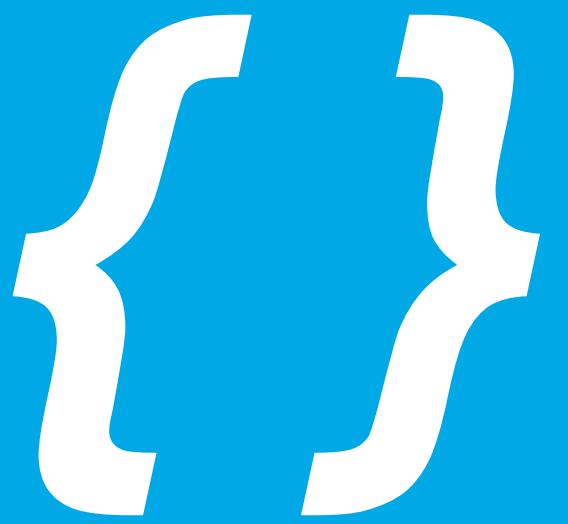




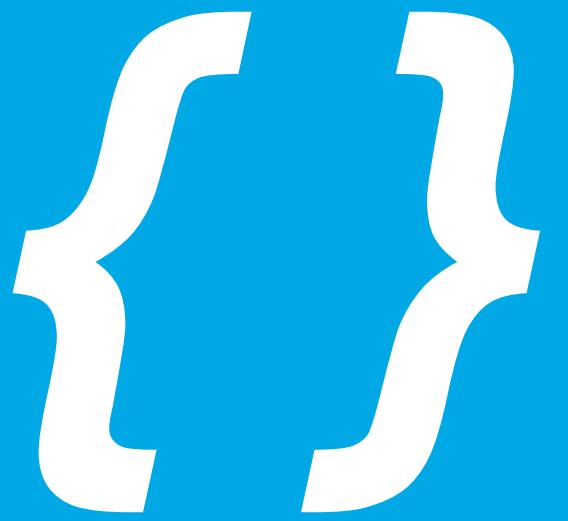
Elasticsearch for Apache Hadoop is an open-source, stand-alone, self-contained, small library that allows Hadoop jobs to interact with Elasticsearch.



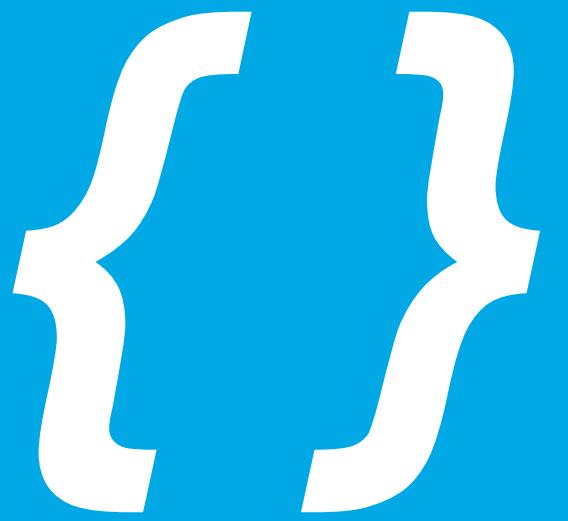
*Elasticsearch for Apache Hadoop is an **open-source**, stand-alone, self-contained, small library that allows Hadoop jobs to interact with Elasticsearch.*



*Elasticsearch for Apache Hadoop is an open-source,
stand-alone, self-contained, small library that allows
Hadoop jobs to interact with Elasticsearch.*

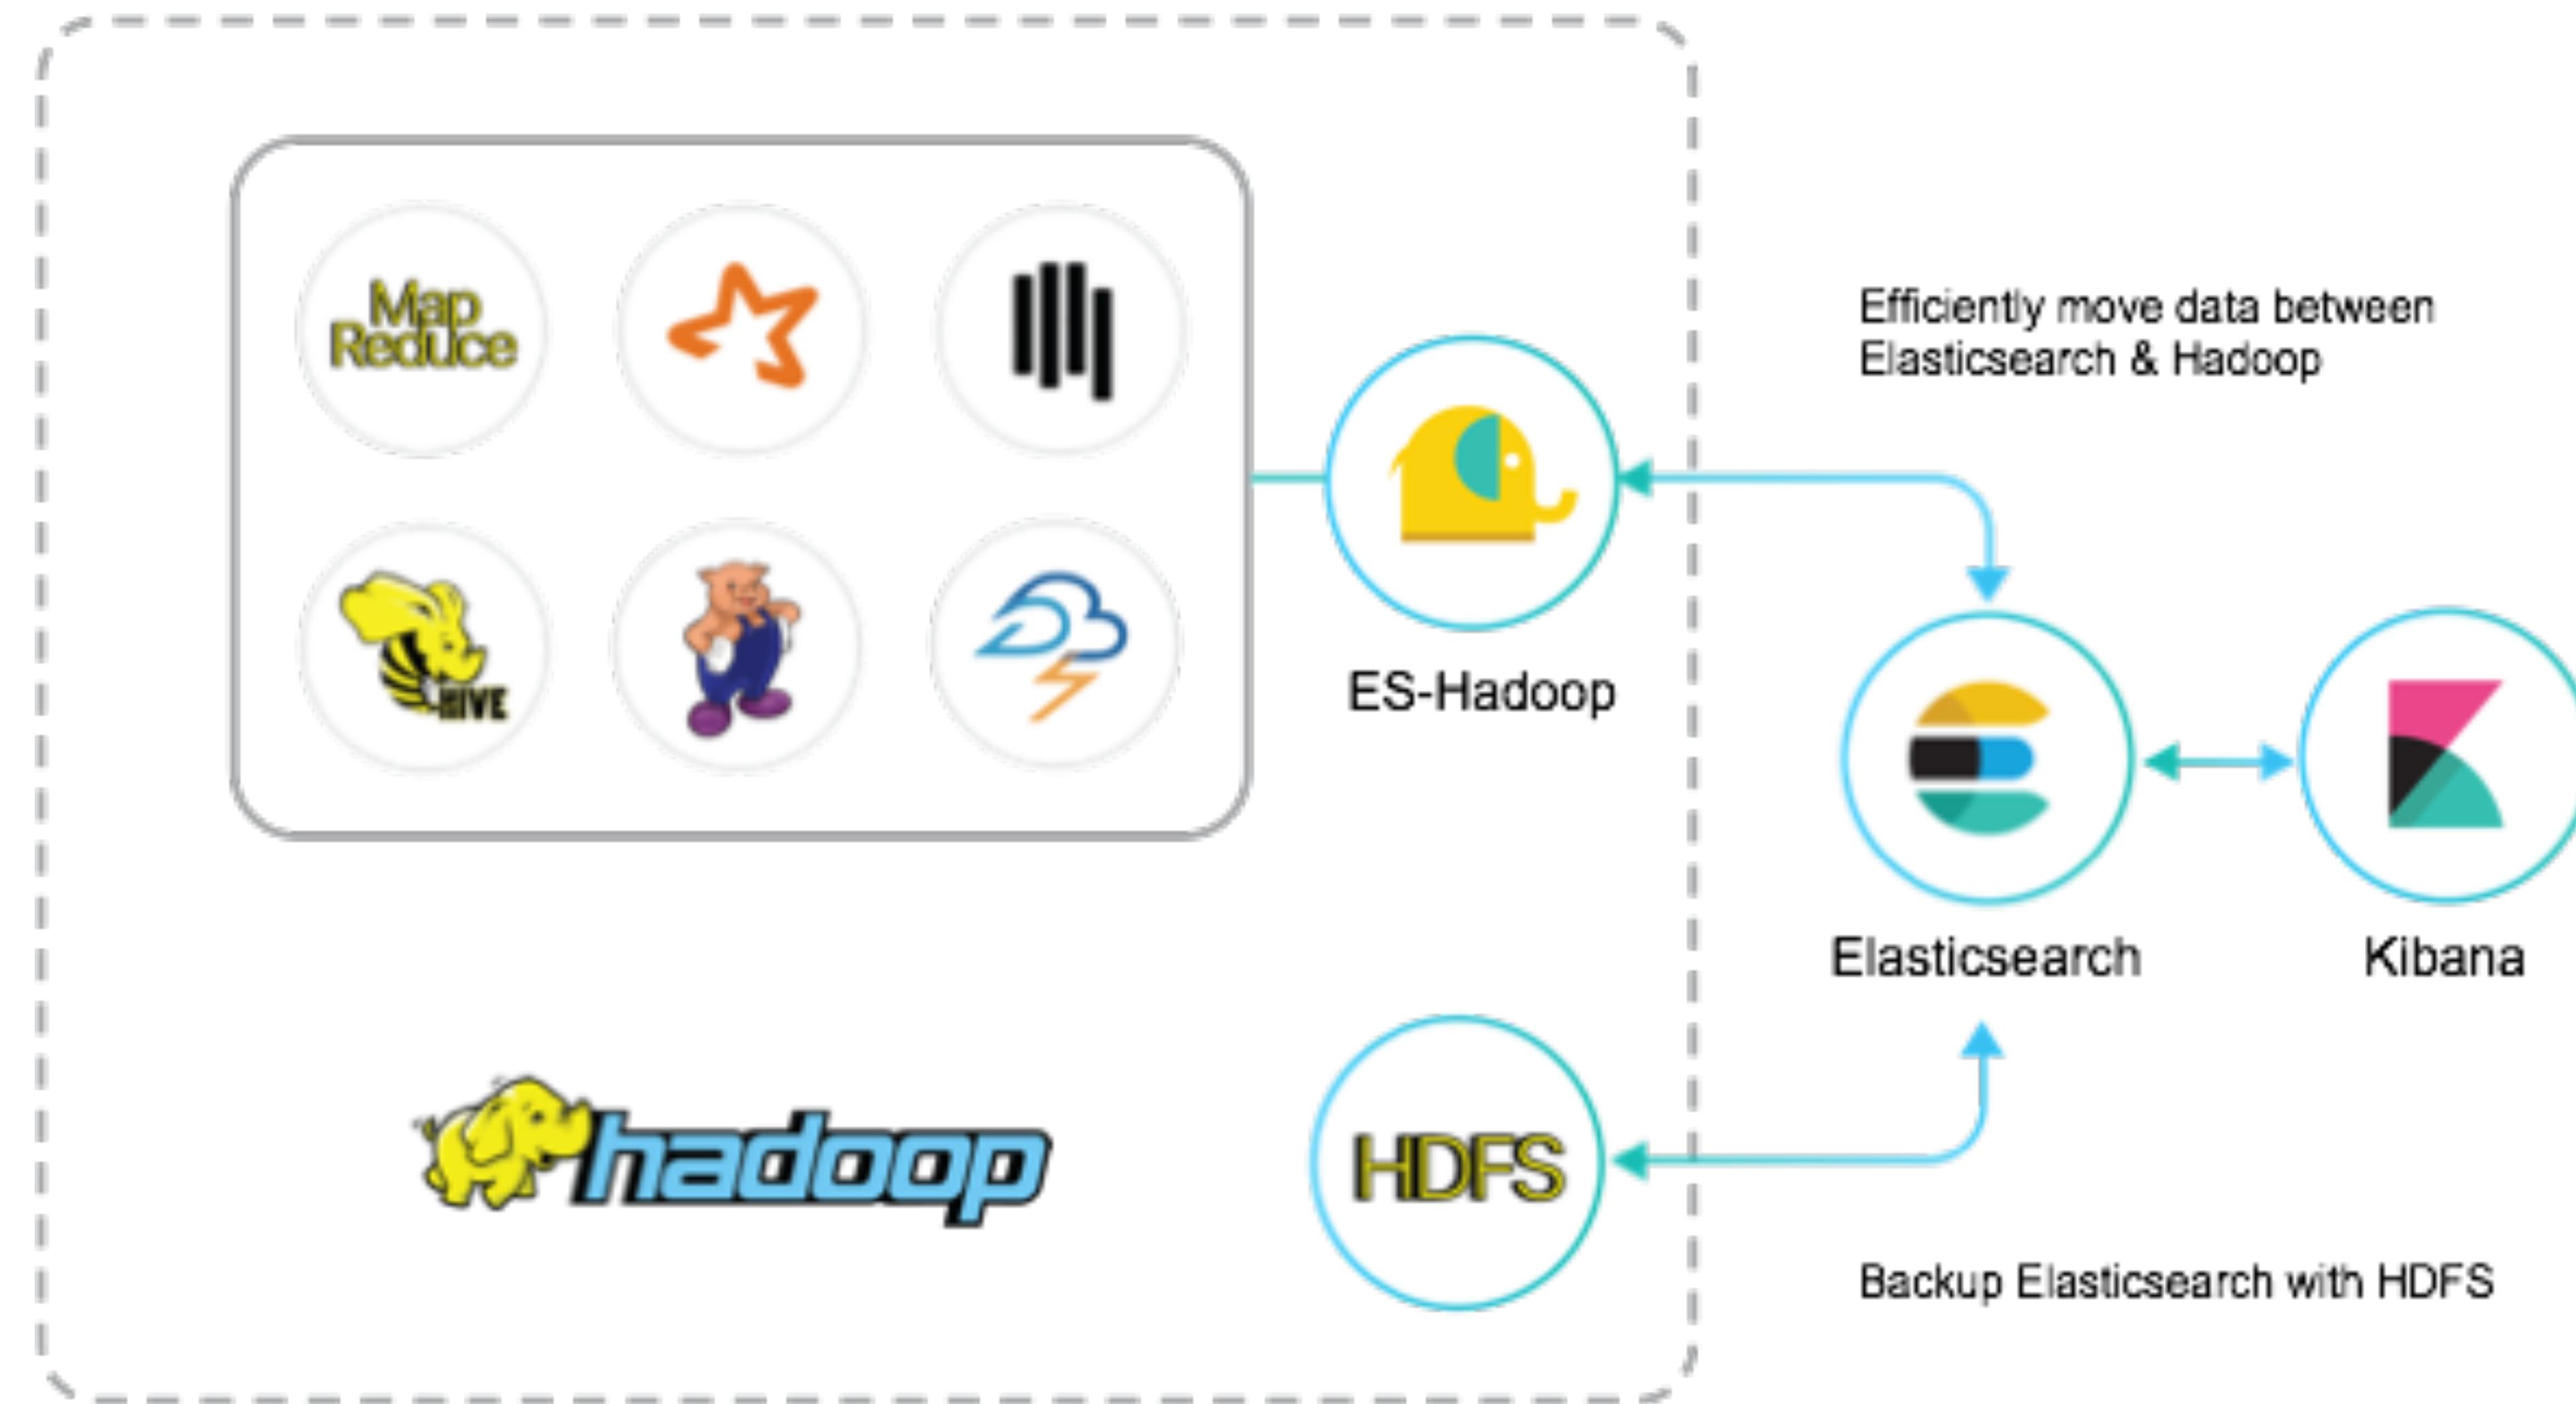


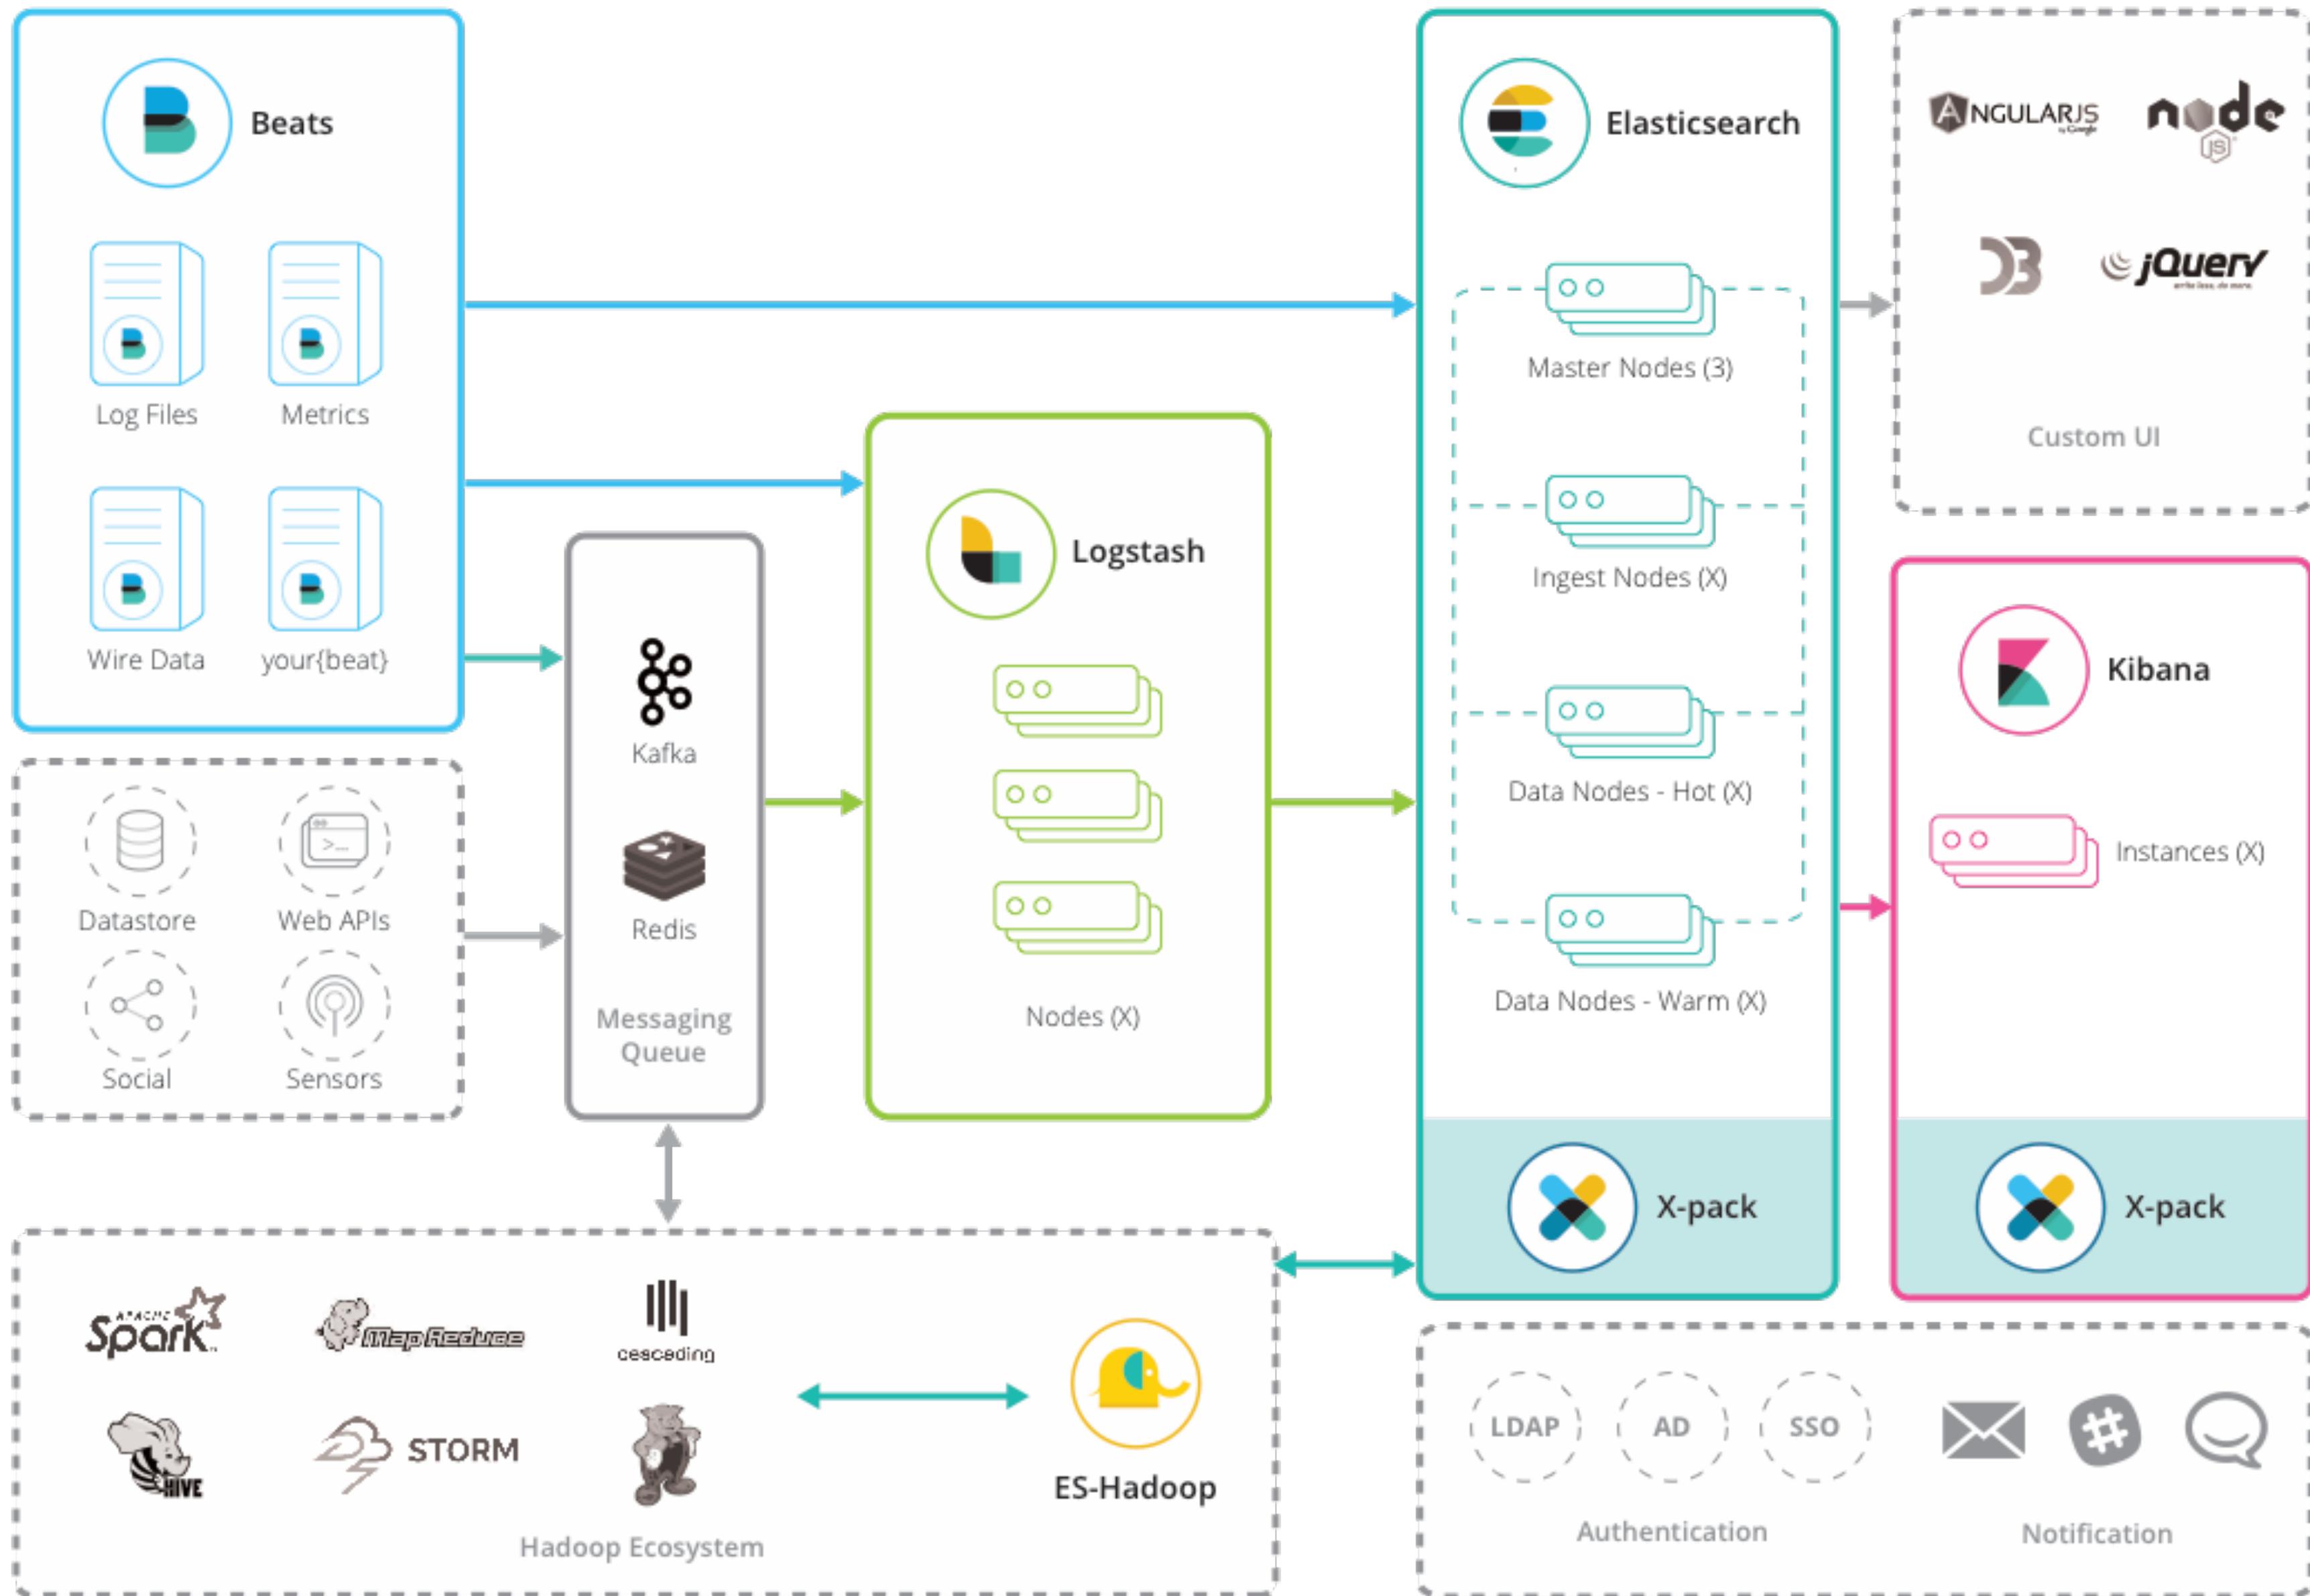
*Elasticsearch for Apache Hadoop is an open-source, stand-alone, self-contained, **small library** that allows Hadoop jobs to interact with Elasticsearch.*



Elasticsearch for Apache Hadoop is an open-source, stand-alone, self-contained, small library that allows Hadoop jobs to interact with Elasticsearch.

Birds Eye View





ES-Hadoop Integrations

Library / API	ES-Hadoop Exposed As
MapReduce	 Input/OutputFormat
Cascading	 Tap/Sink
Apache Pig	 Storage (Load and Store)
Apache Hive	 EXTERNAL Table
Apache Storm	 Spout/Bolt
Apache Spark	 RDD, DStream, Dataframe, Dataset, DataSource

ES-Hadoop Features

- Latest Spark 2.1 support
- Adaptive I/O for error handling, re-routing, backpressure
- Push-down processing on either platform
- Co-location, rack awareness
- Elastic security compatible - basic authentication, SSL/TLS, PKI
- Hadoop Kerberos security compatible
- Hadoop distribution agnostic



Why?

Peeling Back the Layers

ES-Hadoop Functional Deep Dive

Pushdown

Projections and Predicates

ID : Long	First Name : String	Age : Integer	Profile : Text
347	Martha	32	Has a fancy vineyard
348	Mary	37	Friends with Peter and Paul
349	Geoff	25	Spells name with a 'G'
350	Travis	23	Loves jazz music
351	Jeremiah	39	Bullfrog, Good Friend
352	Mark	42	Hates cold spaghetti

Predicates

Predicate: ID > 347 && ID <= 350

ID : Long	First Name : String	Age : Integer	Profile : Text
347	Martha	32	Has a fancy vineyard
348	Mary	37	Friends with Peter and Paul
349	Geoff	25	Spells name with a 'G'
350	Travis	23	Loves jazz music
351	Jeremiah	39	Bullfrog, Good Friend
352	Mark	42	Hates cold spaghetti

Projections

Projection: Select (ID, FirstName, Age)

ID : Long	FirstName : String	Age : Integer	Profile : Text
347	Martha	32	Has a fancy vineyard
348	Mary	37	Friends with Peter and Paul
349	Geoff	25	Spells name with a 'G'
350	Travis	23	Loves jazz music
351	Jeremiah	39	Bullfrog, Good Friend
352	Mark	42	Hates cold spaghetti

Final View

Select (ID, FirstName, Age) where ID > 347 && ID <= 350

ID : Long	First Name : String	Age : Integer	Profile : Text
347	Martha	32	Has a fancy vineyard
348	Mary	37	Friends with Peter and Paul
349	Geoff	25	Spells name with a 'G'
350	Travis	23	Loves jazz music
351	Jeremiah	39	Bullfrog, Good Friend
352	Mark	42	Hates cold spaghetti

Spark SQL Reading

Predicate Pushdown

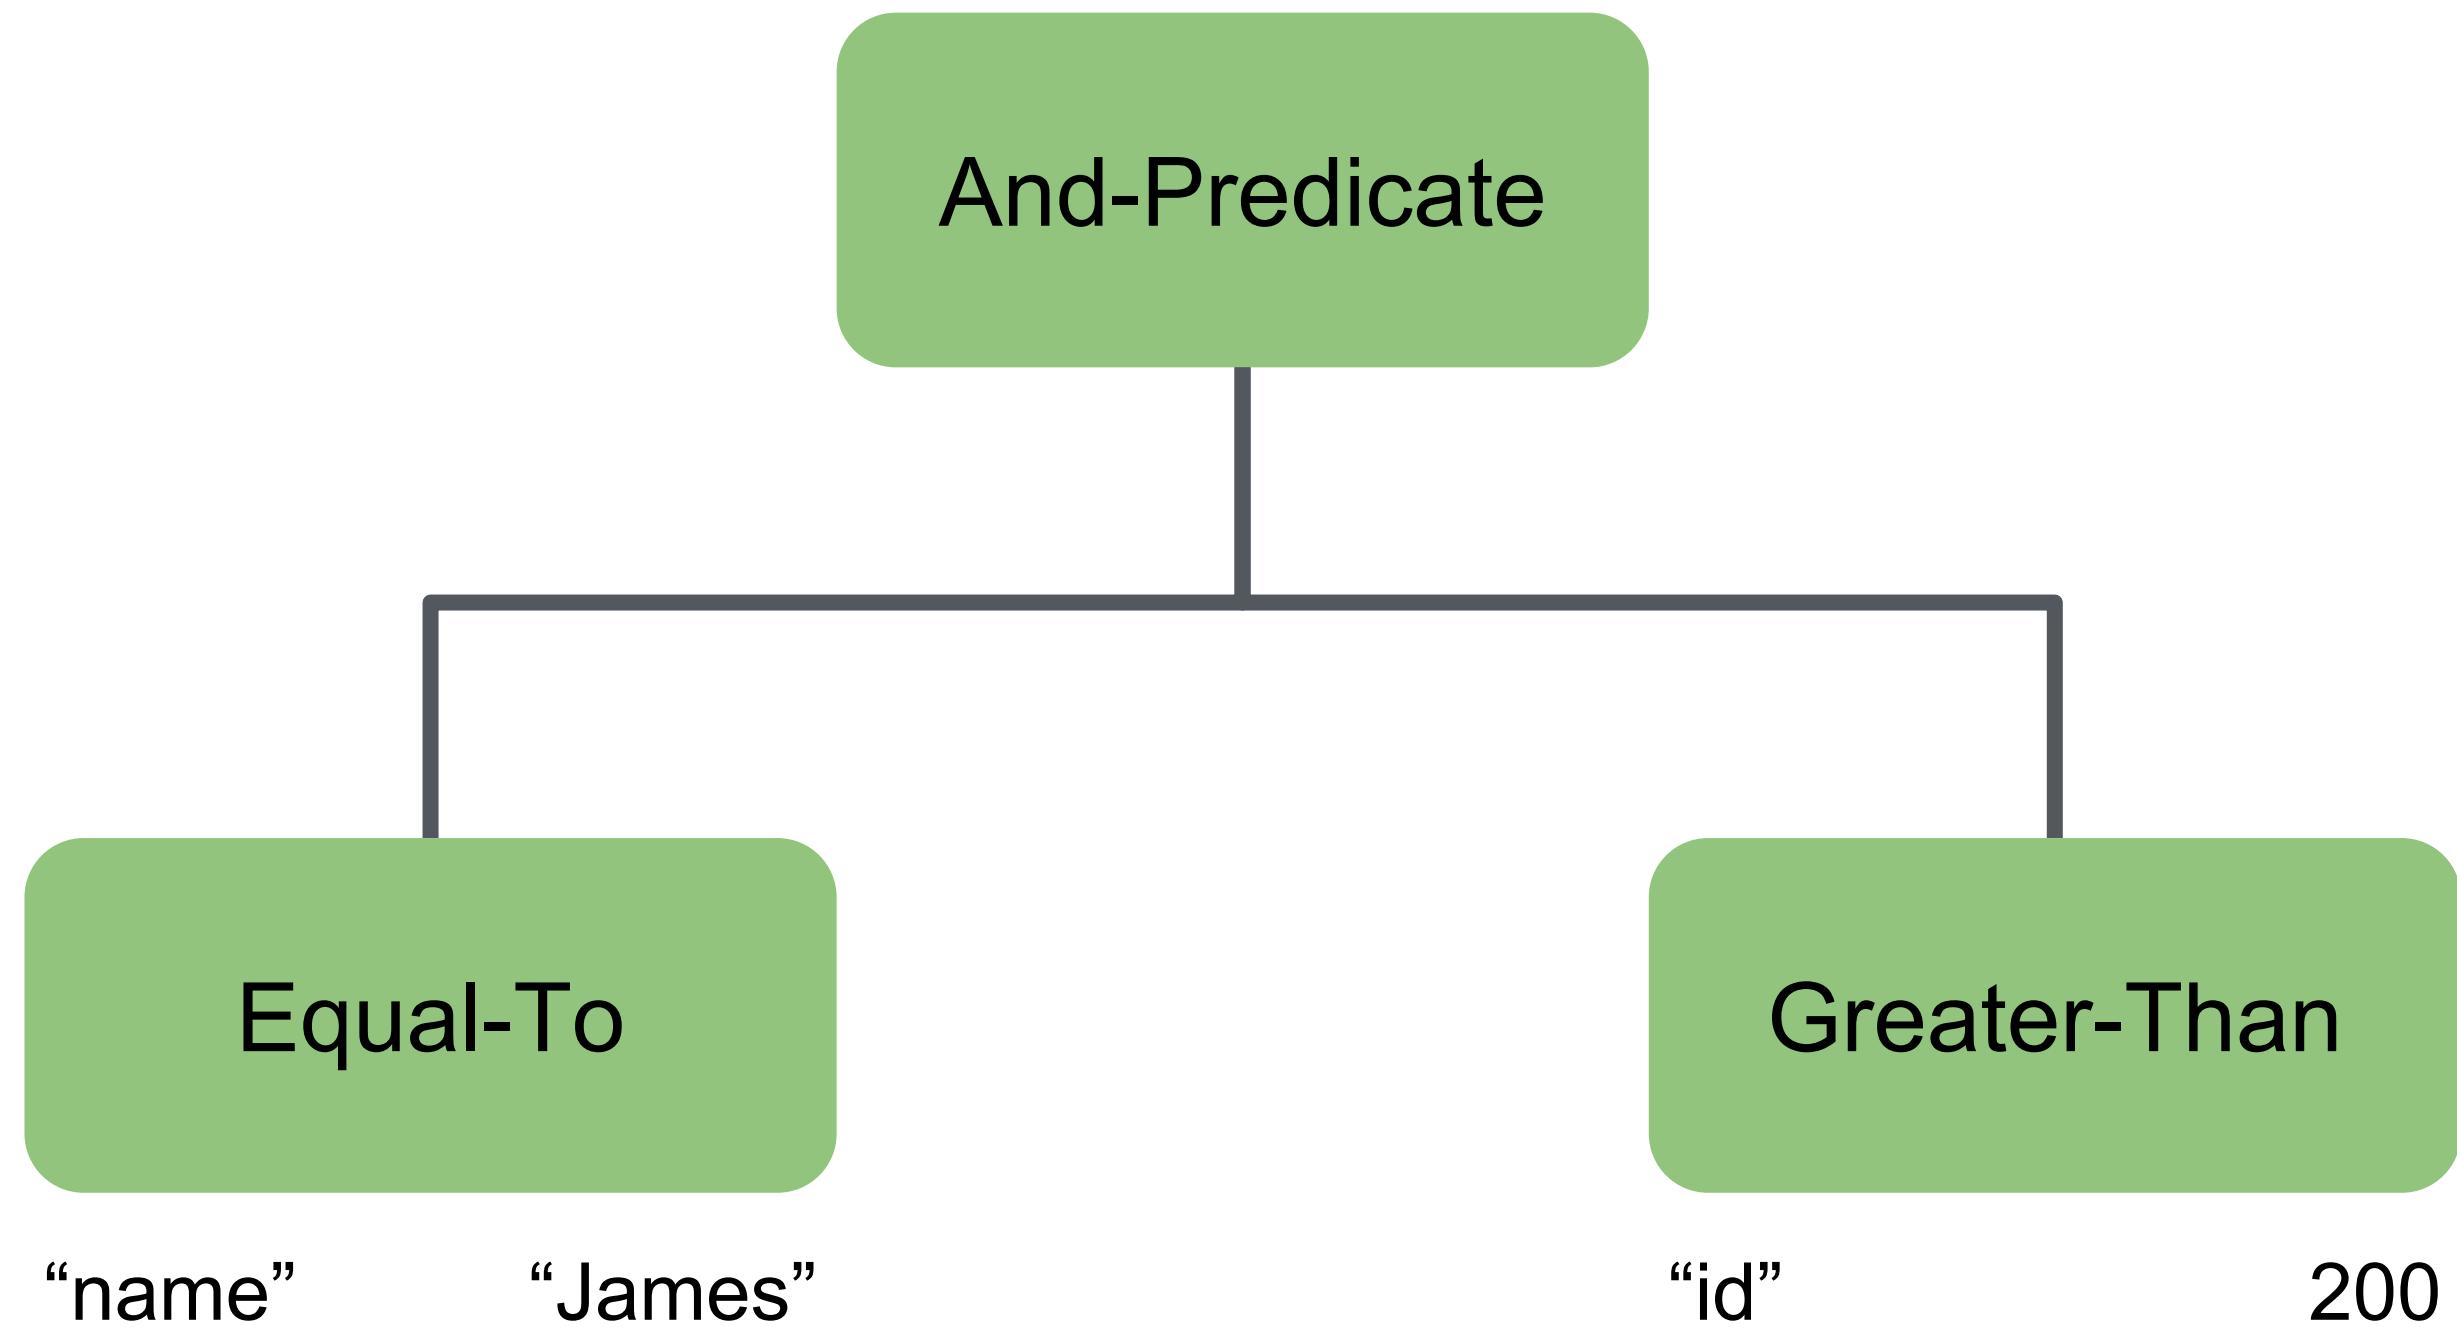
```
val df = sqlContext.read().format("es").load("spark/users")

df.printSchema()
// root
// |-- name: string (nullable = true)
// |-- id: long (nullable = true)
// |-- profile: string (nullable = true)

val filter = df.filter(
  df("name").equalTo("James").and(df("id").gt(200)))
```

Spark SQL Reading

Predicate Pushdown



Spark SQL Reading

Predicate Pushdown

```
{  
  "query" : {  
    "bool" : {  
      "must" : [ { "match_all" : {} } ],  
      "filter" : [ {  
        "bool": {"filter": [  
          {"match" : { "name" : "James"}},  
          {"range" : { "id" : { "gt" : 200 } }}]  
        ]  
      }  
    }  
  }  
}
```

Spark SQL Reading

Projection Pushdown

```
val df = sqlContext.read().format("es").load("spark/users")
df.createOrReplaceTempView("myIndex")

df.printSchema()
// root
// |-- name: string (nullable = true)
// |-- id: long (nullable = true)
// |-- profile: string (nullable = true)

val names = sqlContext.sql("SELECT name, id FROM myIndex")
```

Spark SQL Reading

Projection Pushdown

```
{  
  "_source": [ "name", "id" ],  
  "query" : {  
    "bool":{  
      "must" : [ { "match_all" : {} } ]  
    }  
  }  
}
```

Spark SQL Reading

Projection Pushdown + Predicate Pushdown

```
val df = sqlContext.read().format("es").load("spark/users")
df.createOrReplaceTempView("myIndex")

df.printSchema()
// root
// |-- name: string (nullable = true)
// |-- id: long (nullable = true)
// |-- profile: string (nullable = true)

val names = df.sqlContext.sql(
    "SELECT name FROM myIndex WHERE id >=1 AND id <= 10")
```

Spark SQL Reading

Projection Pushdown + Predicate Pushdown

```
{  
  "_source": [ "name", "id" ],  
  "query" : {  
    "bool":{  
      "must" : [ { "match_all" : {} } ],  
      "filter" : [ {  
        "bool":{ "filter": [  
          { "range" : { "id" : { "gte" : 1 } } },  
          { "range" : { "id" : { "lte" : 10 } } }  
        ]}  
      }]  
    }]  
  }  
}
```

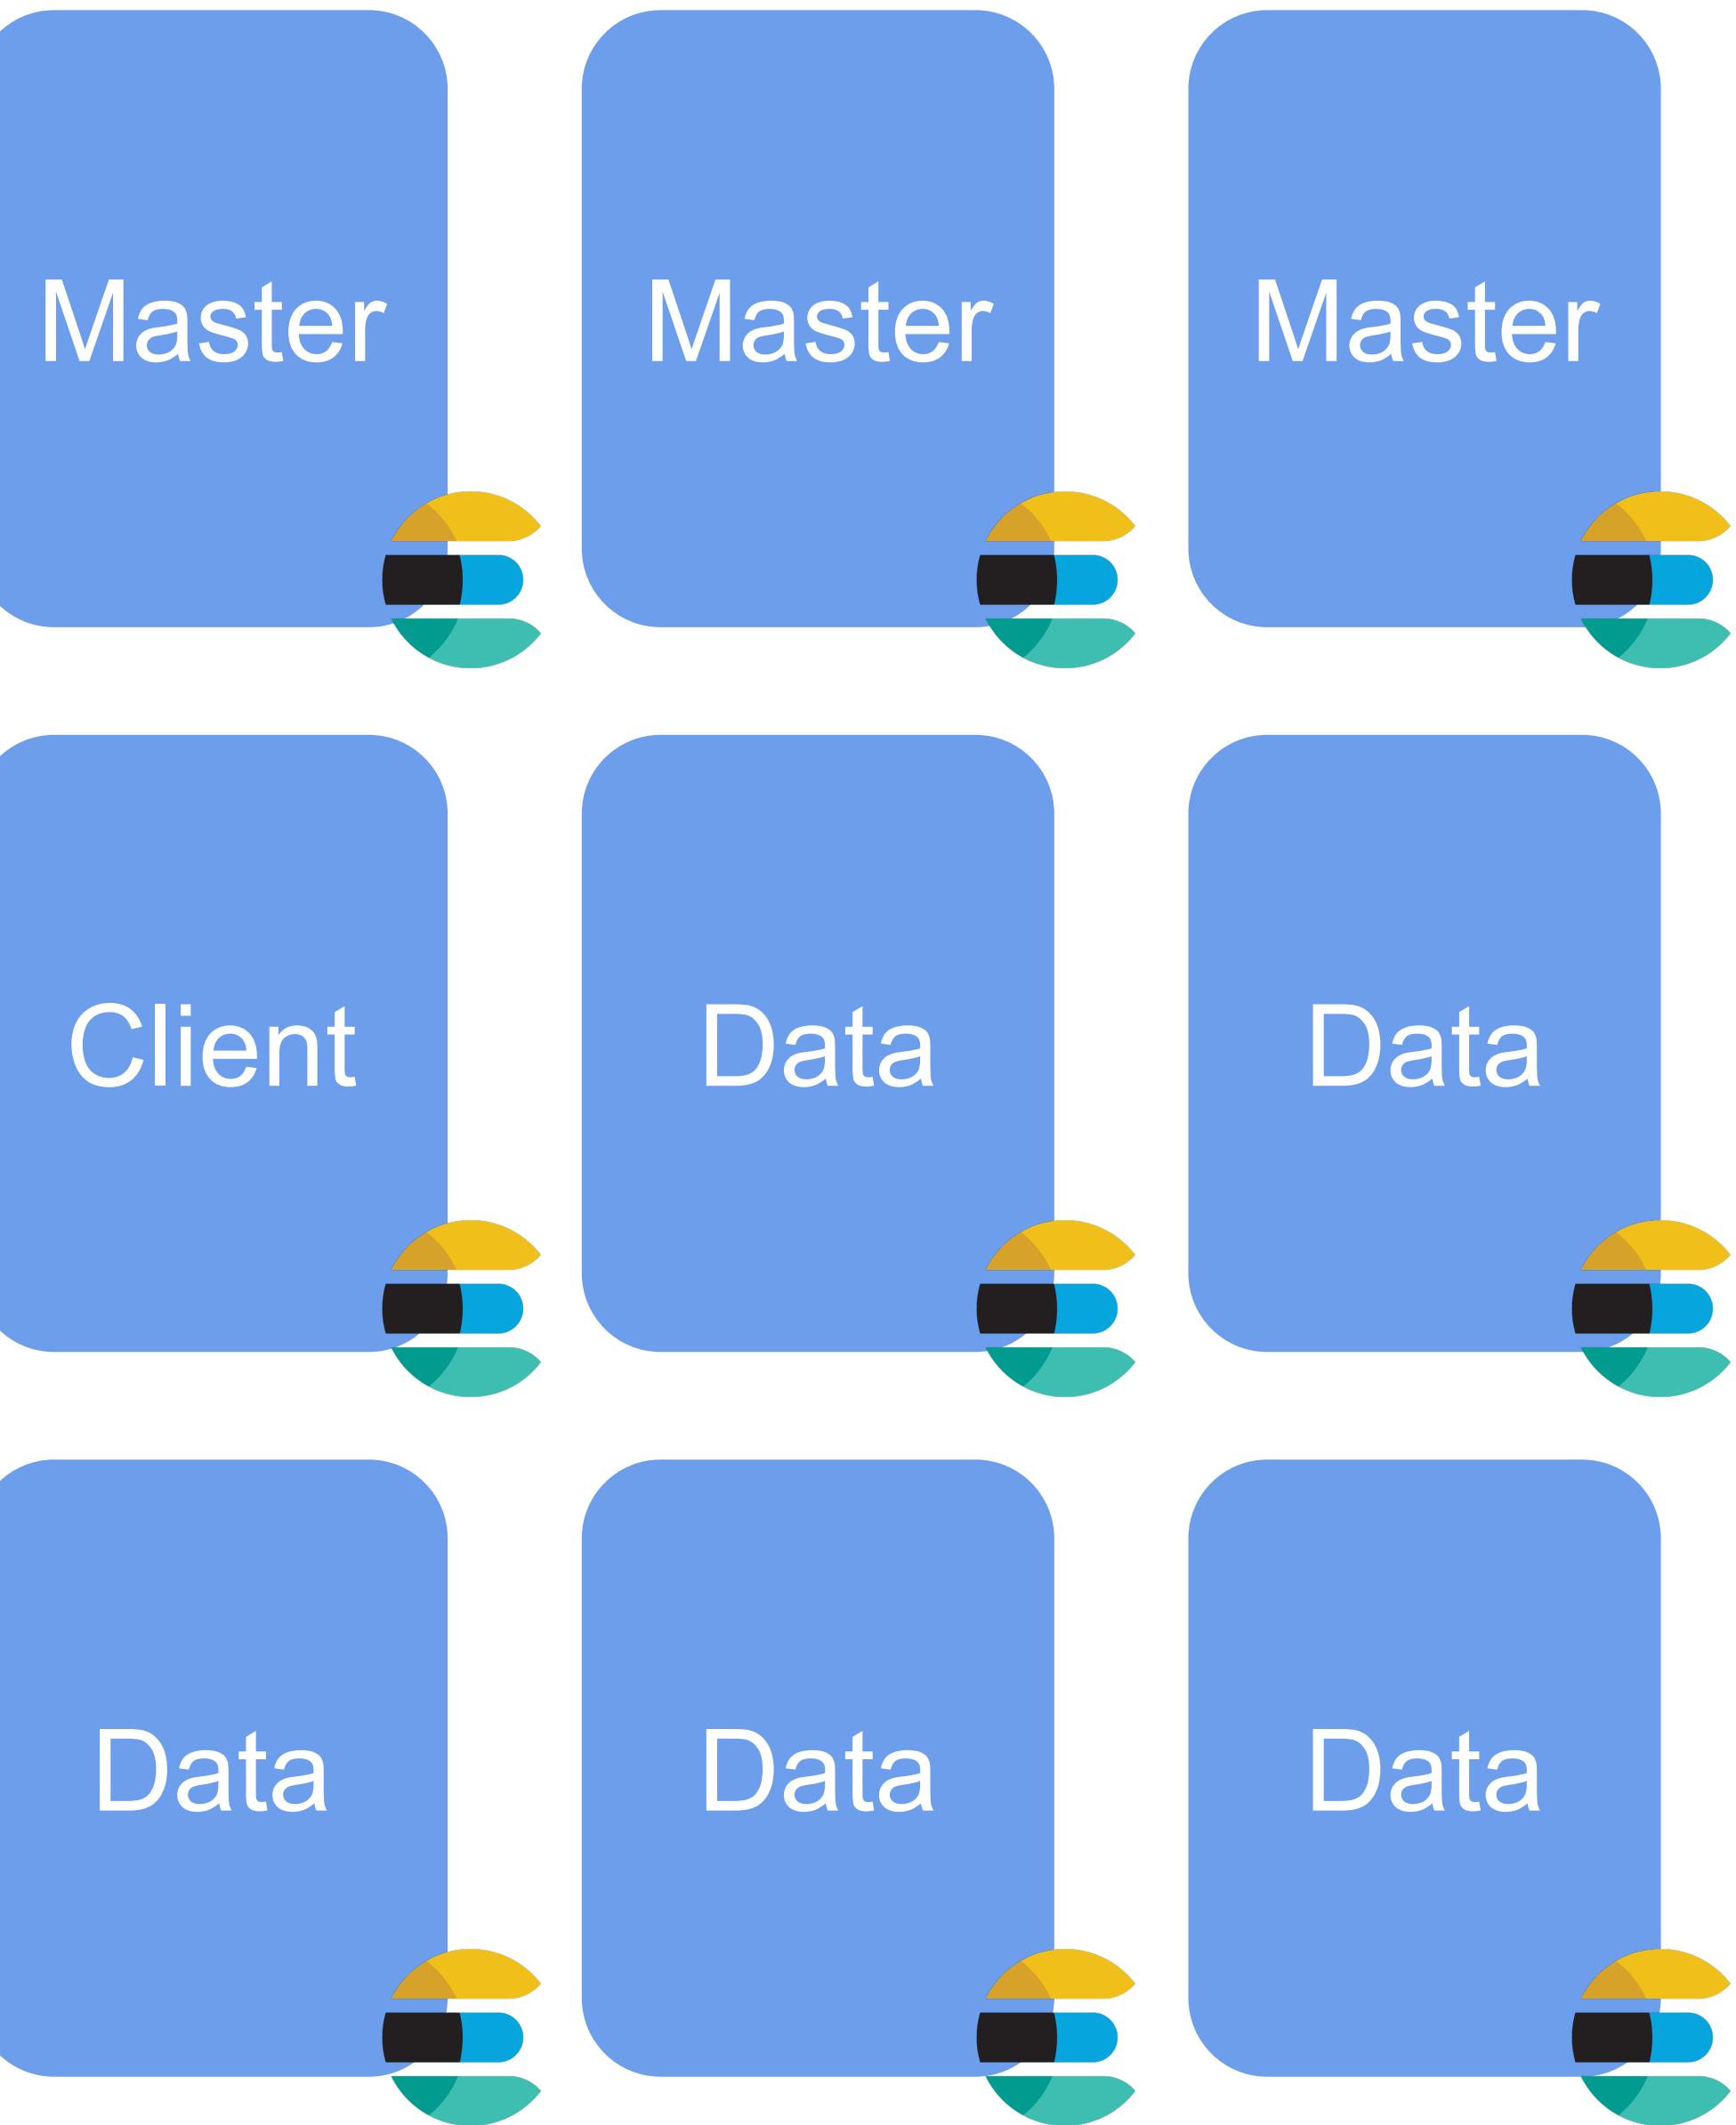
Automatic Pushdown Support in ES-Hadoop

Library / API	Projection	Predicate
MapReduce	 Manual	Manual
Cascading	 Automatic	Manual
Apache Pig	 Automatic	Manual
Apache Hive	 Automatic	Manual
Apache Storm	 Manual	Manual
Apache Spark	(SQL) Automatic (RDD) Manual	(SQL) Automatic (RDD) Manual

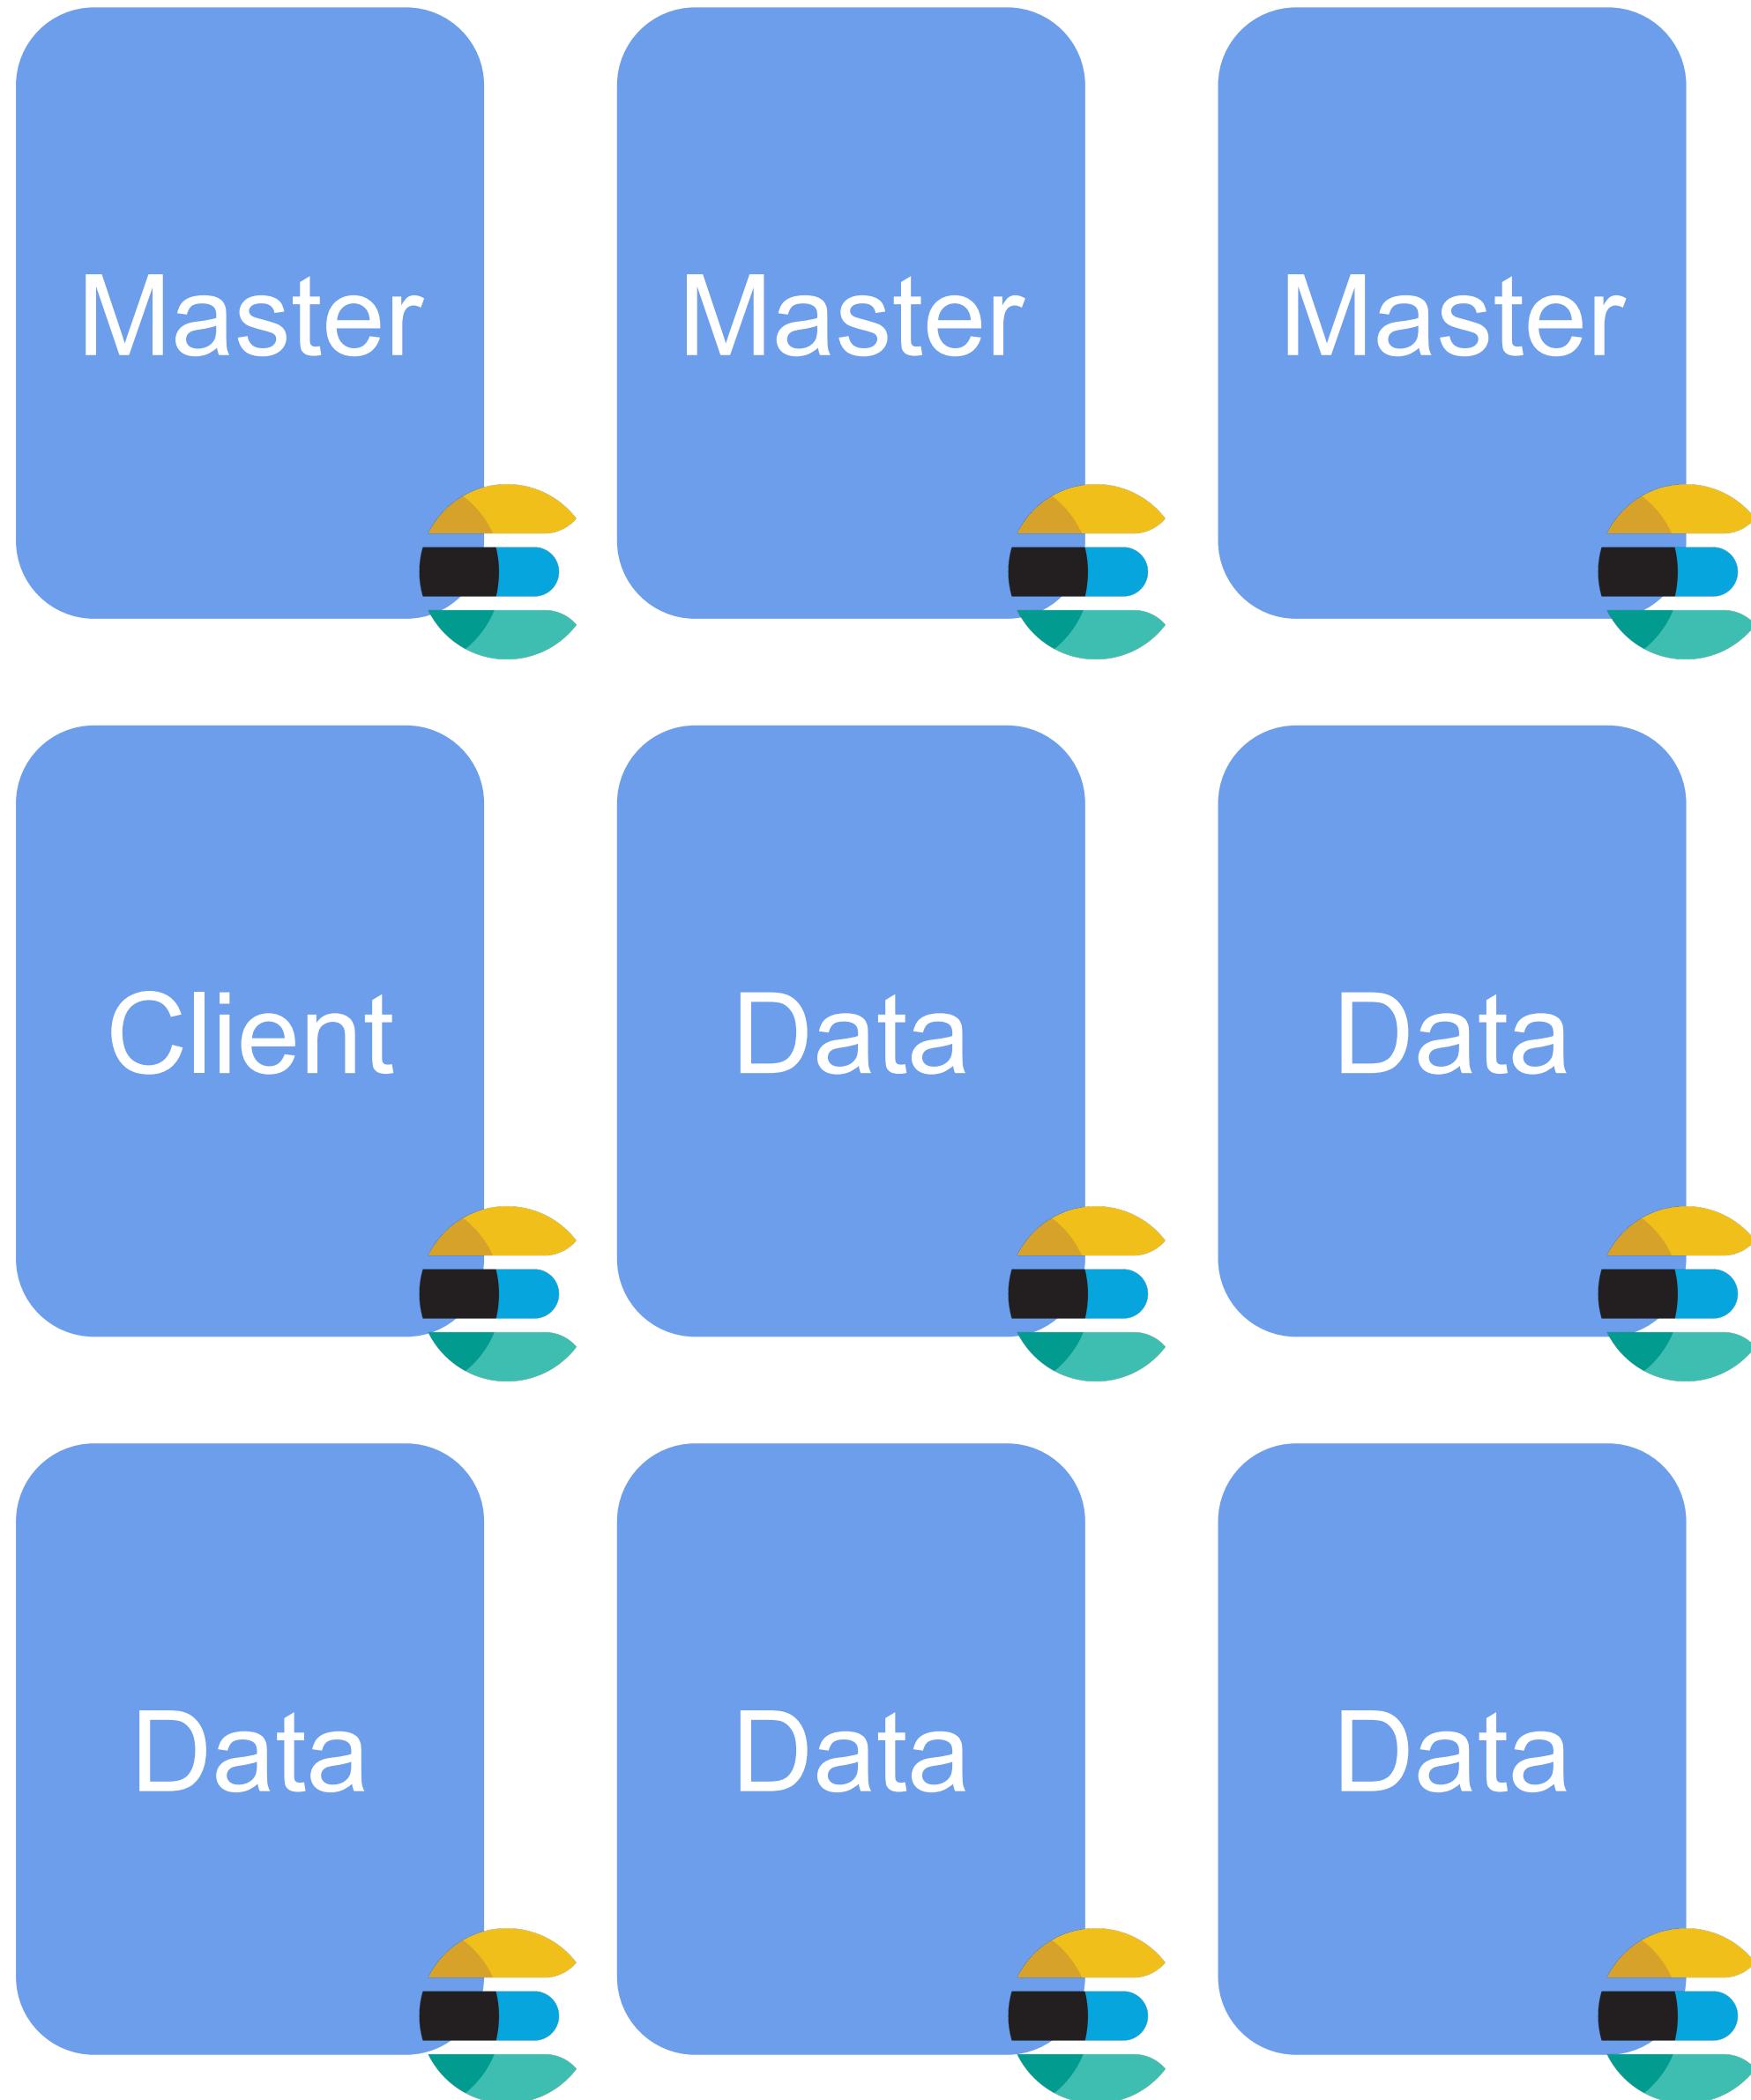
Node Discovery



elasticsearch

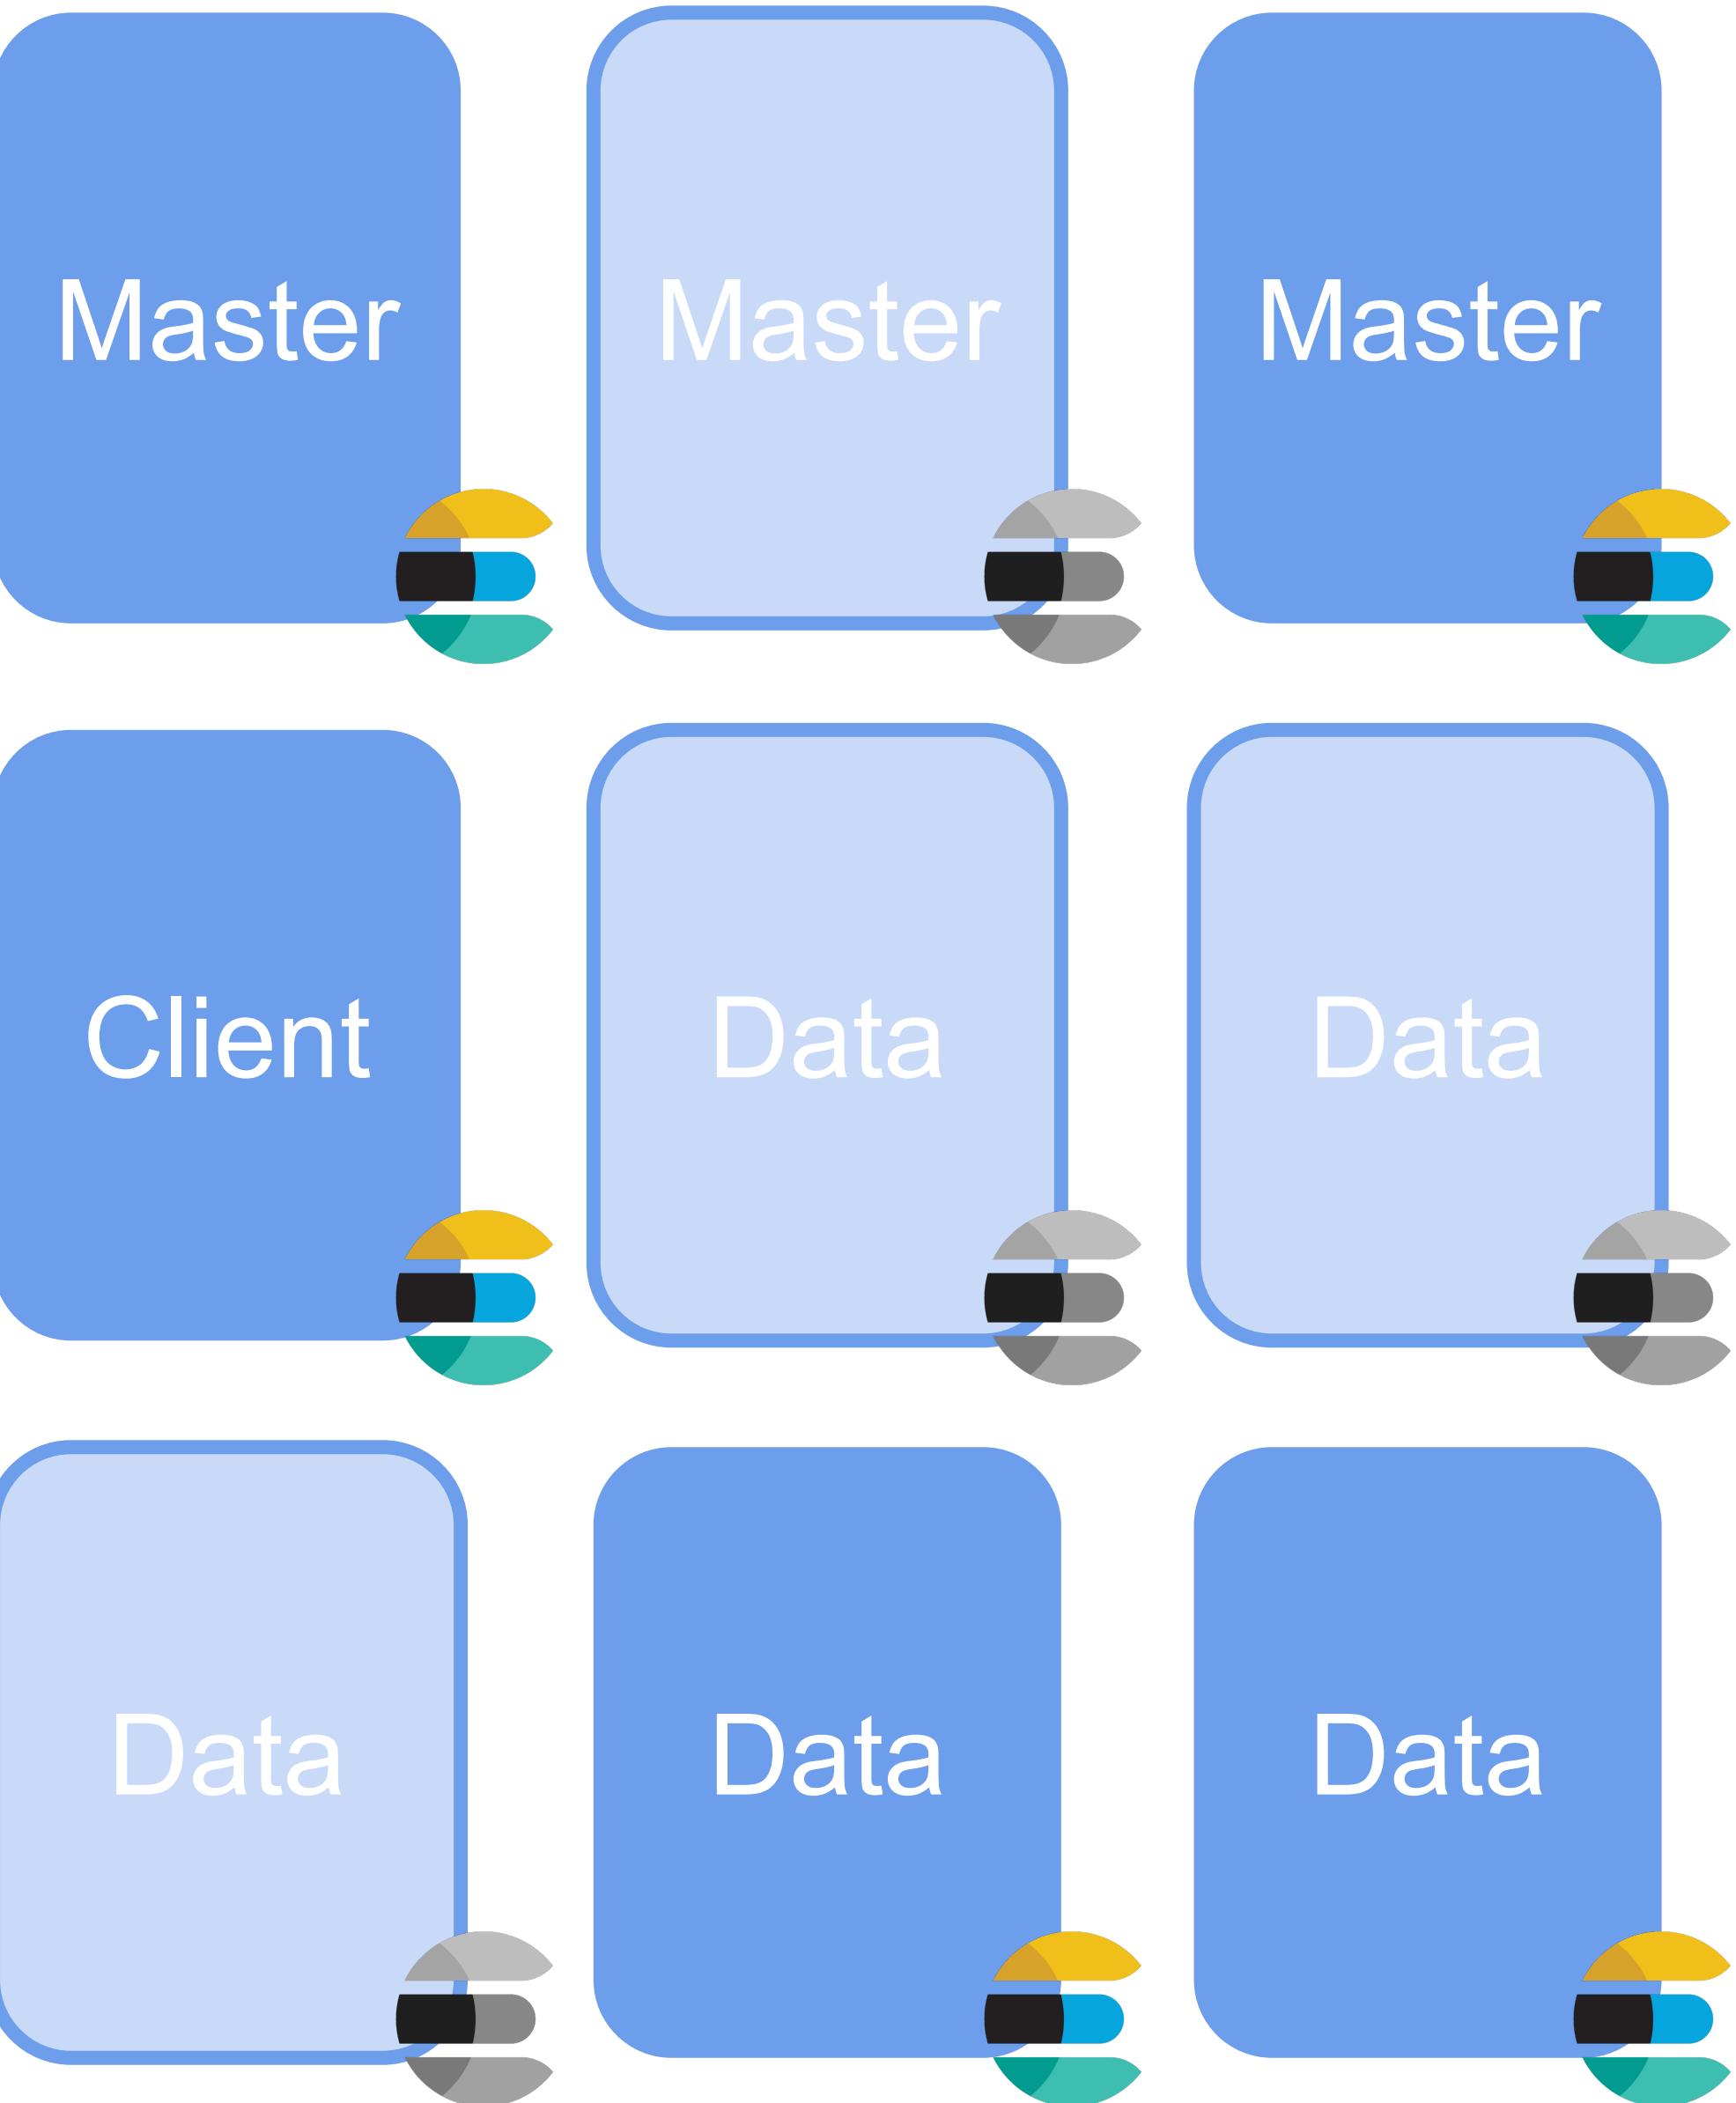


Driver

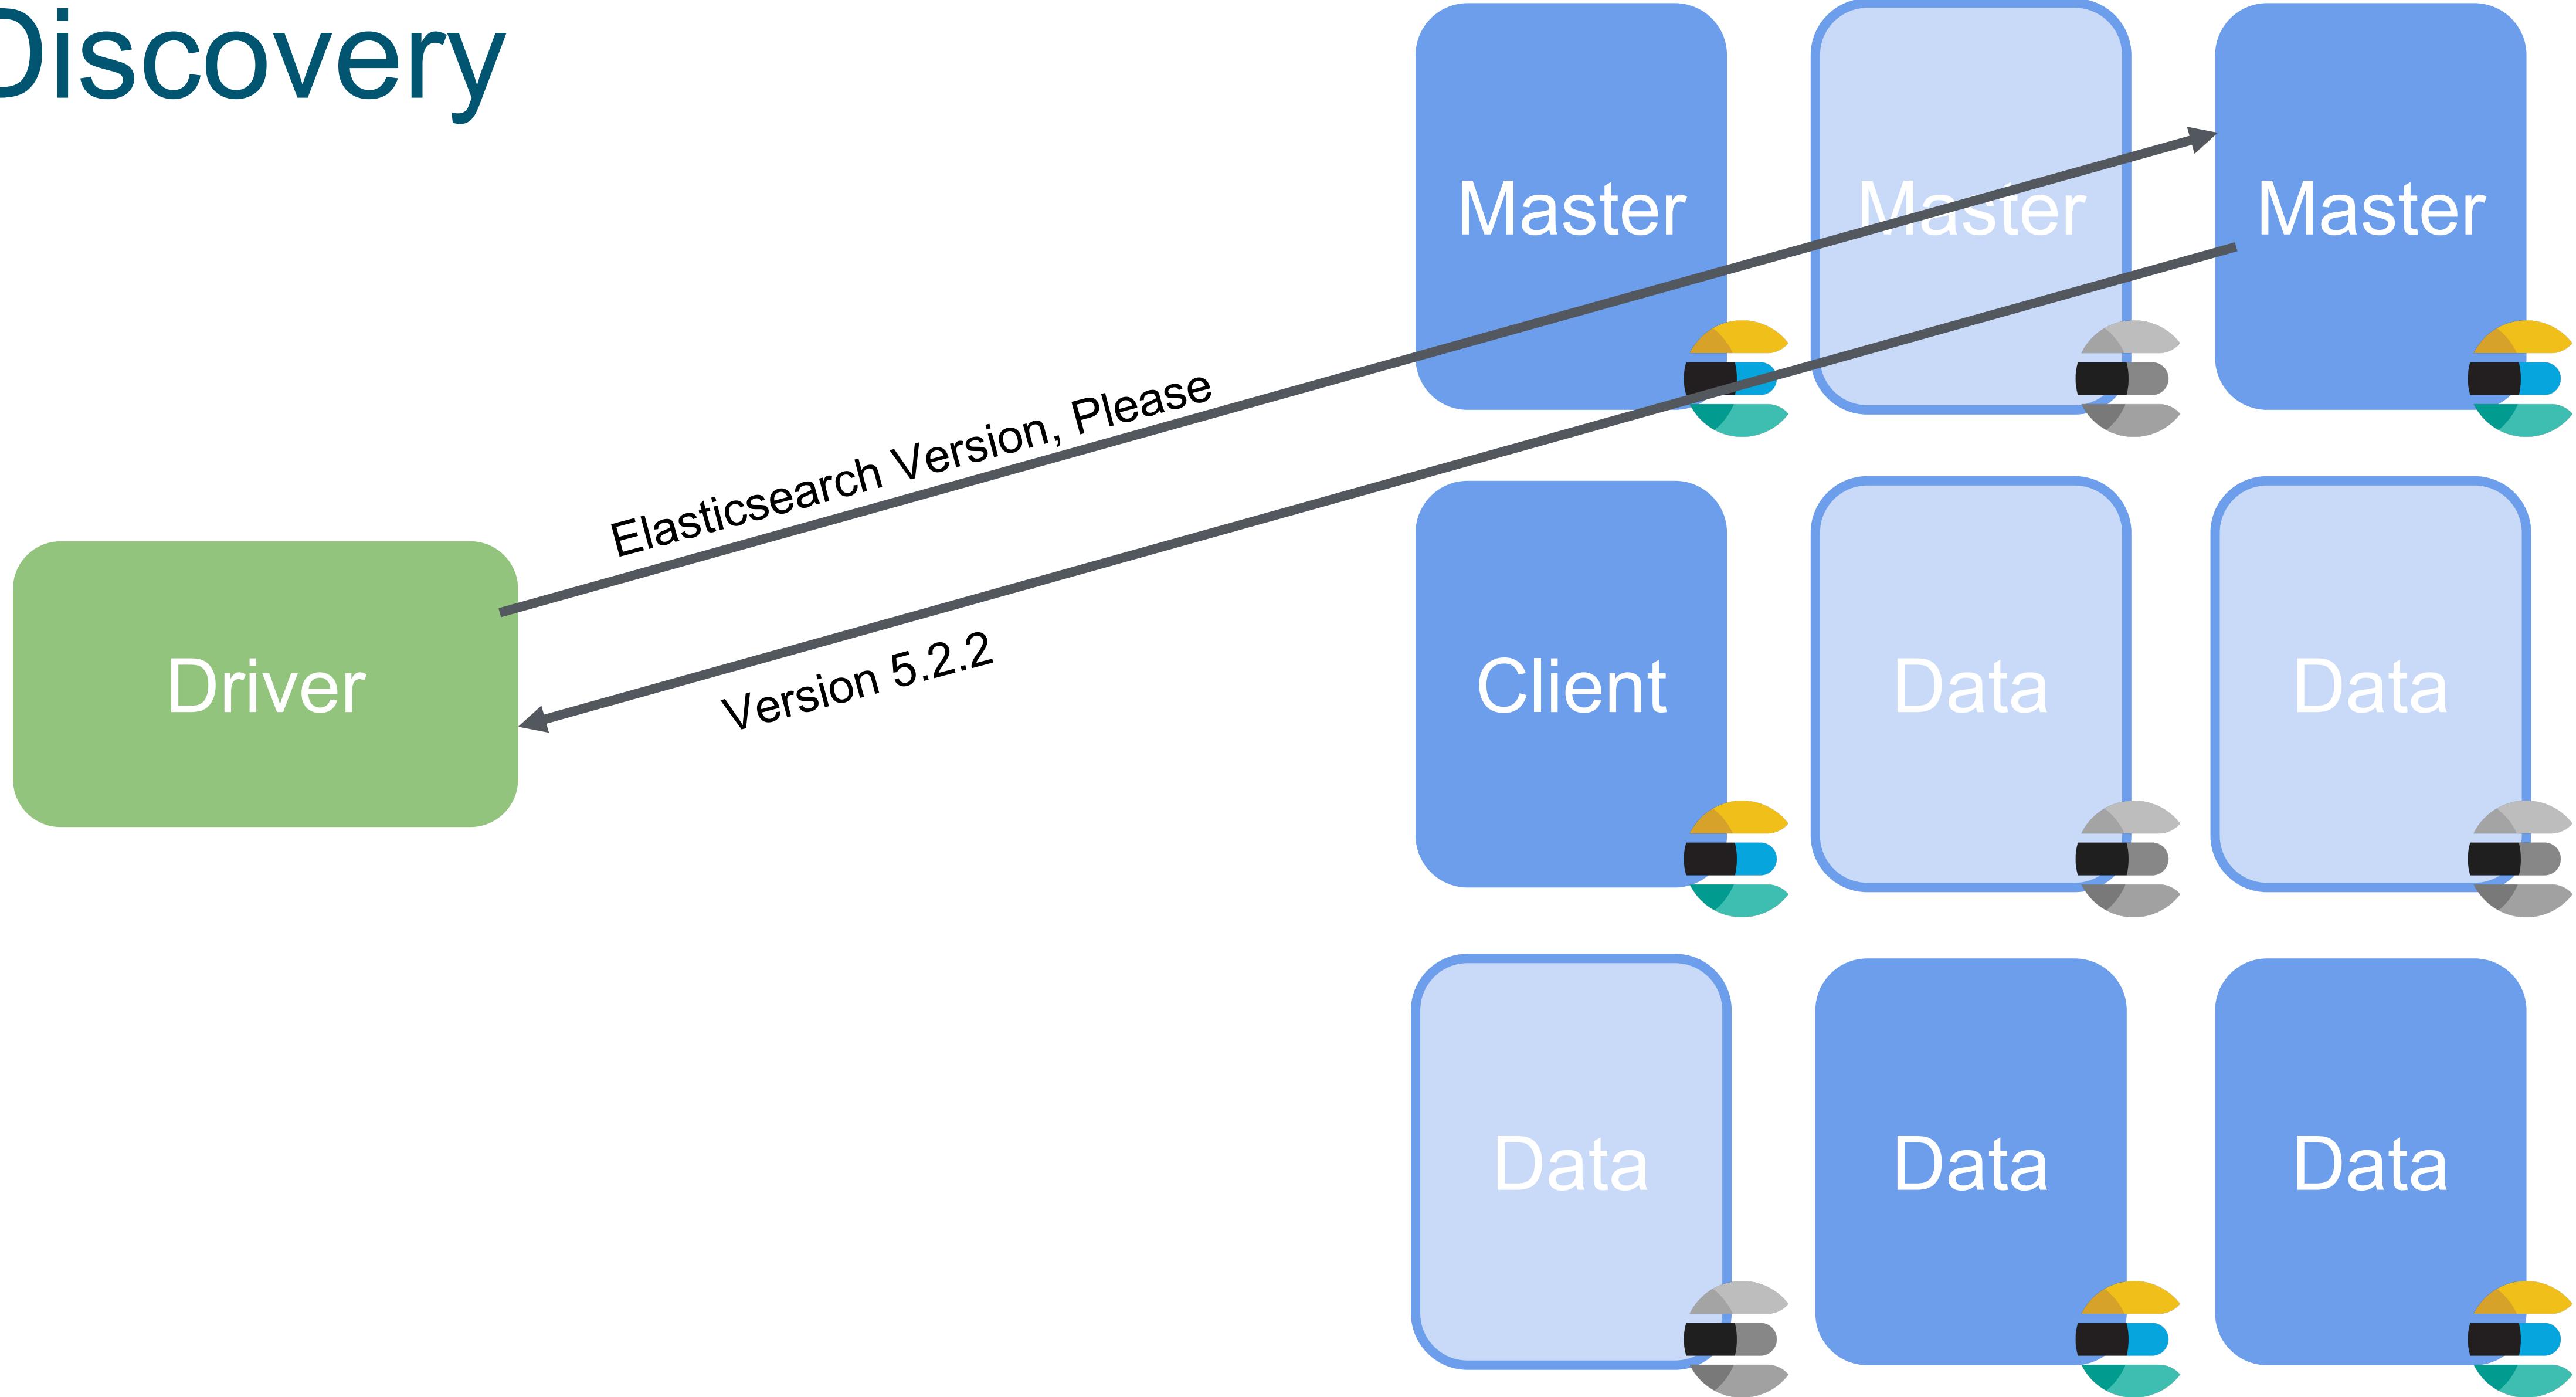


Initial Nodes Set

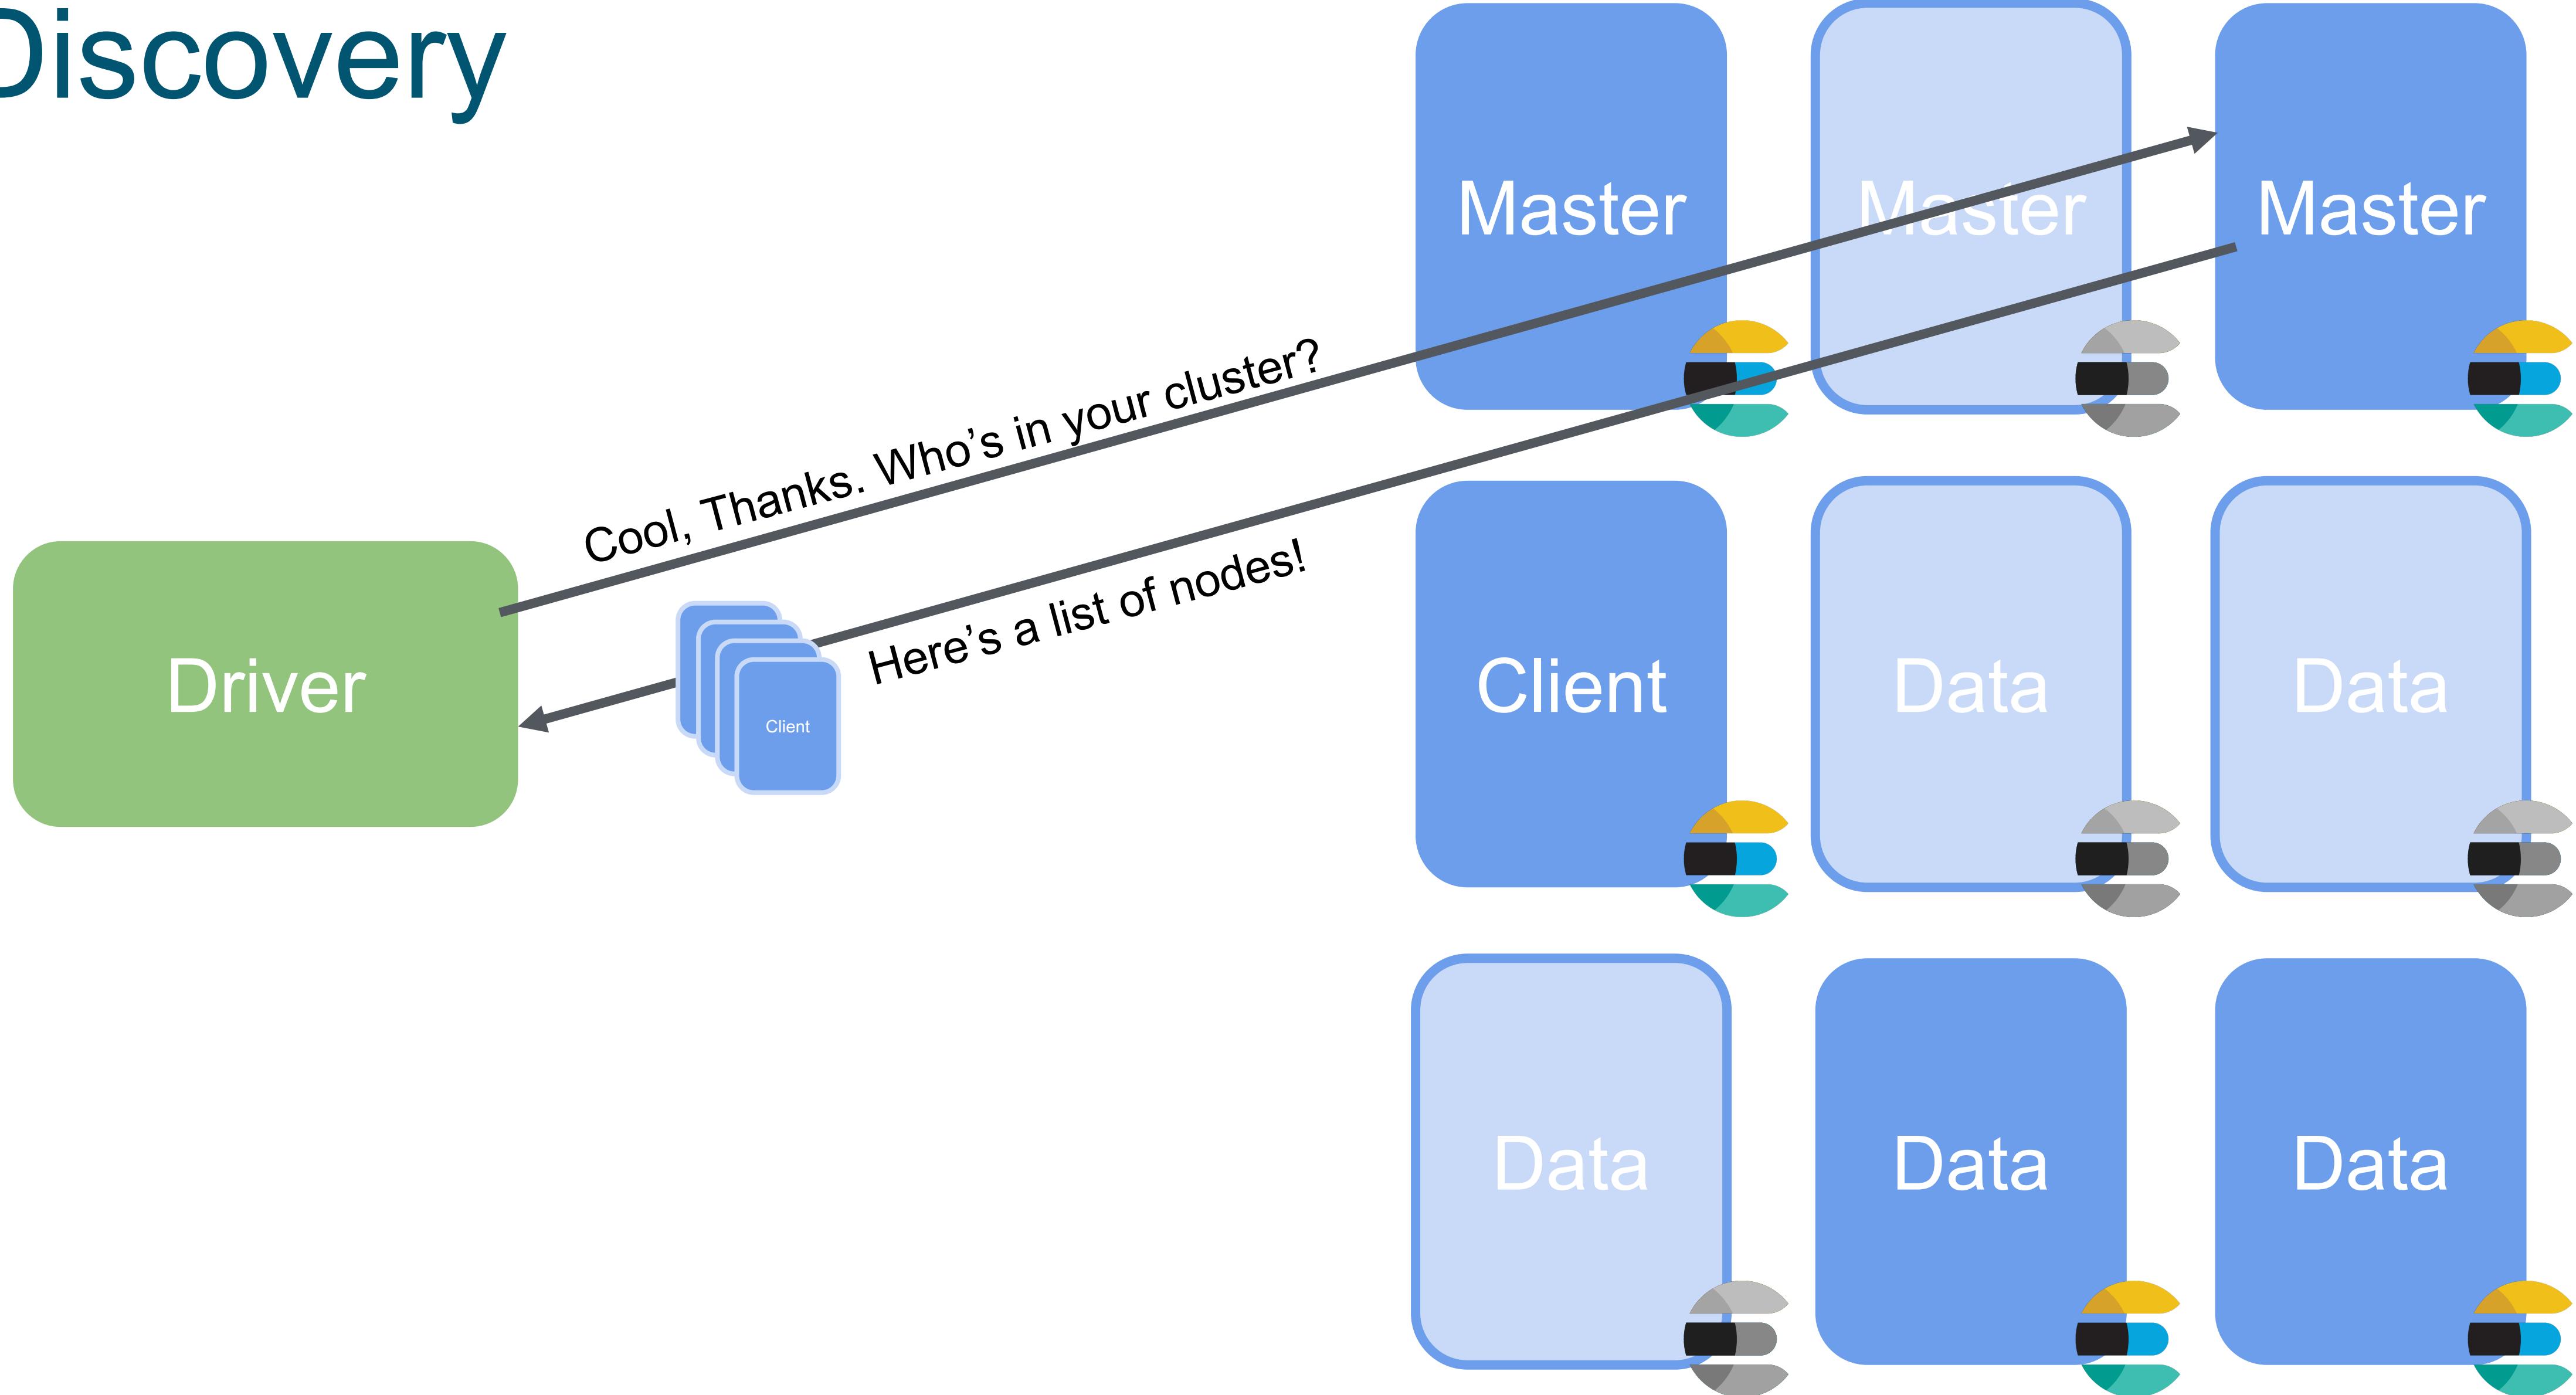
Driver



Discovery



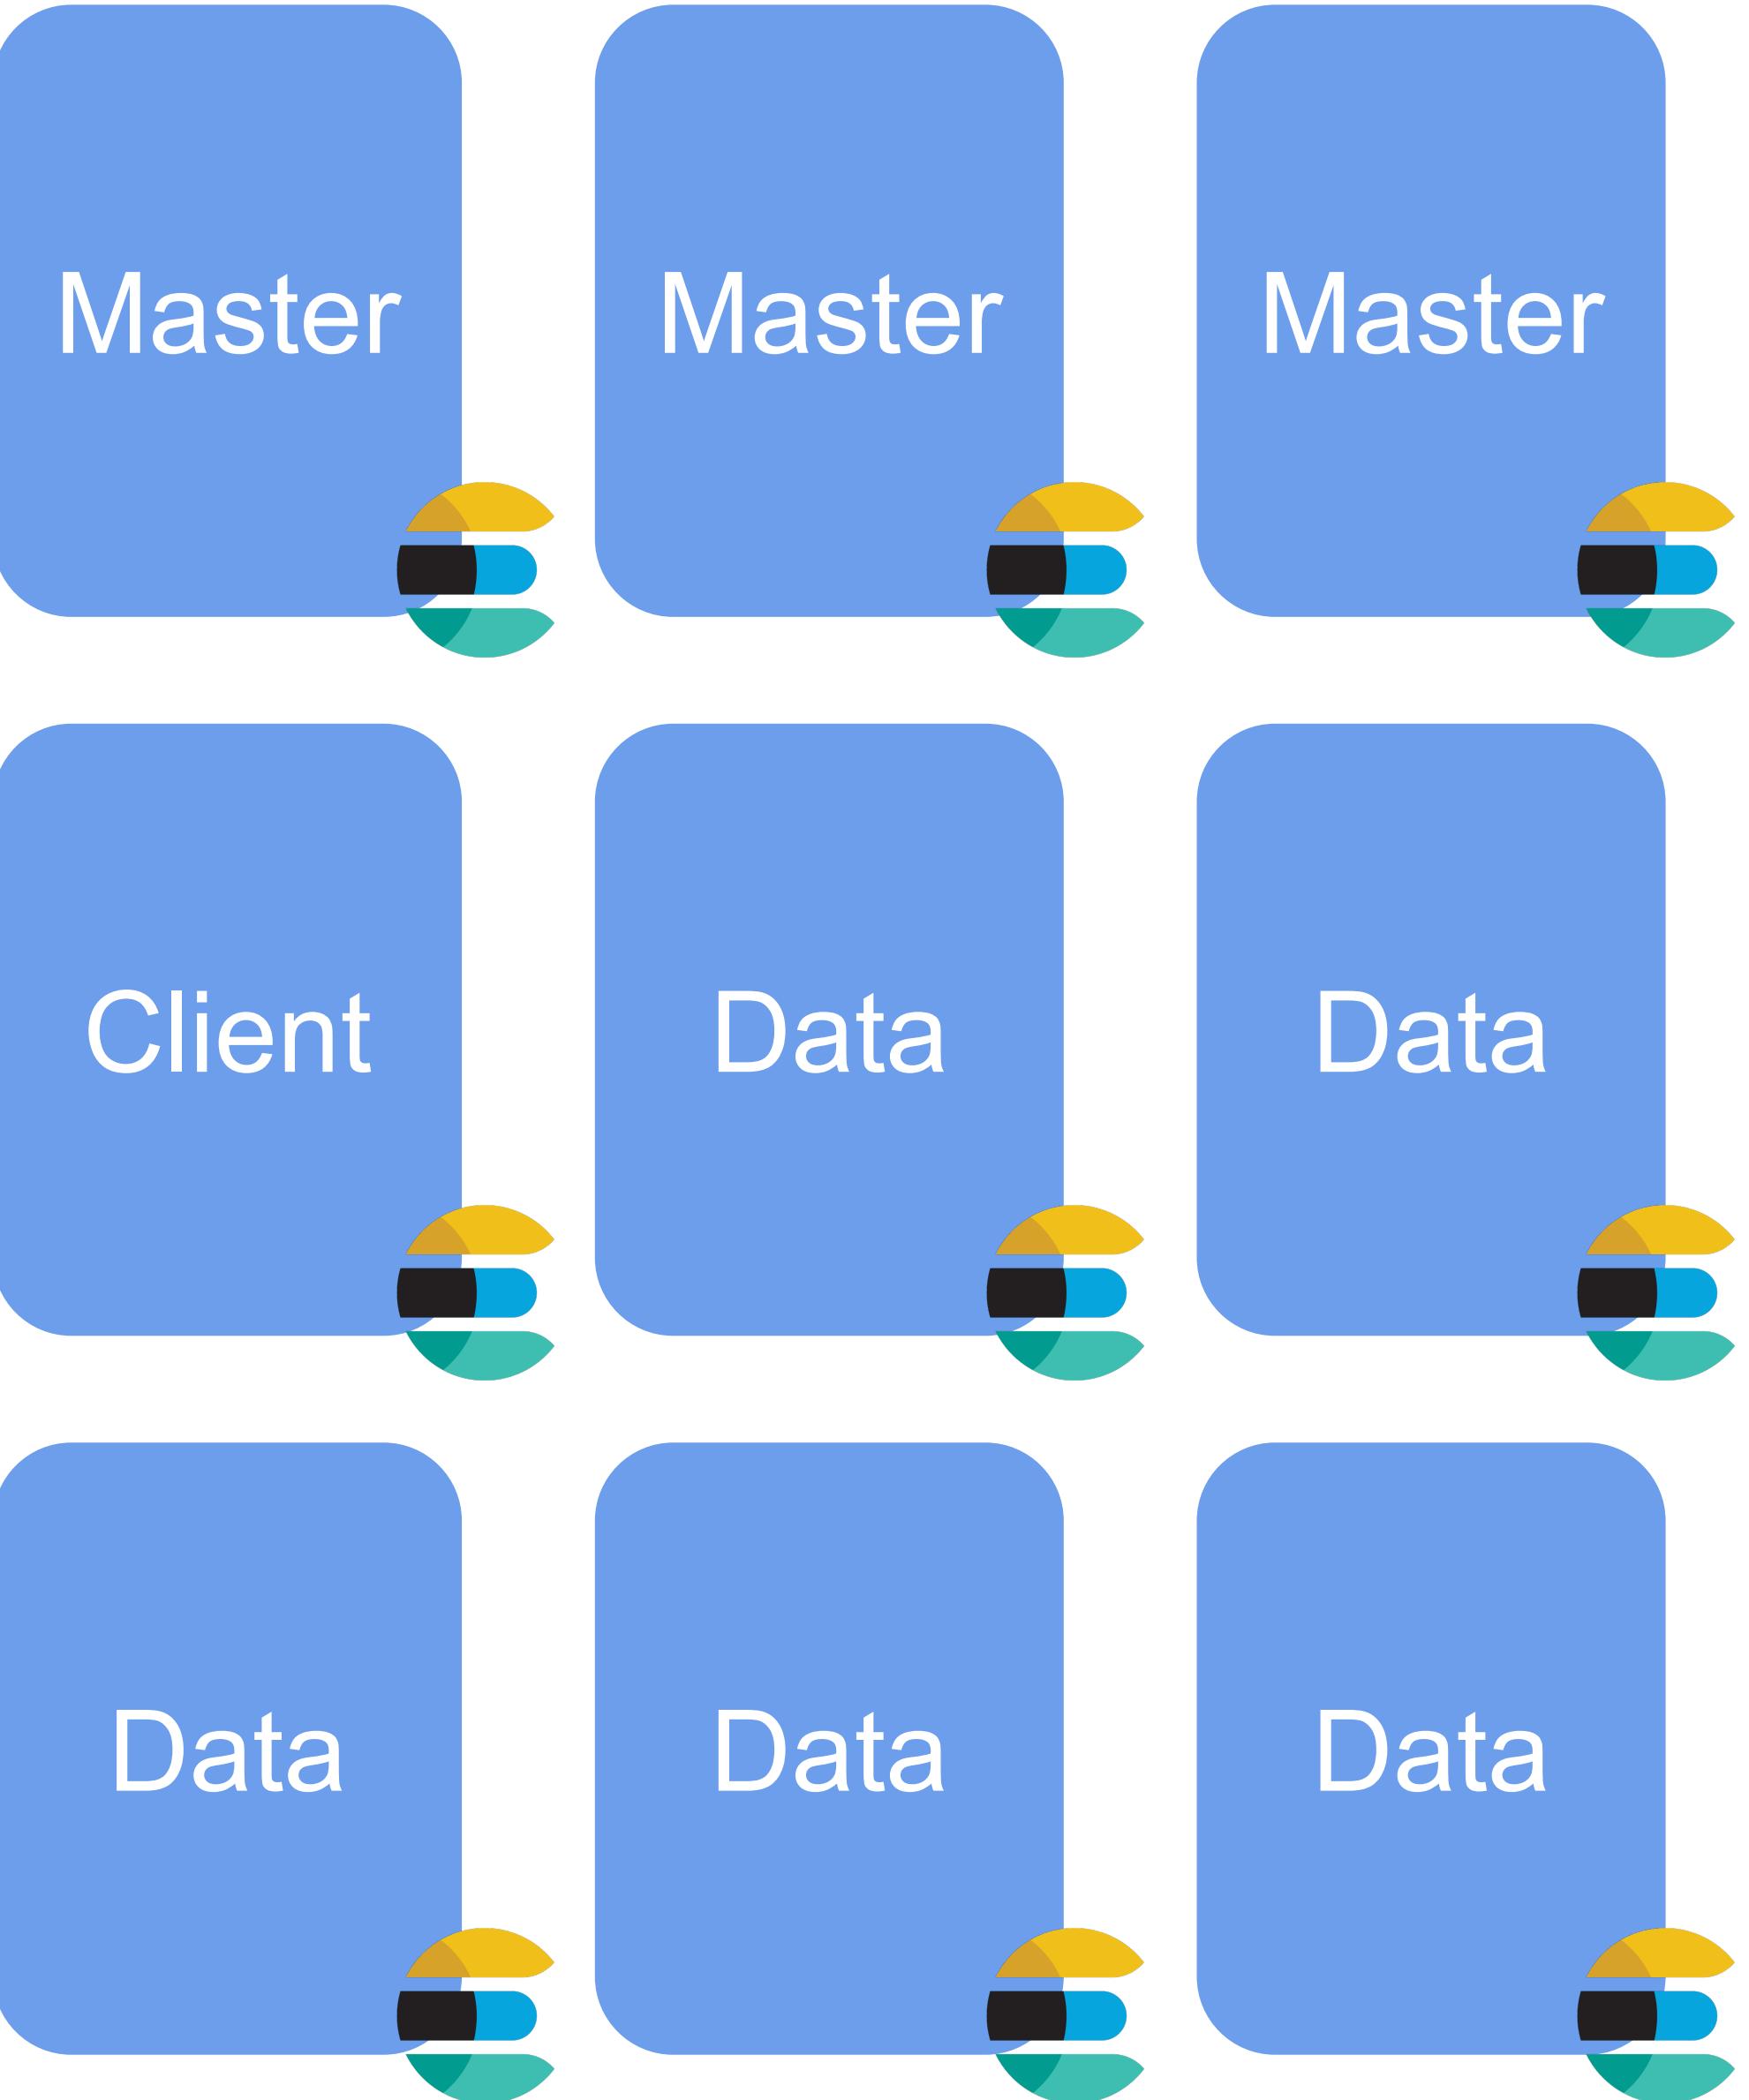
Discovery



Discovery



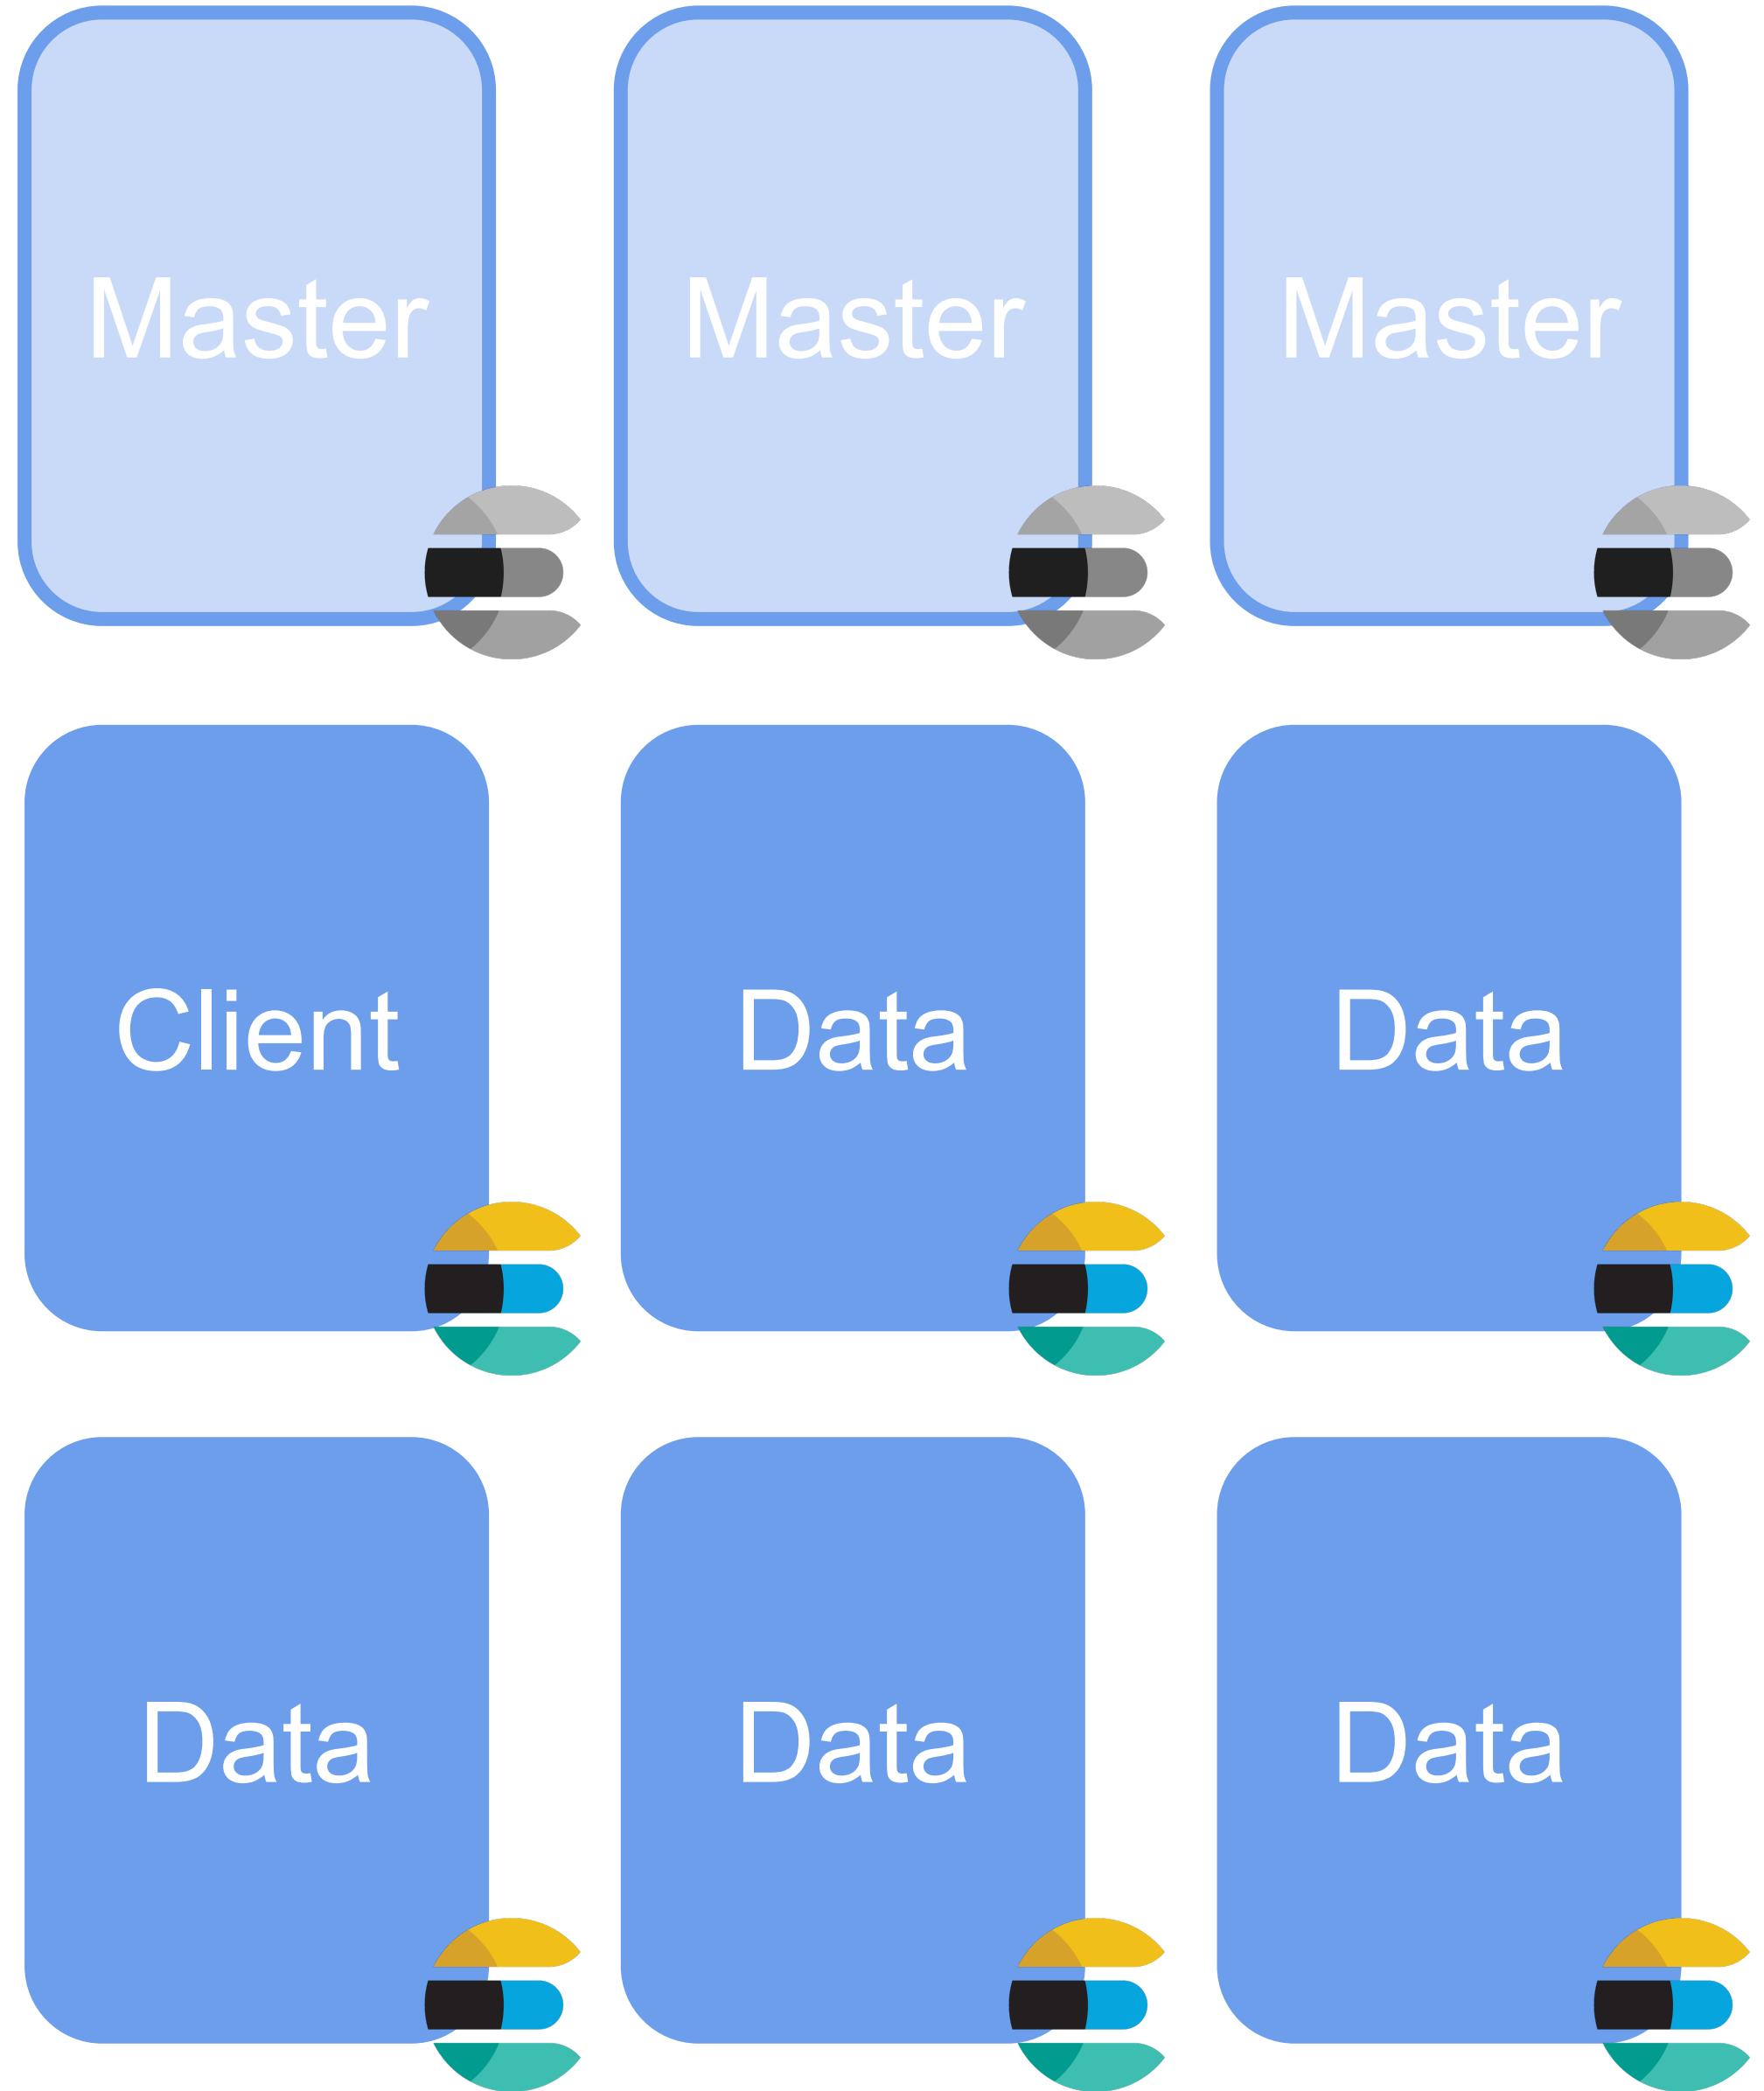
Wow, look at all these nodes.



Discovery

Driver

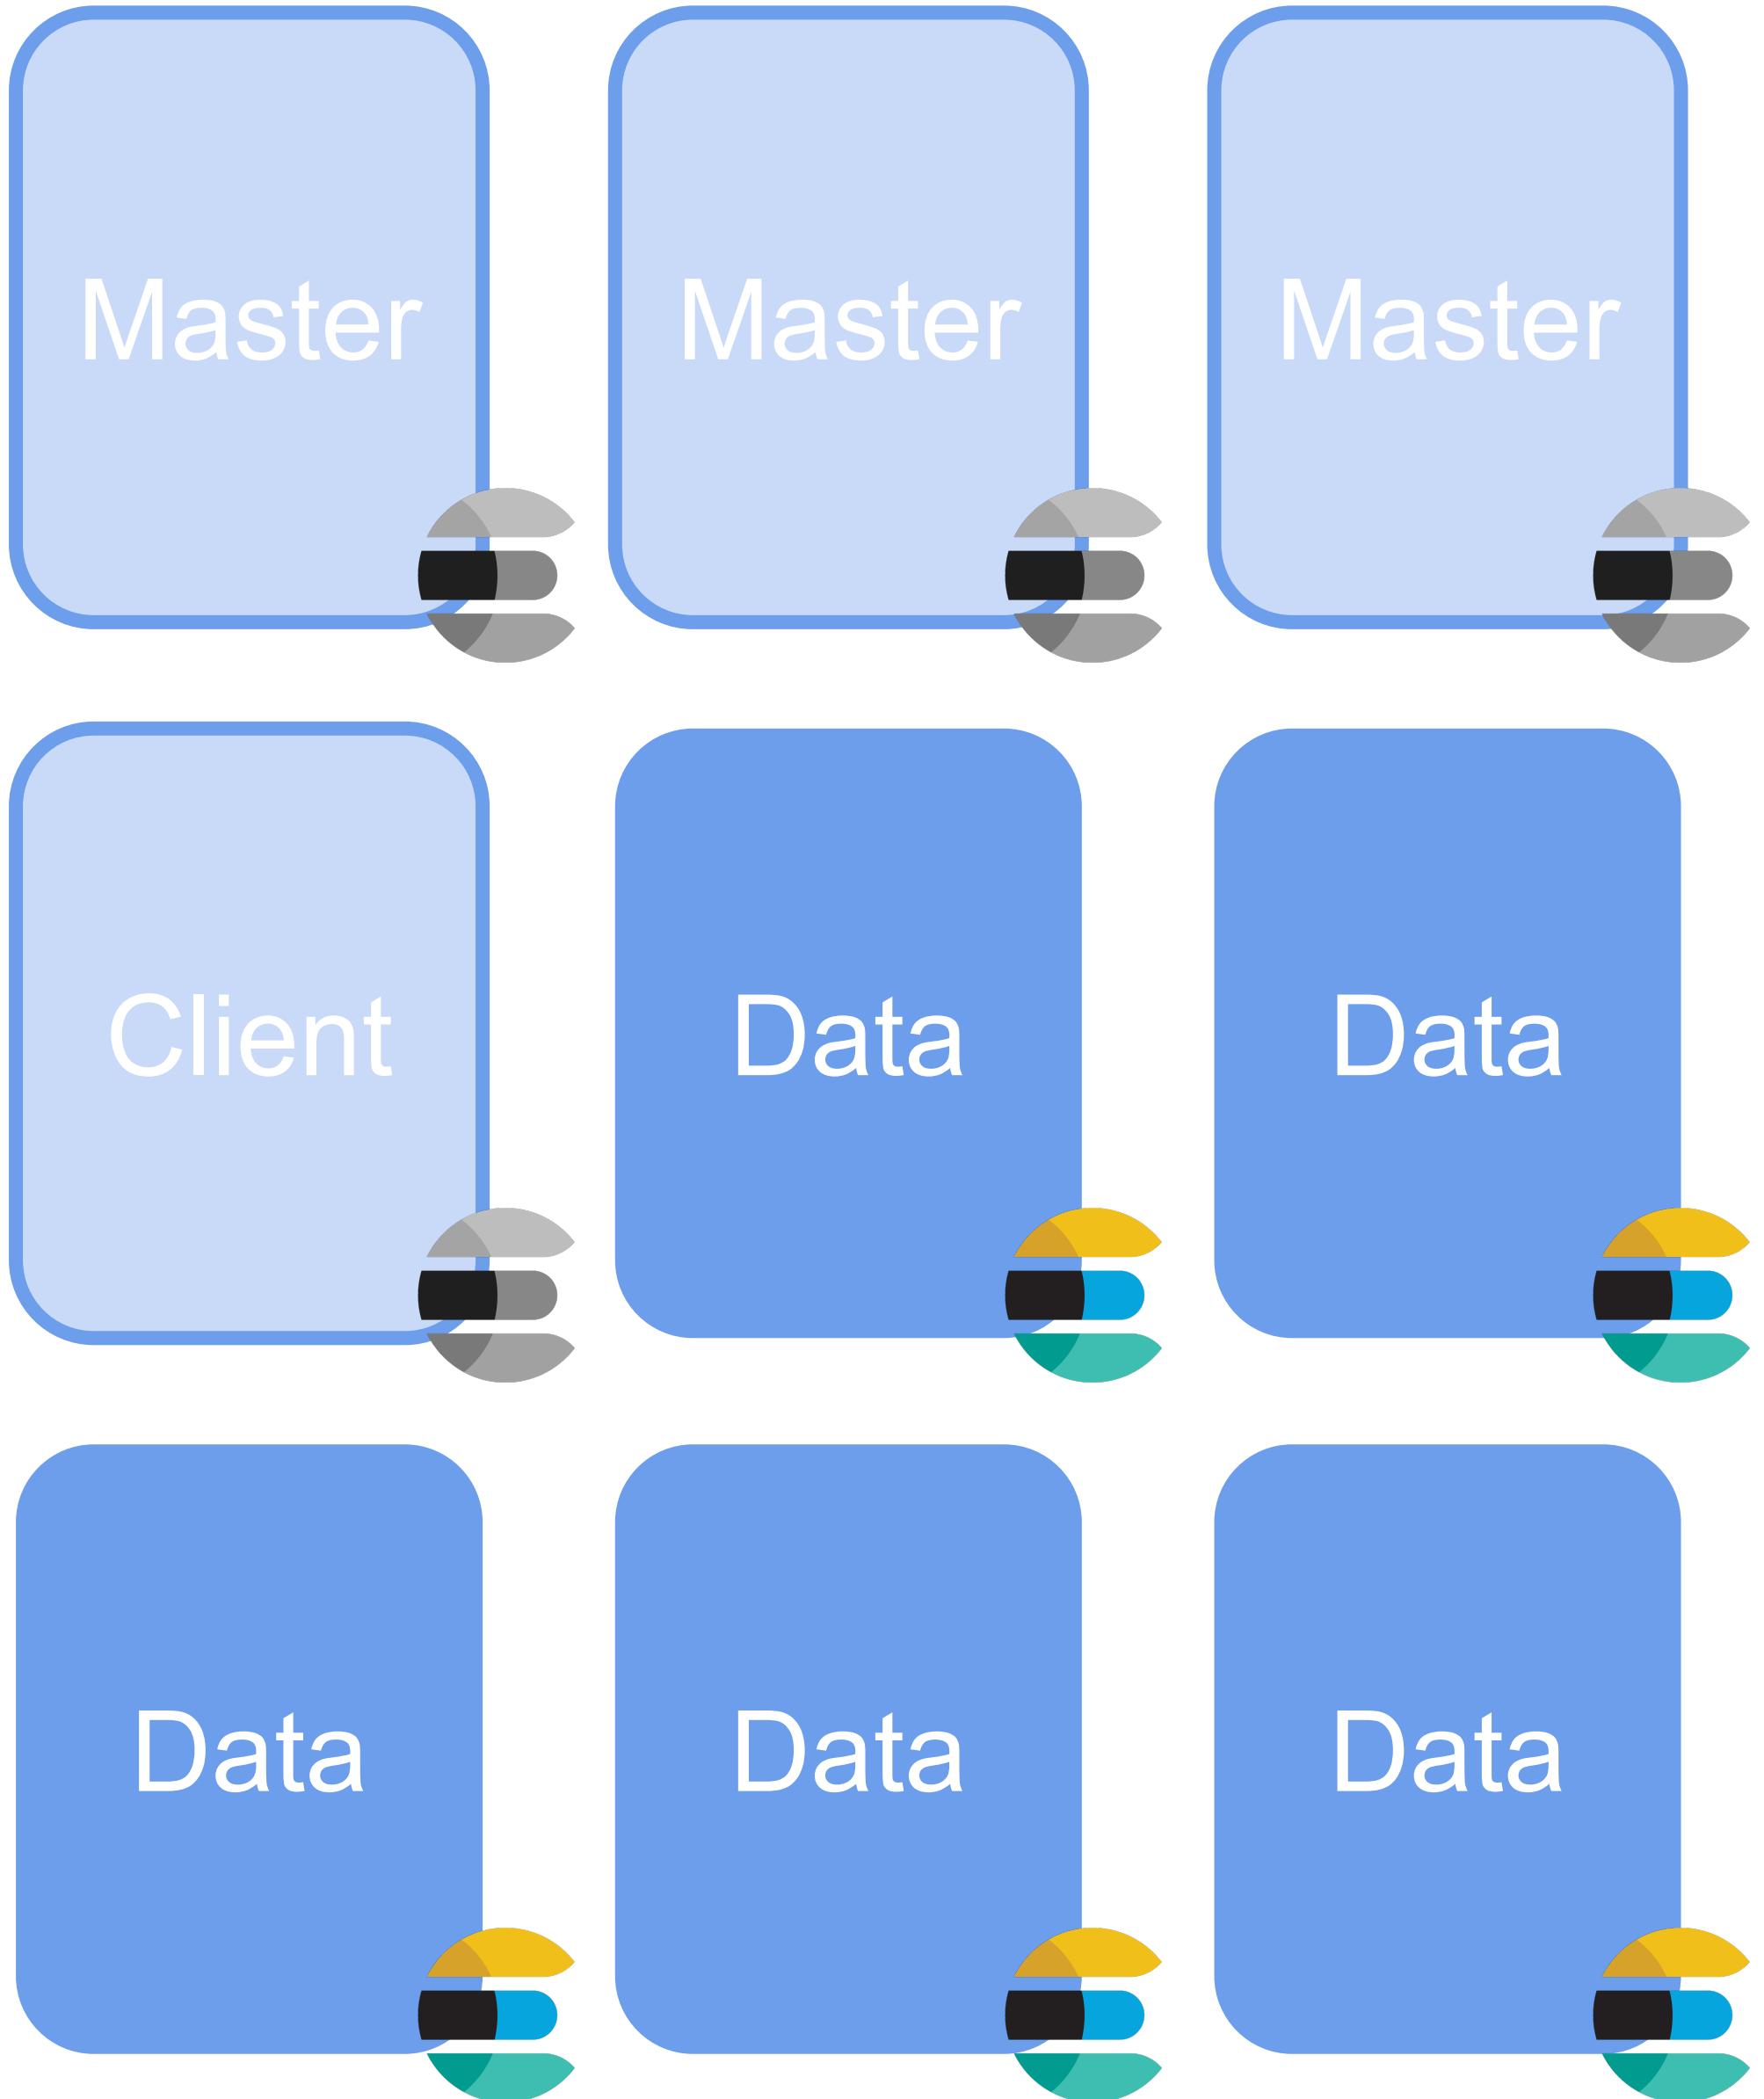
Don't need to talk to these...



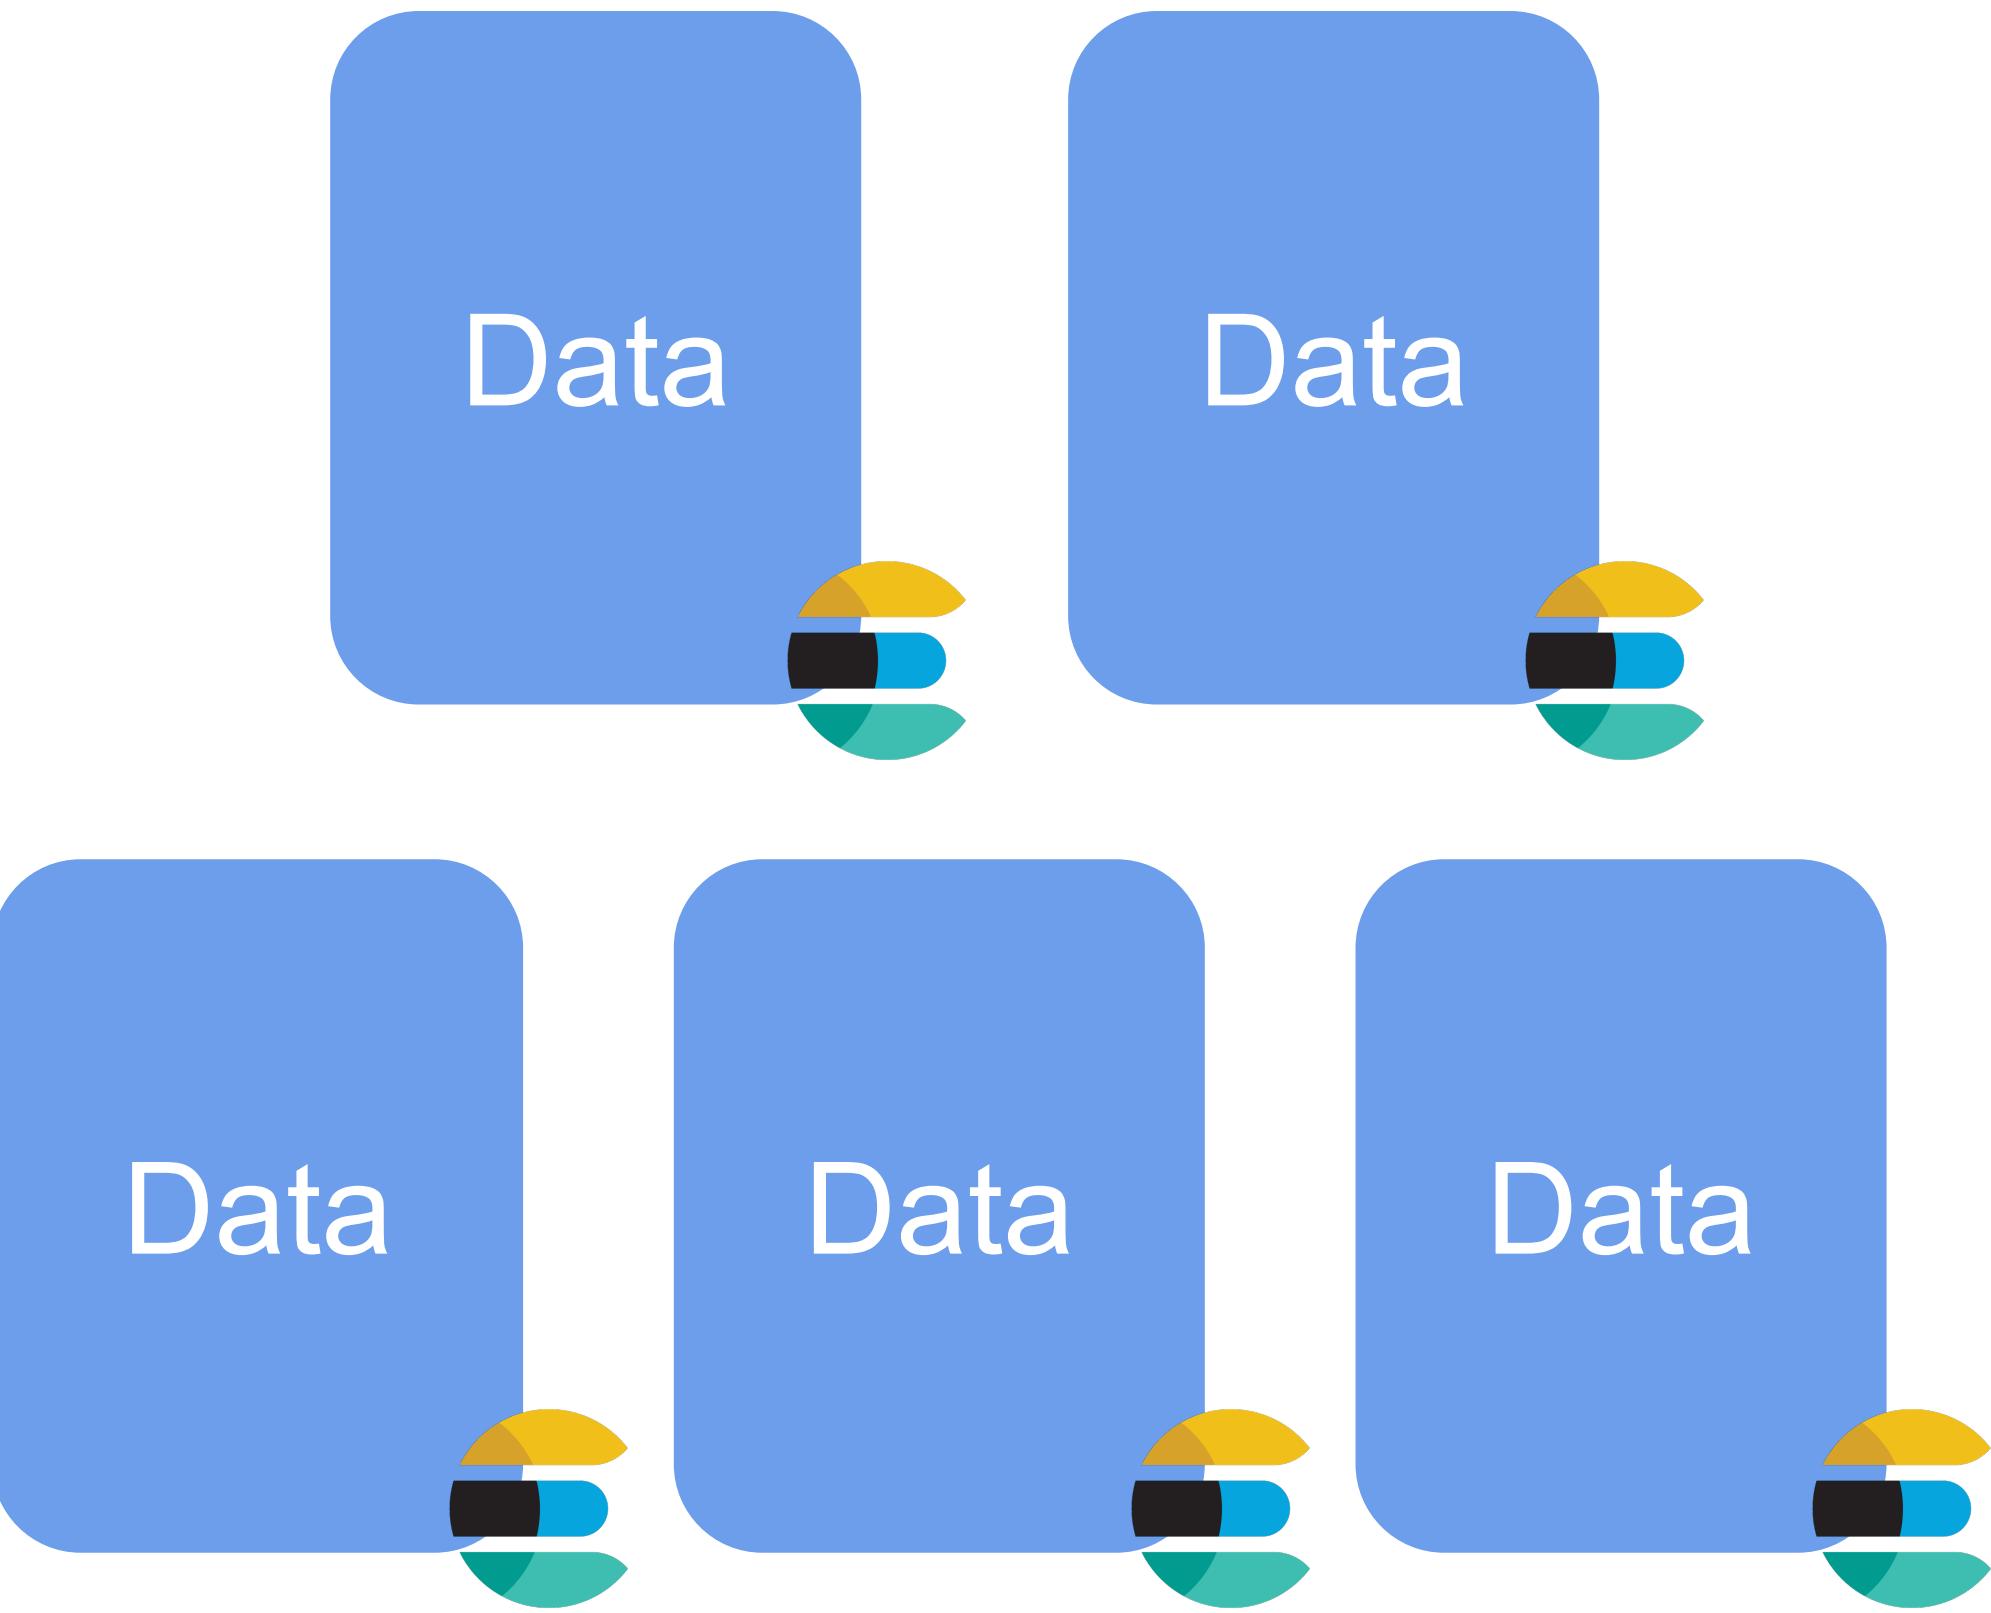
Discovery

Driver

Nope, don't need this one either...



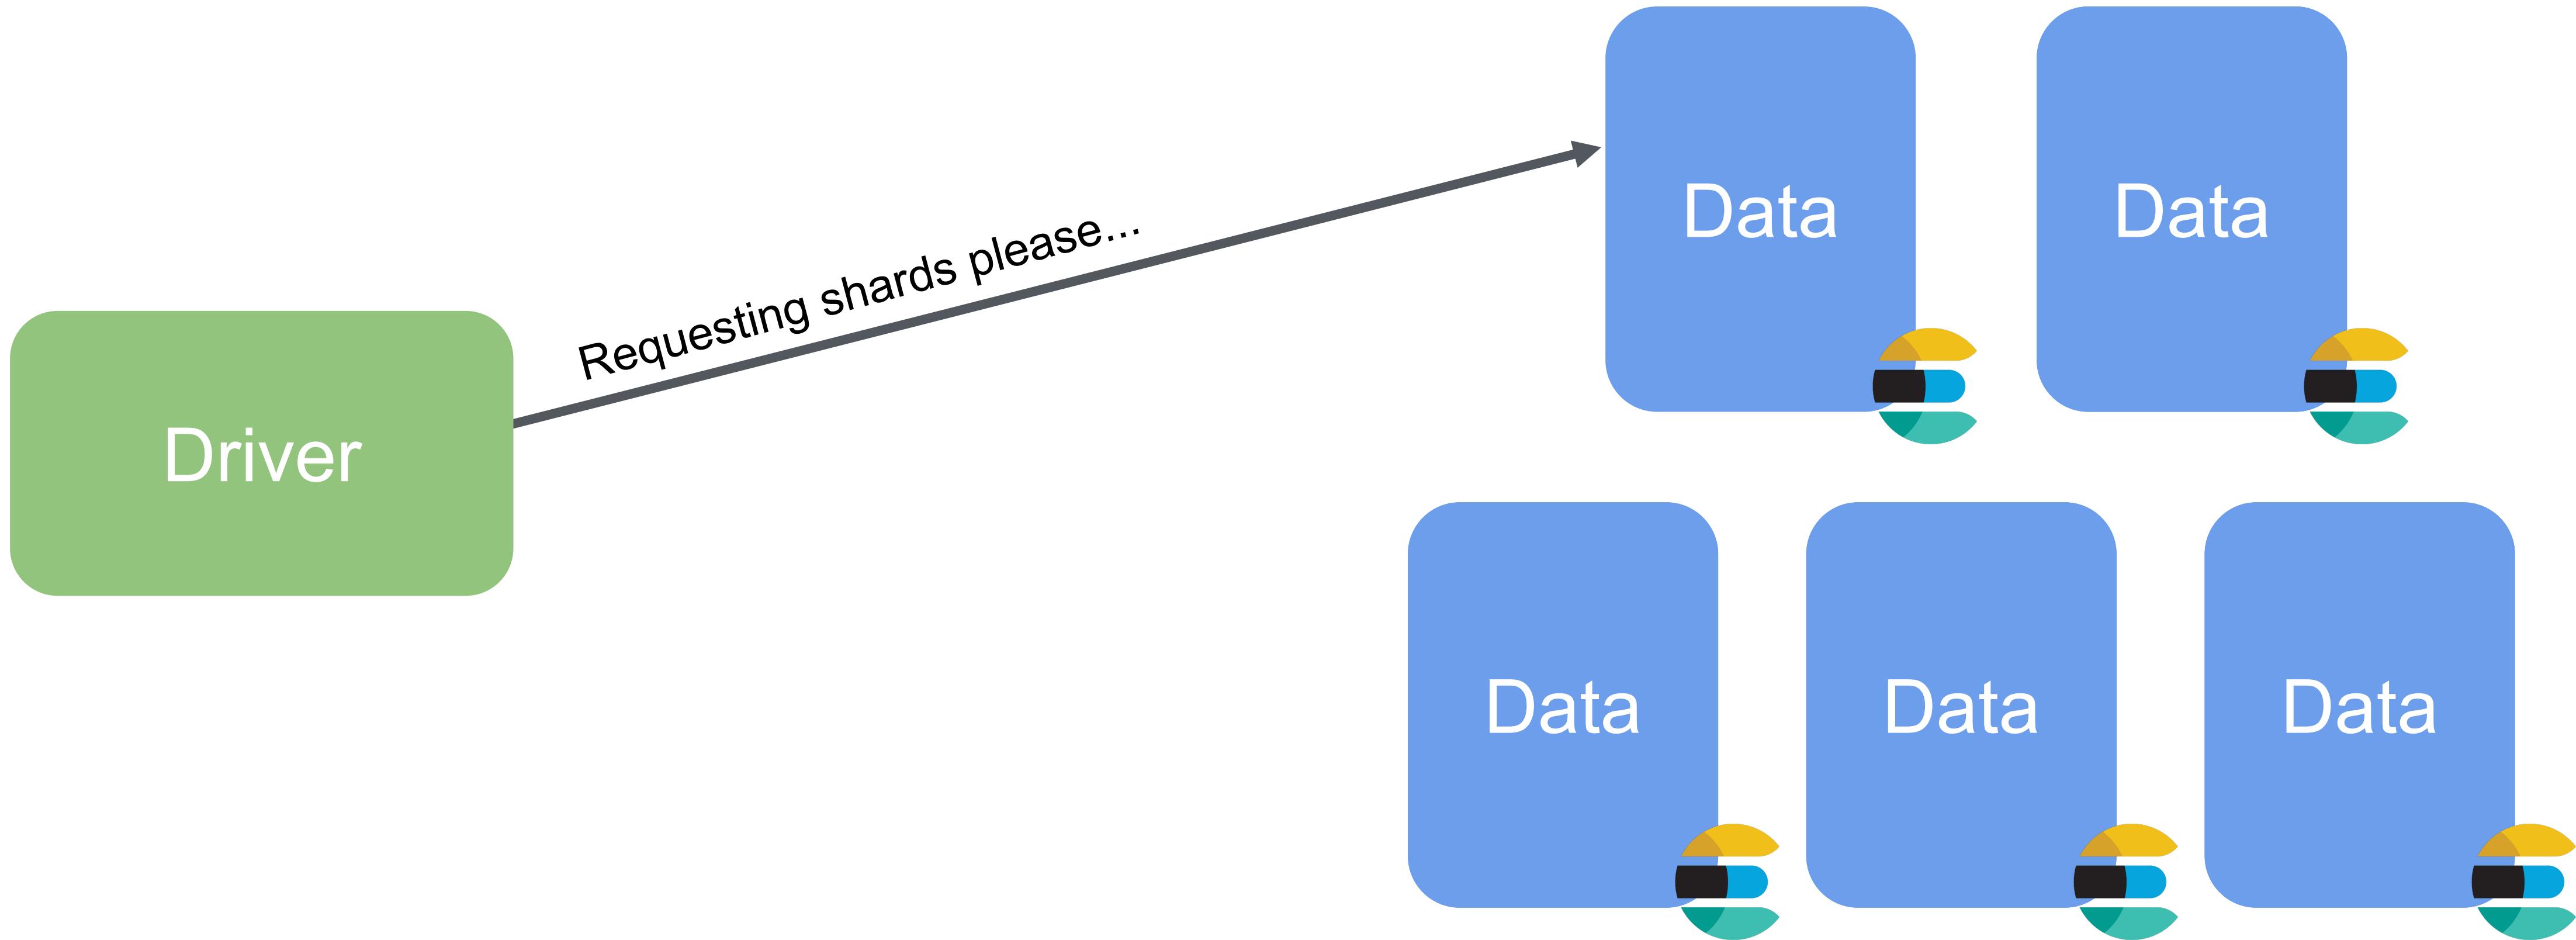
Driver



Reading from Elasticsearch

Finding Partitions

Reading index "logs/data"...



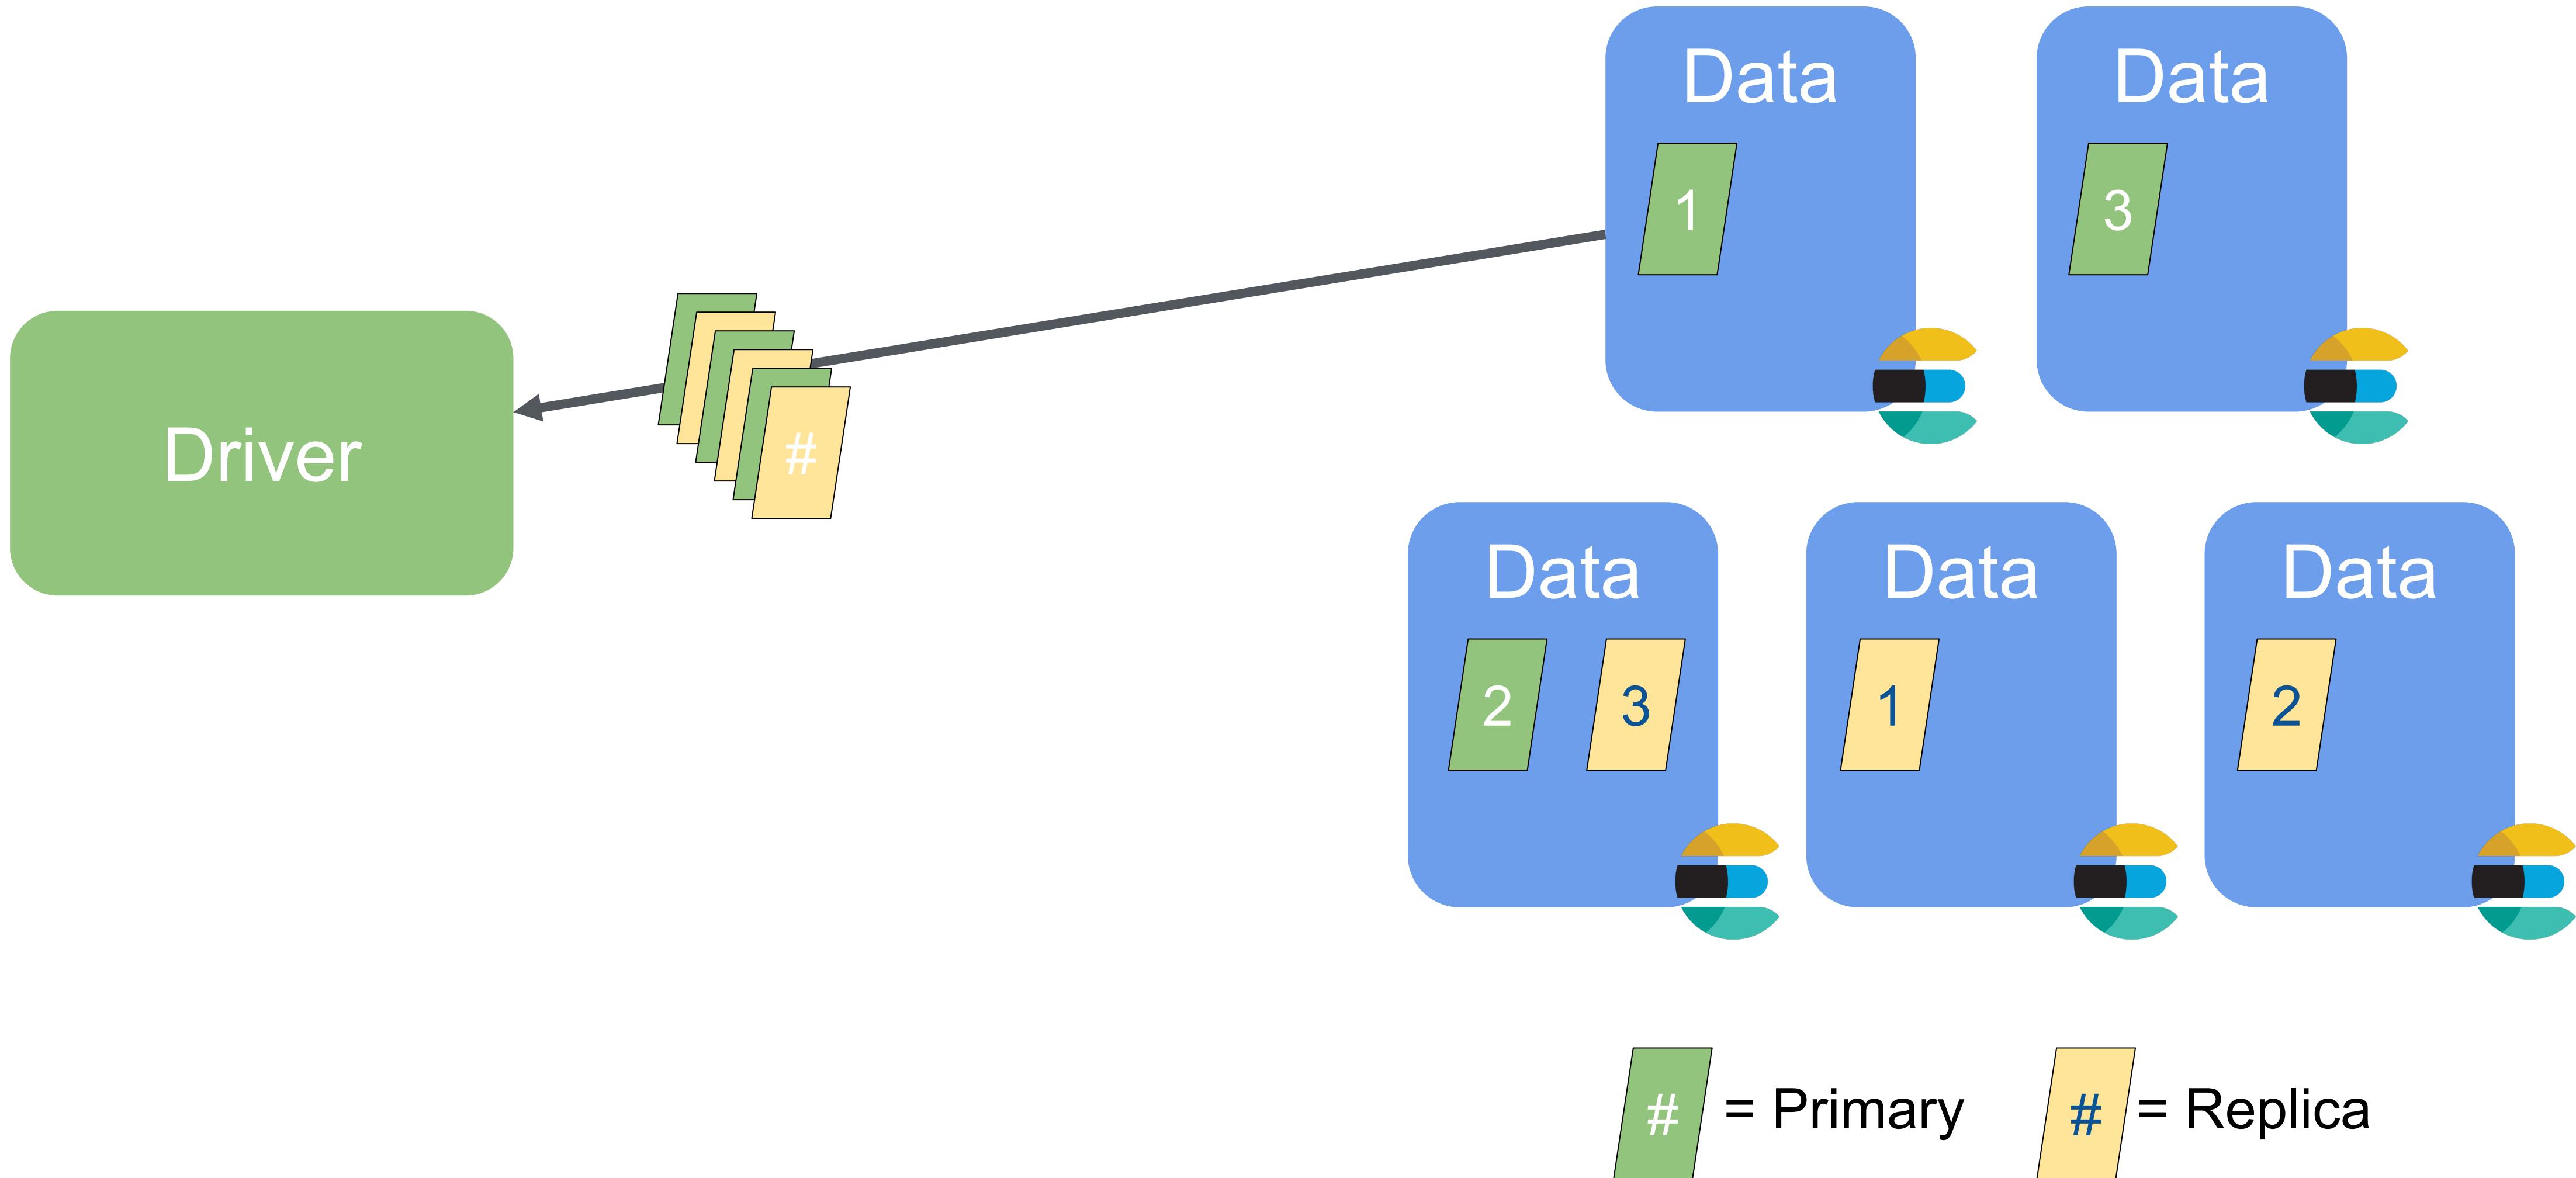
Finding Partitions

Reading index "logs/data"...



Finding Partitions

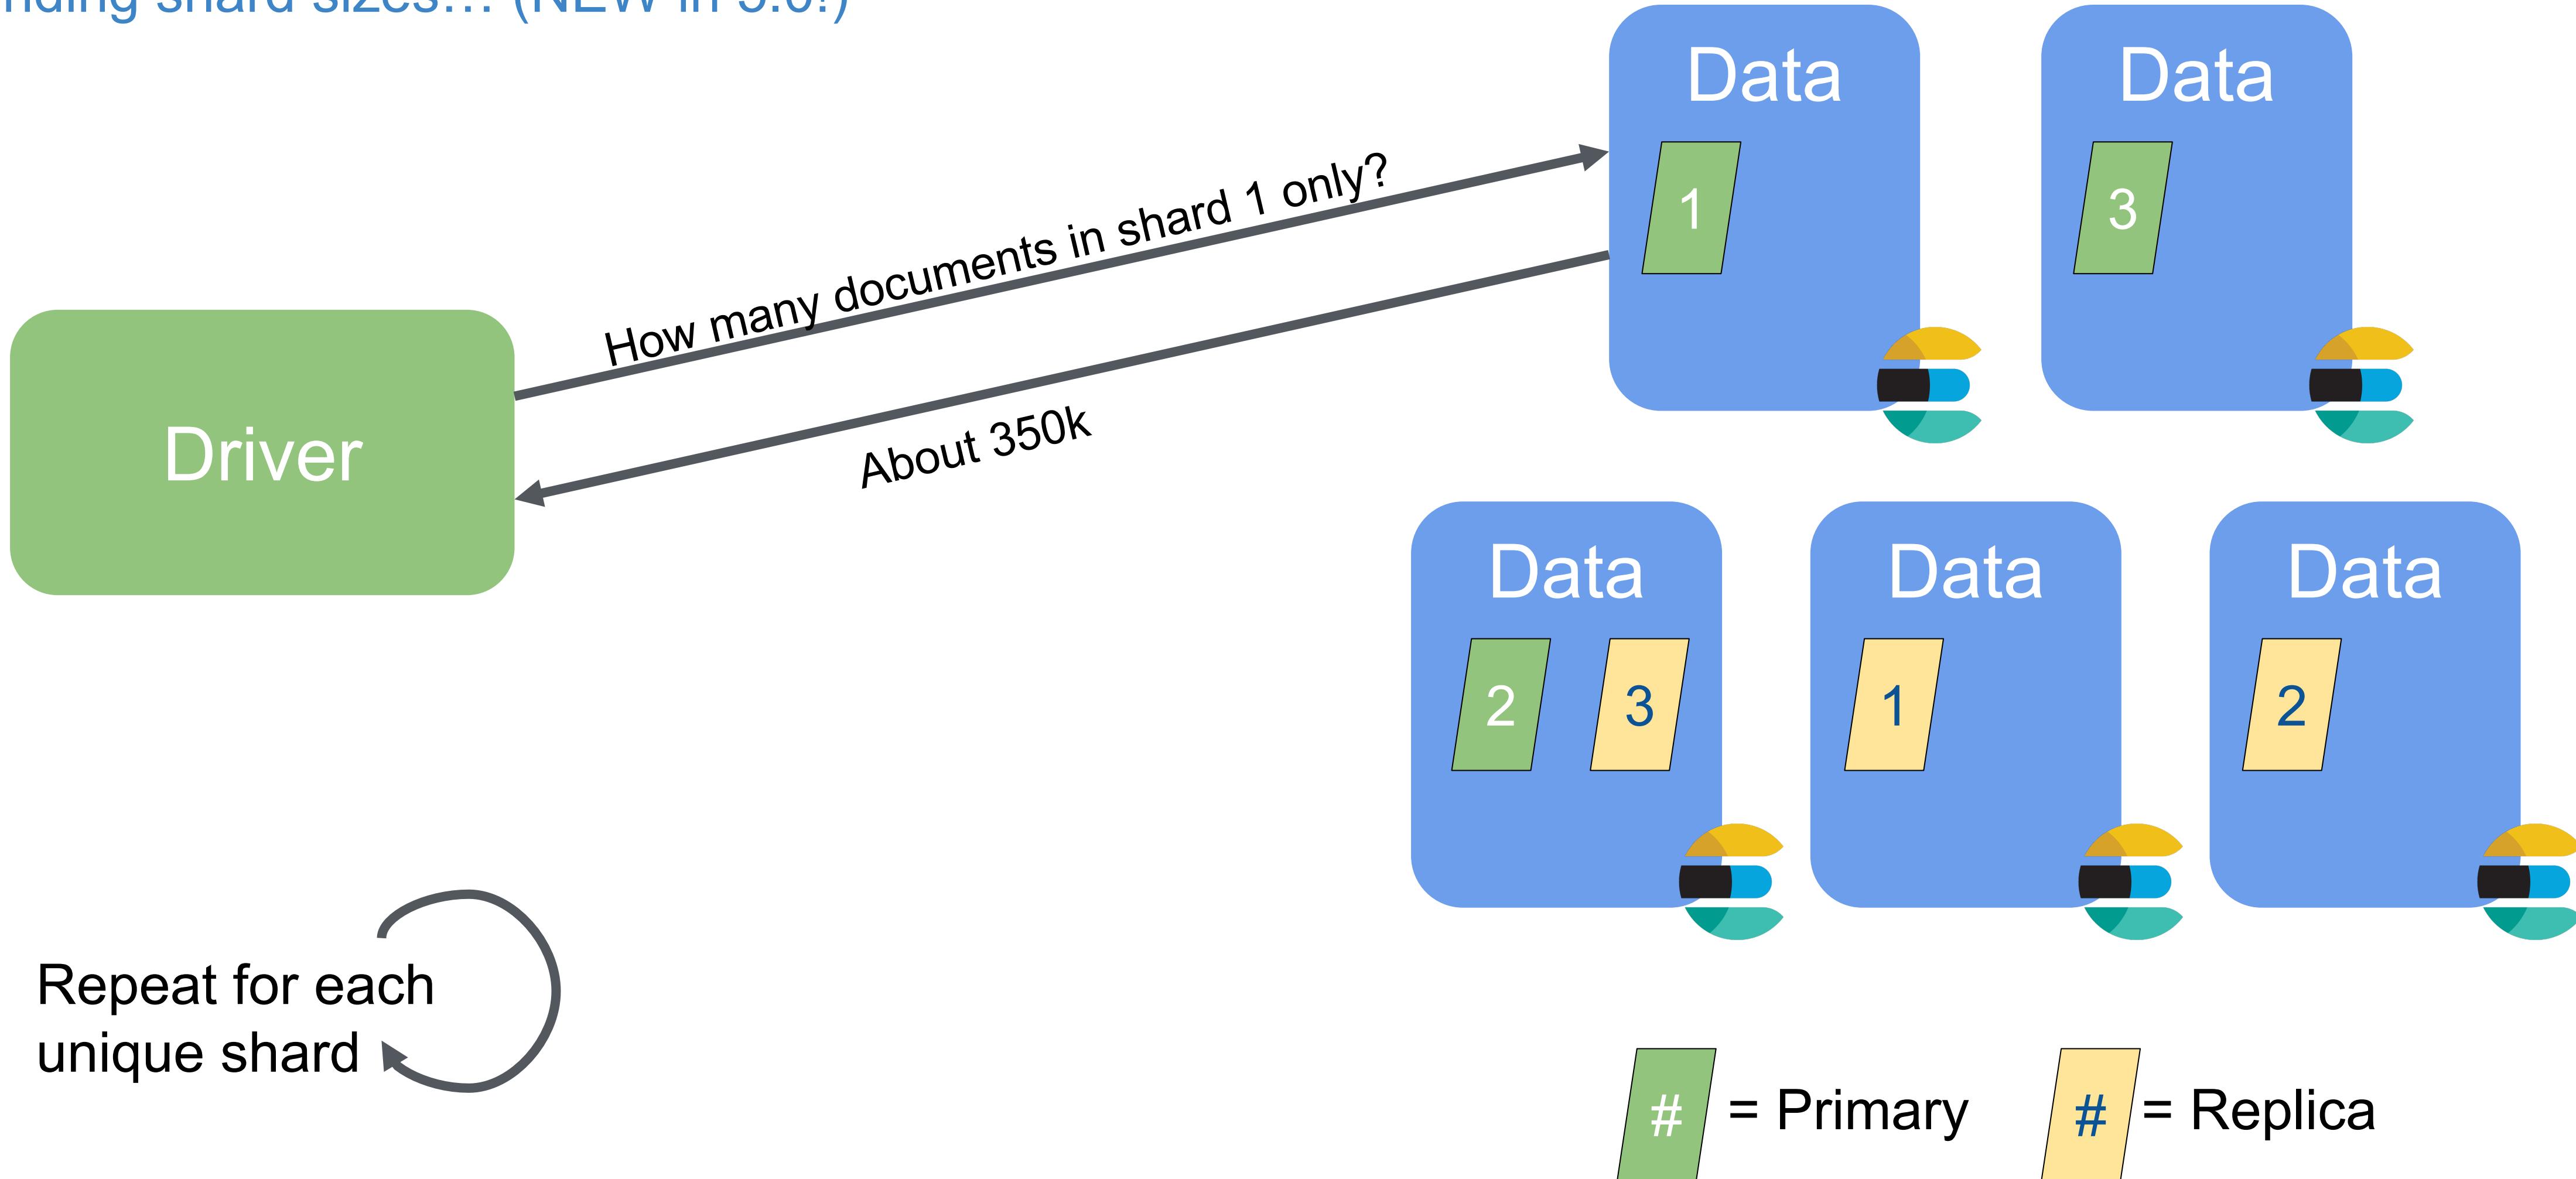
Reading index "logs/data"...



Finding Partitions

Reading index "logs/data"...

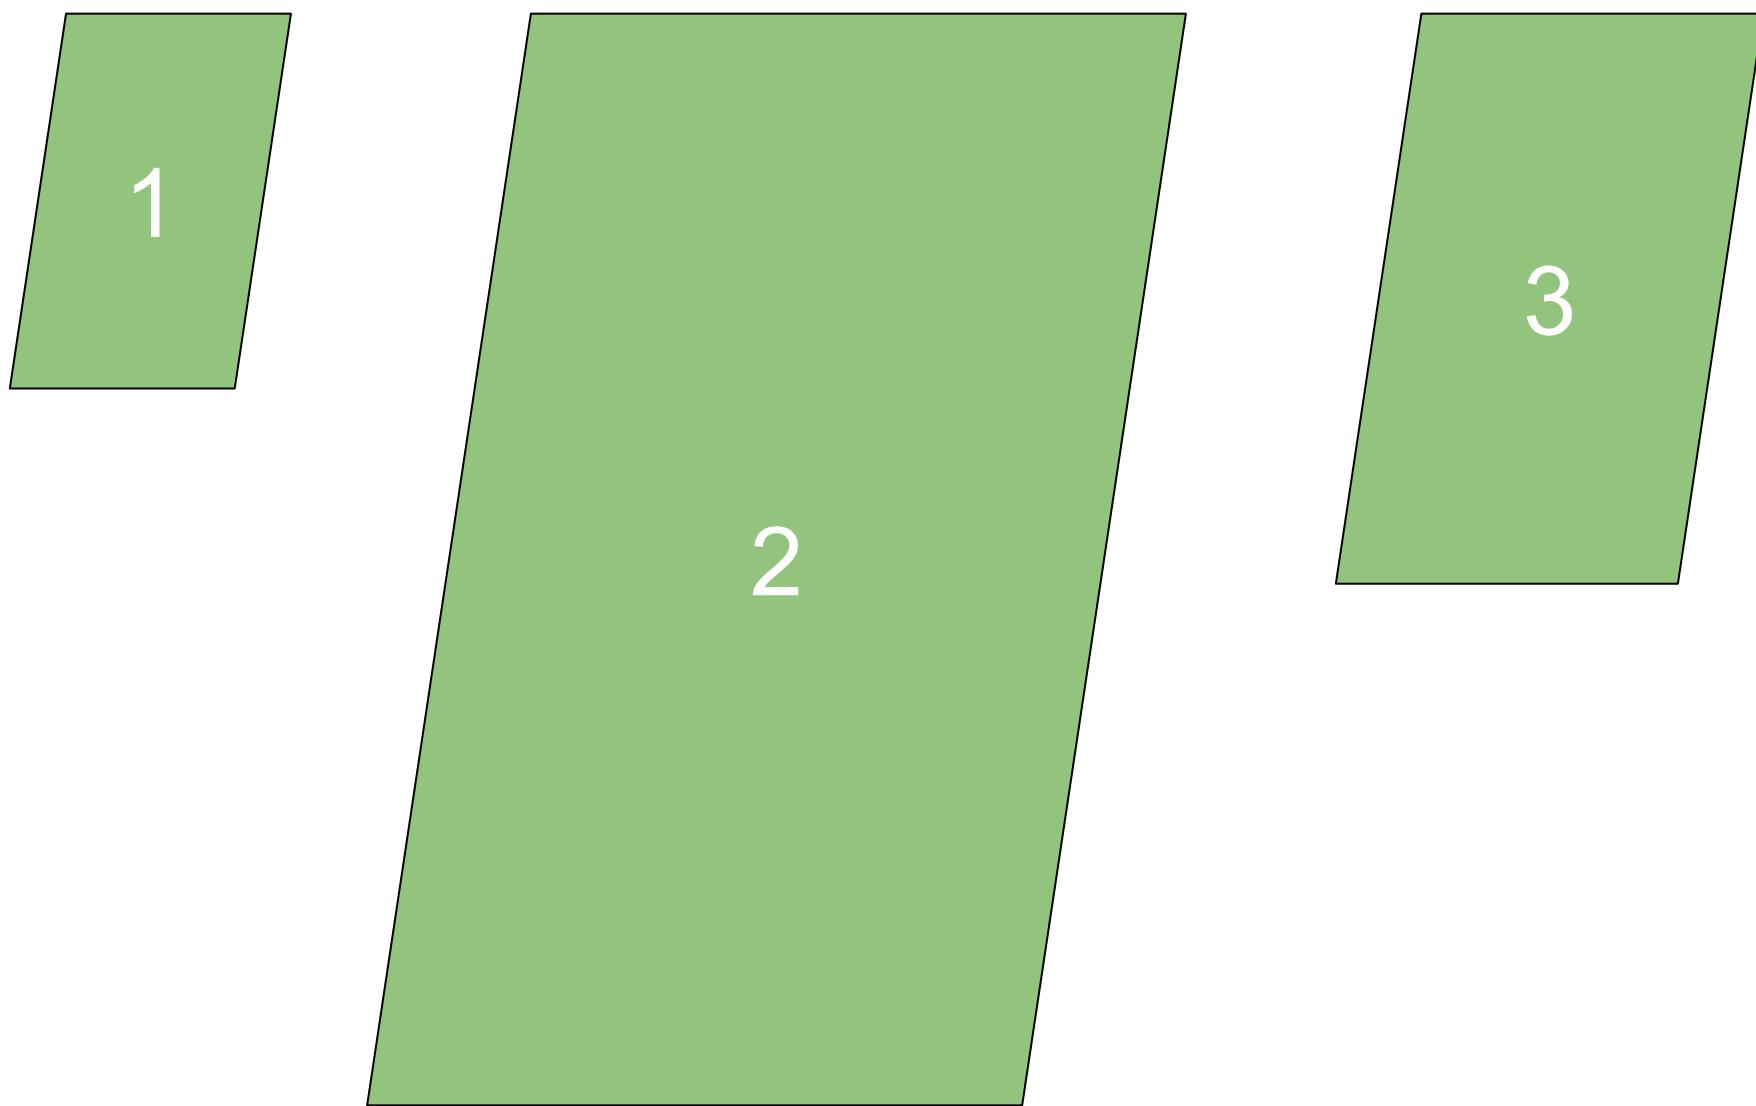
Finding shard sizes... (NEW in 5.0!)



Finding Partitions

Reading index "logs/data"...

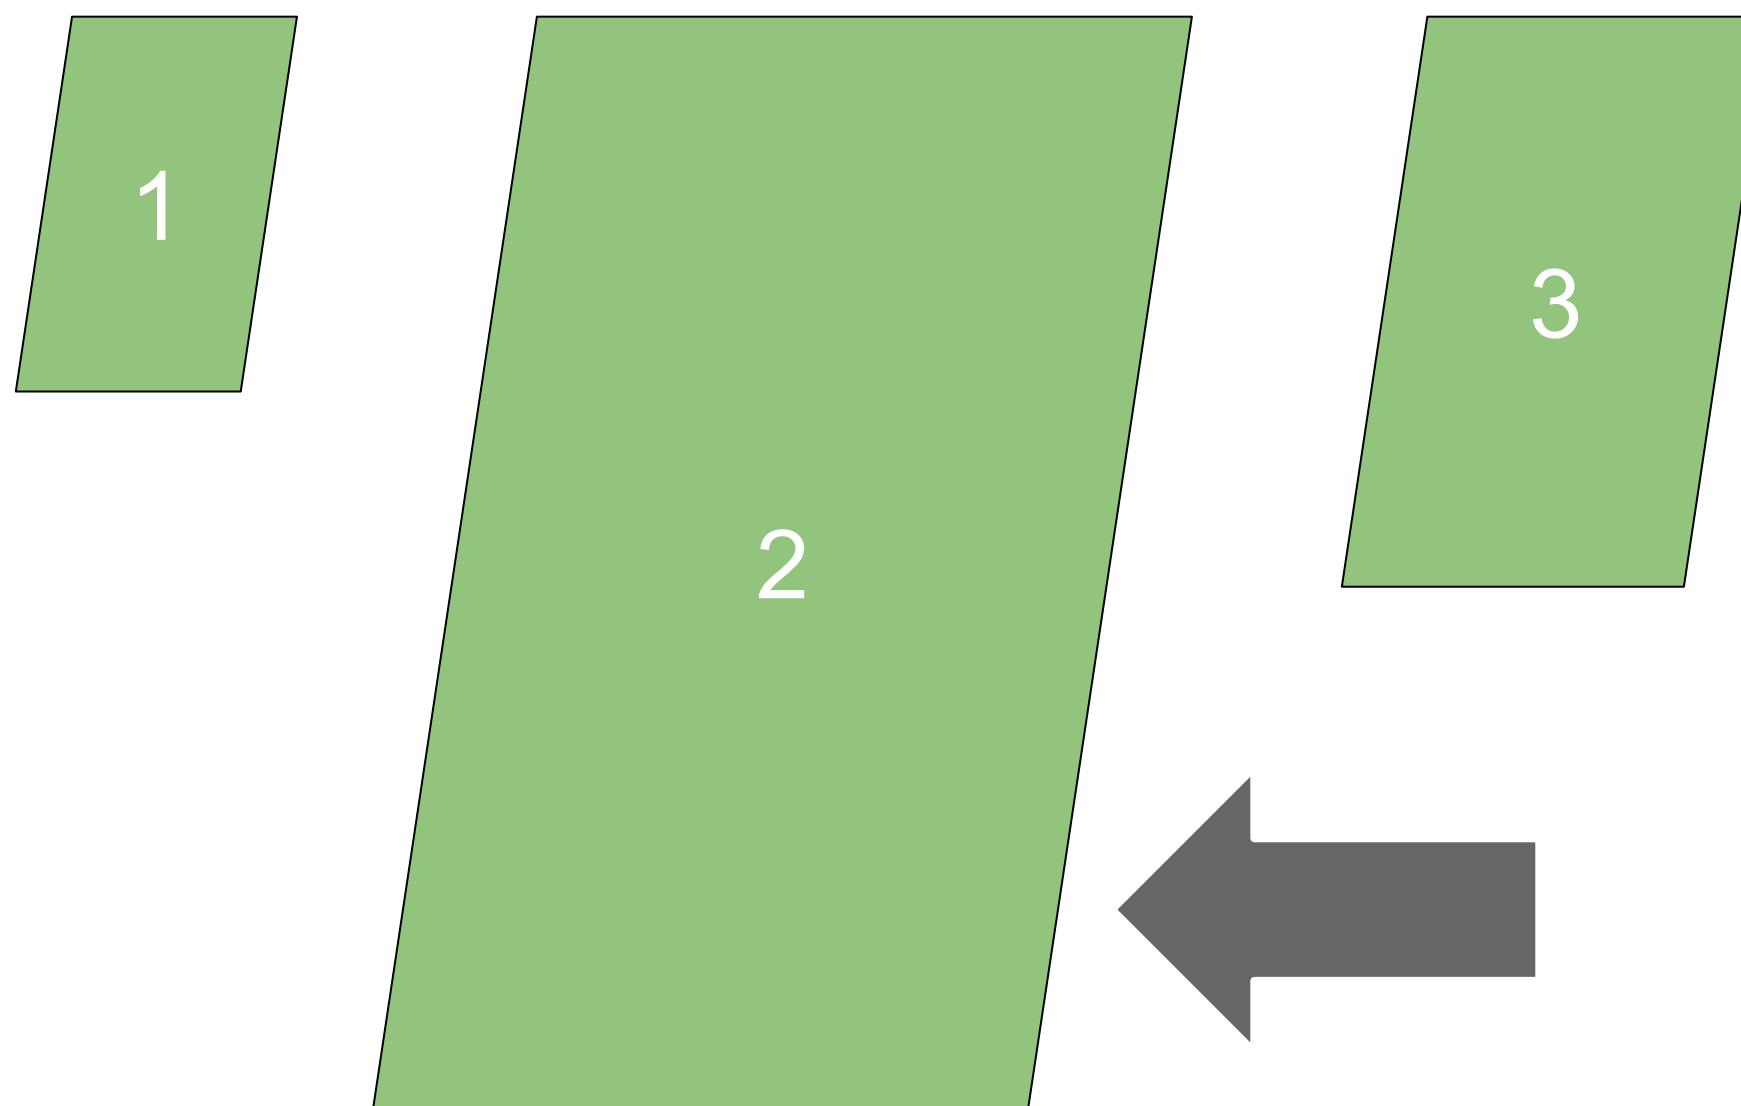
Subdividing shards... (NEW in 5.0!)



Finding Partitions

Reading index "logs/data"...

Subdividing shards... (NEW in 5.0!)



NEW IN 5.0 OF ELASTICSEARCH

Sliced Scrolls

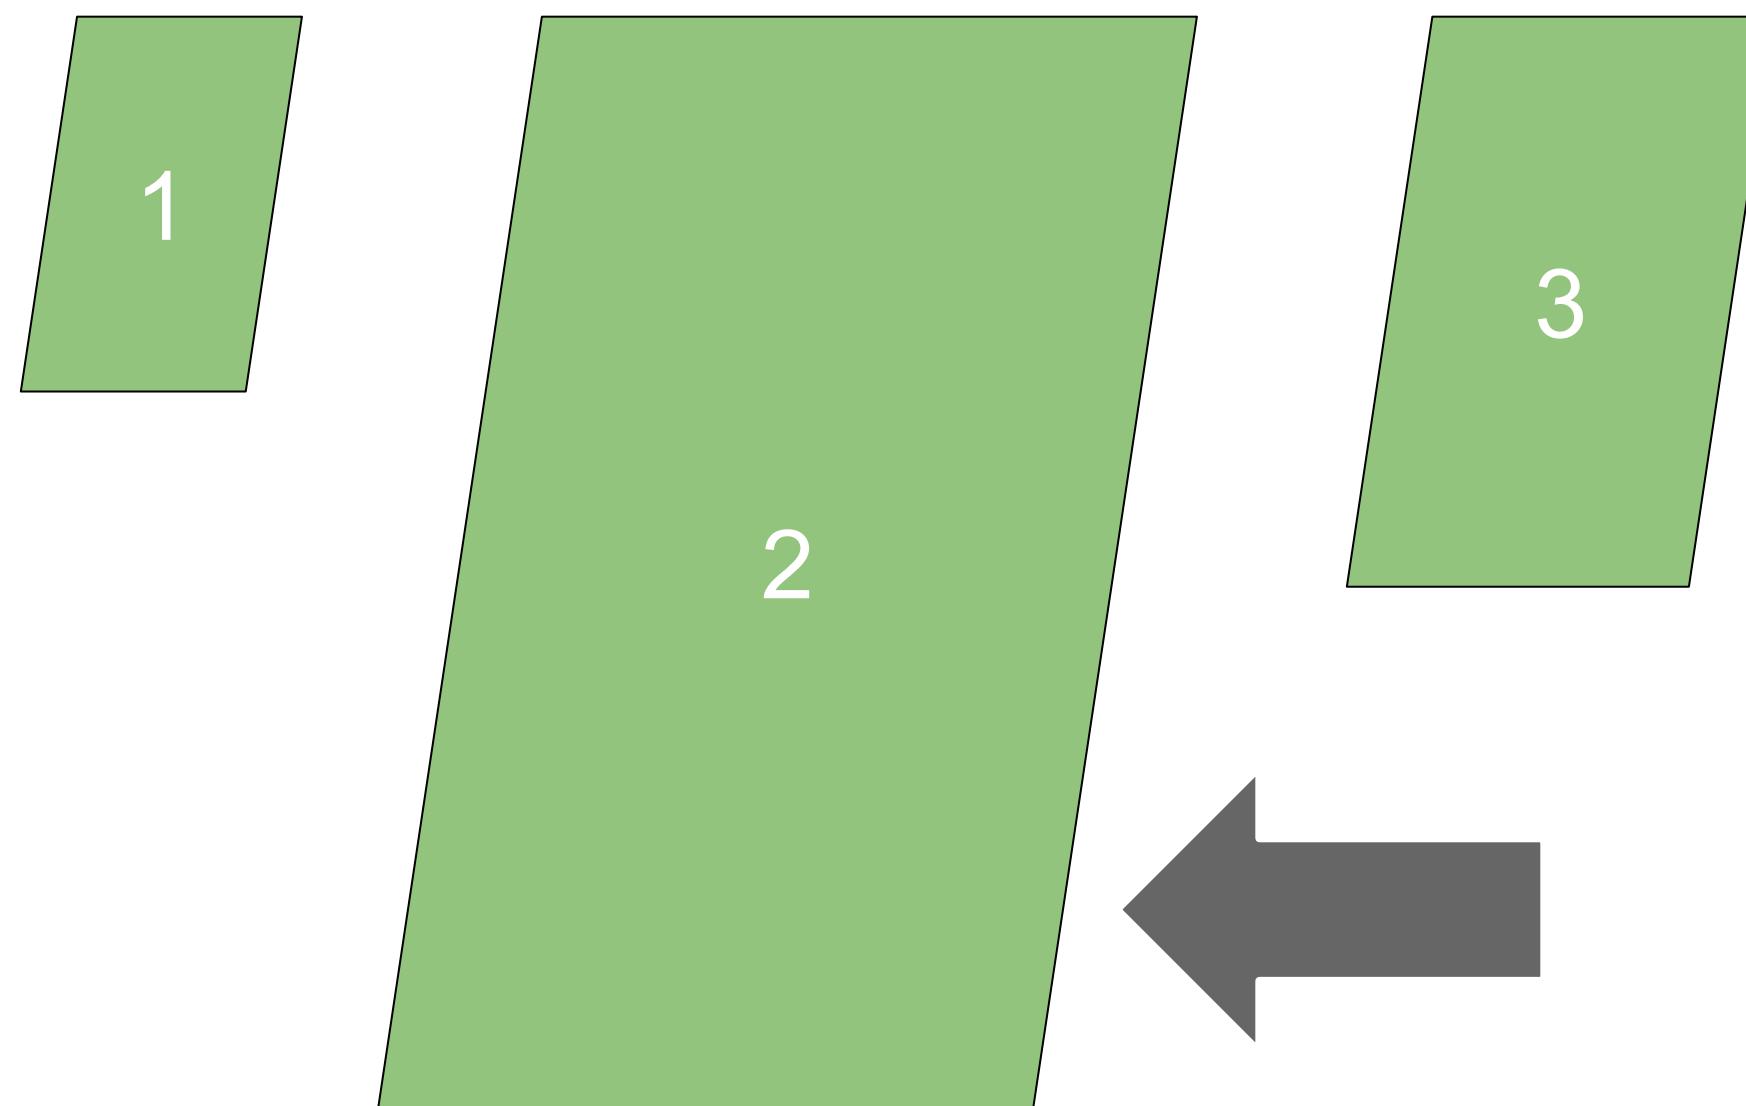
New in Elasticsearch 5.0

```
curl -XGET localhost:9200/idx/t/_search?scroll=1m -d' {  
  "slice": {  
    "id": 0,  
    "max": 2  
  },  
  "query": {  
    "match" : {  
      "title" : "elasticsearch"  
    }  
  }  
}'
```

Finding Partitions

Reading index "logs/data"...

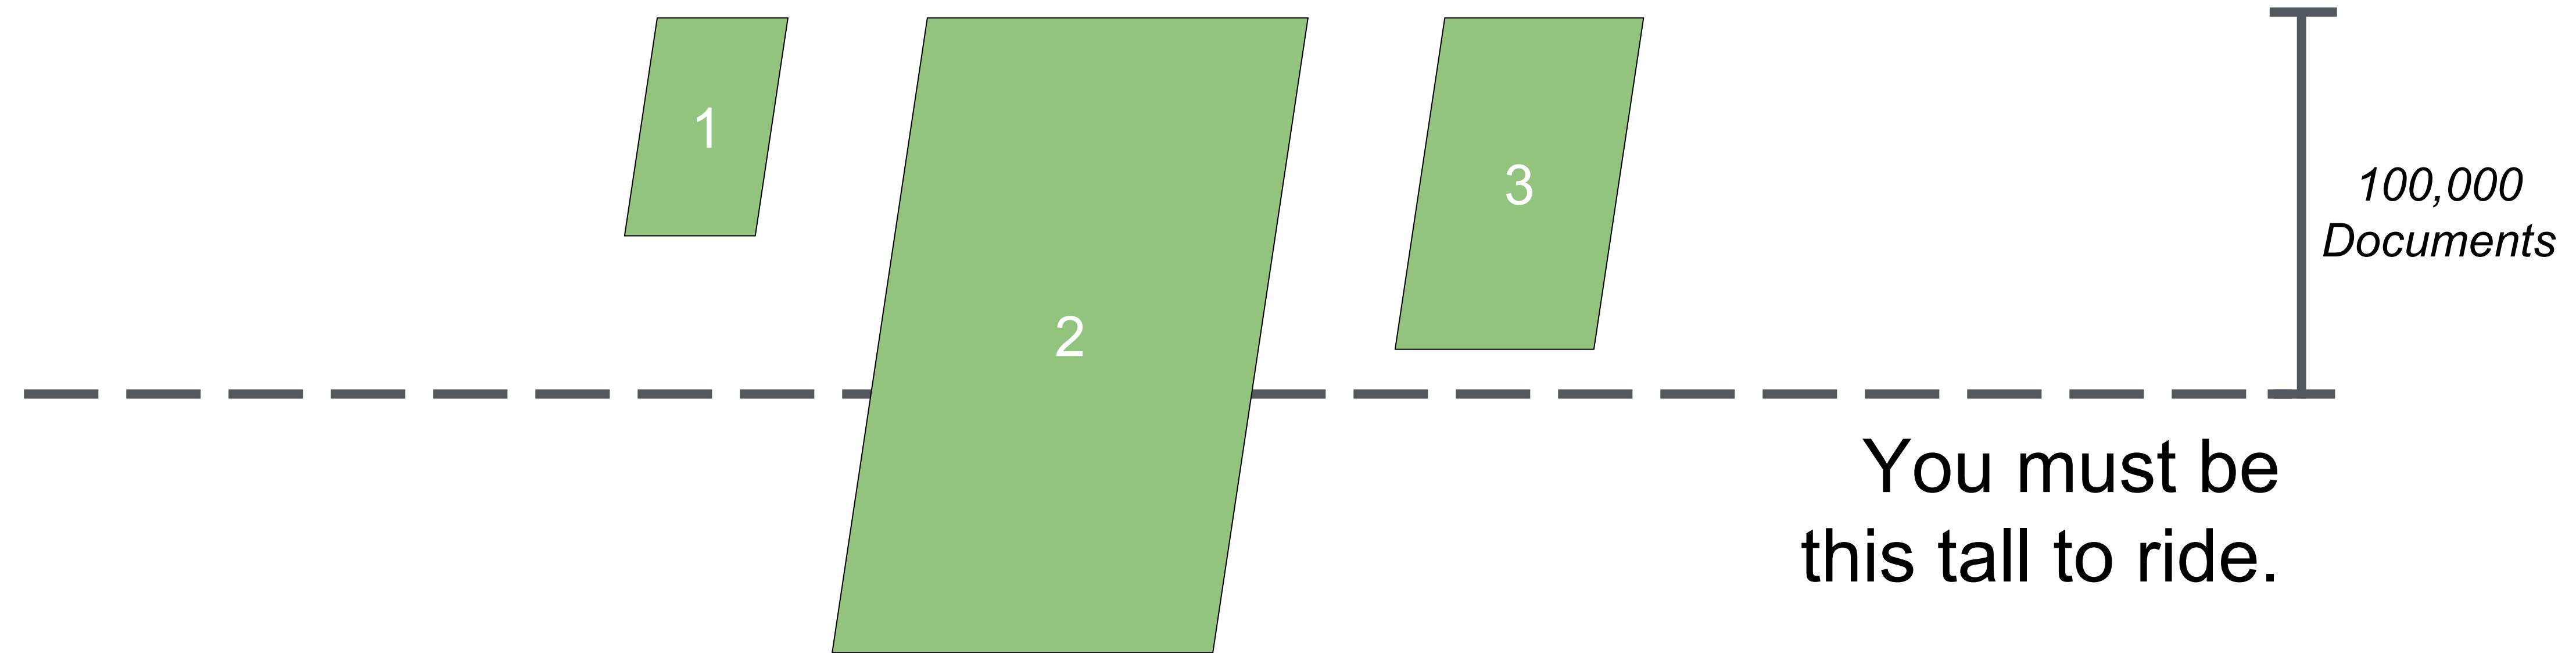
Subdividing shards... (NEW in 5.0!)



Finding Partitions

Reading index "logs/data"...

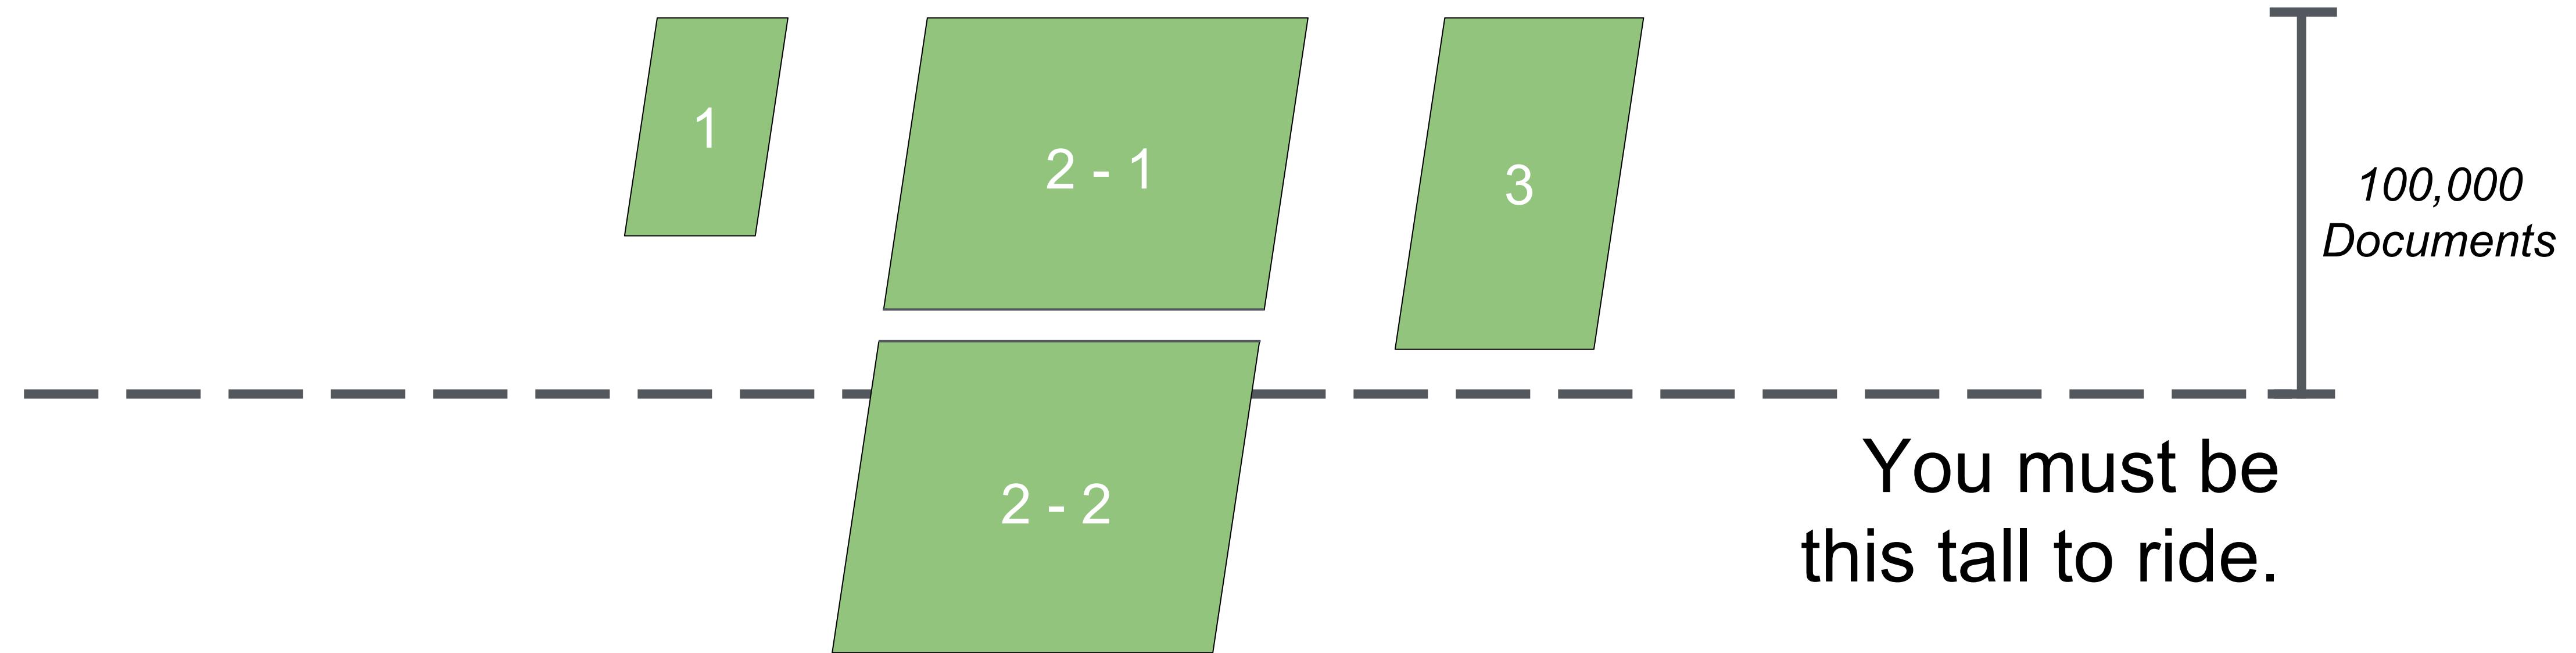
Subdividing shards... (NEW in 5.0!)



Finding Partitions

Reading index "logs/data"...

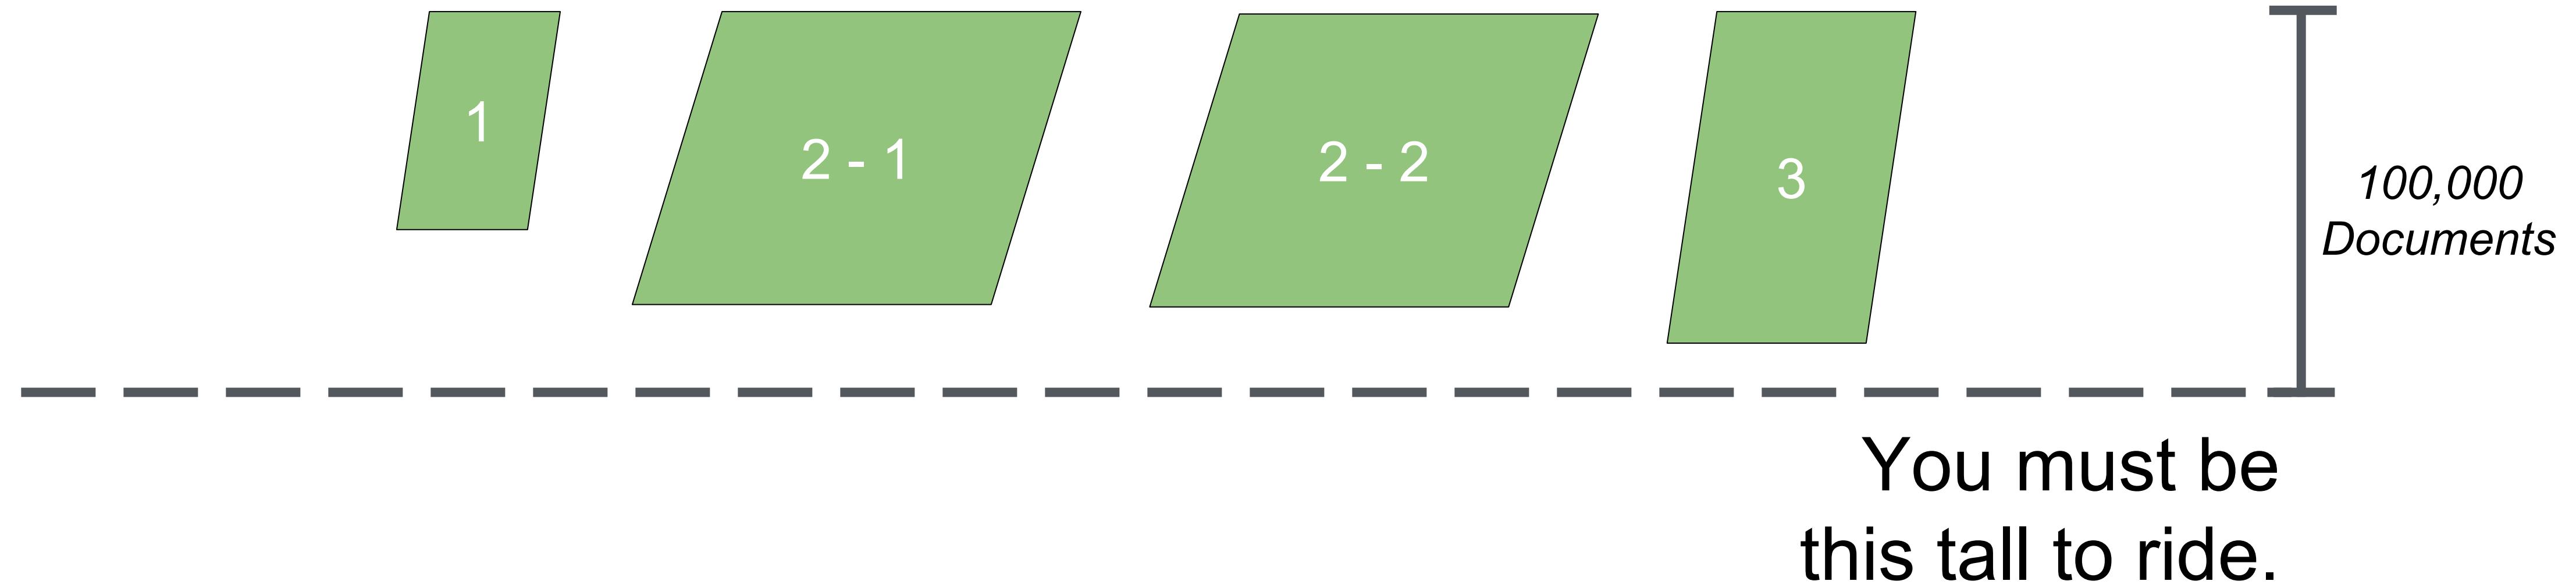
Subdividing shards... (NEW in 5.0!)



Finding Partitions

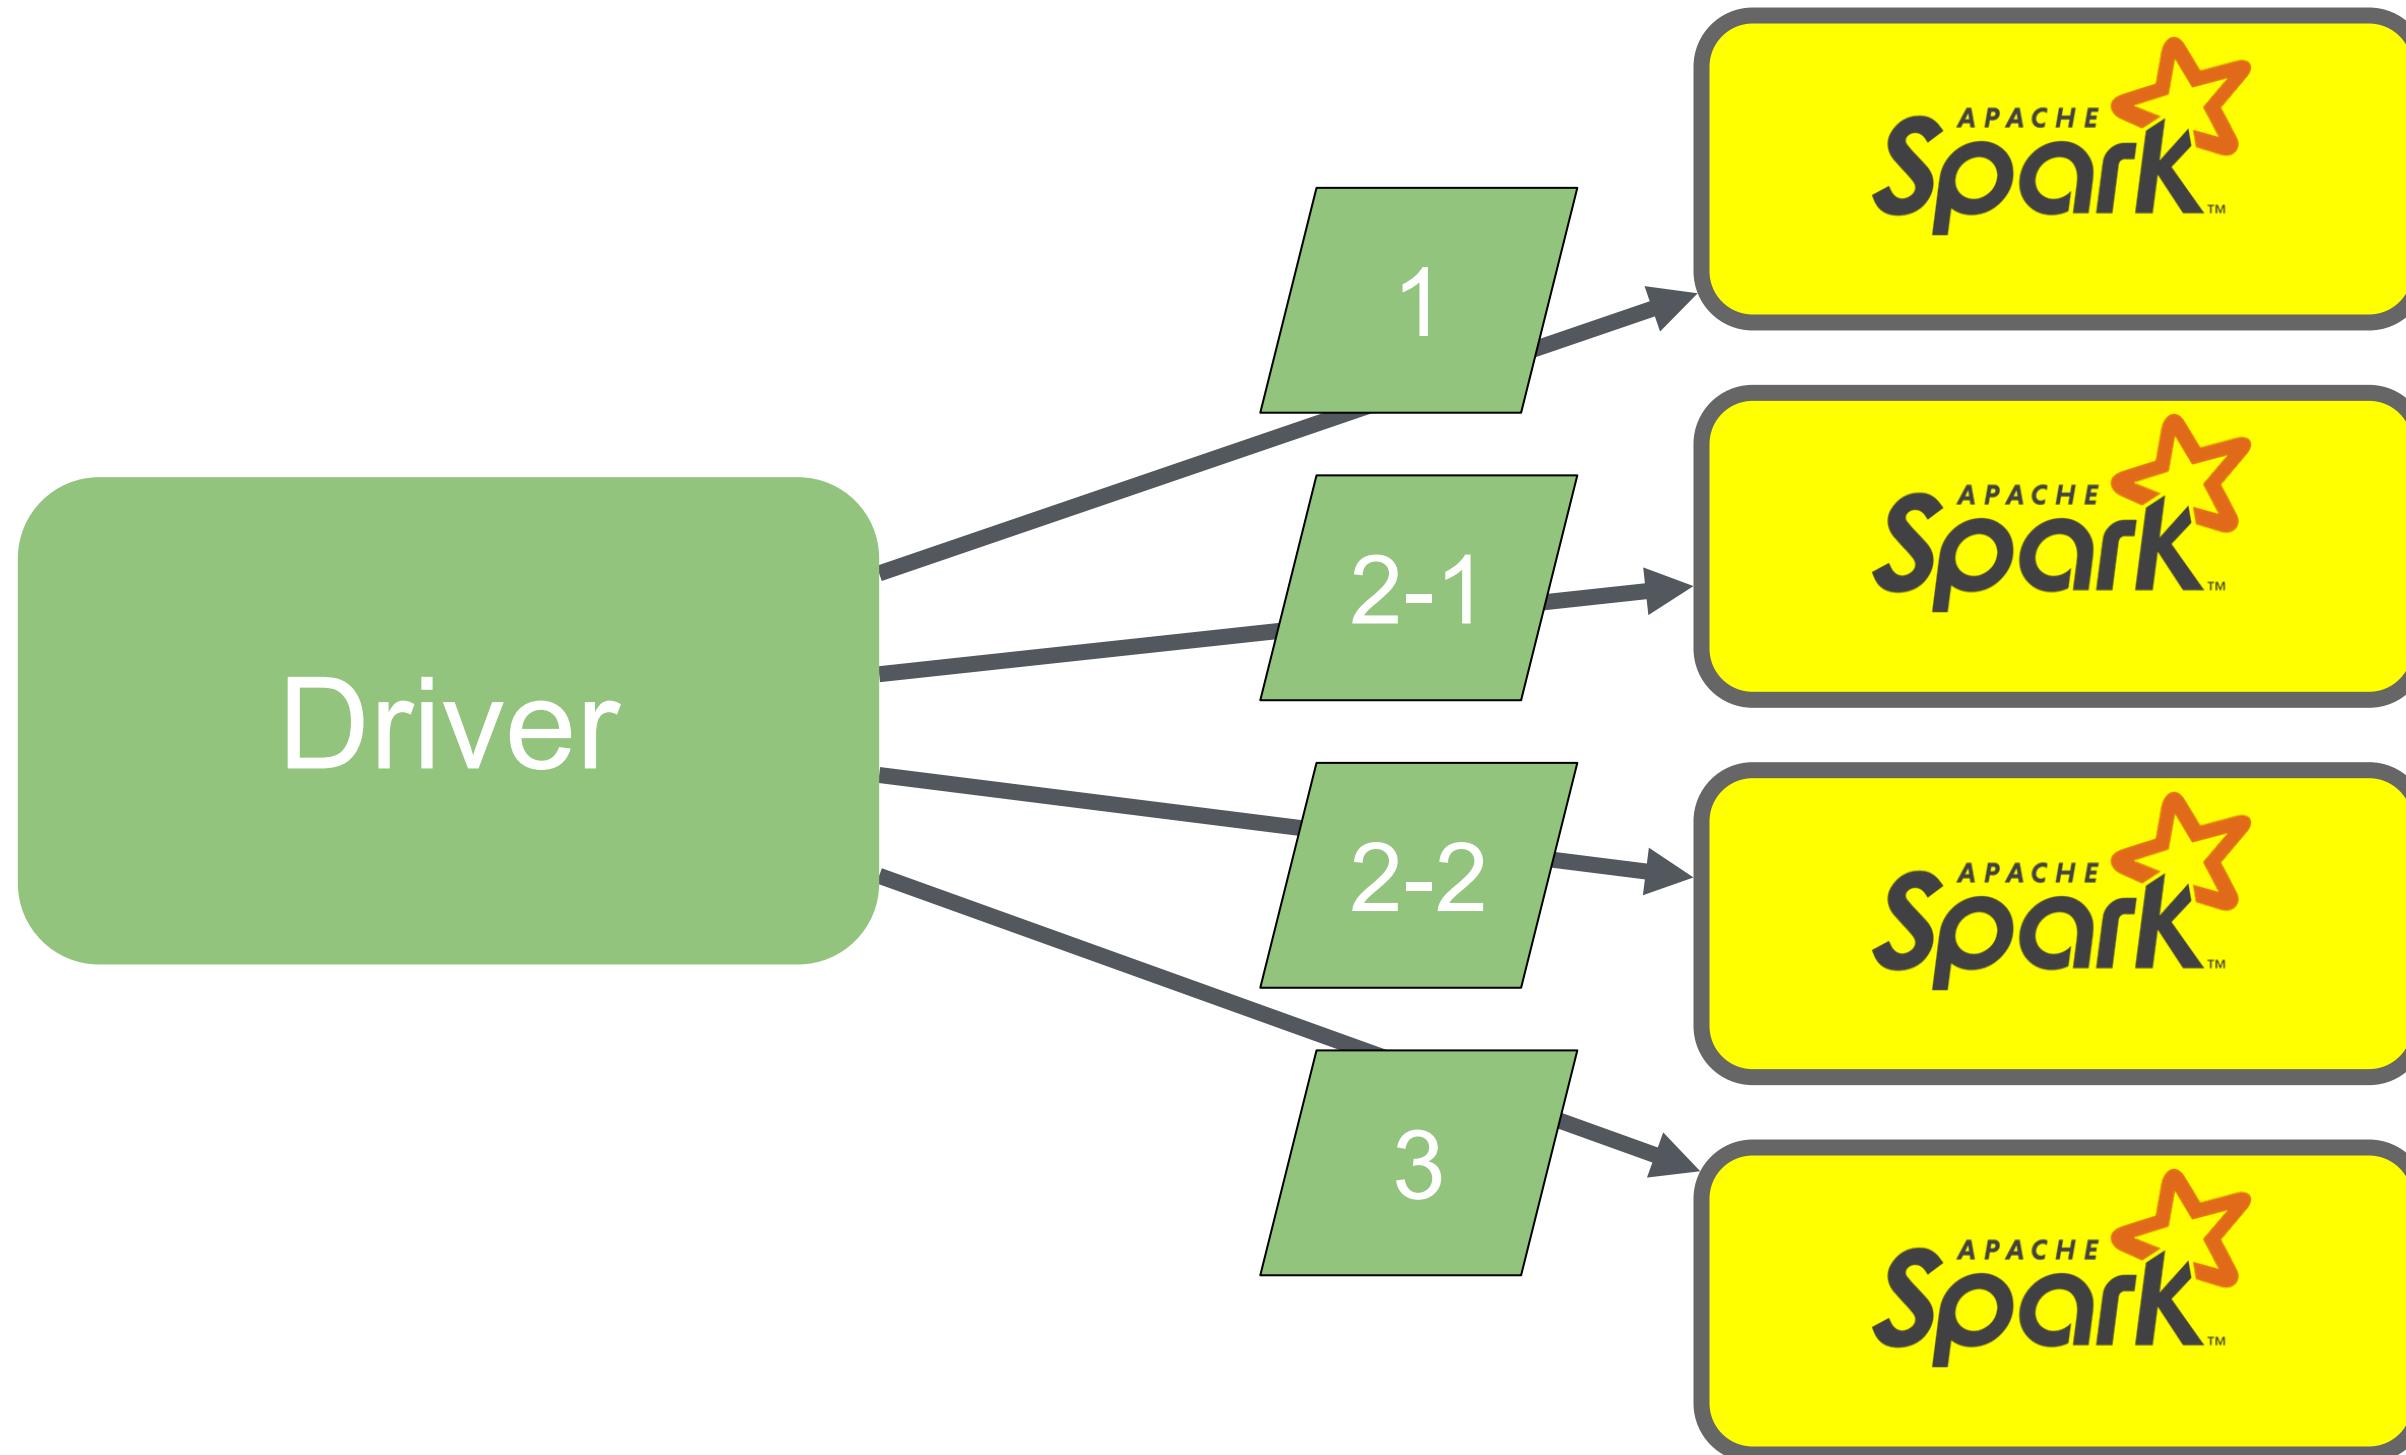
Reading index "logs/data"...

Subdividing shards... (NEW in 5.0!)

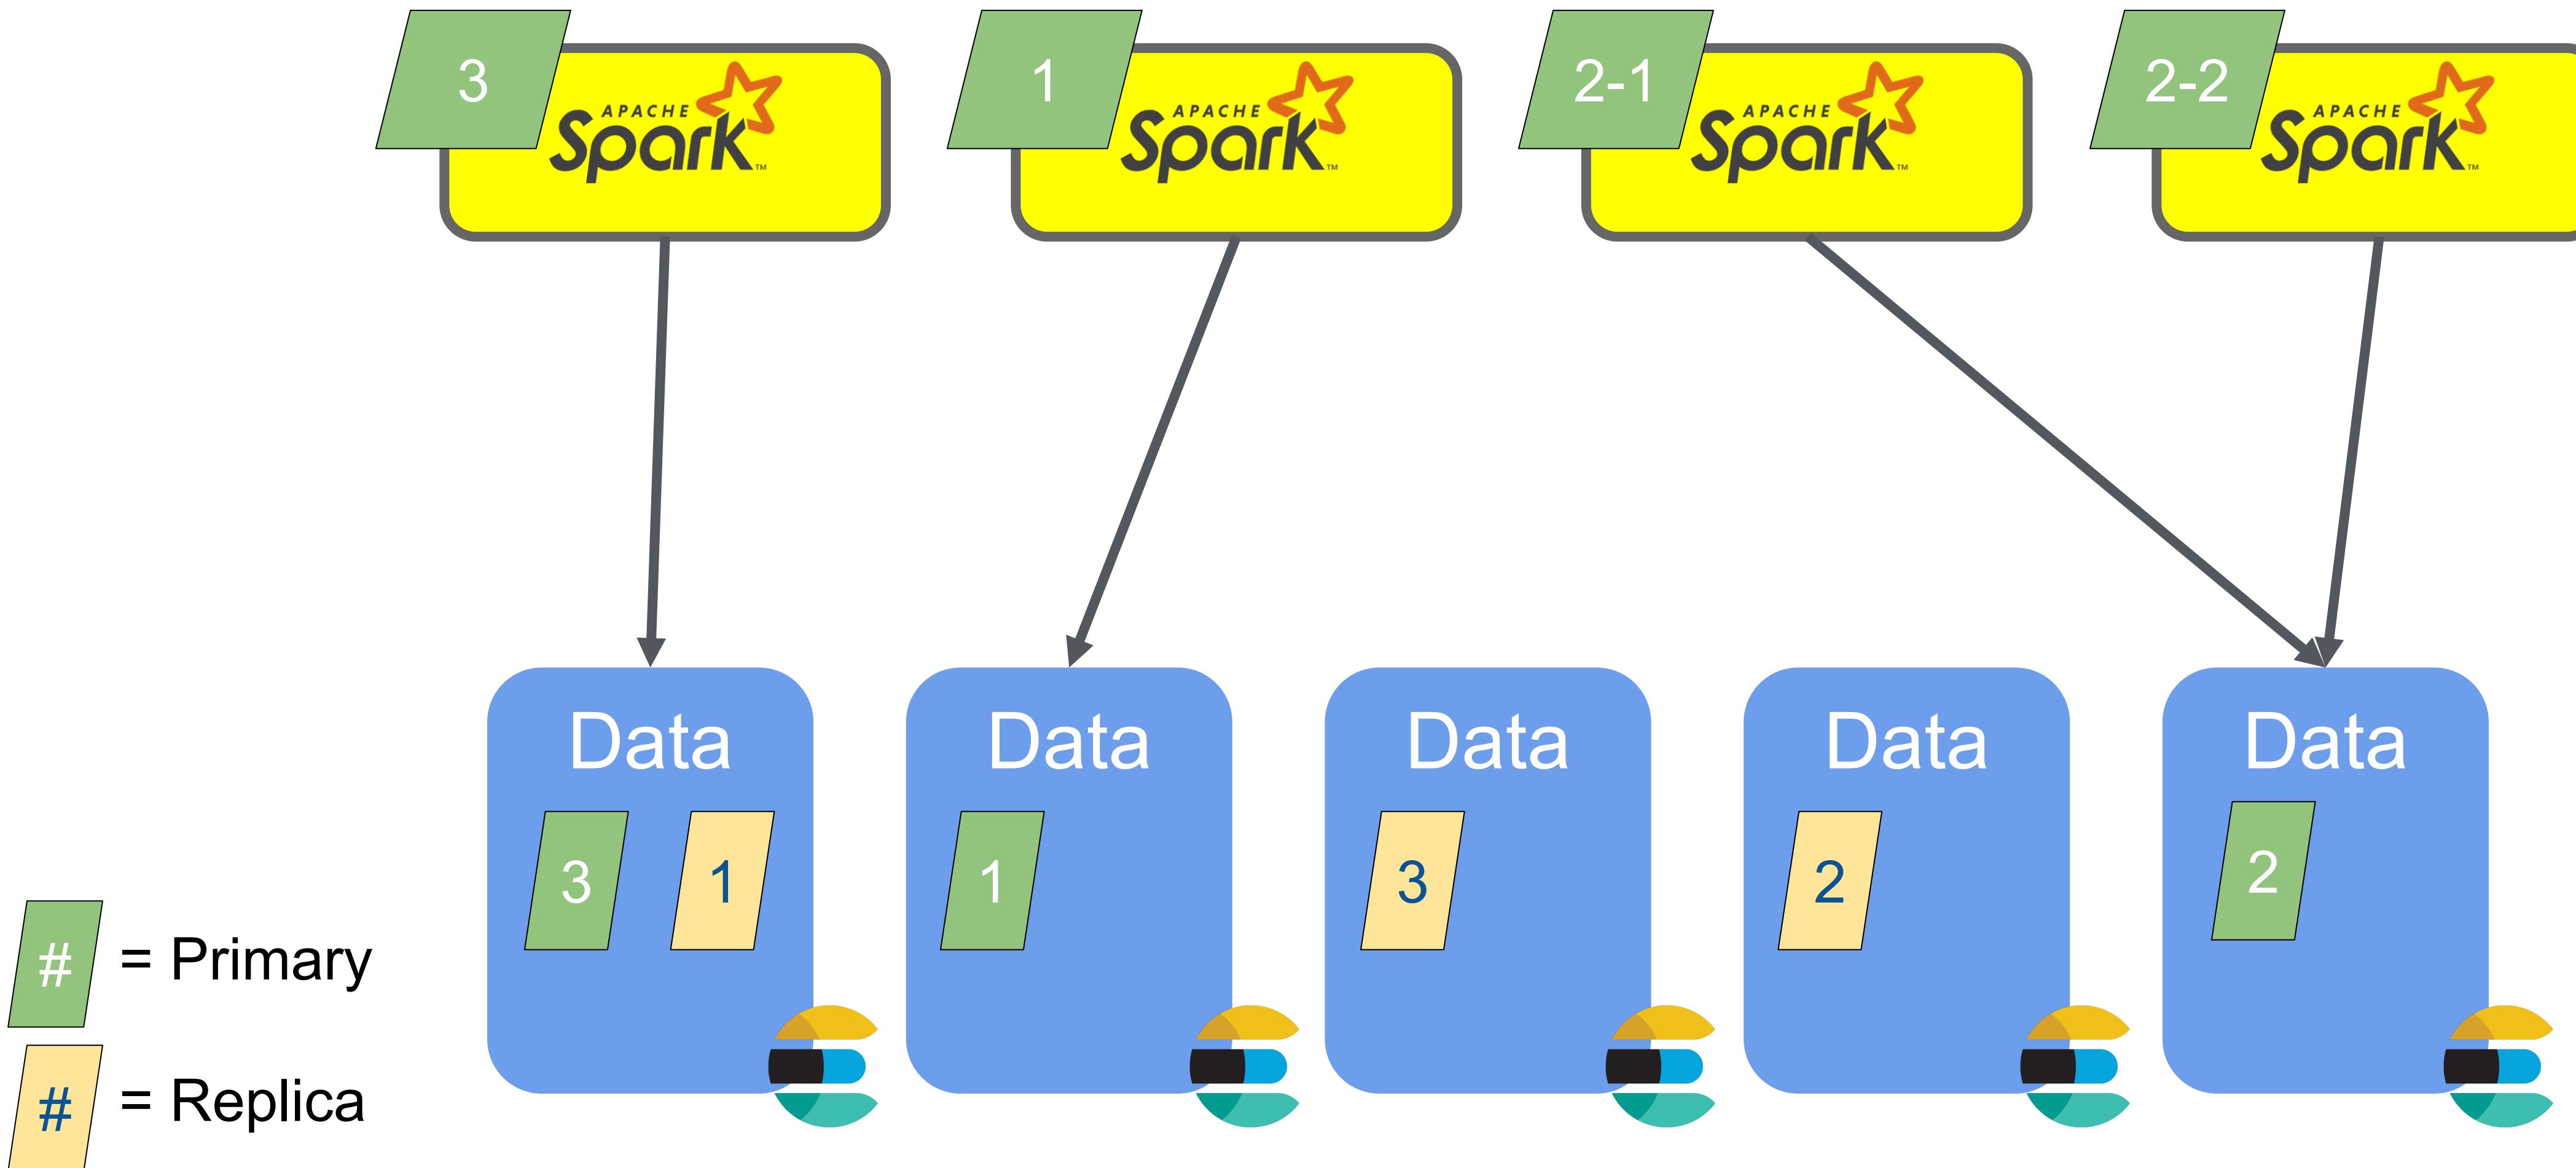


Job Execution

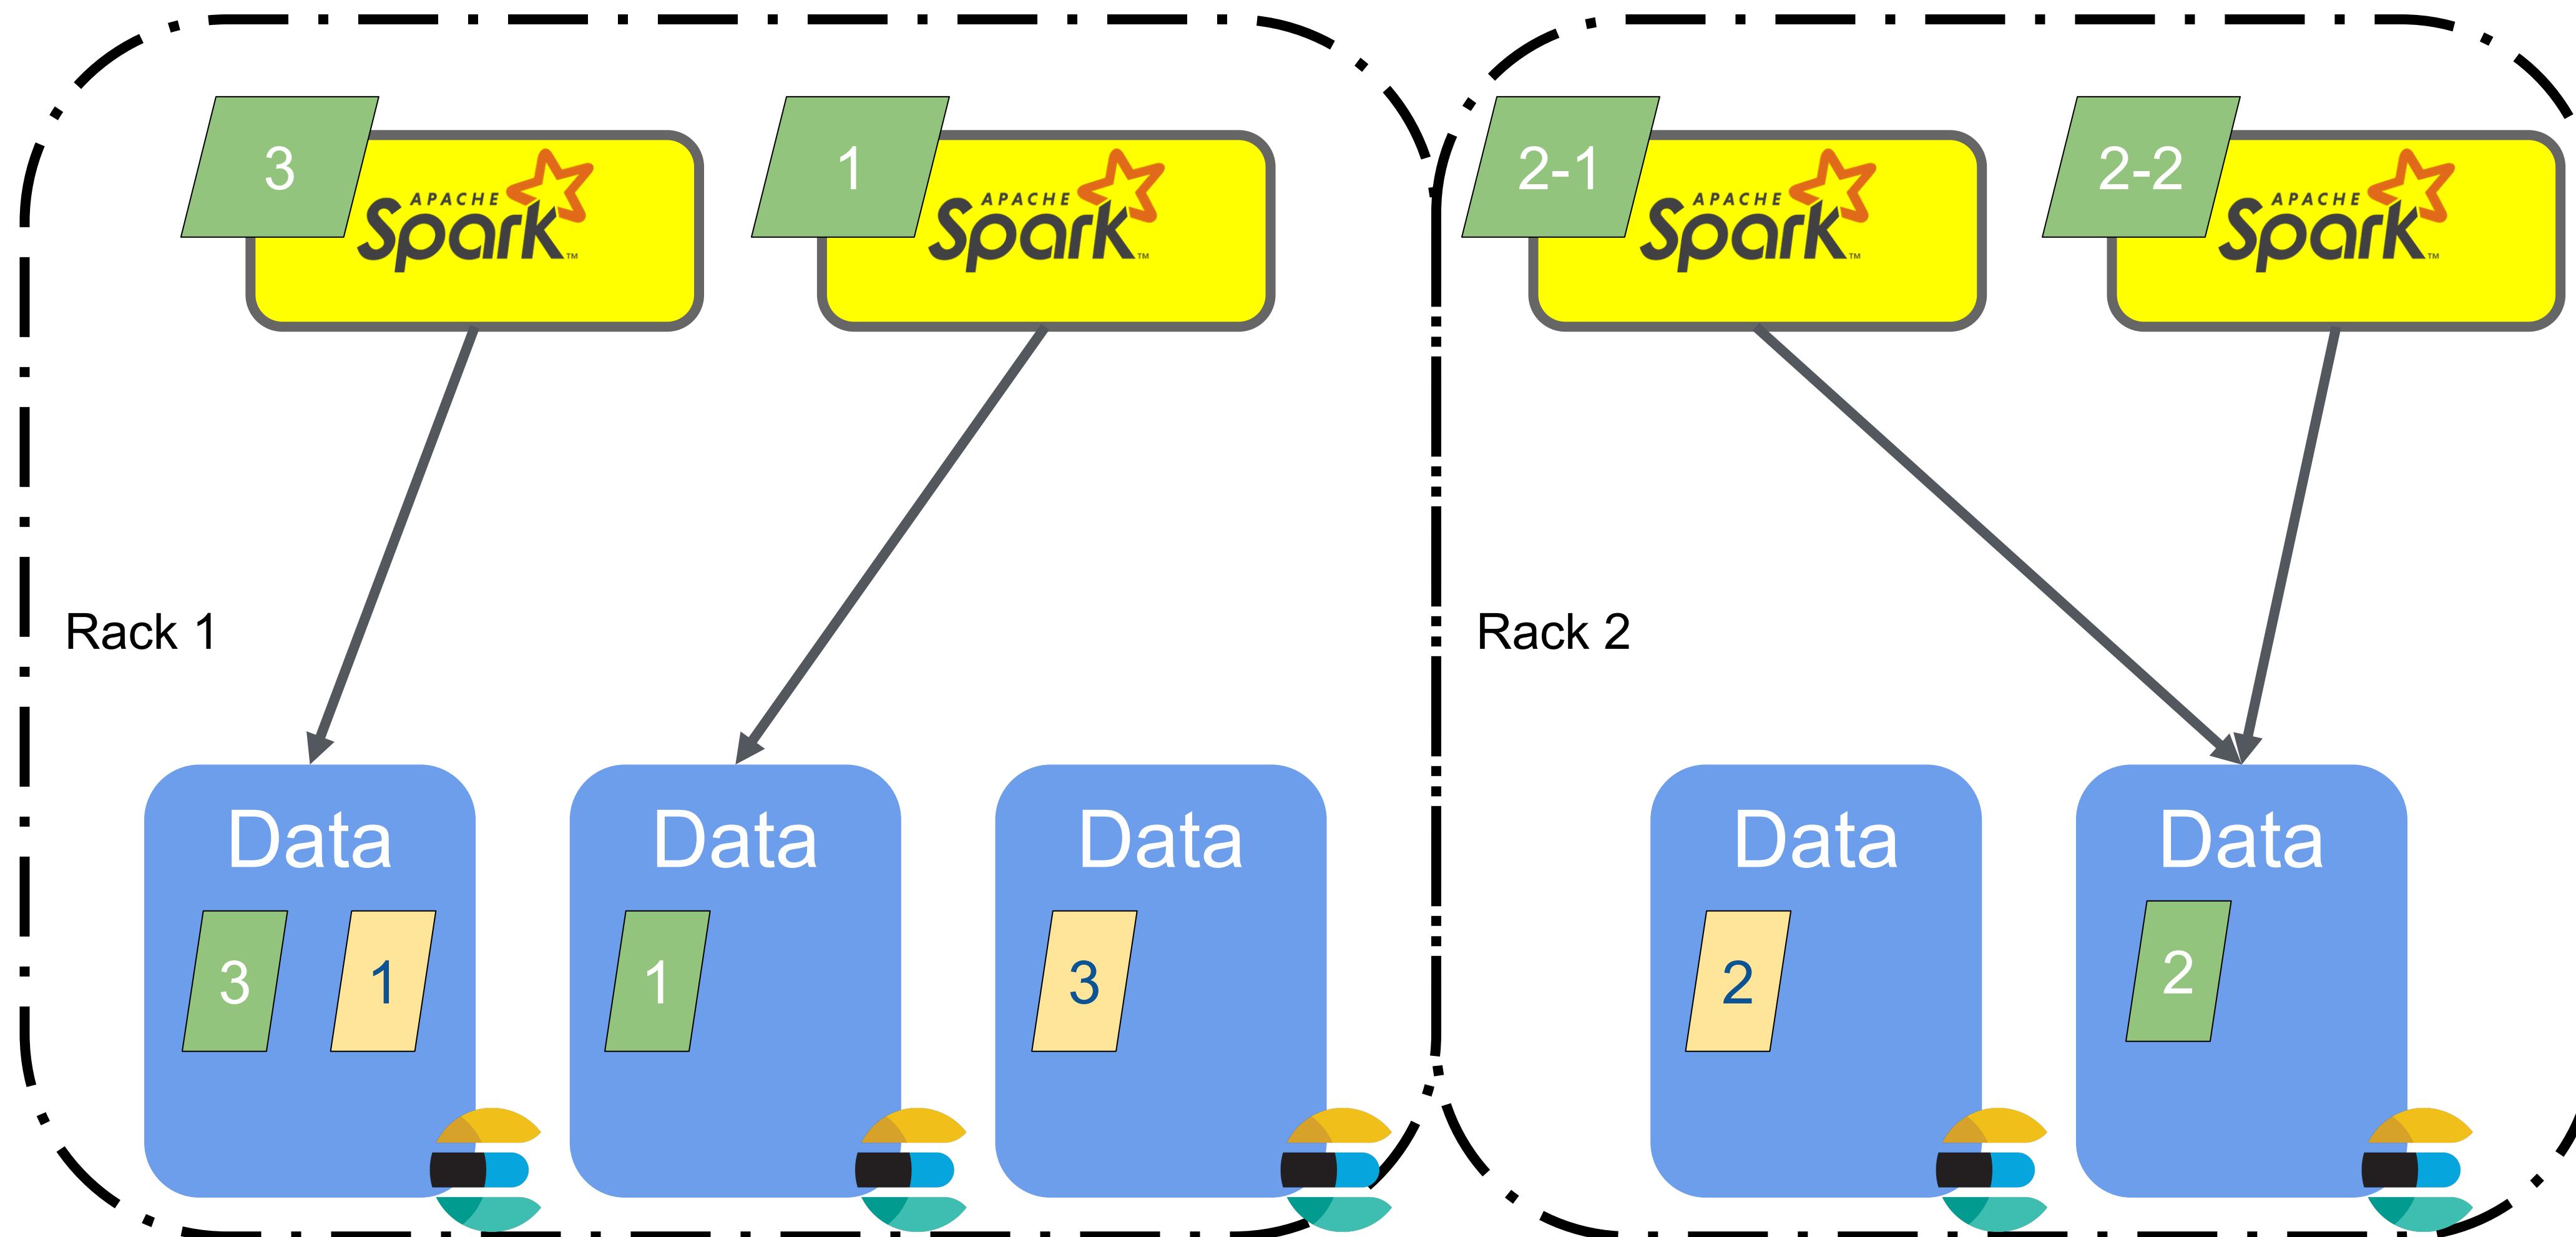
Launching Tasks...



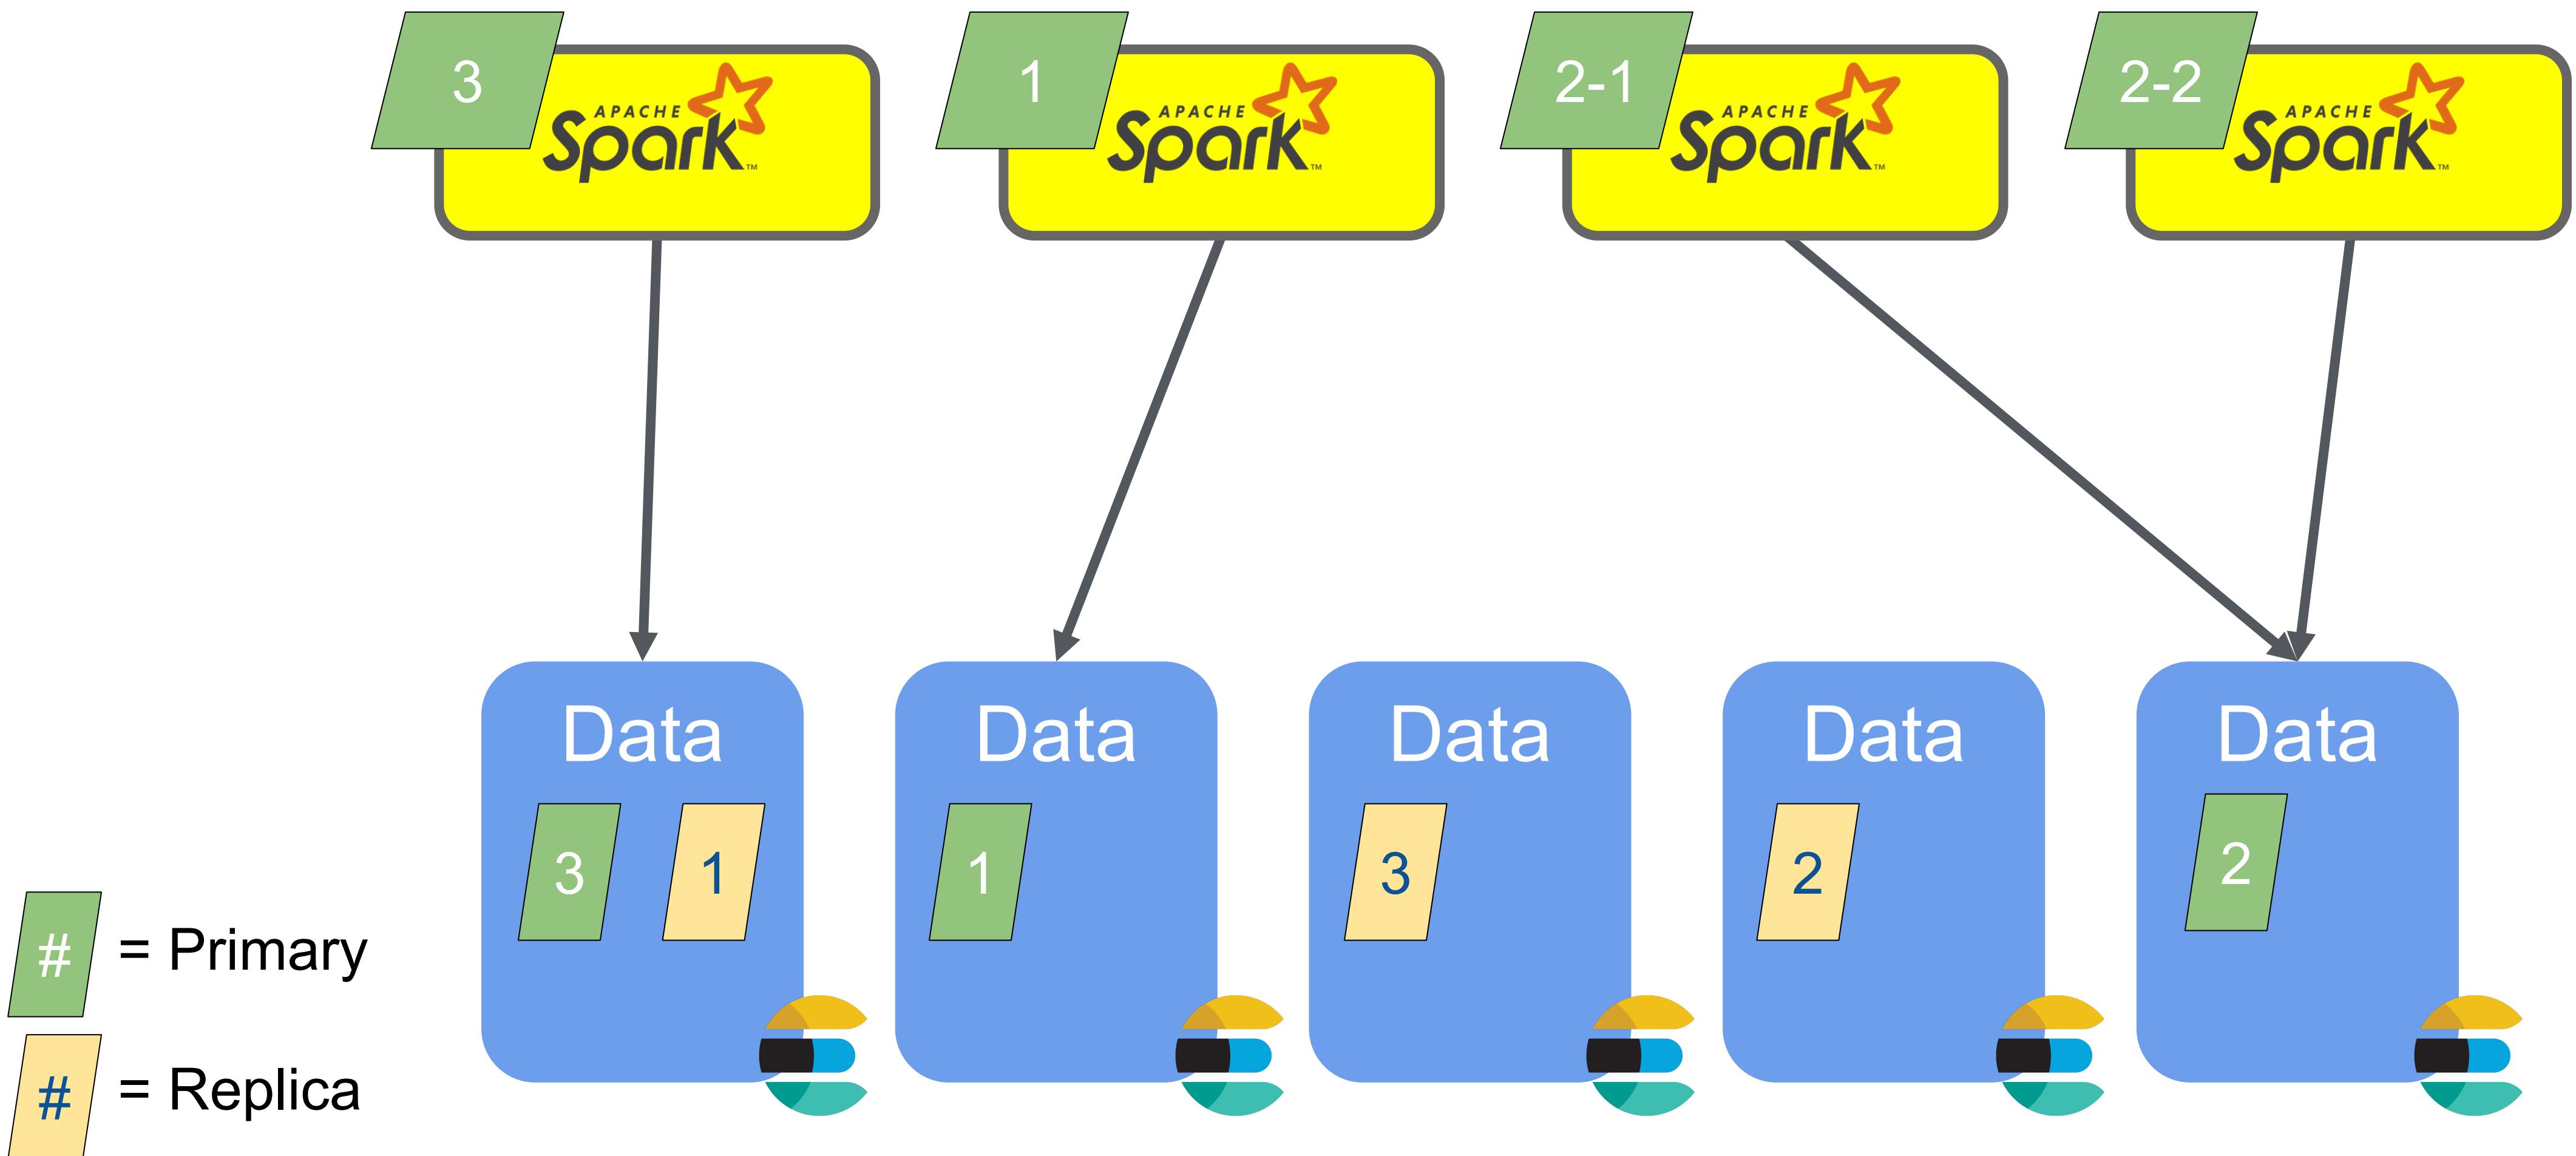
Job Execution



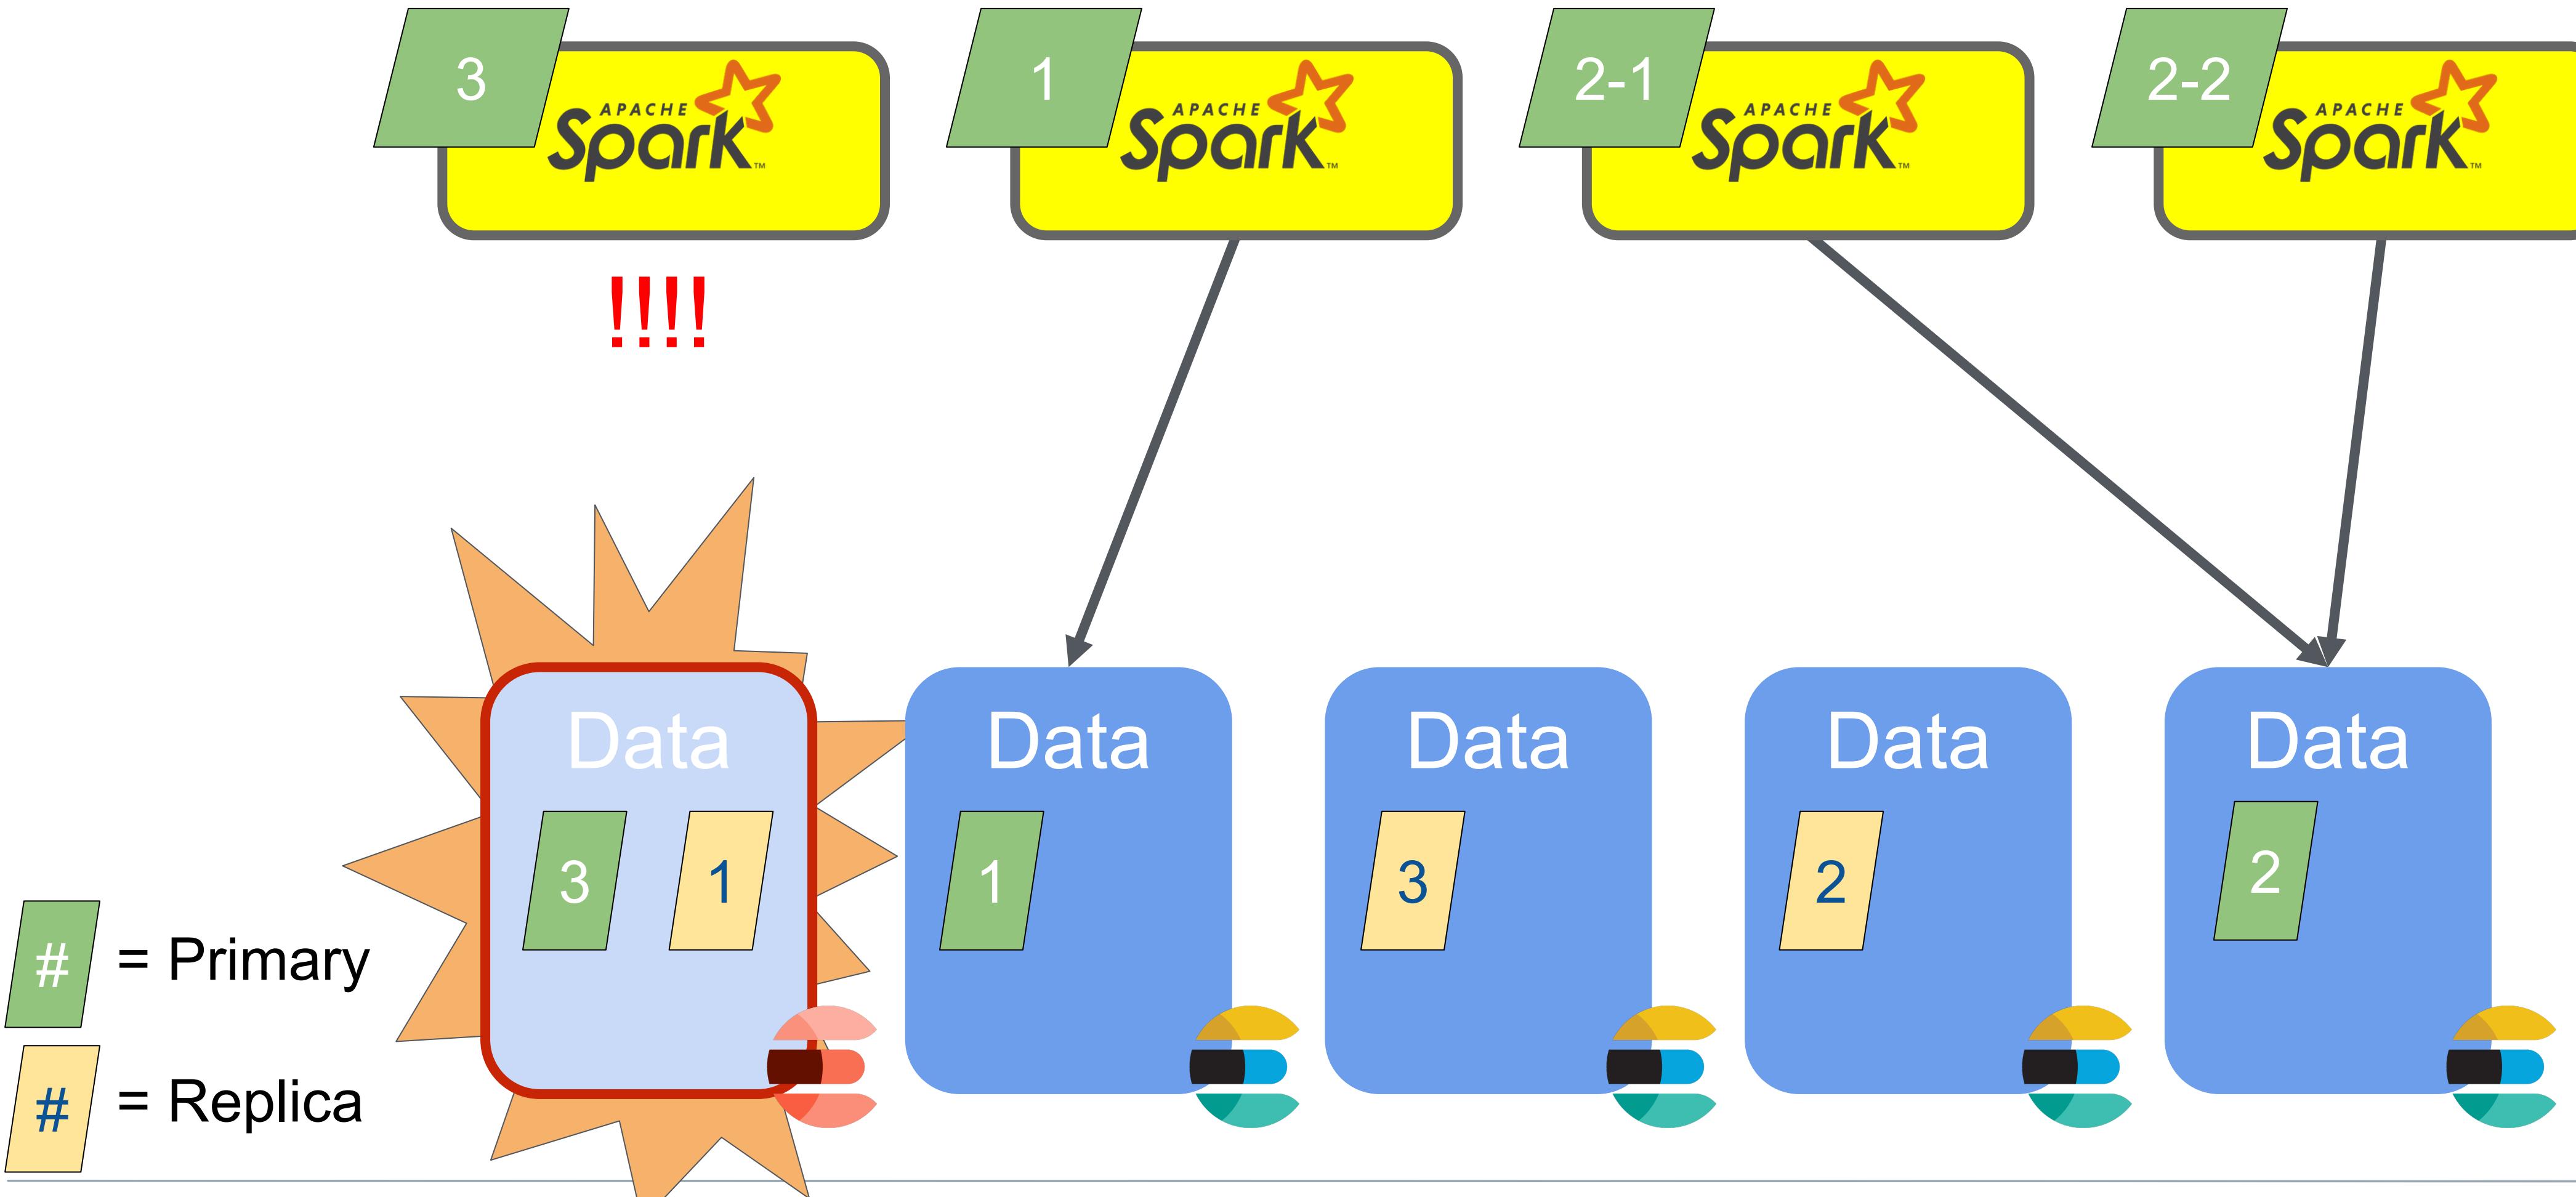
Job Execution



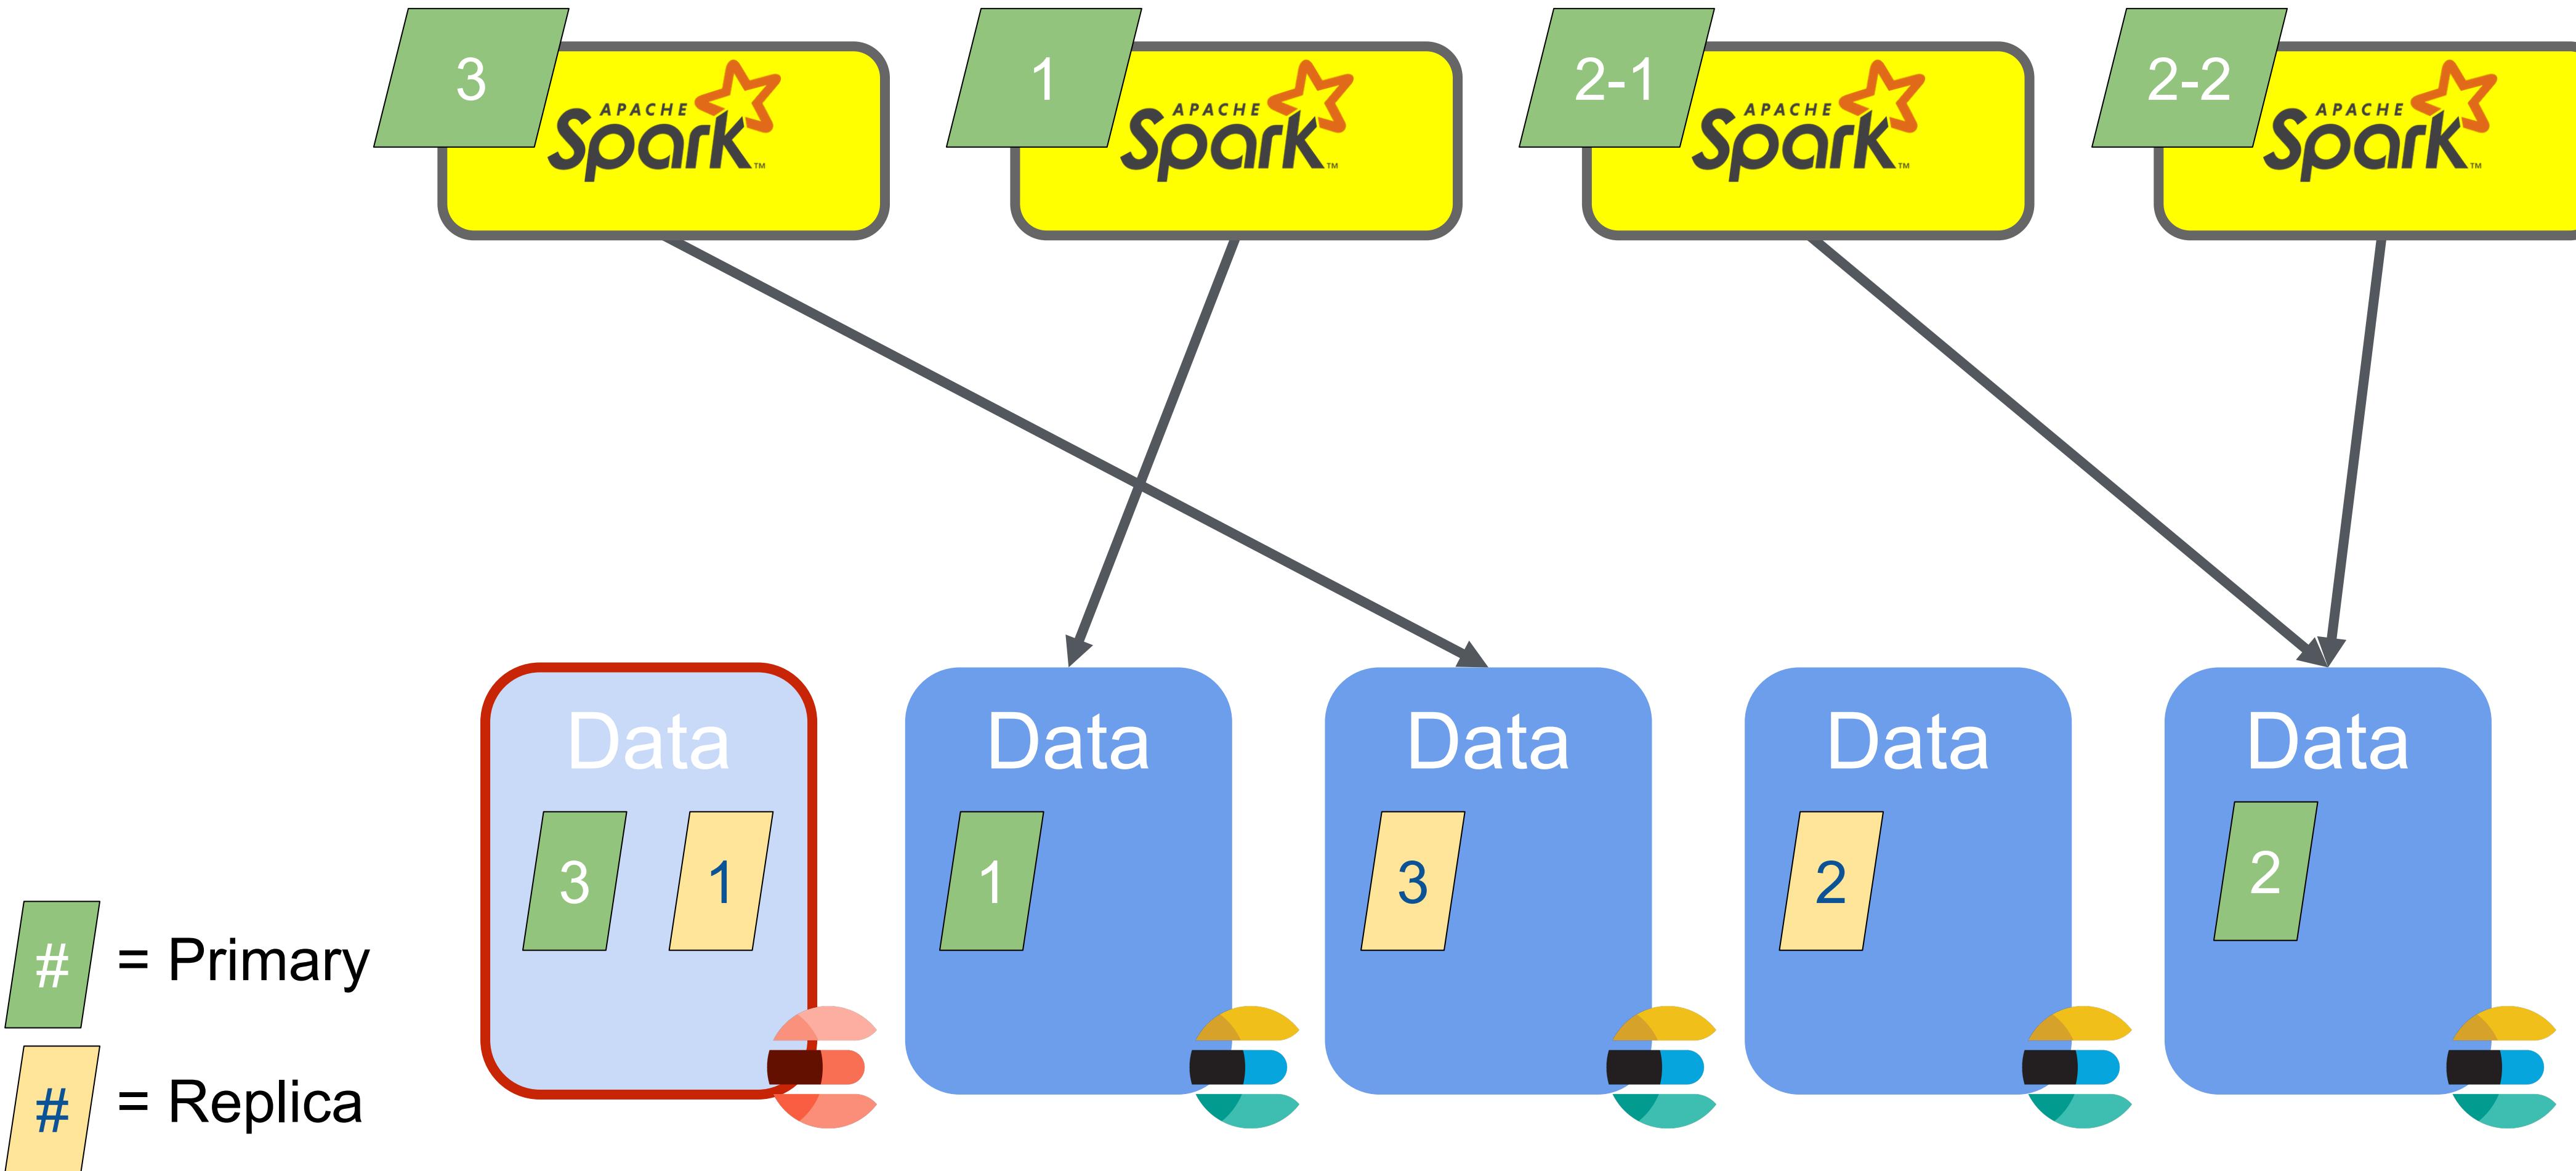
Job Execution



Communication Failure



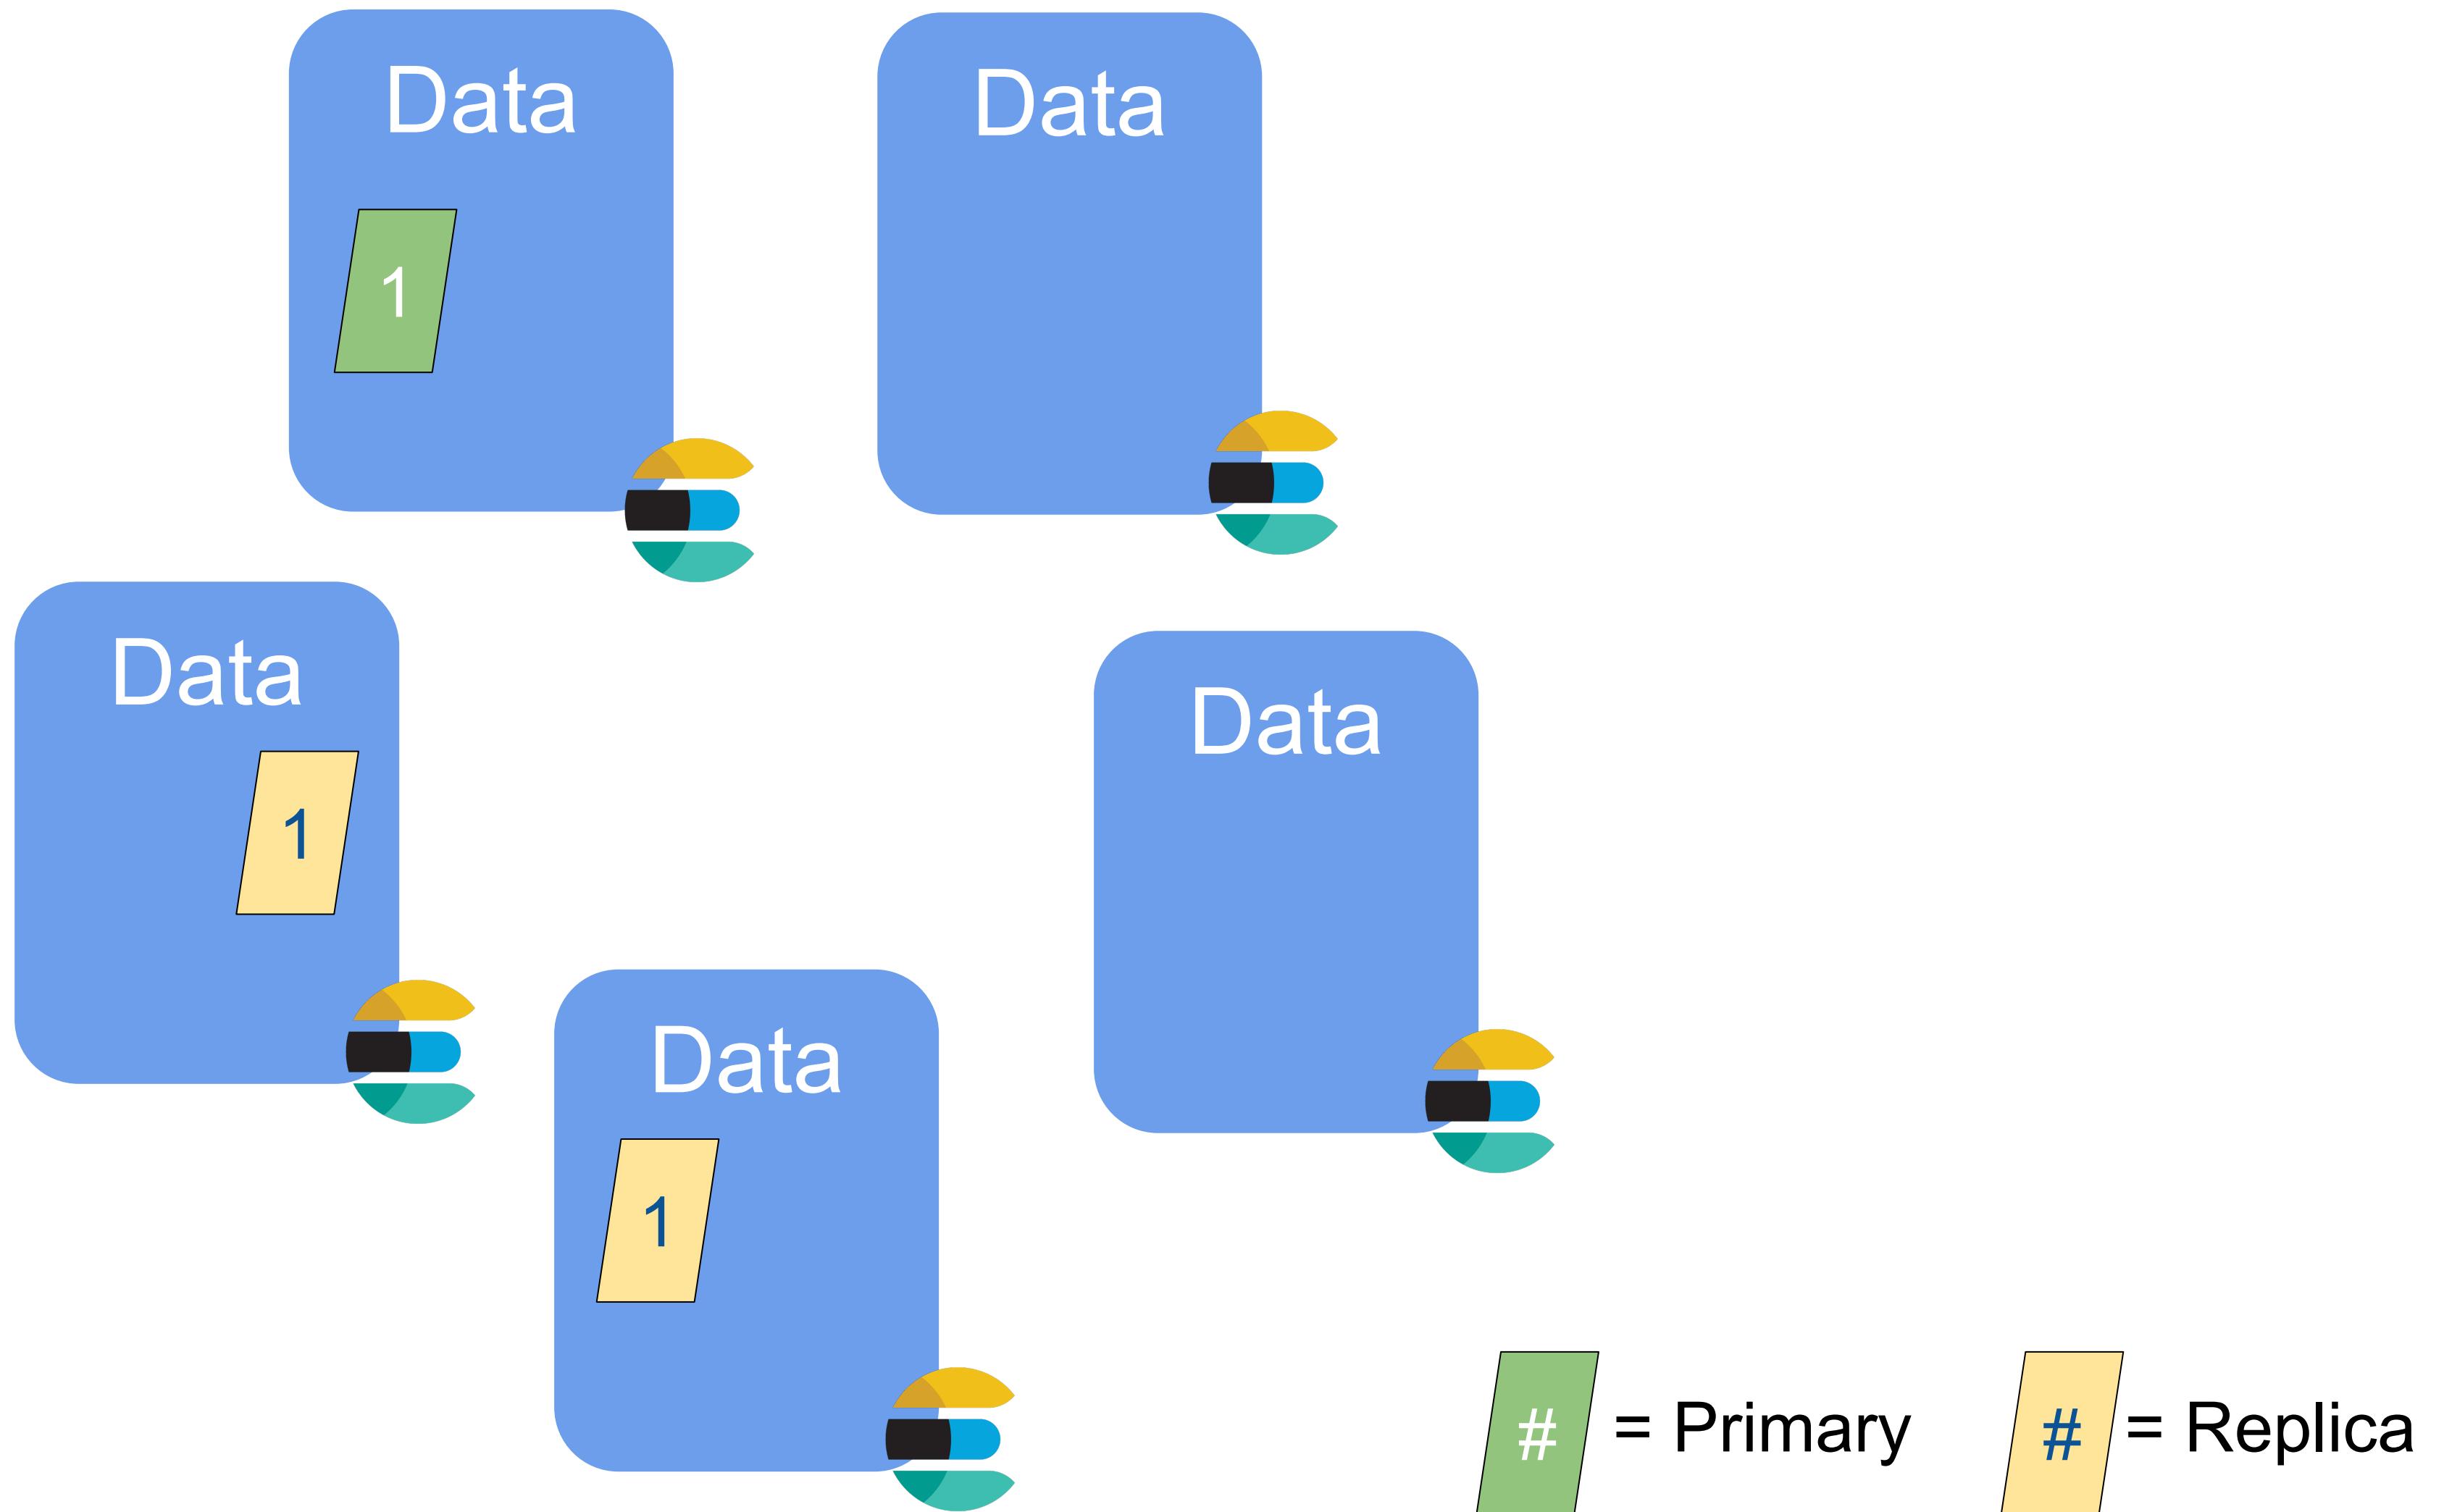
Communication Failure



Writing to Elasticsearch

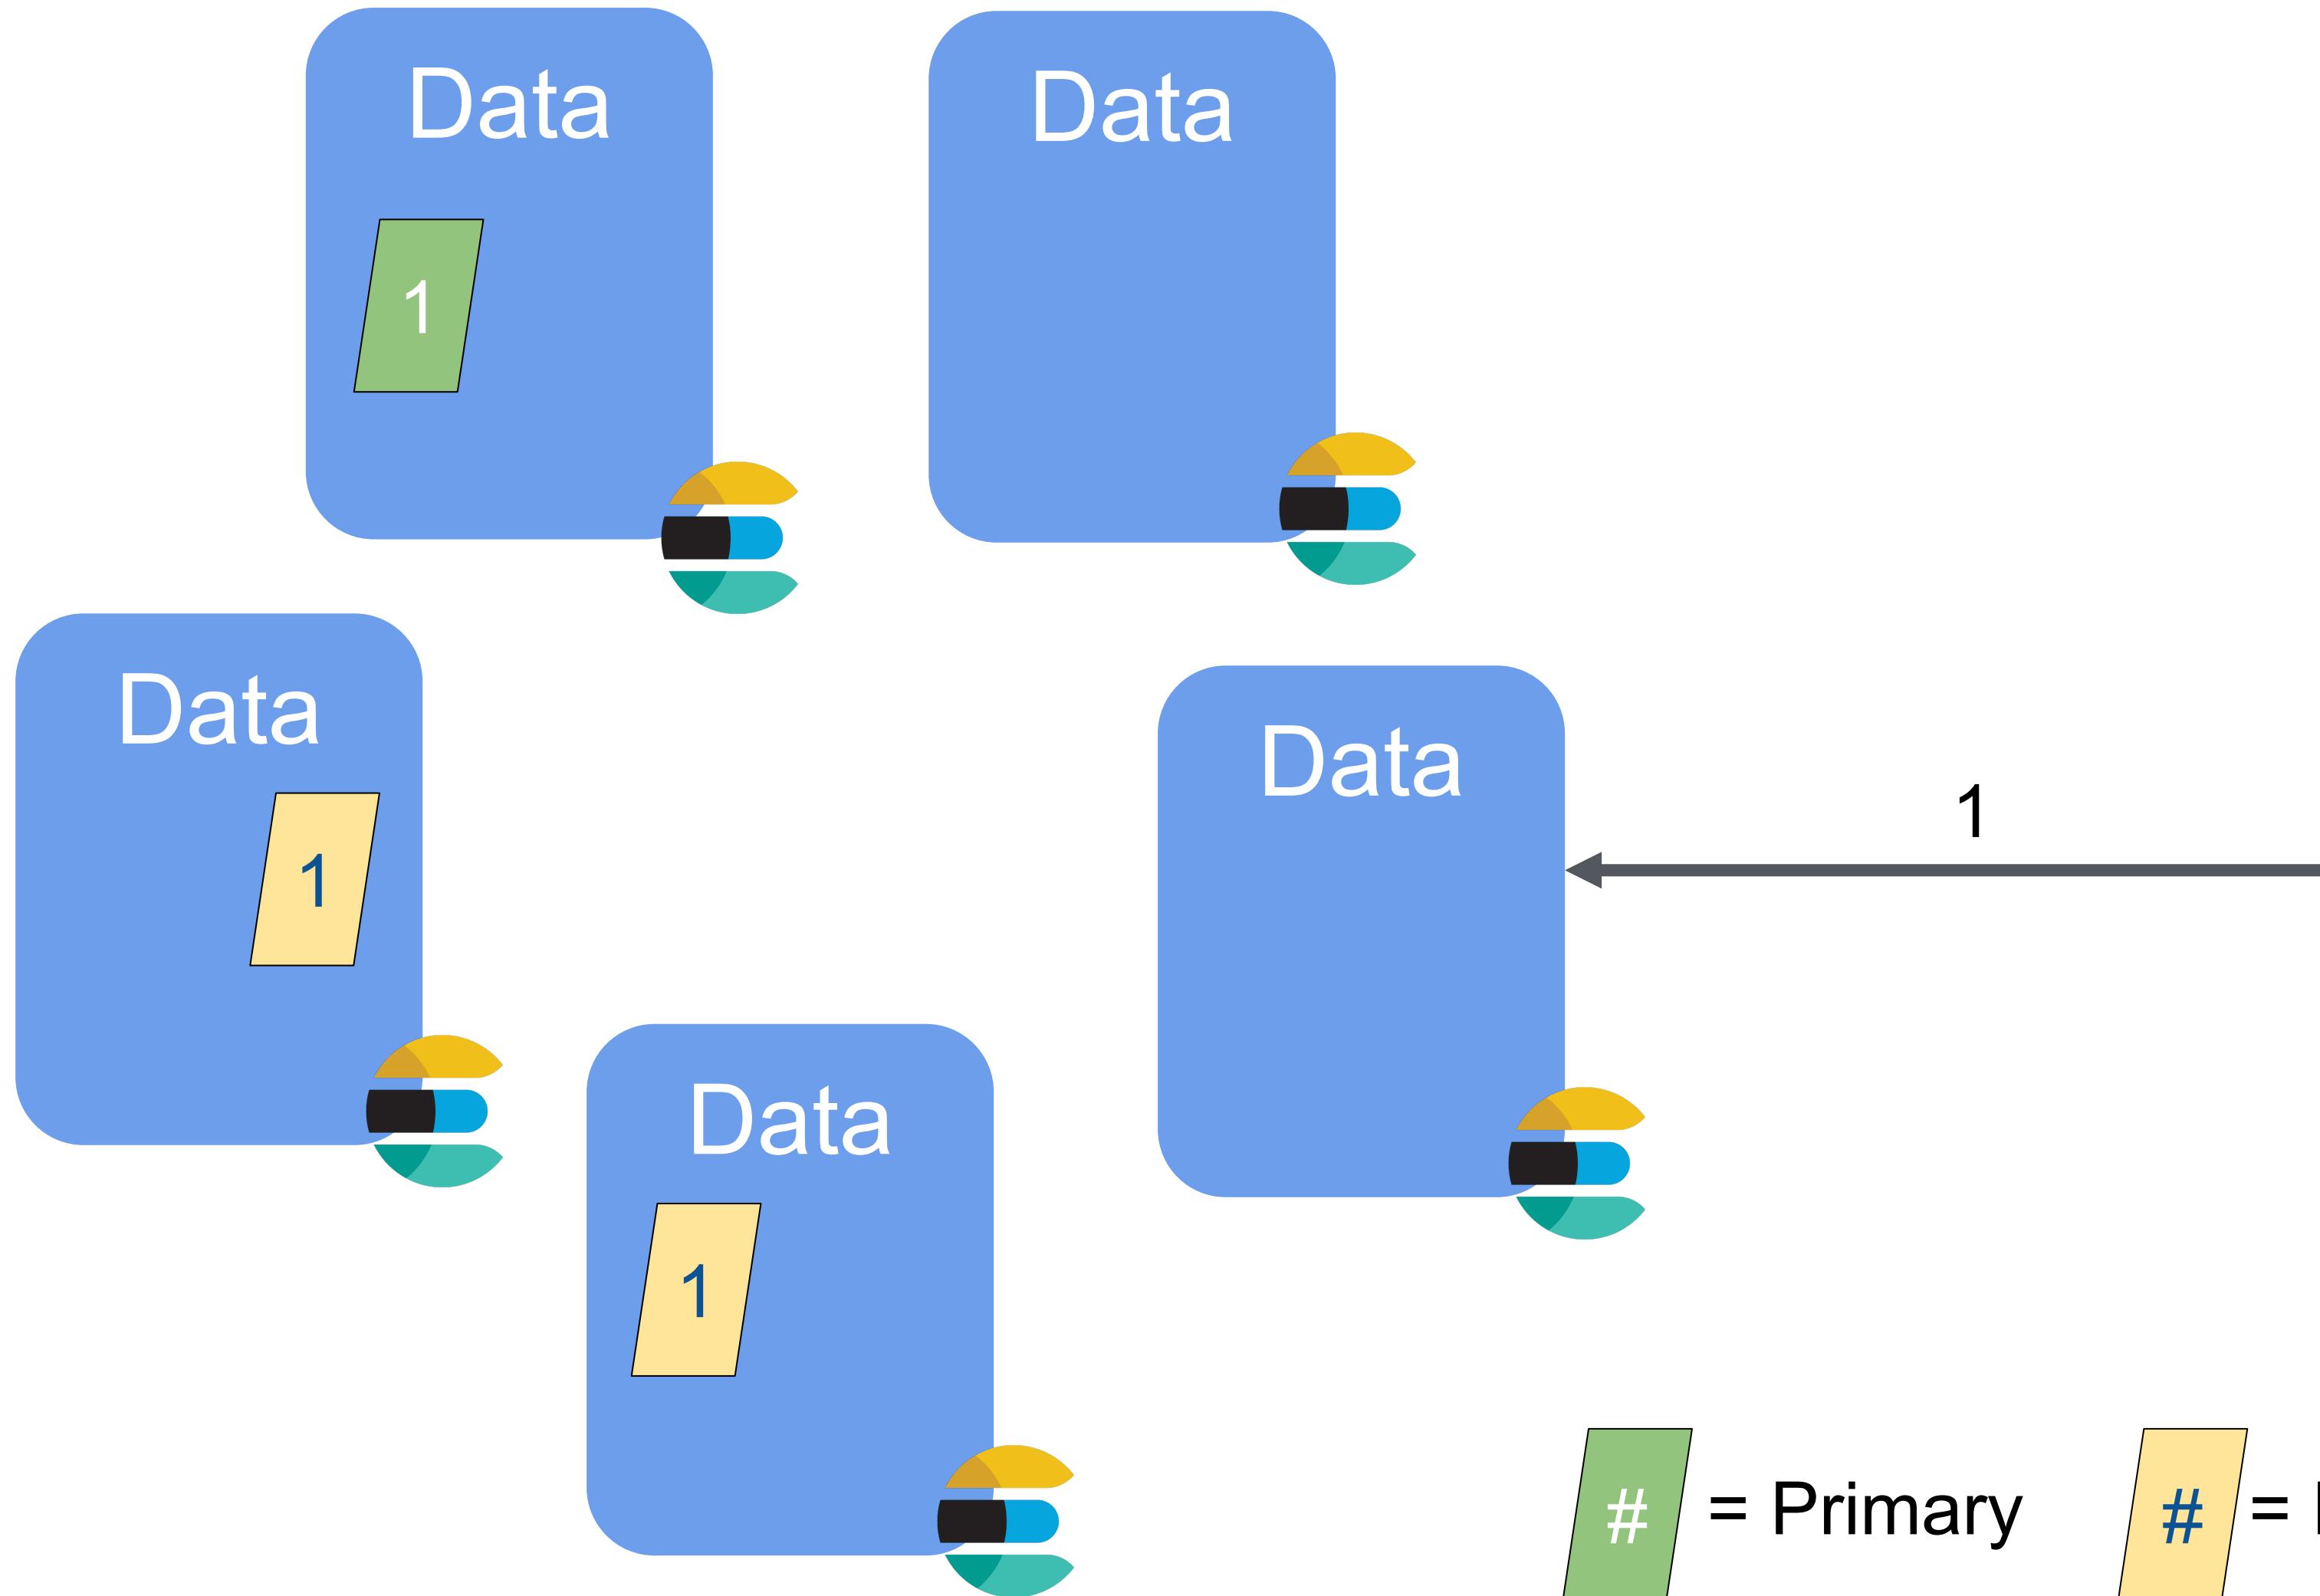
Finding Partitions

Writing to index ...



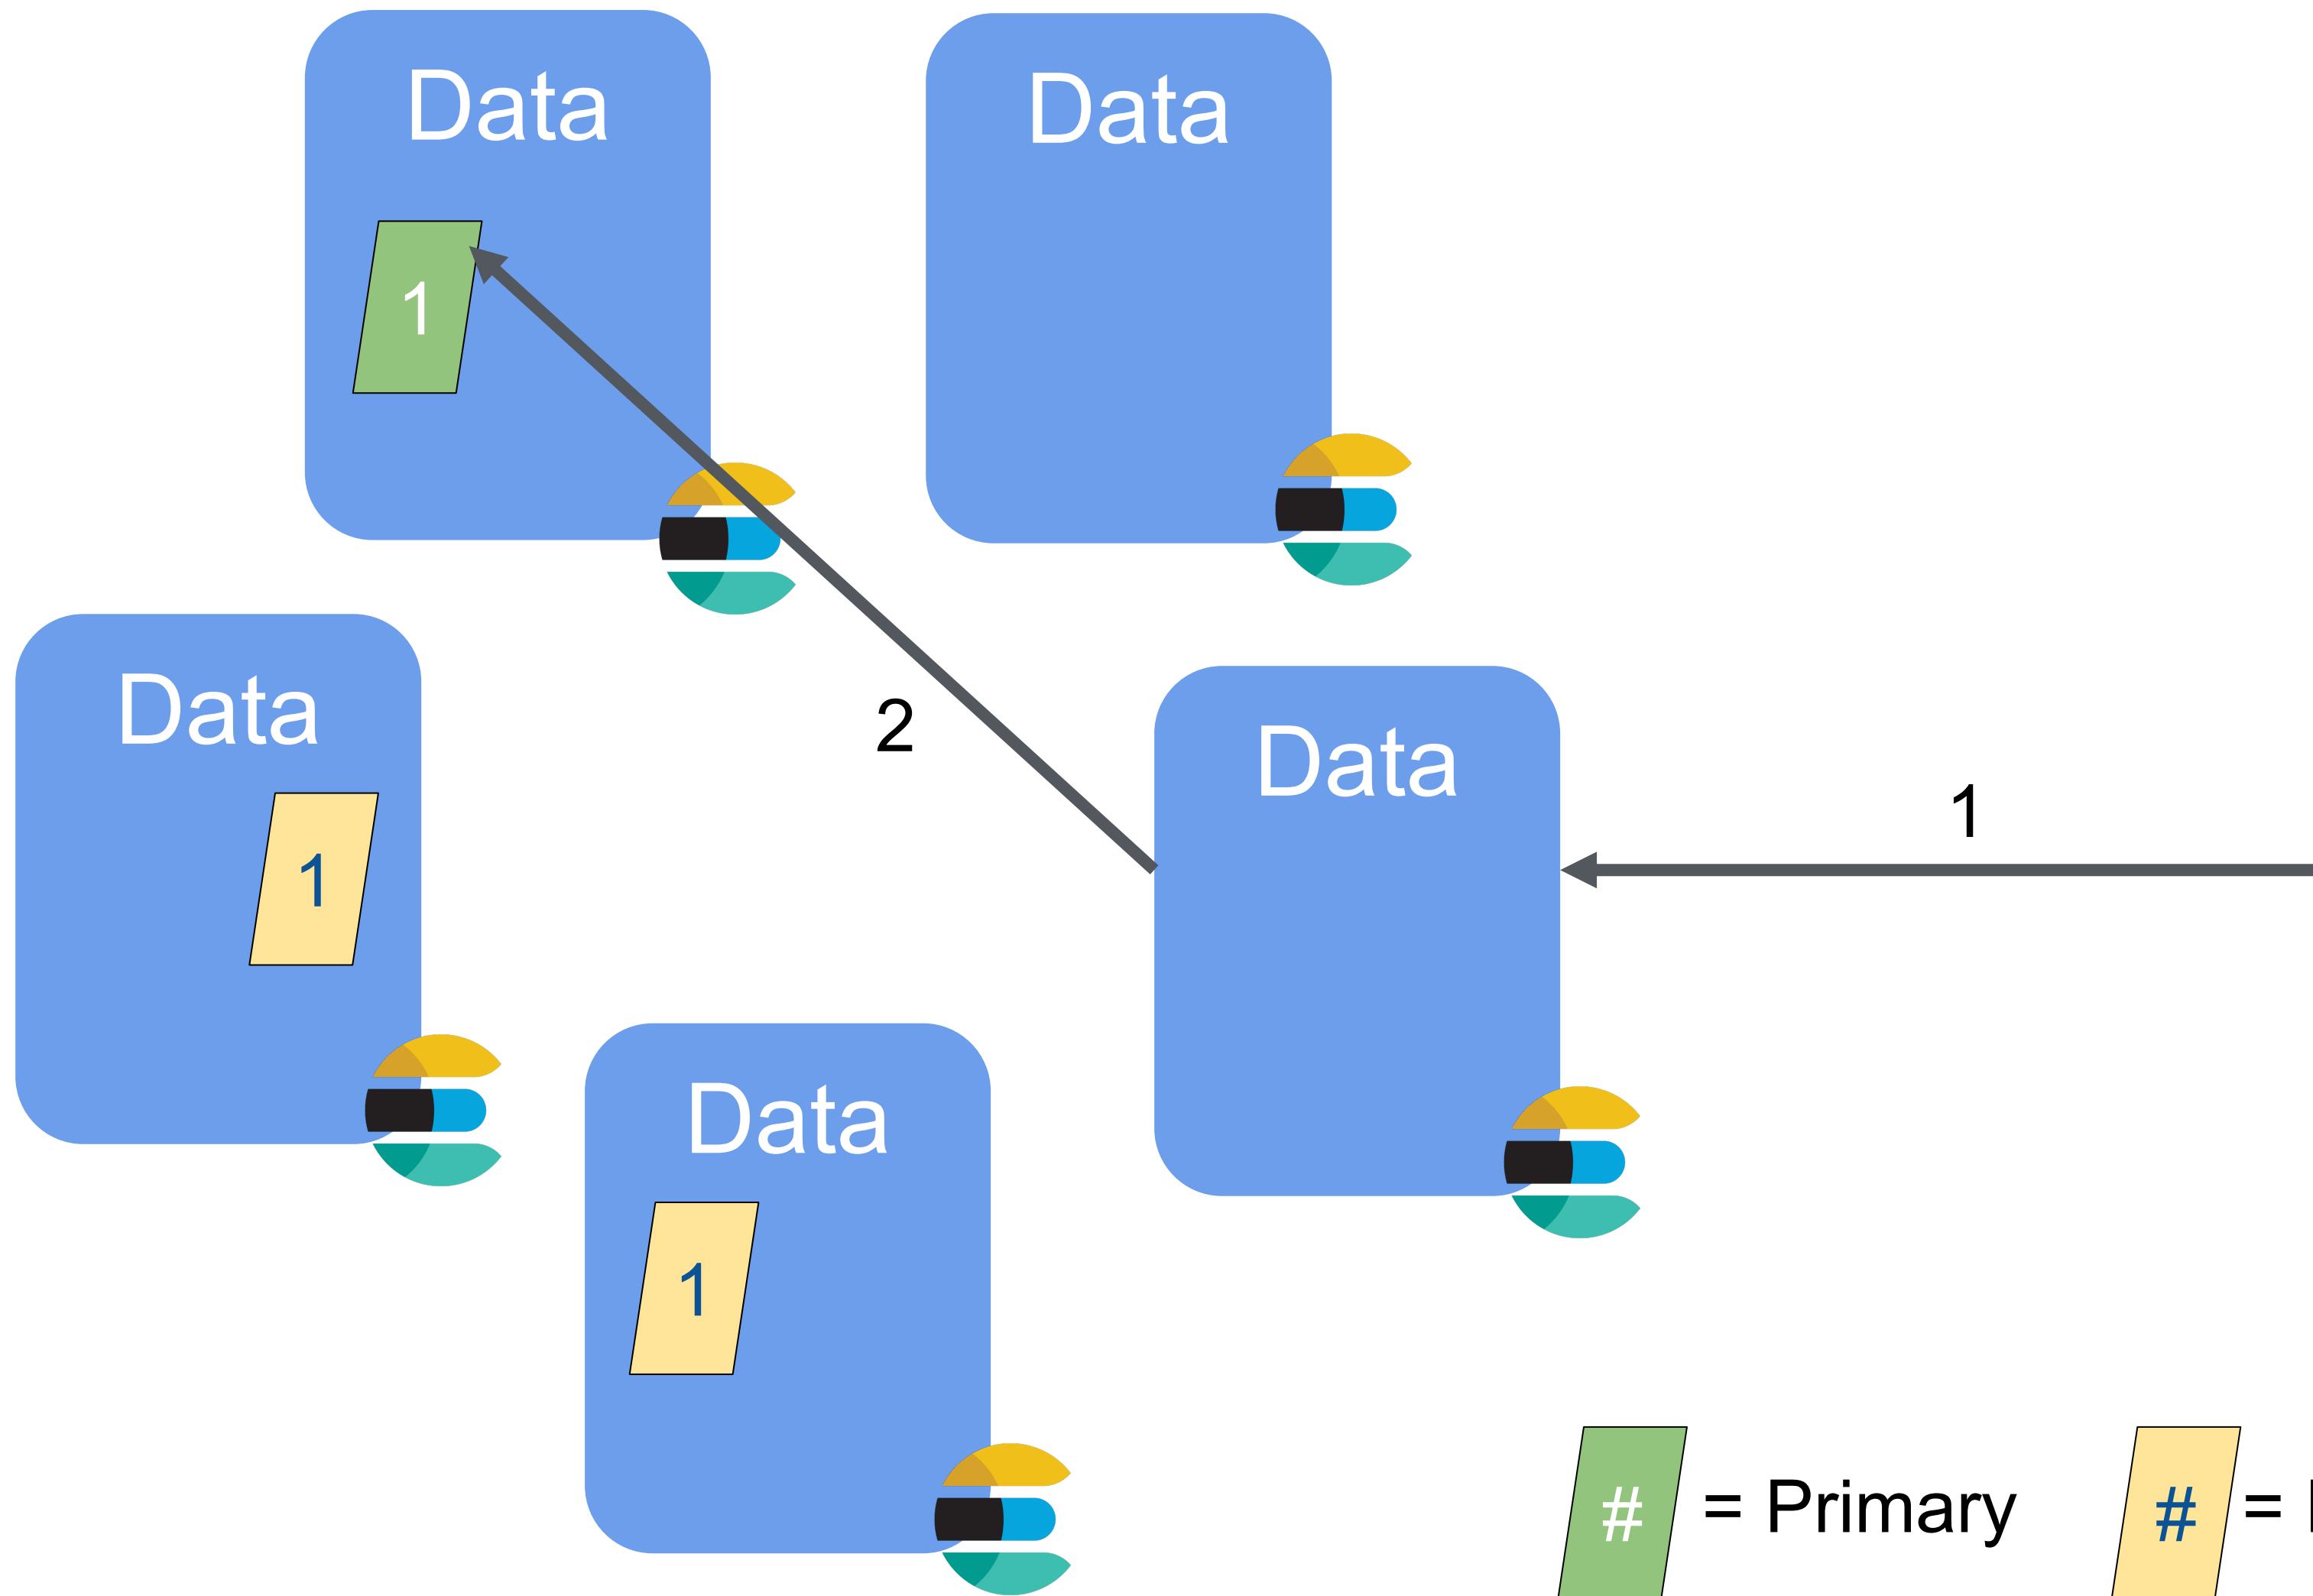
Finding Partitions

Writing to index ...



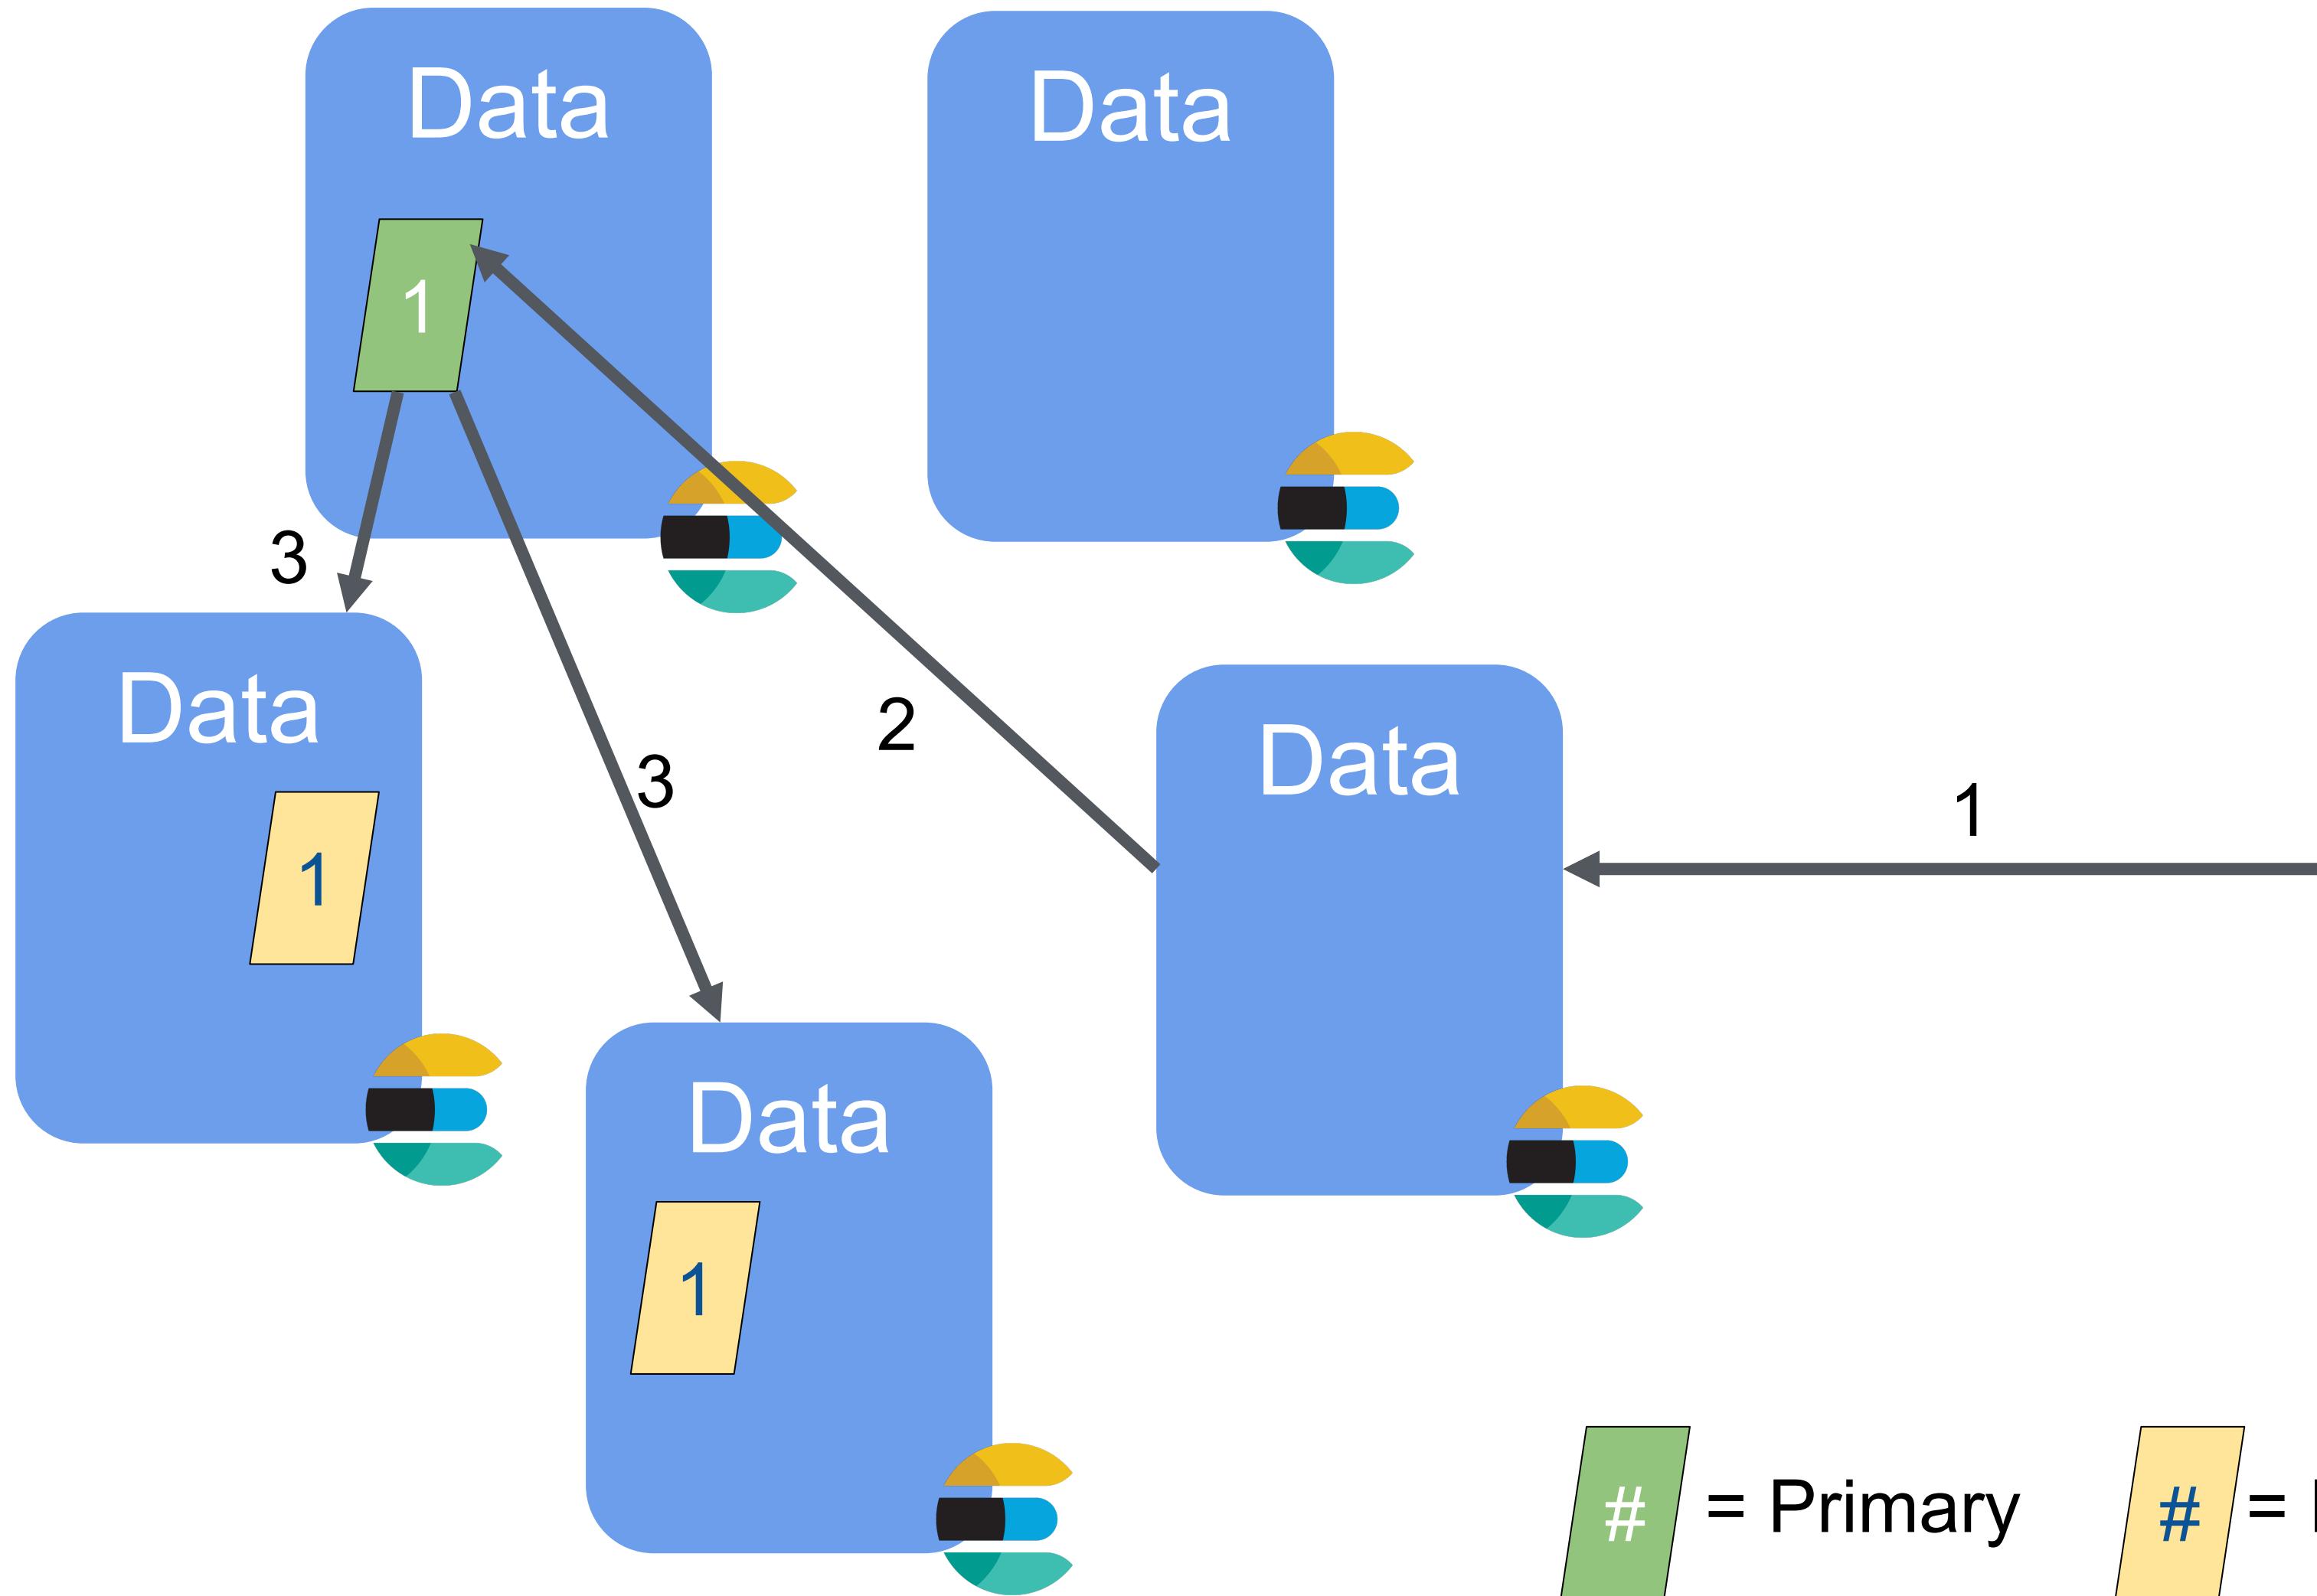
Finding Partitions

Writing to index ...



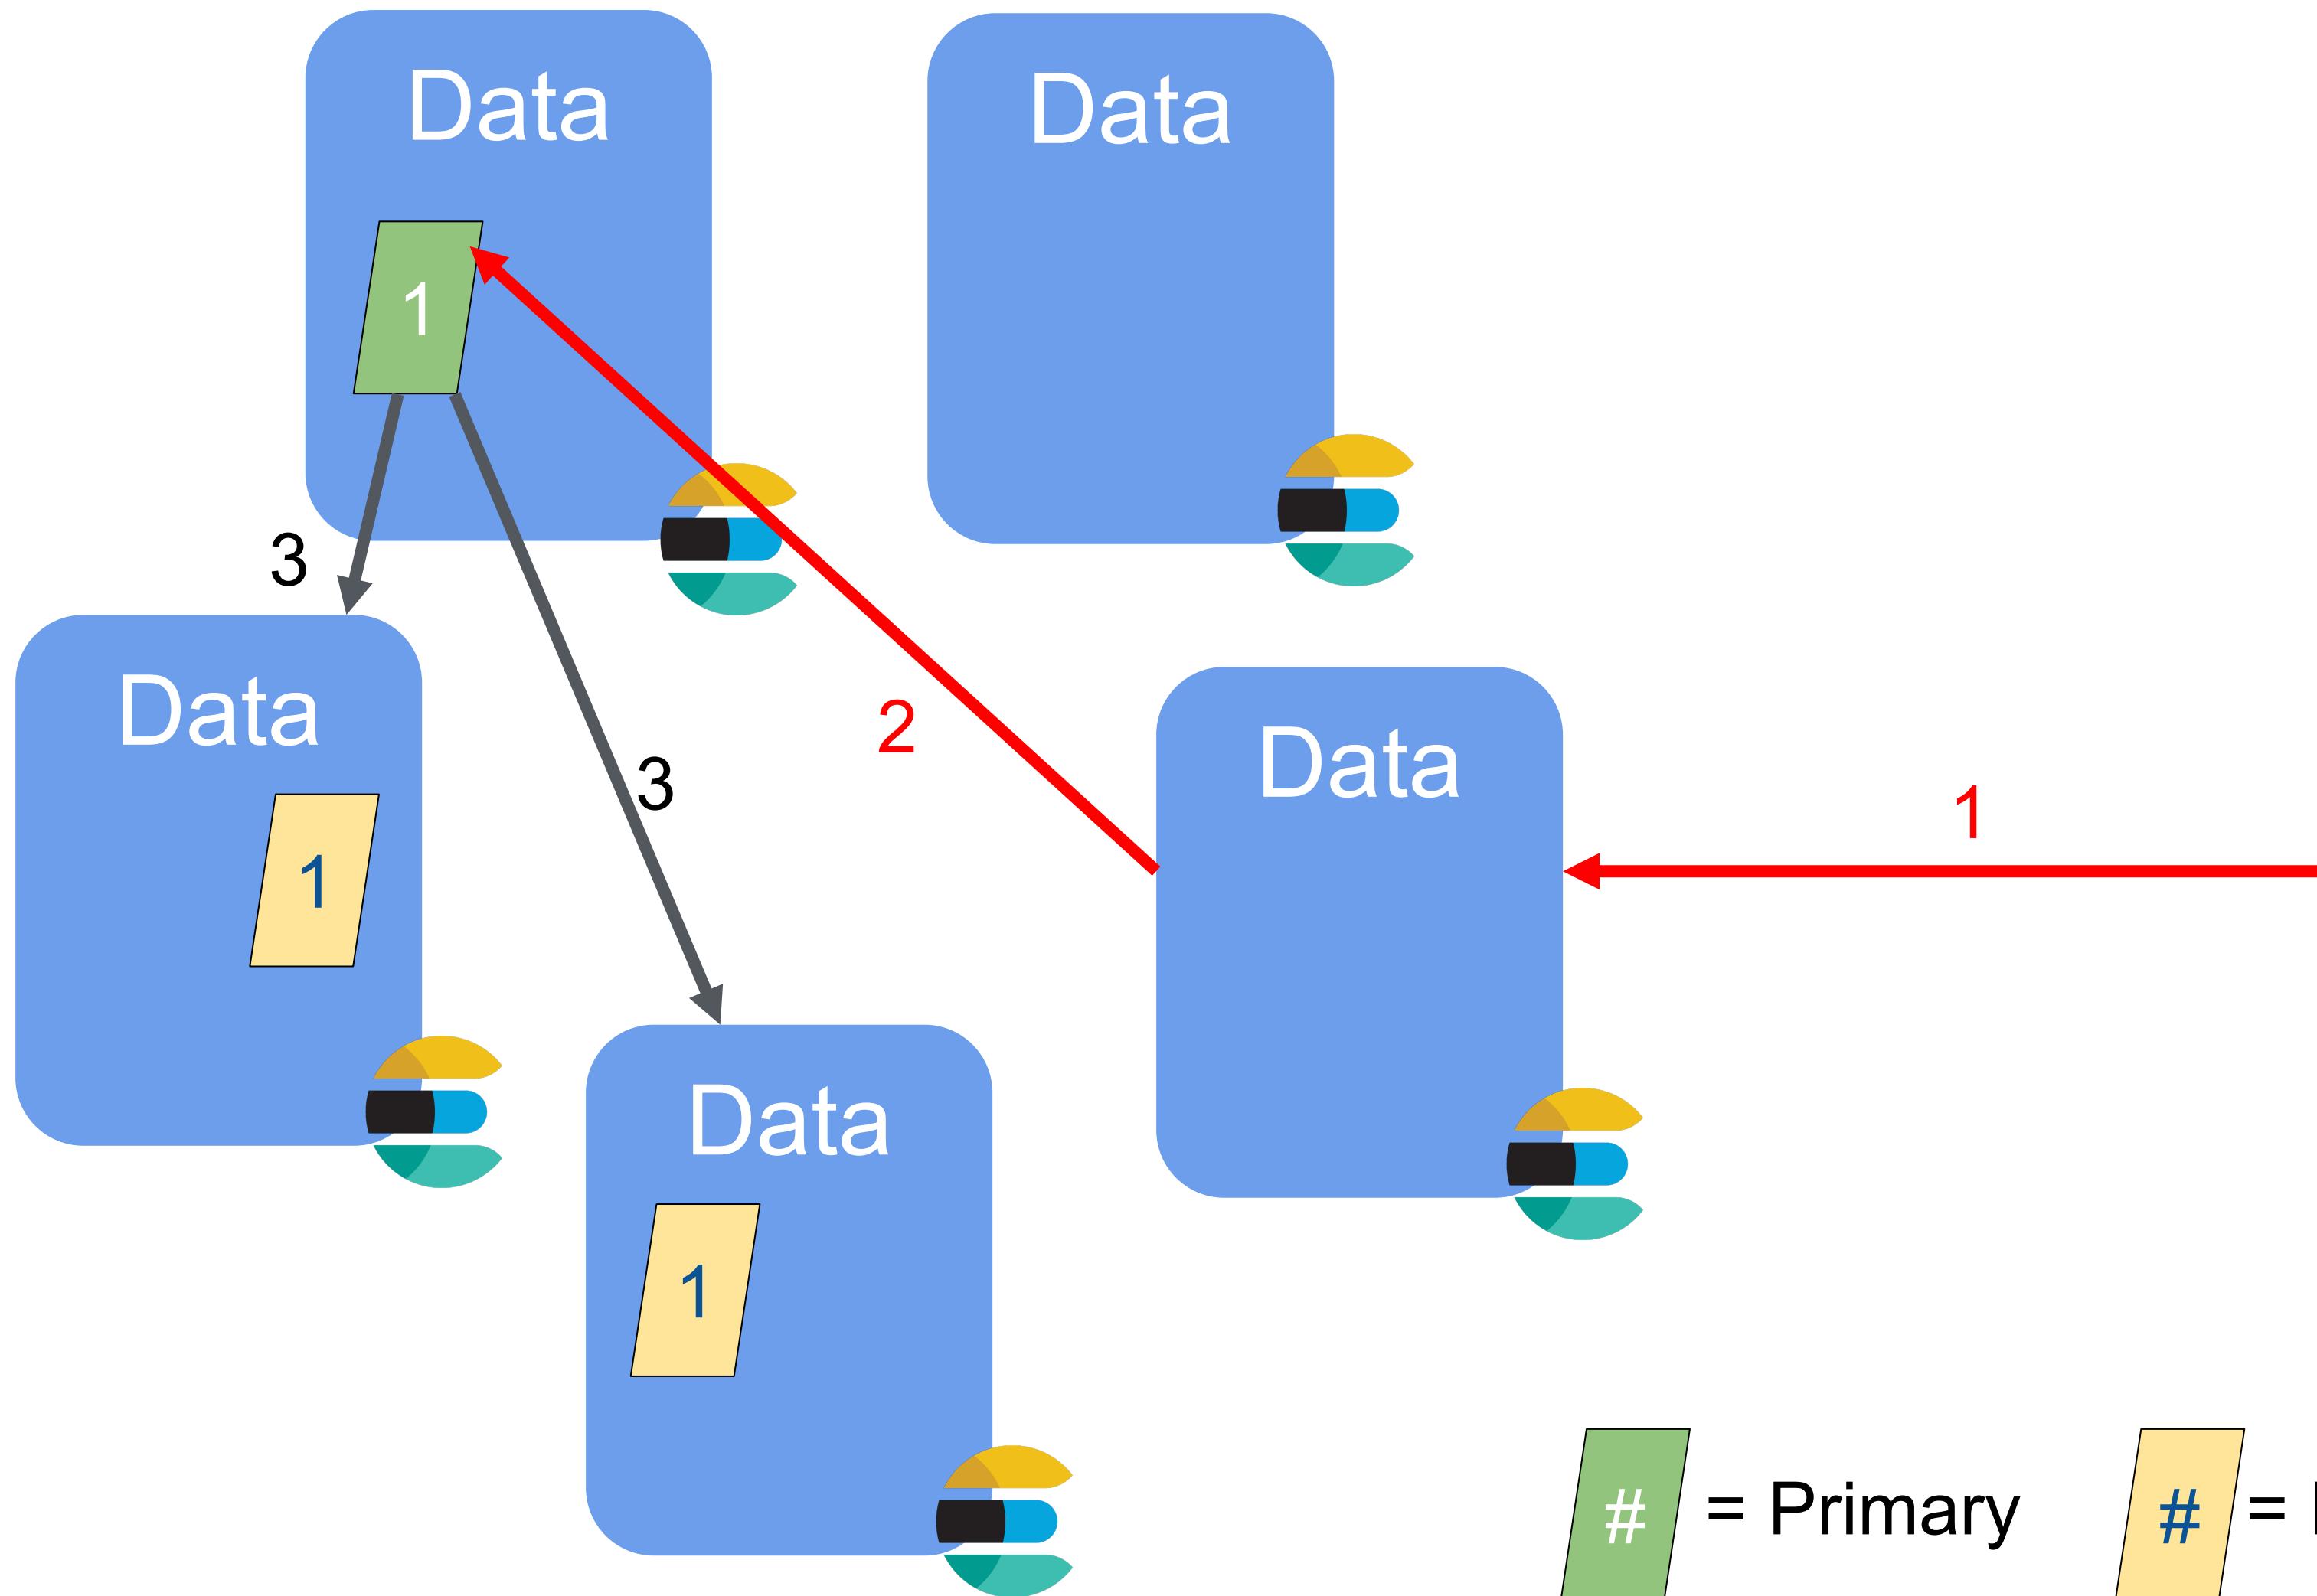
Finding Partitions

Writing to index ...



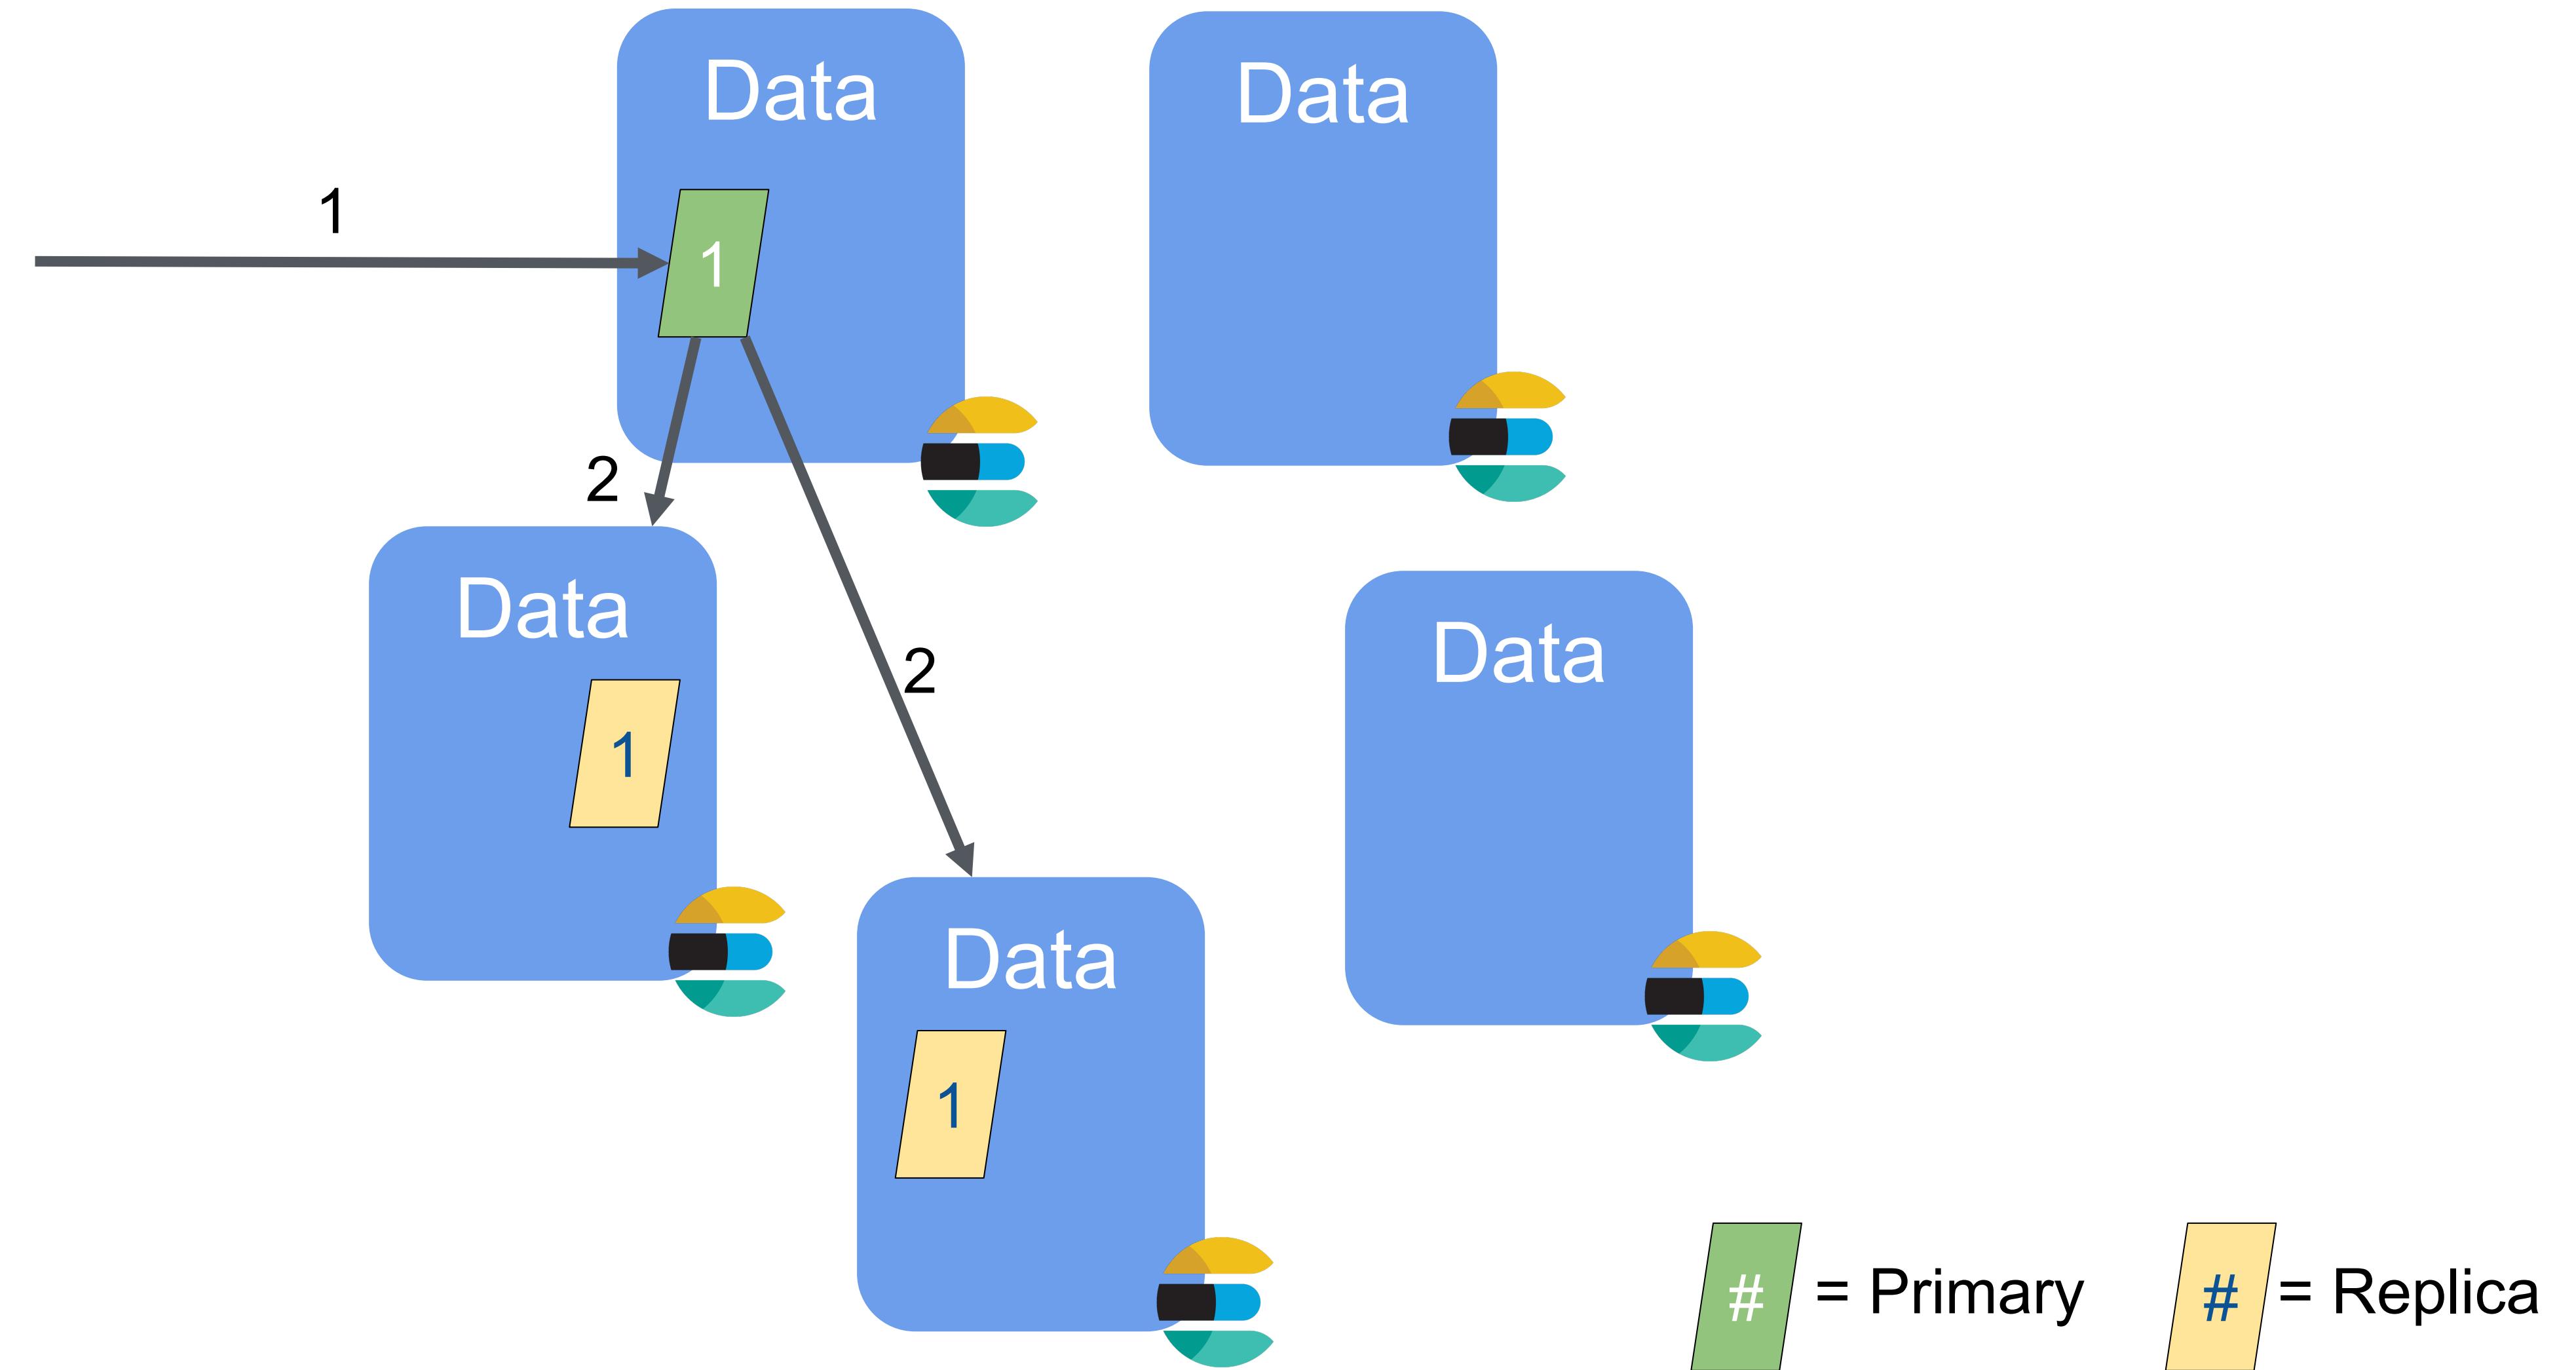
Finding Partitions

Writing to index ...



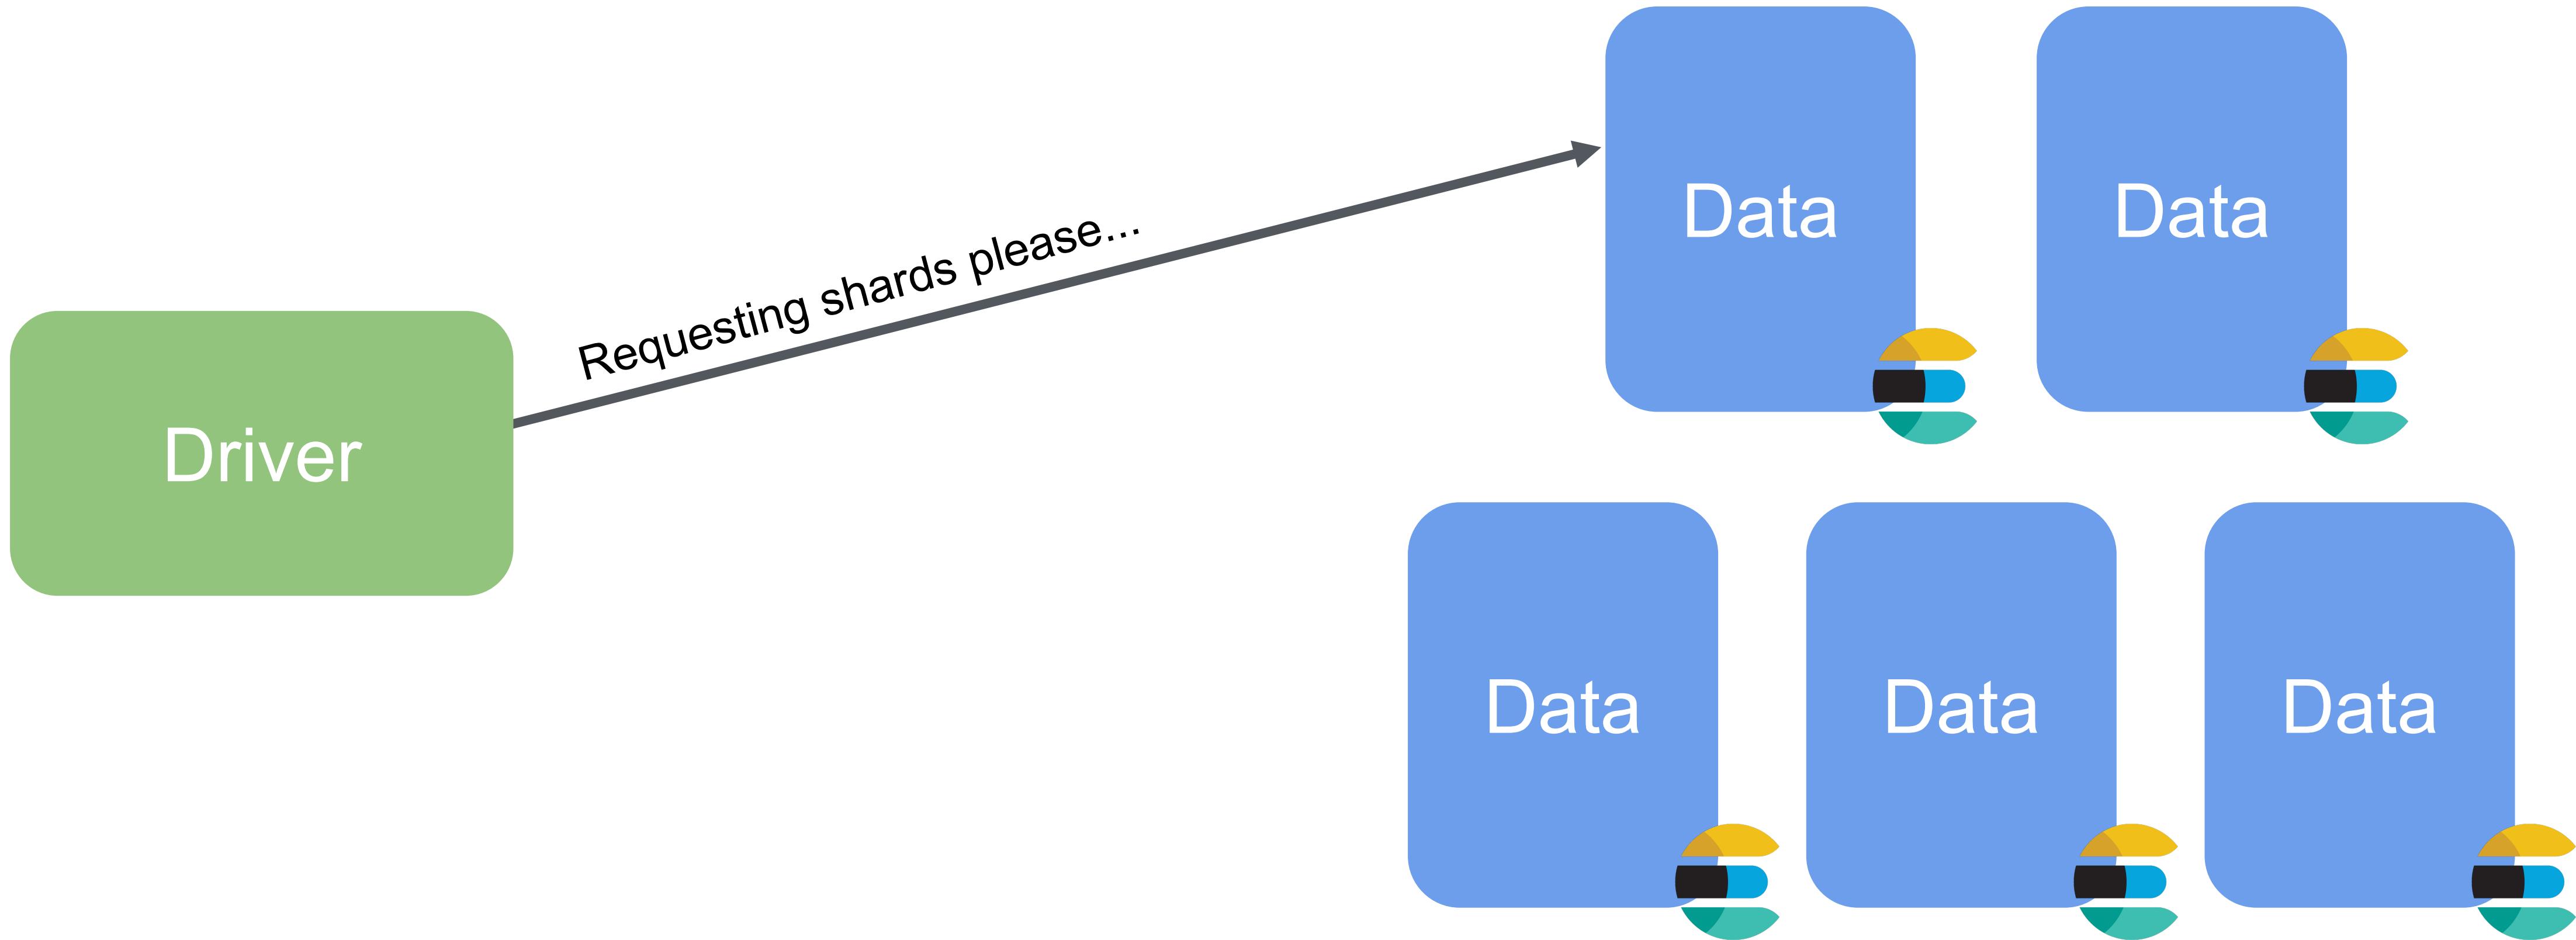
Finding Partitions

Writing to index ...



Finding Partitions

Writing to index ...



Finding Partitions

Writing to index ...

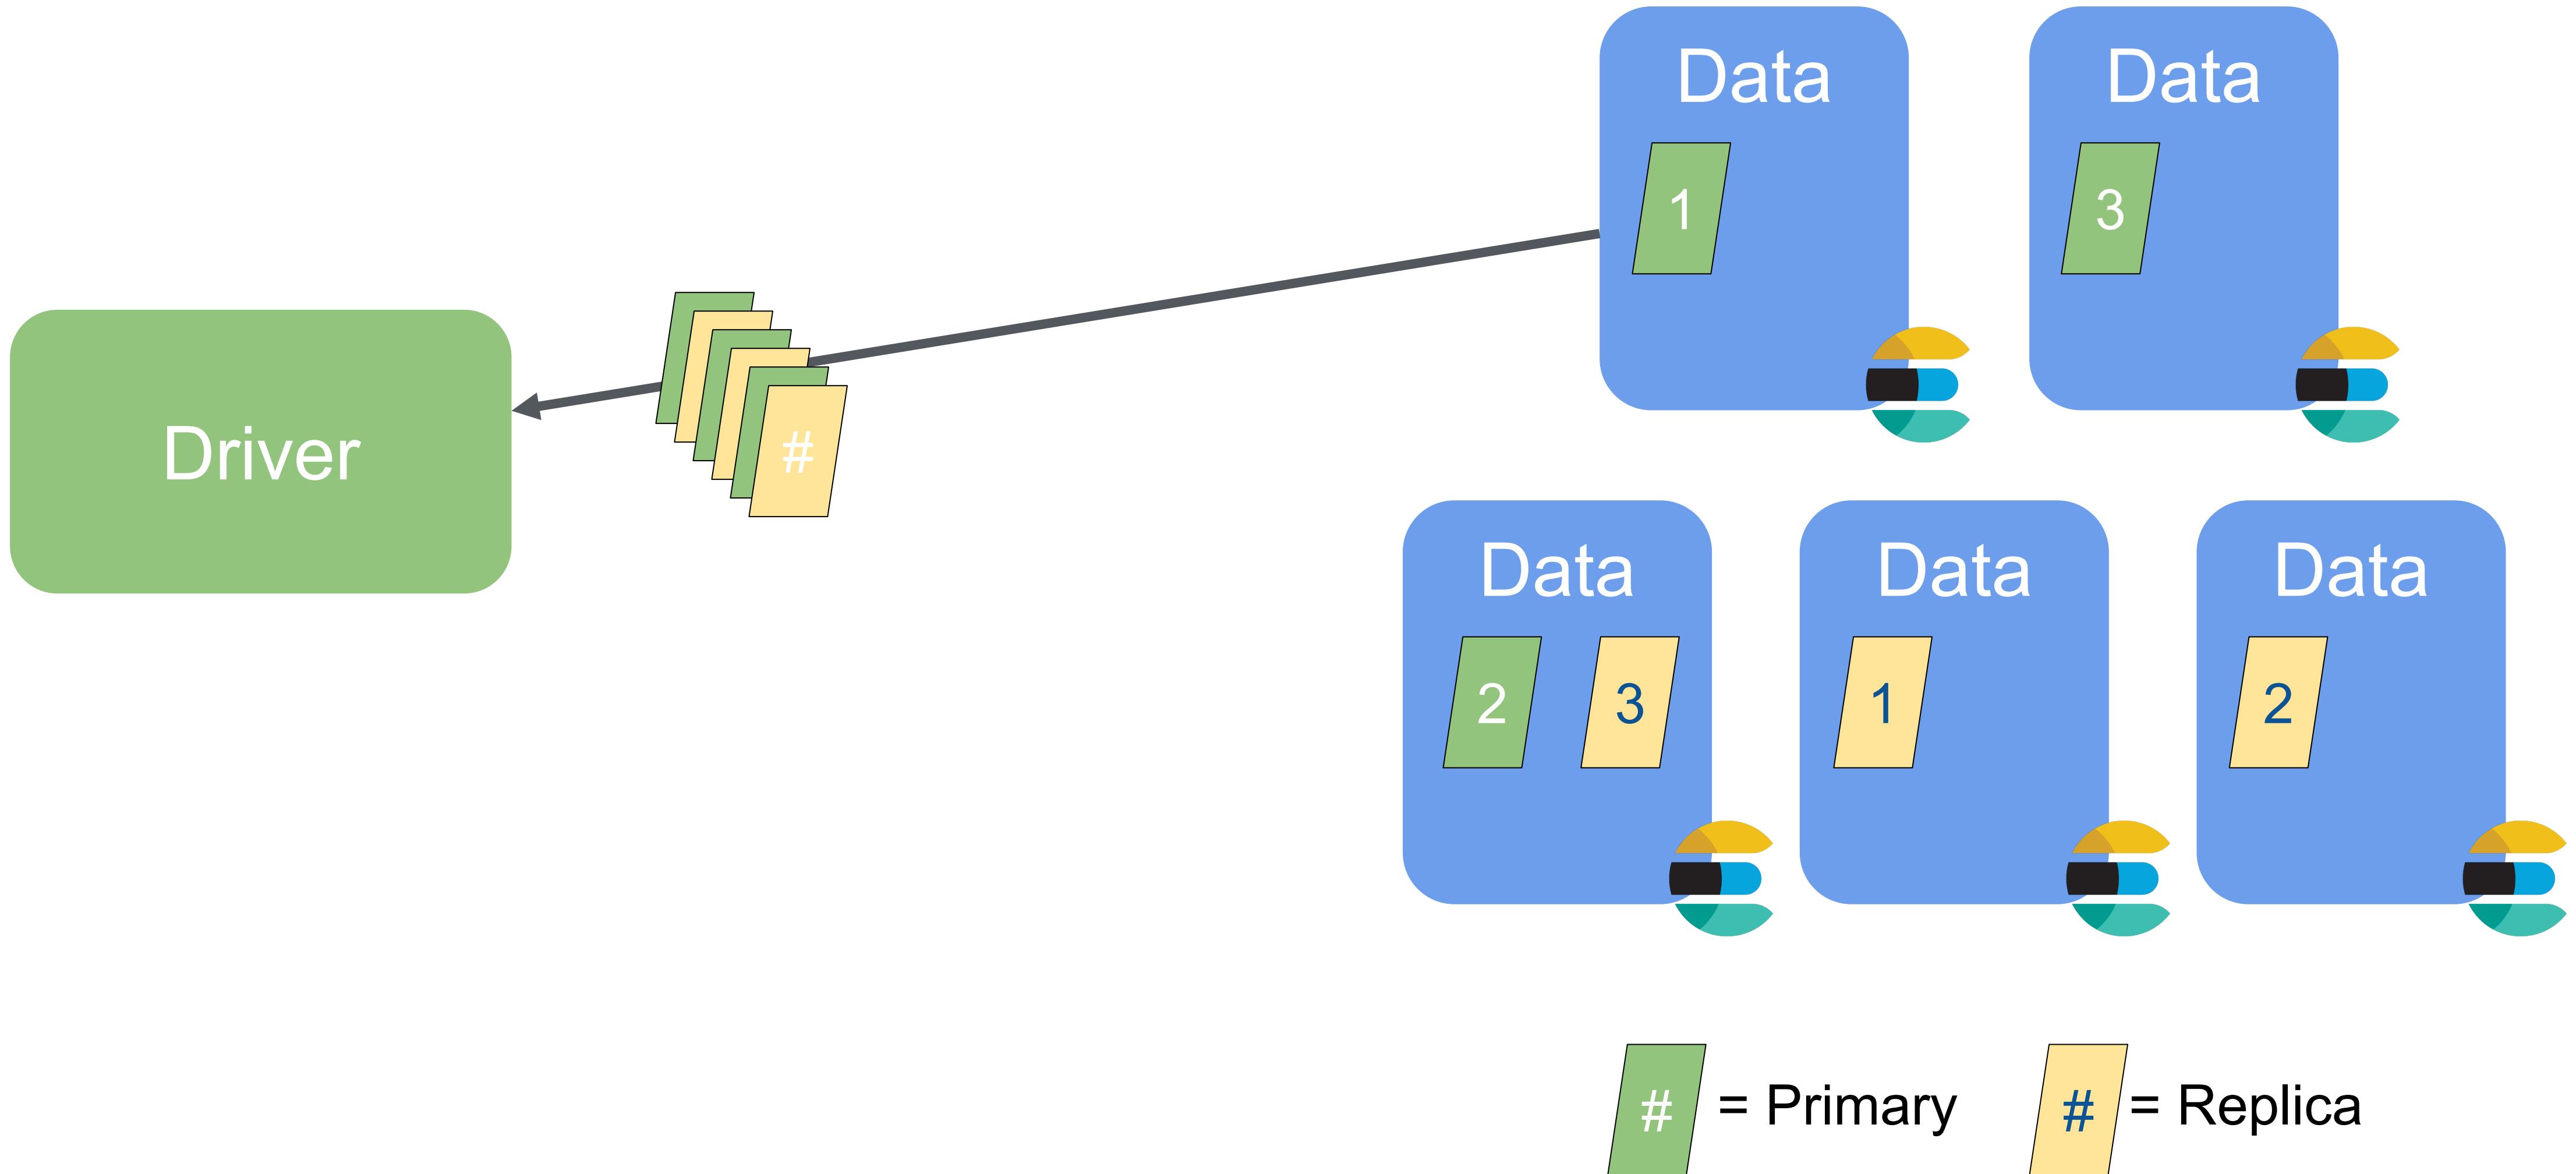


= Primary

= Replica

Finding Partitions

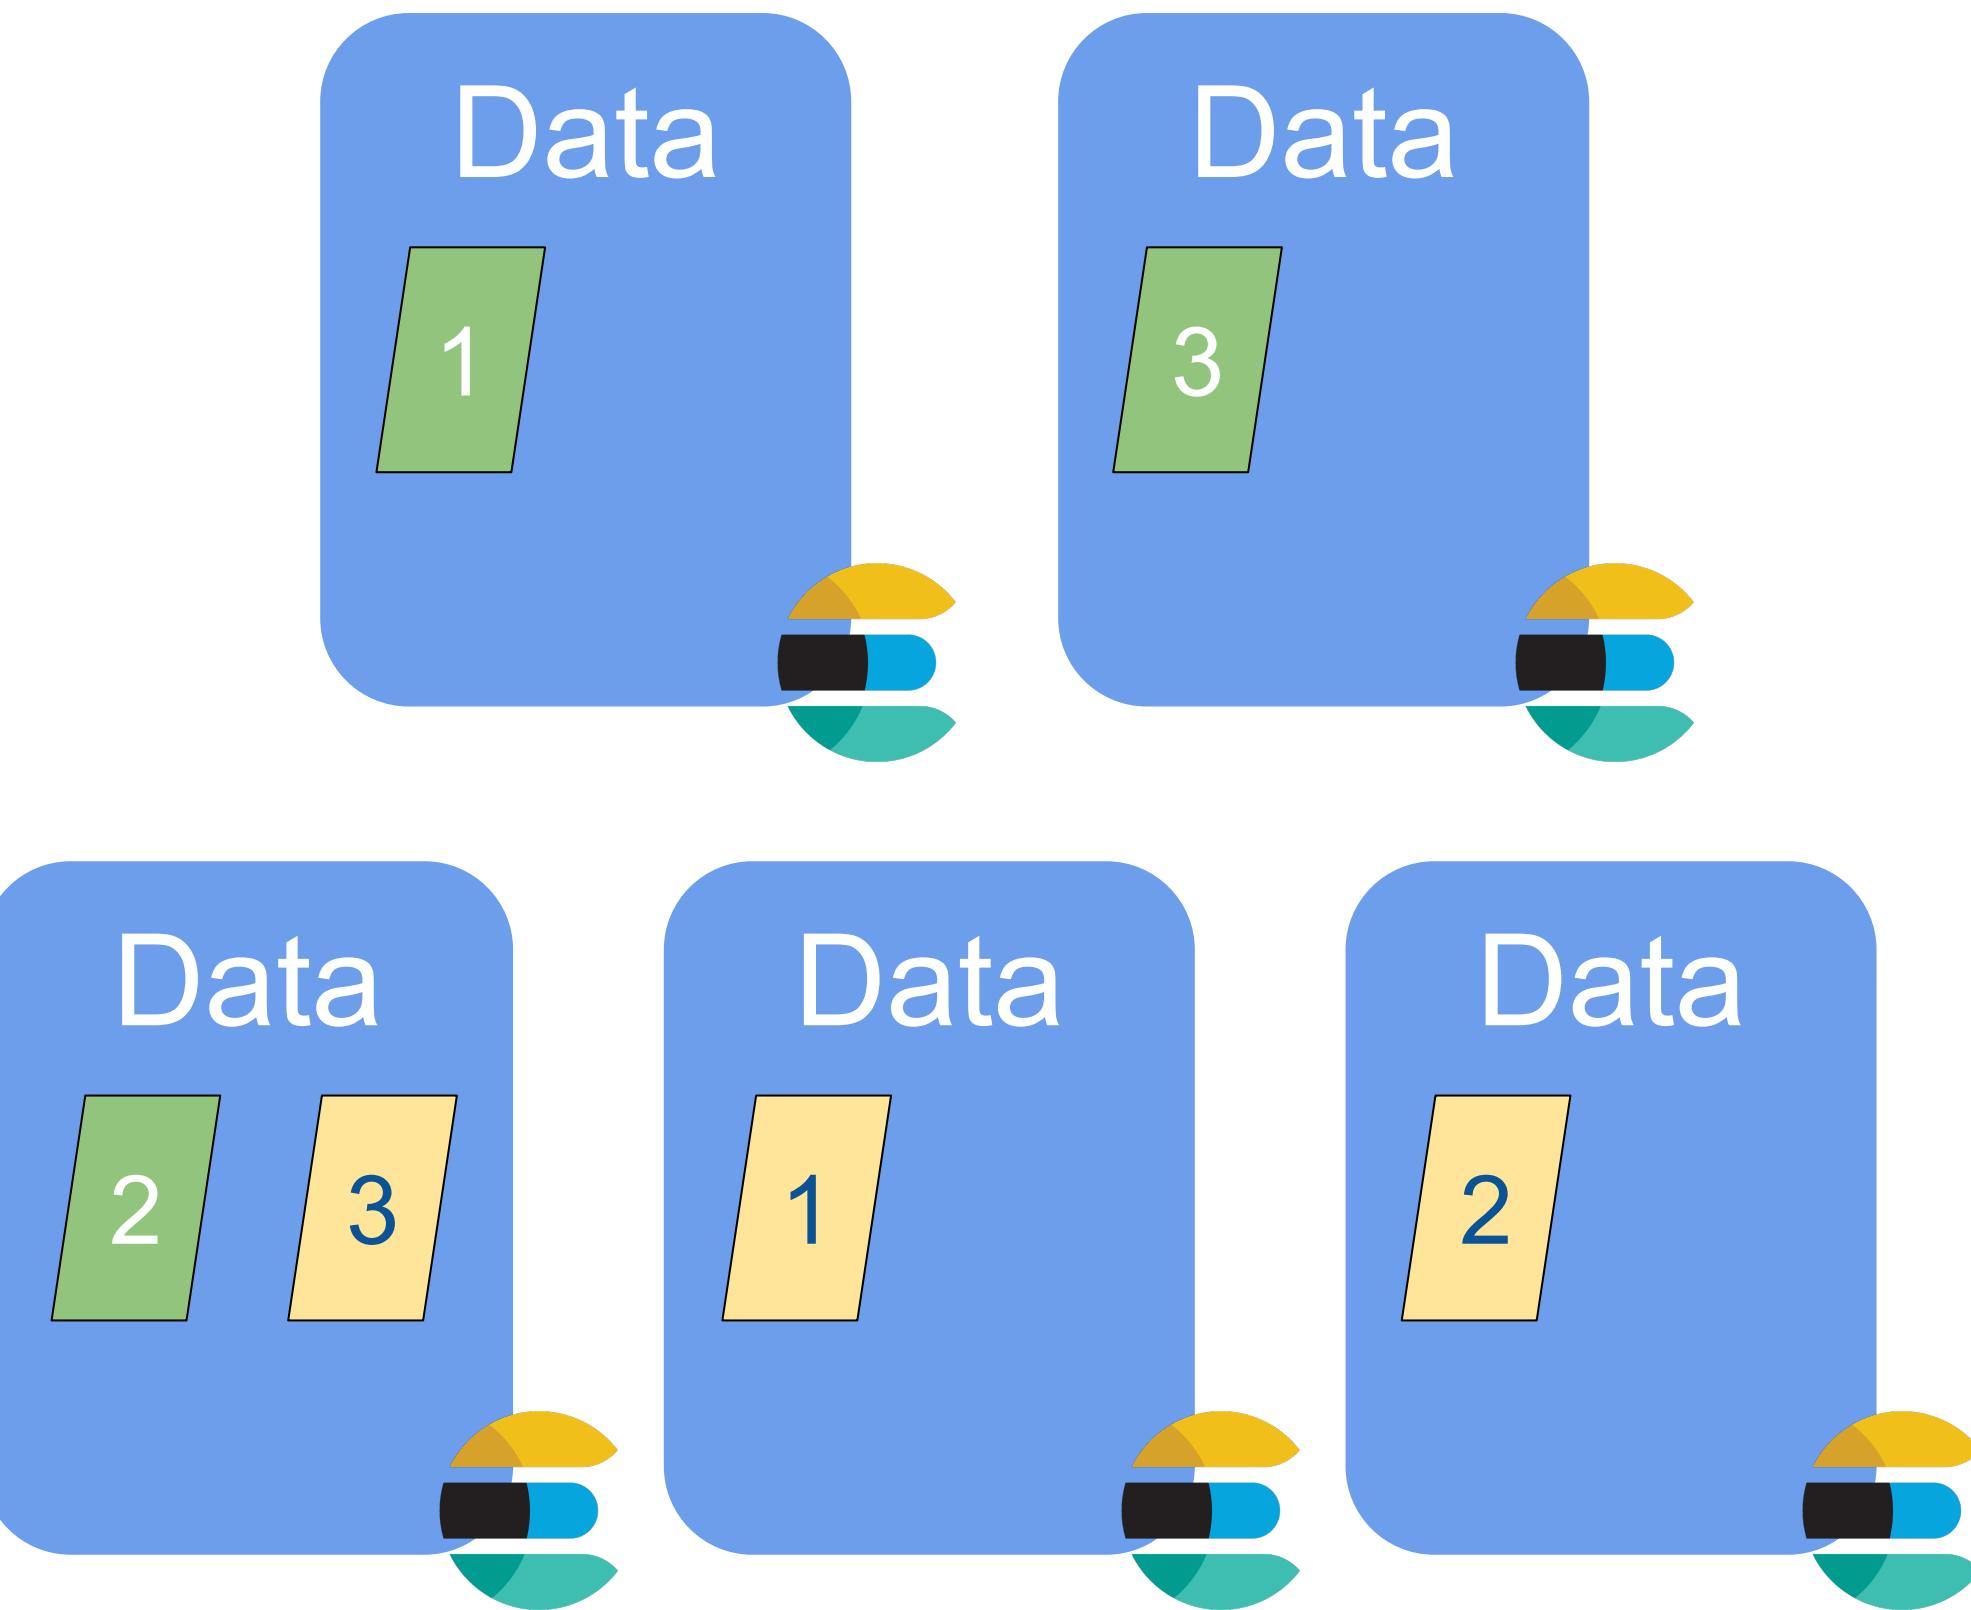
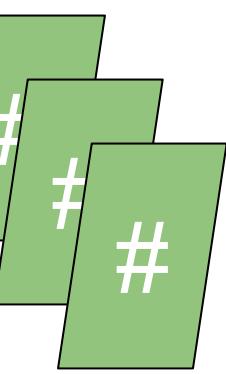
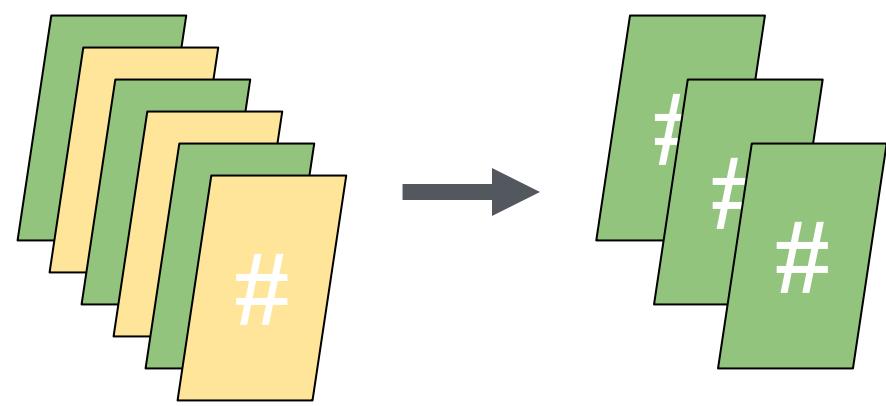
Writing to index ...



Finding Partitions

Writing to index ...

Driver

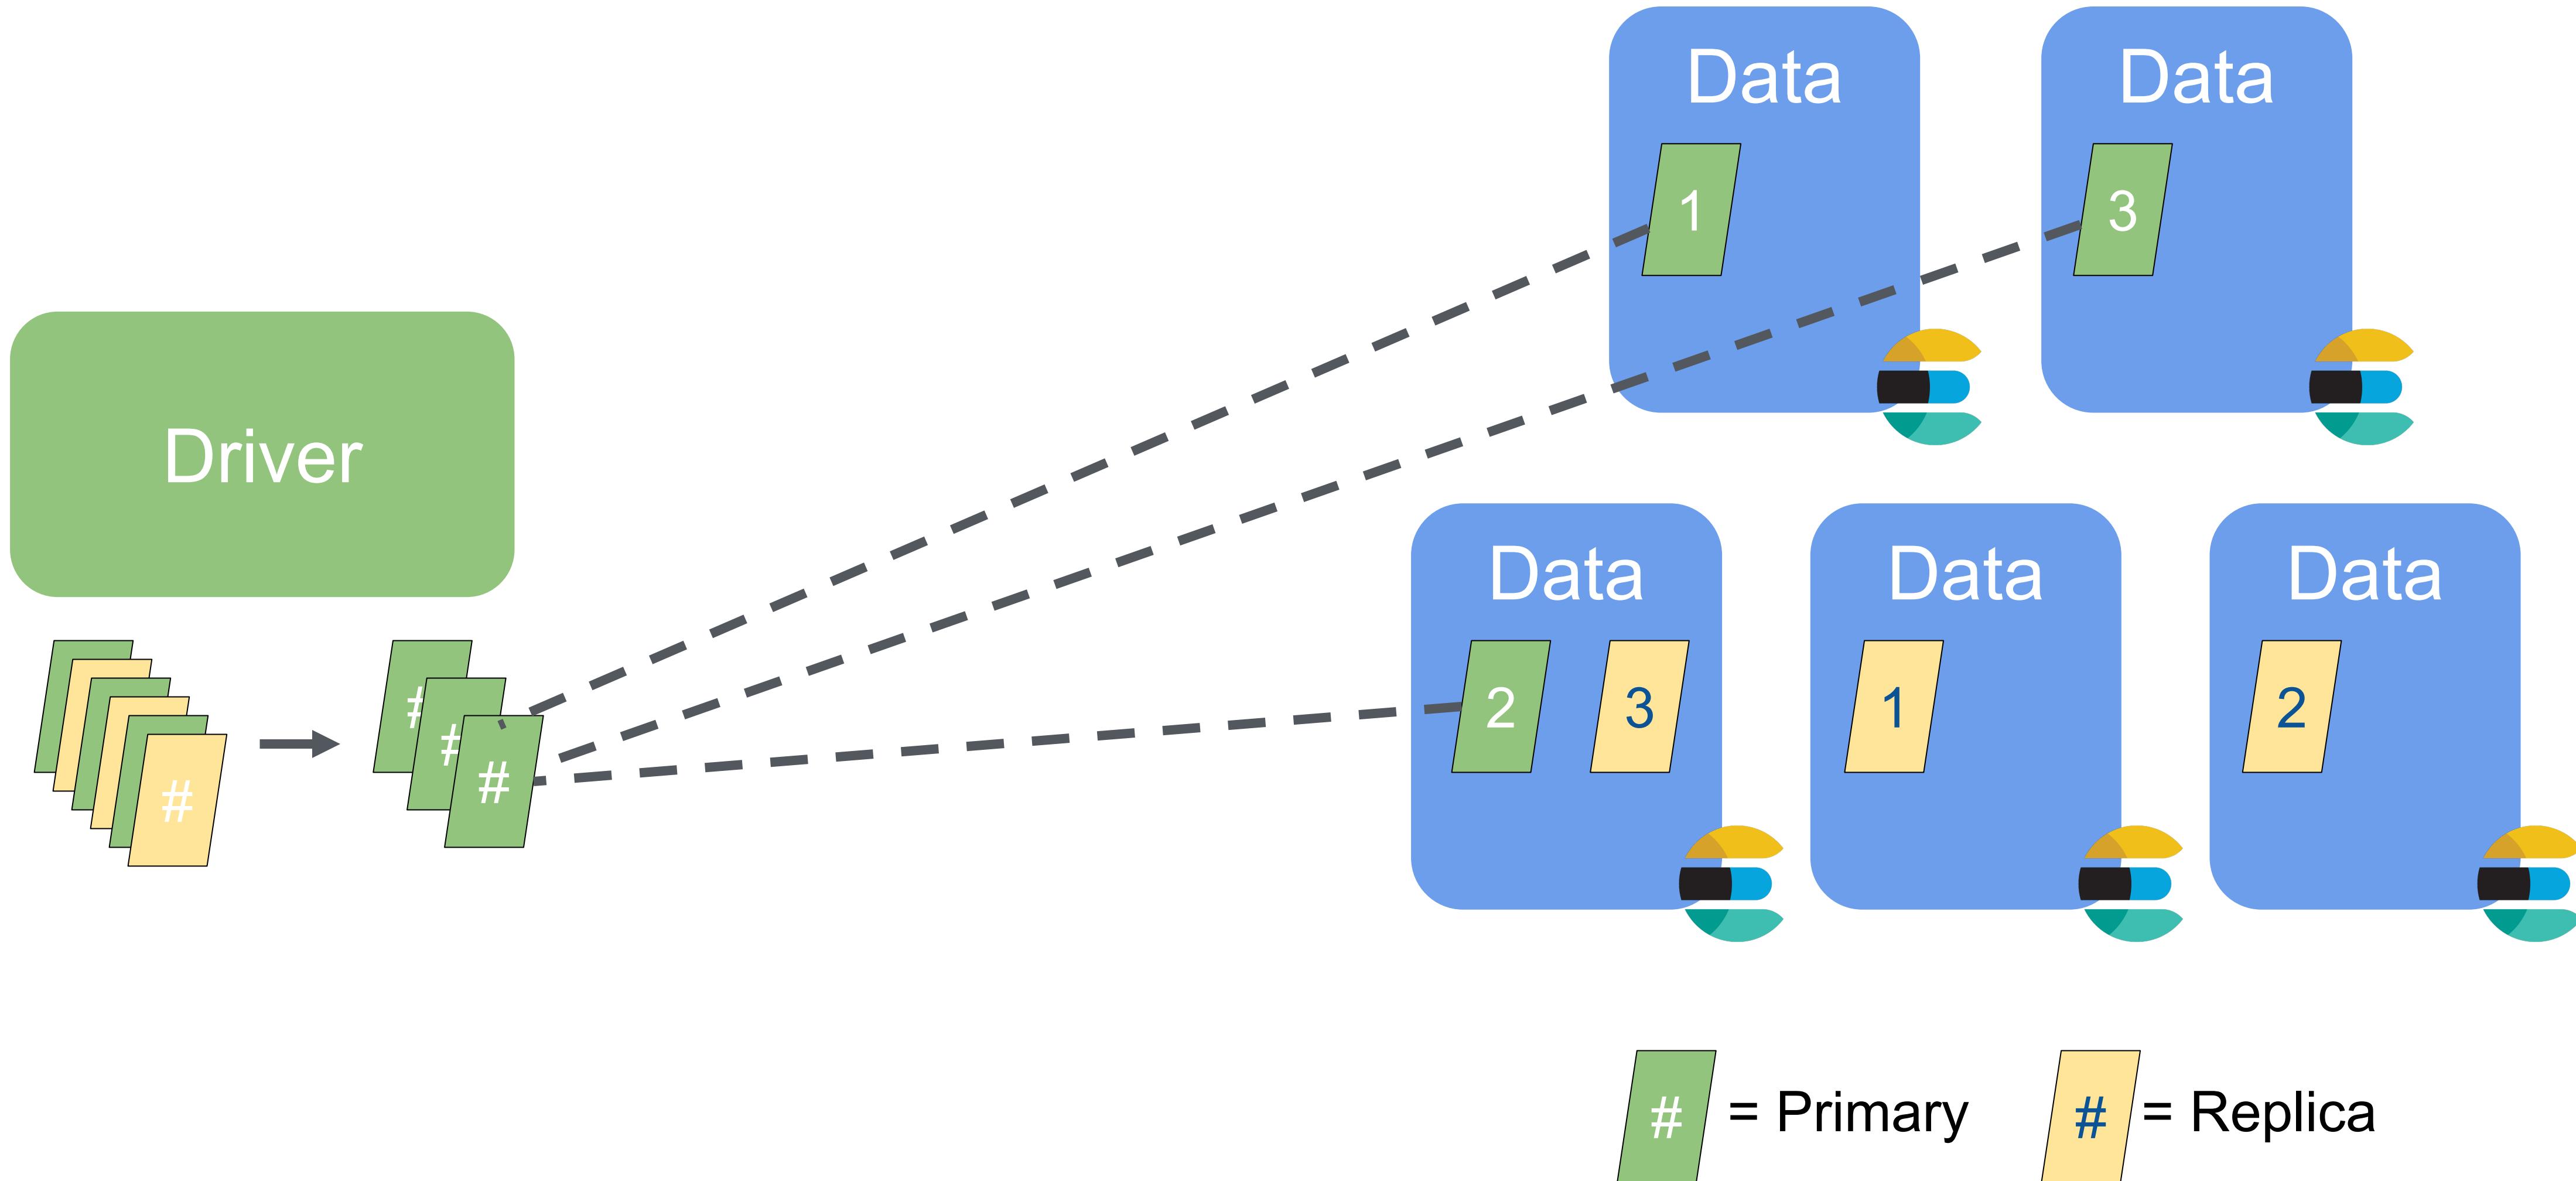


= Primary

= Replica

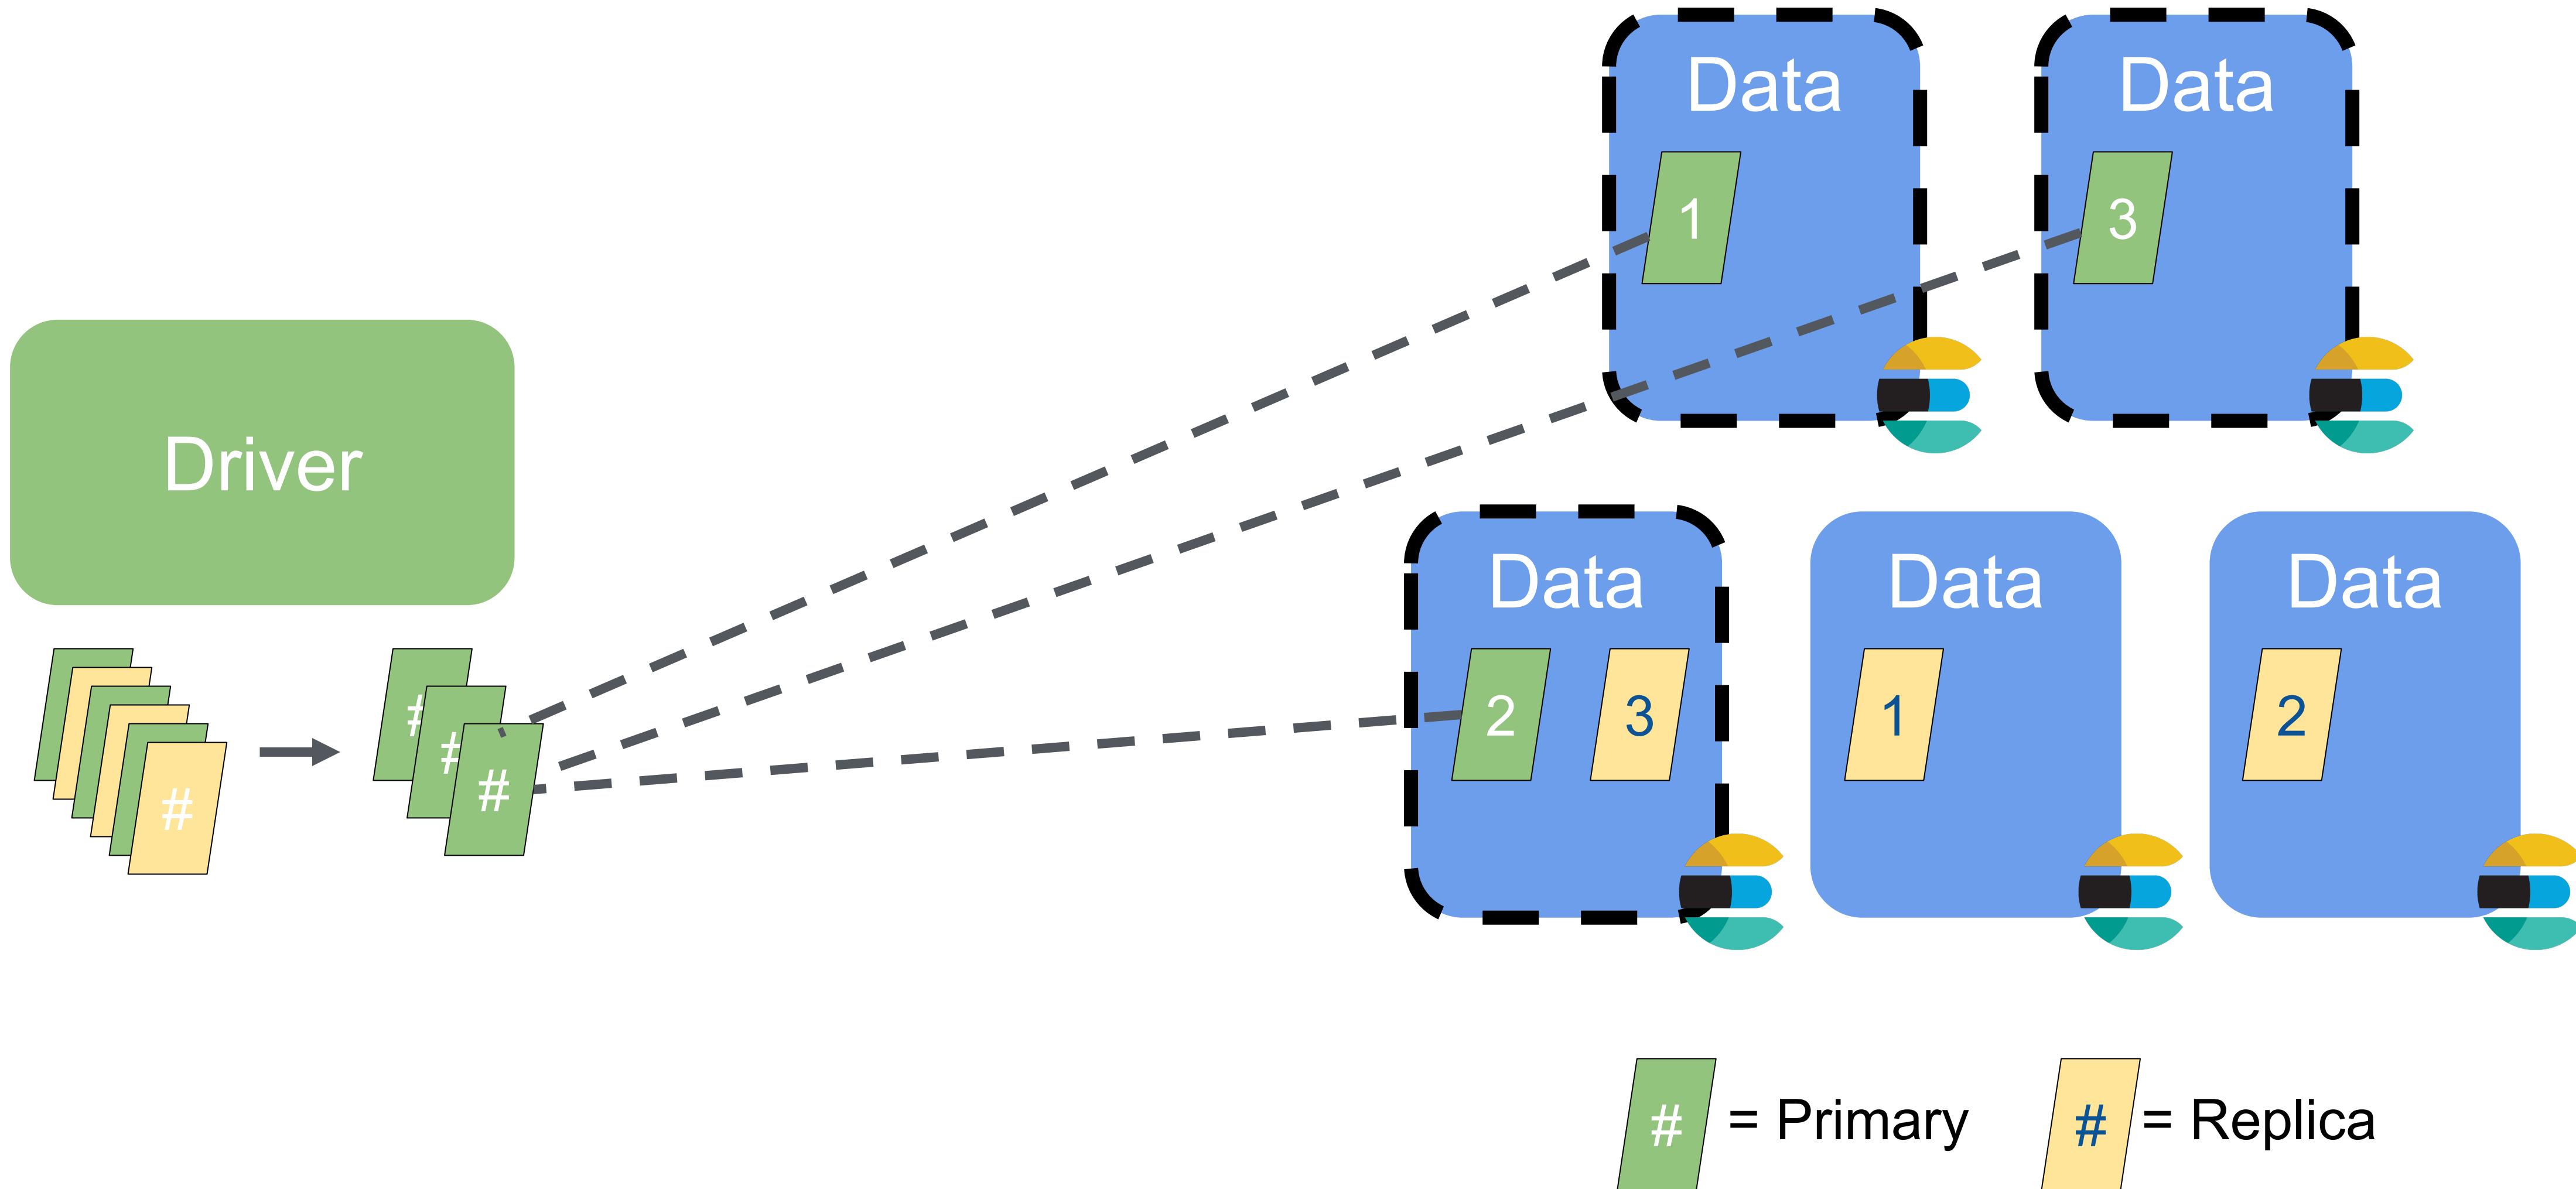
Finding Partitions

Writing to index ...



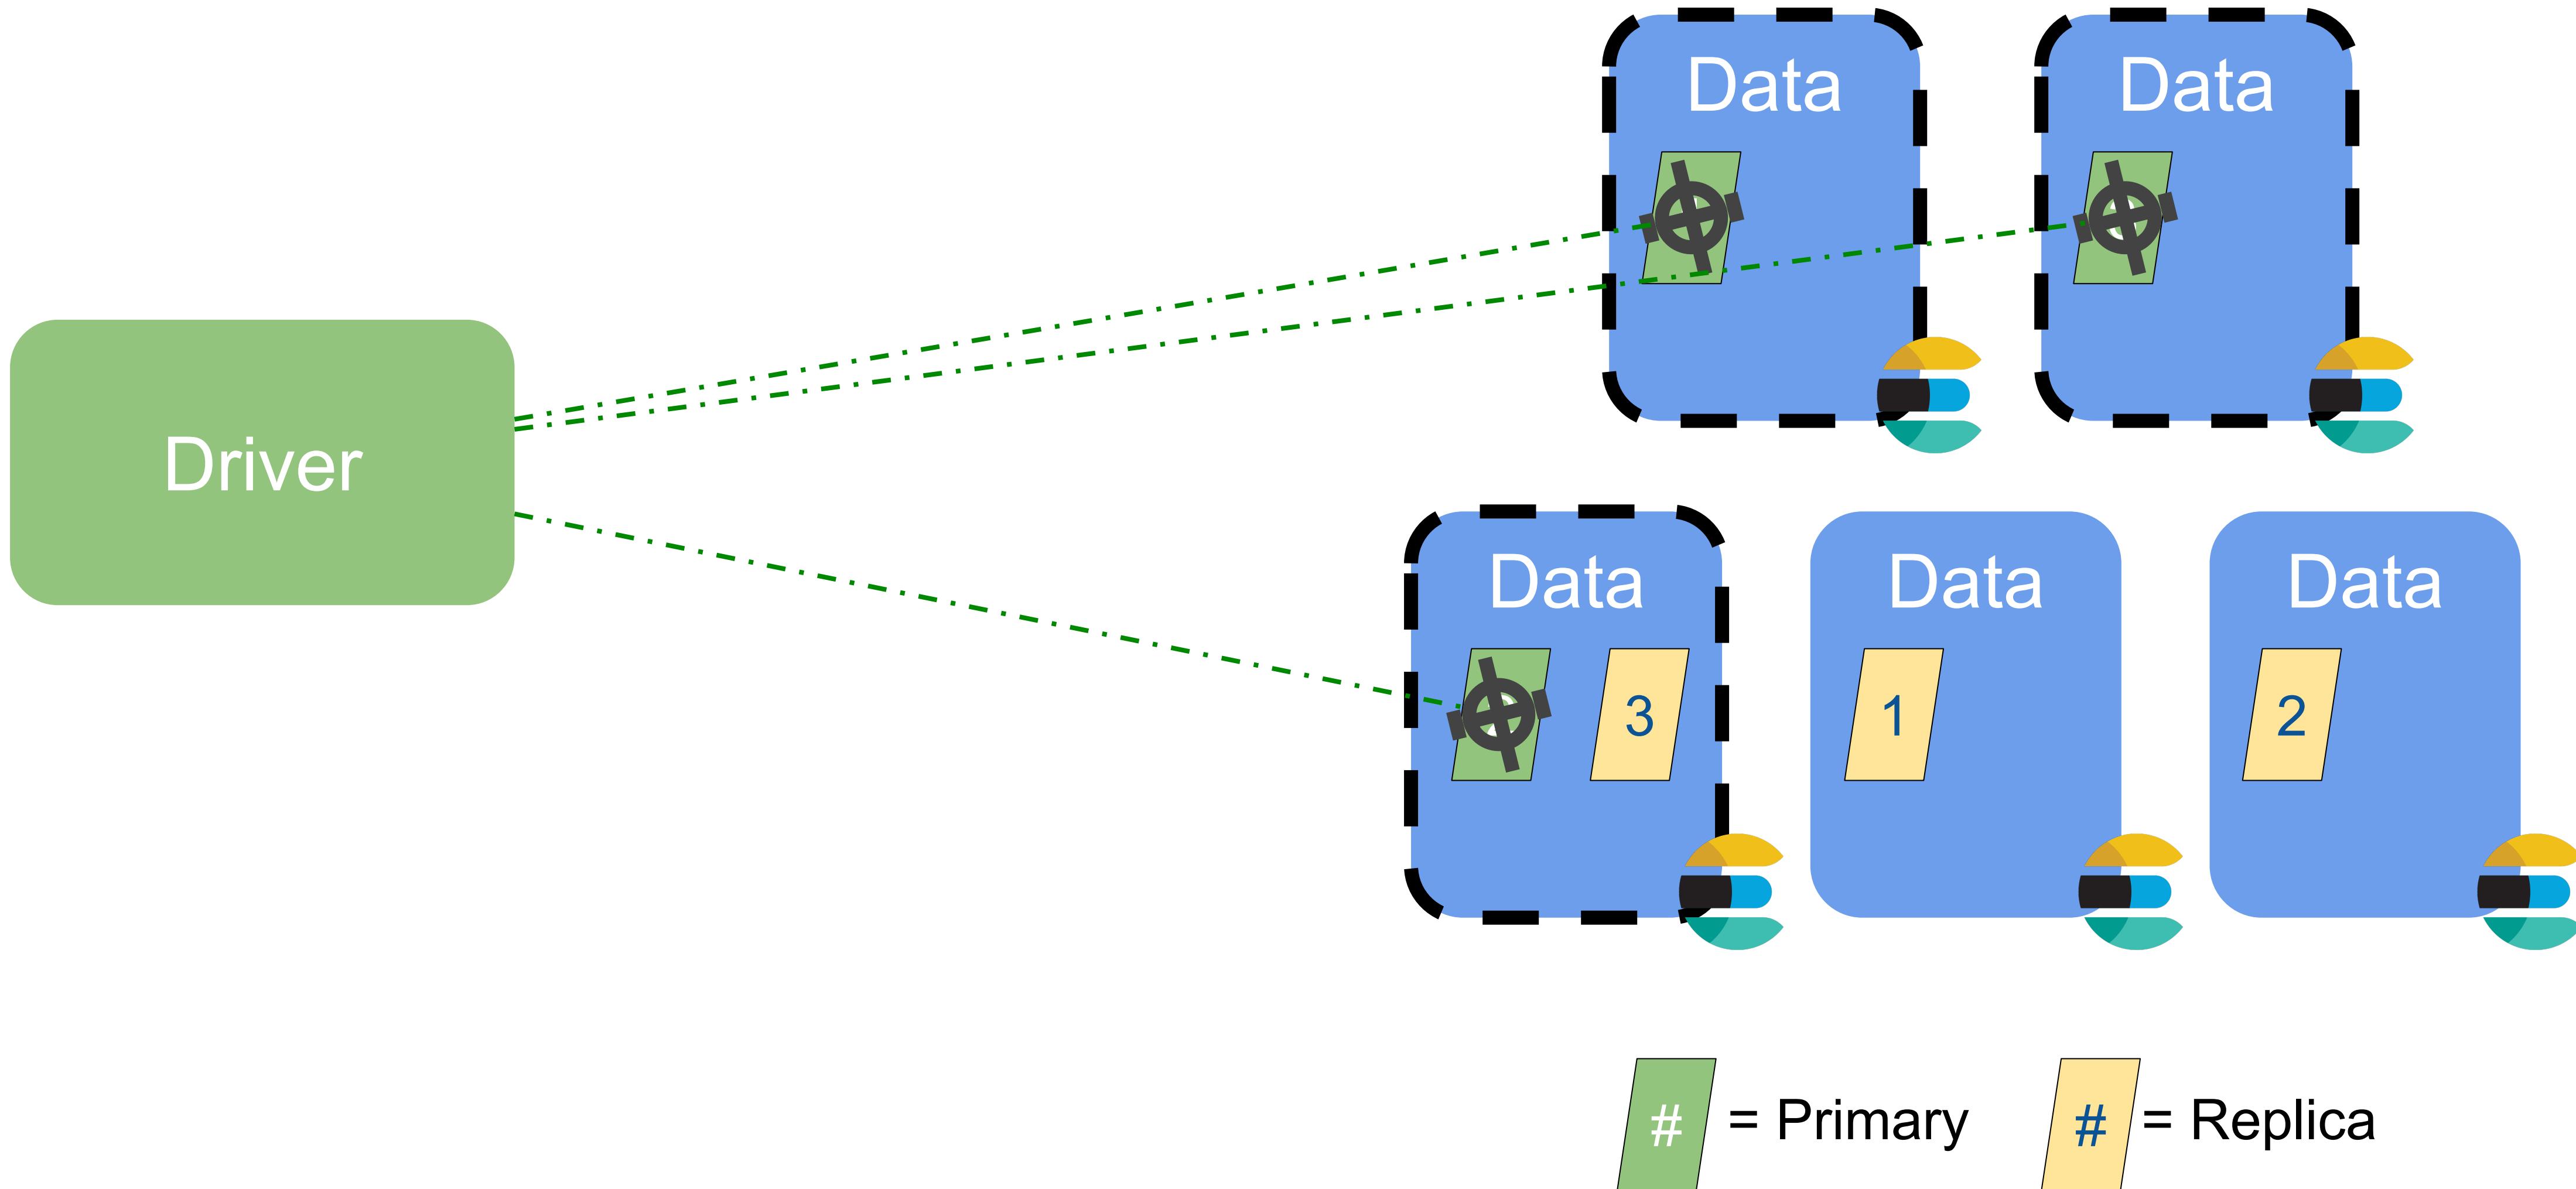
Finding Partitions

Writing to index ...

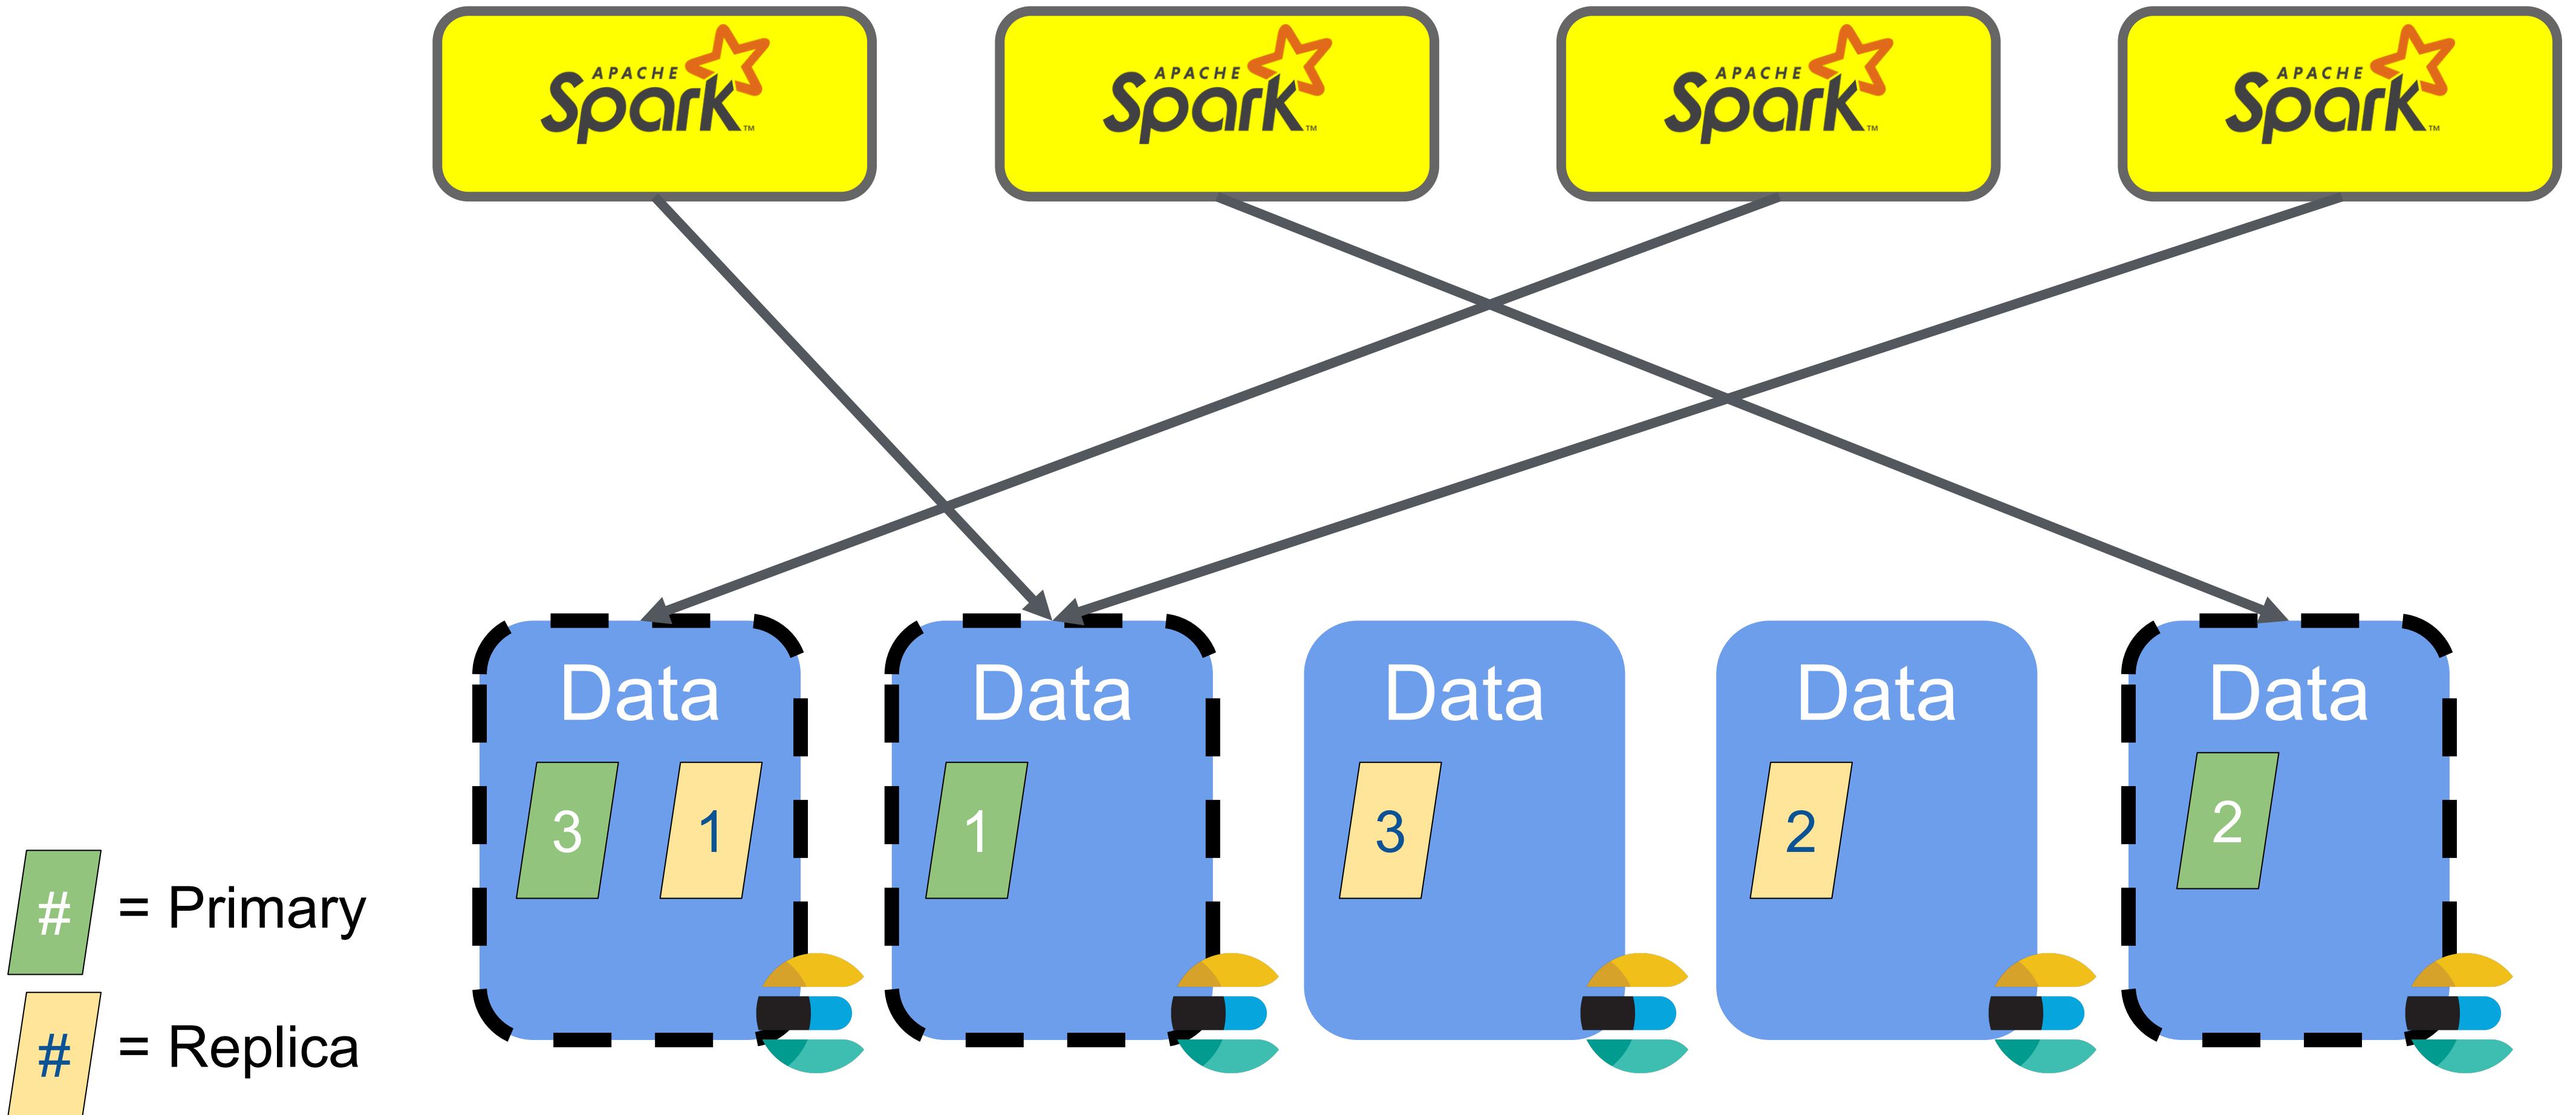


Finding Partitions

Writing to index ...



Job Execution



Feature Tour

Code Examples

Apache Hive

Reading from Elasticsearch

```
CREATE EXTERNAL TABLE artists (
    id      BIGINT,
    name    STRING,
    links   STRUCT<url:STRING, picture:STRING>)
STORED BY 'org.elasticsearch.hadoop.hive.EsStorageHandler'
TBLPROPERTIES(
    'es.resource' = 'radio/artists',
    'es.query' = '?q=me*');

-- stream data from Elasticsearch
SELECT * FROM artists;
```

Apache Pig

Writing to Elasticsearch

```
-- load data from HDFS into Pig using a schema
A = LOAD 'src/test/resources/artists.dat' USING PigStorage()
AS (id:long, name, genre, url:chararray, picture:chararray);

-- transform data
B = FOREACH A GENERATE name, genre, TOTUPLE(url, picture) AS
links;

-- save the result to Elasticsearch
STORE B INTO 'radio/{genre}' USING
org.elasticsearch.hadoop.pig.EsStorage('es.mapping.id=name');
```

Apache Spark

Reading from Elasticsearch

```
import org.elasticsearch.spark._

val conf = ...
val sc = new SparkContext(conf)
val rdd = sc.esRDD("radio/artists")

rdd.take(10)
```

Apache Spark Streaming

DStreams: The Old Way

```
val sc = new SparkContext(conf)
val ssc = new StreamingContext(sc, Seconds(1))

ssc.socketTextStream("127.0.0.1", 9999)
    .foreachRDD(EsSpark.saveToEs(_, "netcat/data"))

ssc.start()
```

Apache Spark Streaming

DStream Native Integration (NEW IN 5.0)

```
import org.elasticsearch.spark.streaming._

val sc = new SparkContext(conf)
val ssc = new StreamingContext(sc, Seconds(1))

ssc.socketTextStream("127.0.0.1", 9999)
    .saveJsonToEs("netcat/data")

ssc.start()
```

Ingest Node Support

Configurations Added for Ingest (NEW IN 5.0)

```
import org.elasticsearch.spark.streaming._

val sc = new SparkContext(conf)
val ssc = new StreamingContext(sc, Seconds(1))

val jobConf = Map("es.ingest.pipeline" -> "hadoop_test")

ssc.socketTextStream("127.0.0.1", 9999)
    .saveJsonToEs("netcat/data", jobConf)

ssc.start()
```

Coming Soon™

- Support for Unicode Index/Type Names
- Source Filtering for RDD's and MapReduce
- User Specified HTTP Headers
- Spark Structured Streaming Integration



Common Use Cases

Anoop Sunke



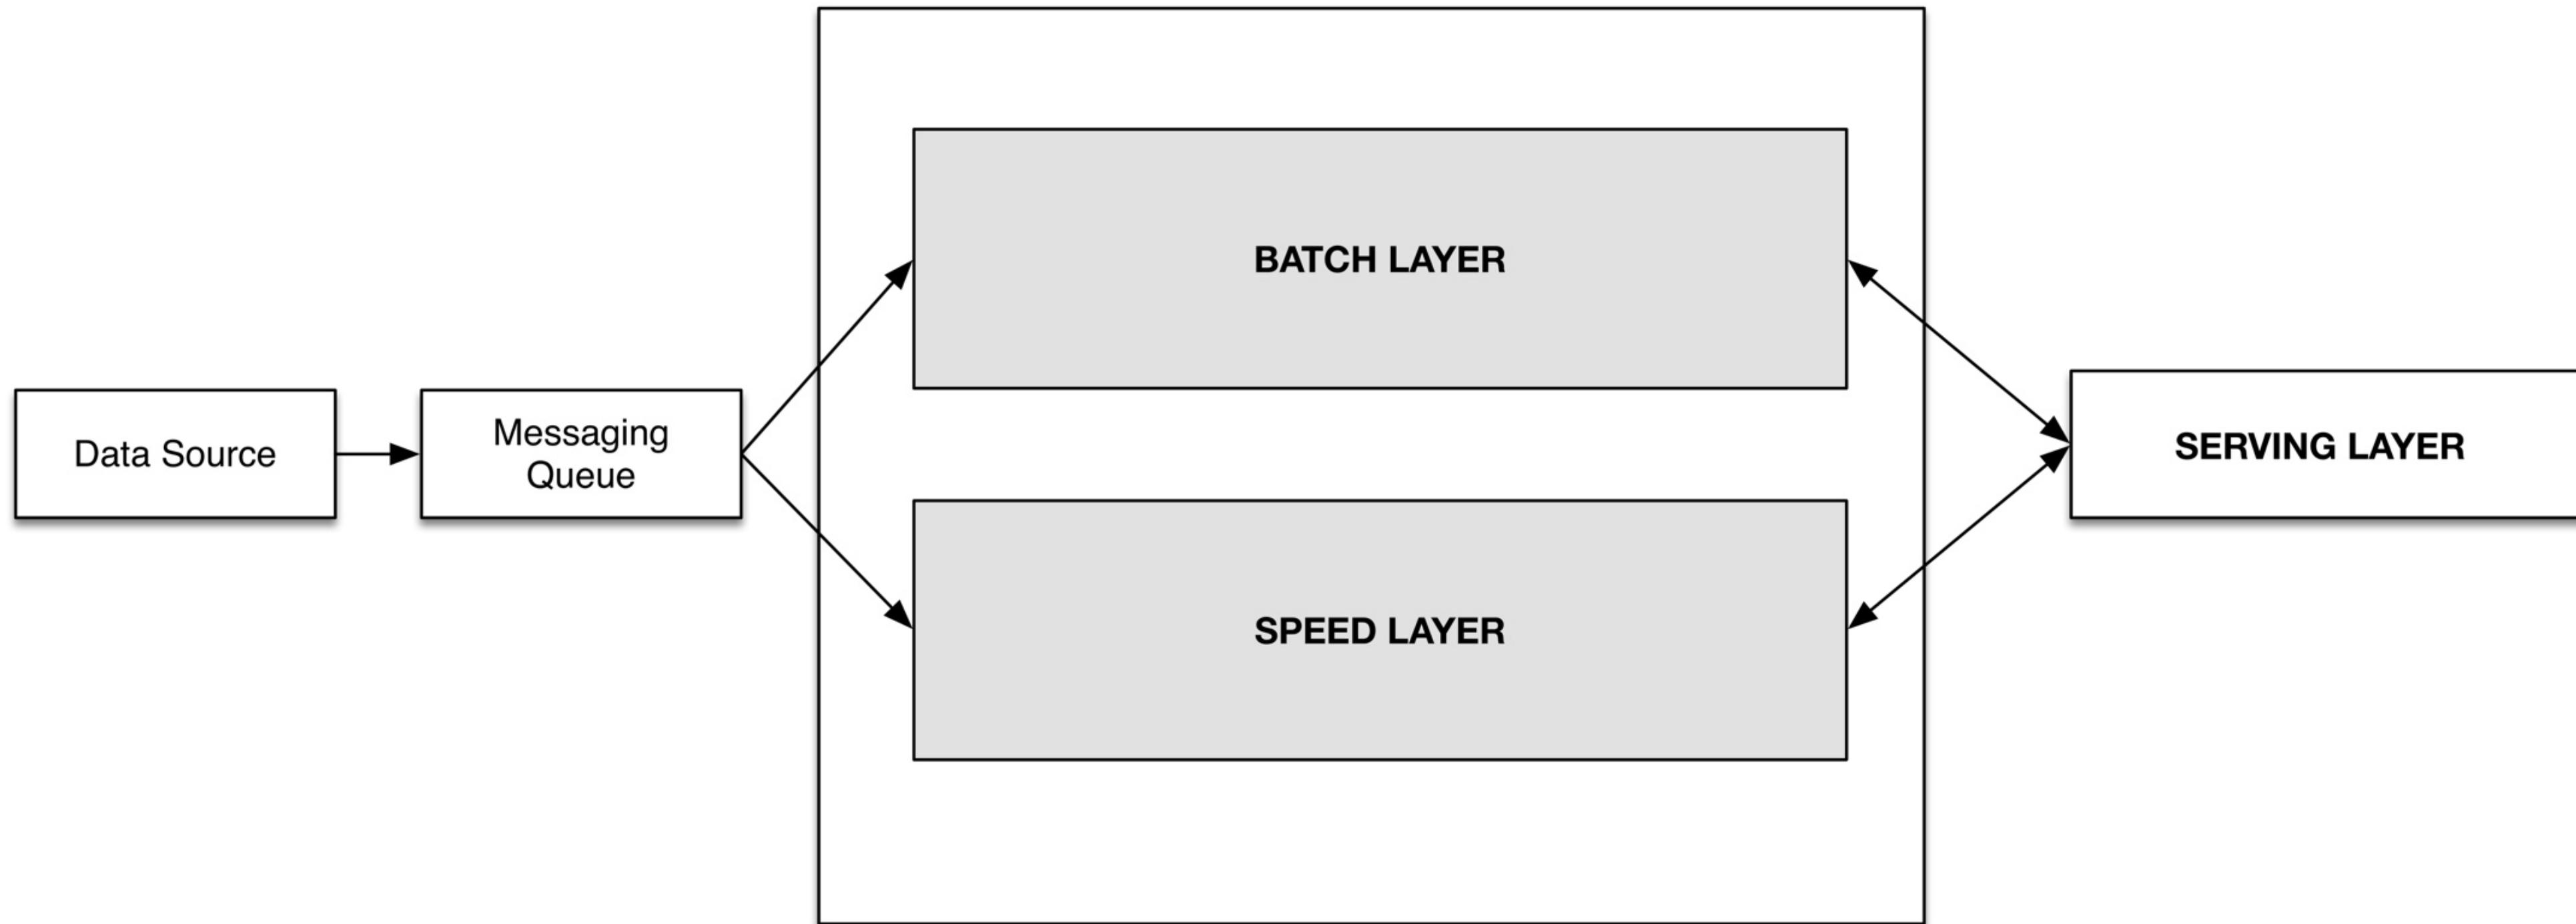
Anoop Sunke

- Solutions Architect @ Elastic
- Ex-Hortonworks, Microsoft
- @apsunke on twitter

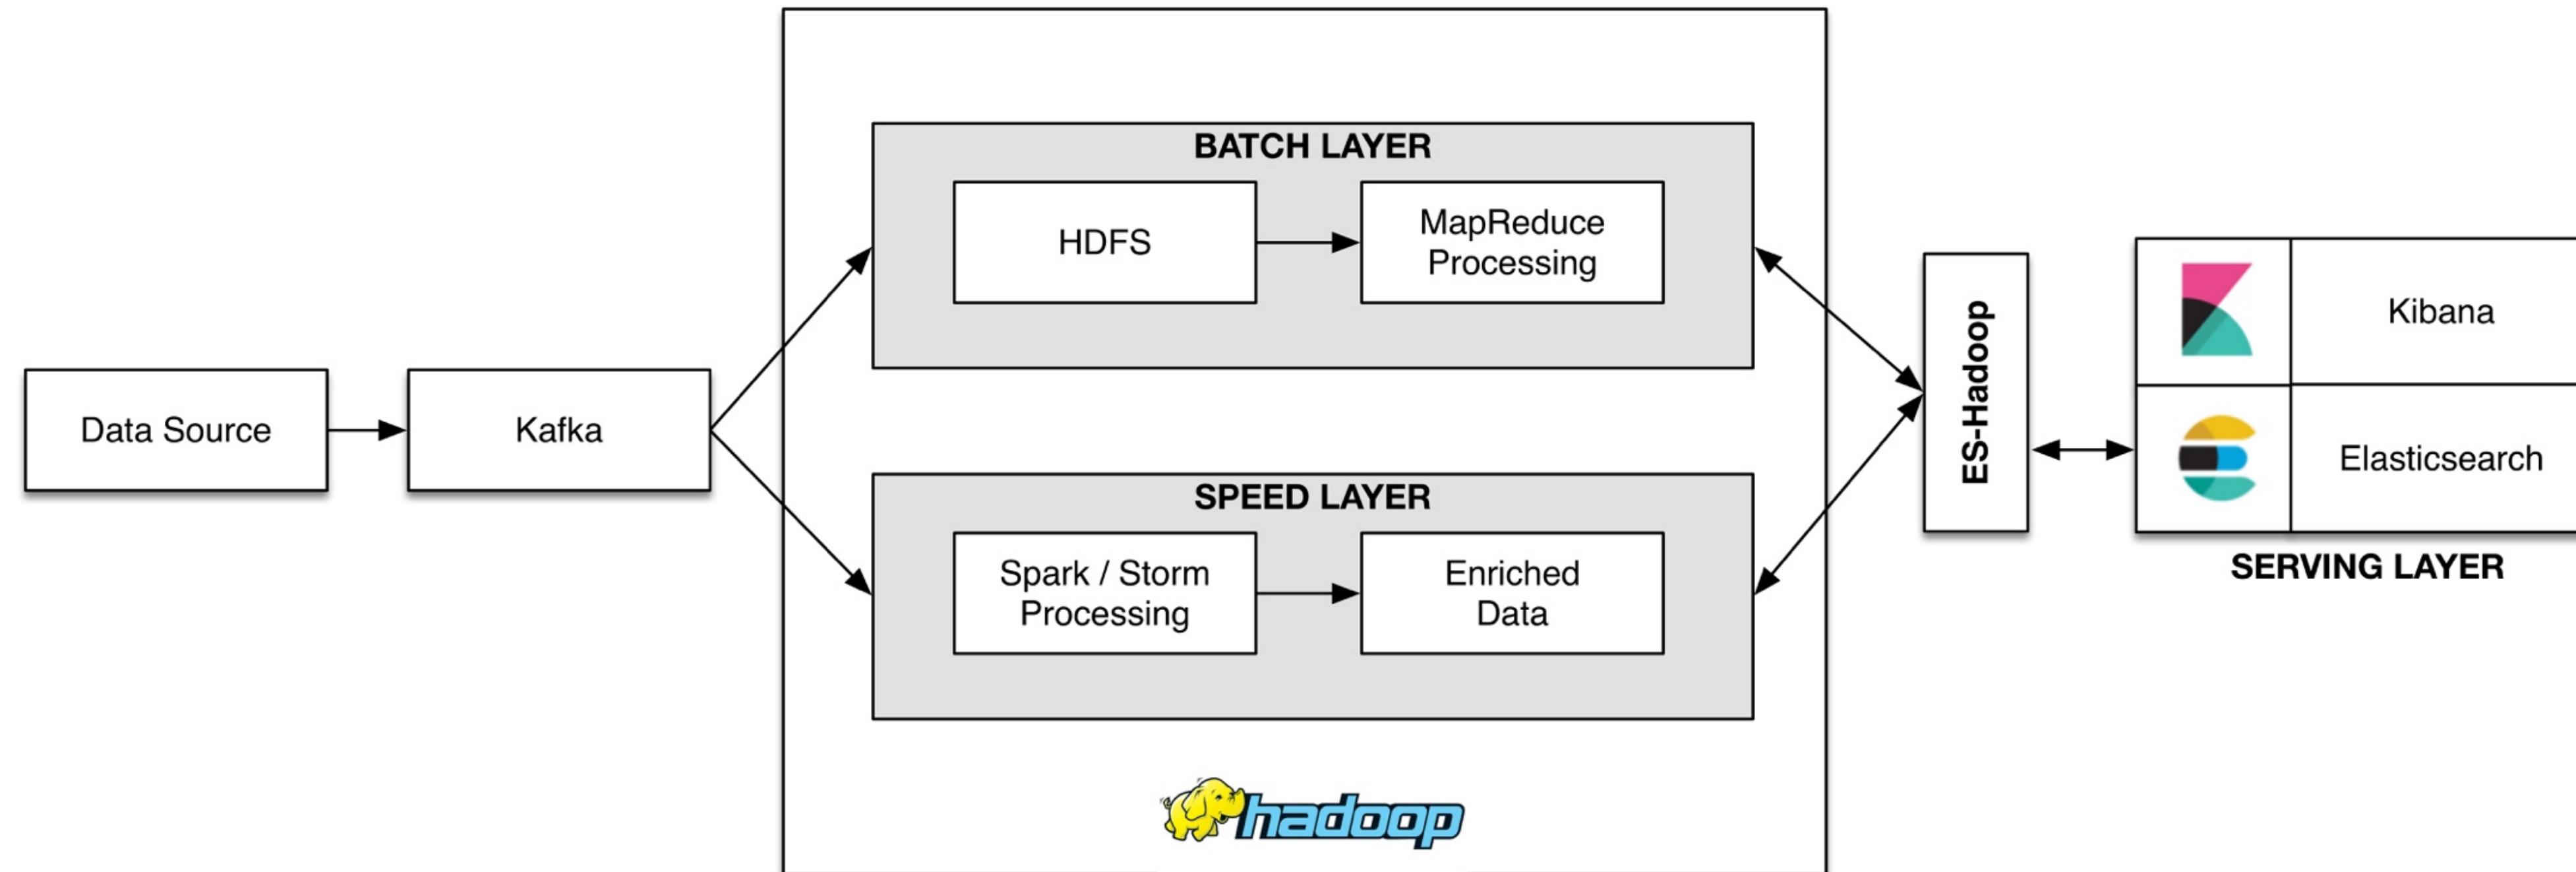
Use Case #1

Lambda Architecture

Lambda Architecture



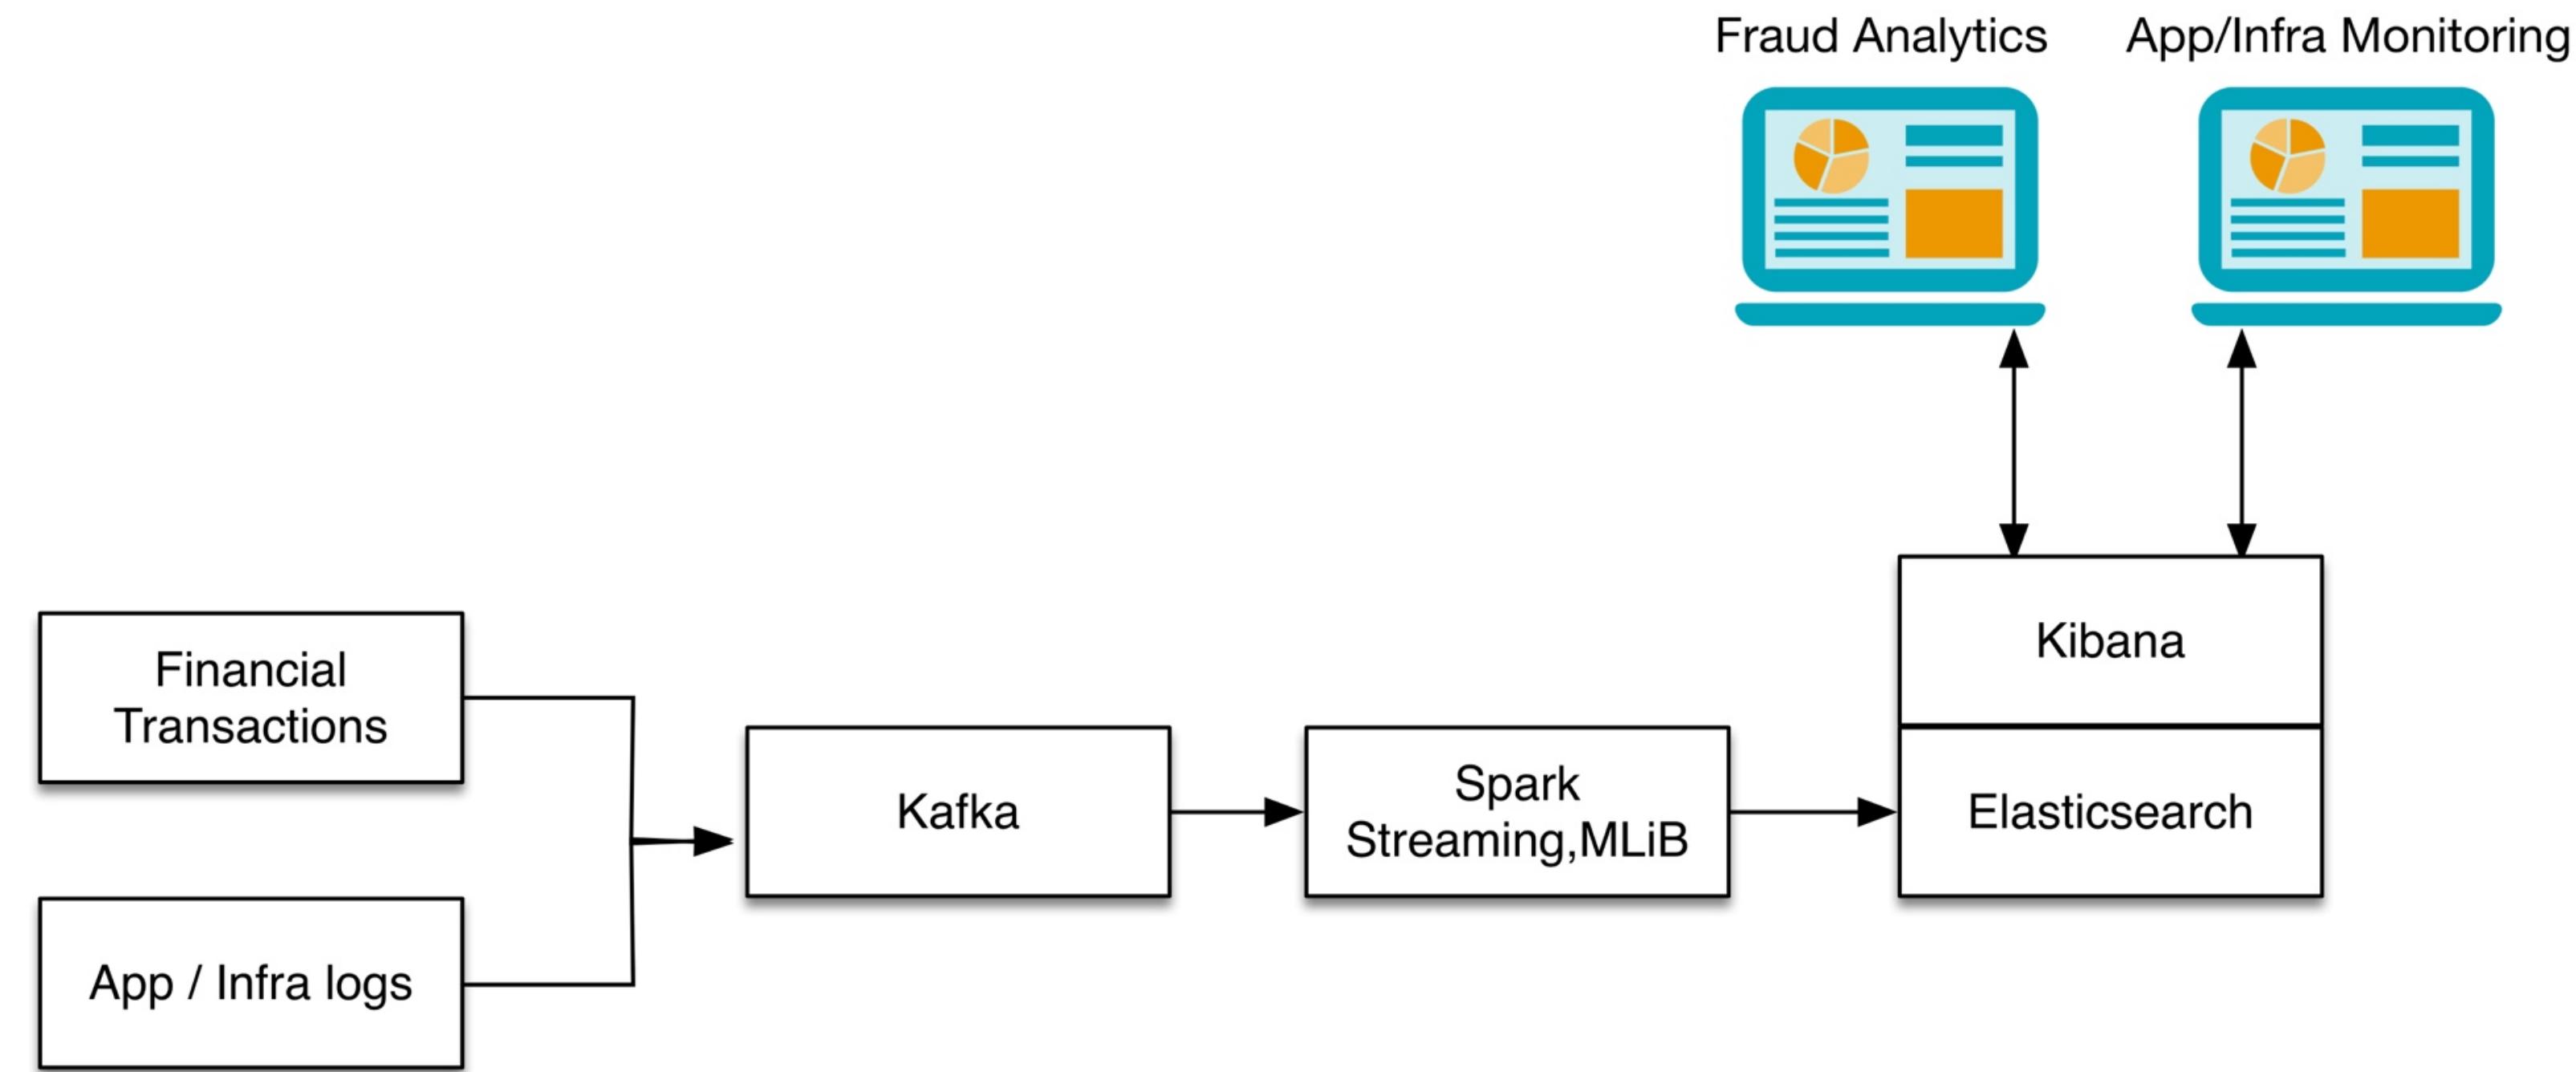
Lambda Architecture



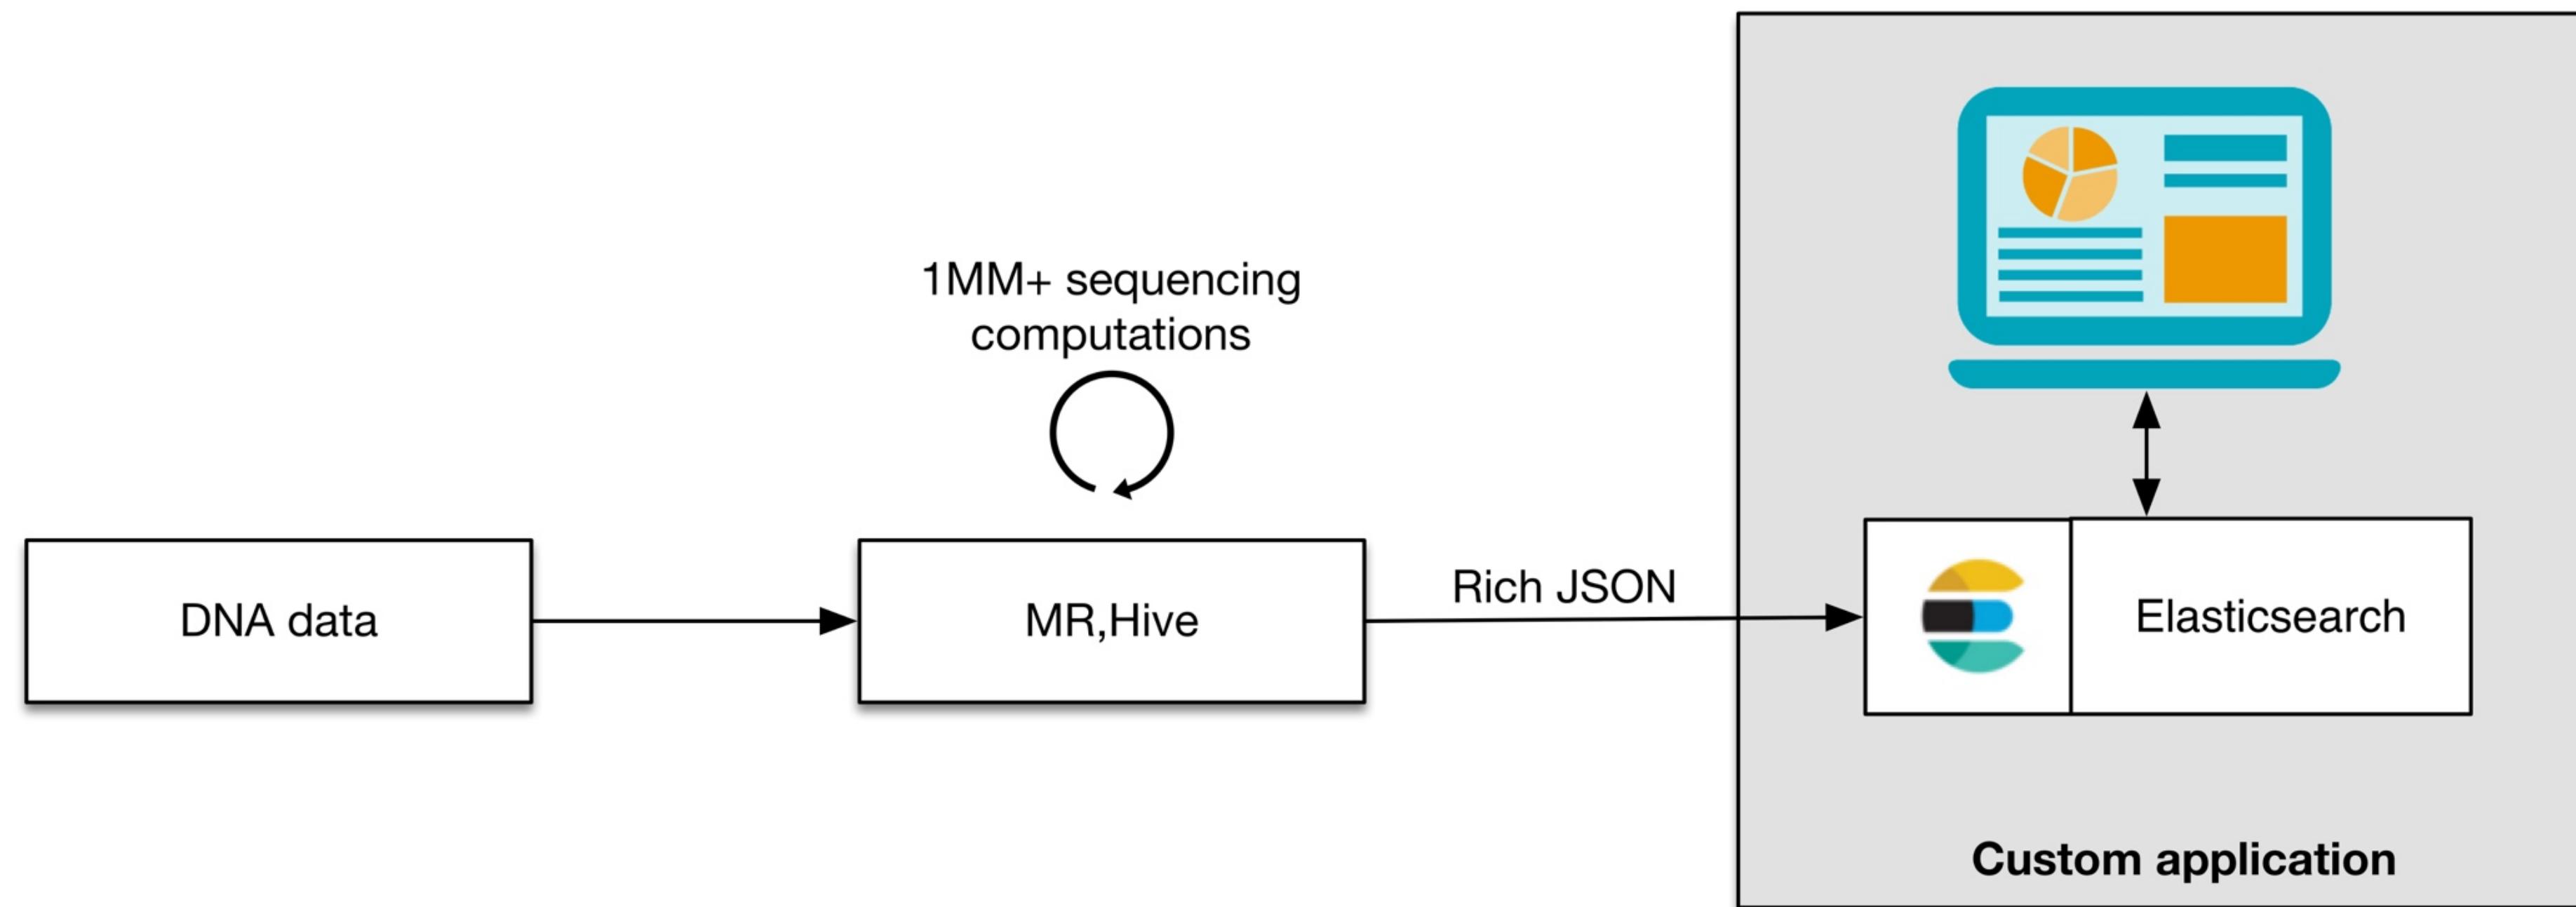
Why ?

- High scale Ingest rate - 1 Million+ Events Per Second
- Data available in 1 second, default refresh interval
- Simple data model – JSON
- Flexible schema - Add fields anytime
- Kibana - Native real time exploration, even at scale
- REST API first data store – Fast Aggregations, Custom UI, Embedding
- Index Index Index !
- Out of the box functionality - sharding, rebalancing, easy scaling

Financial Services company



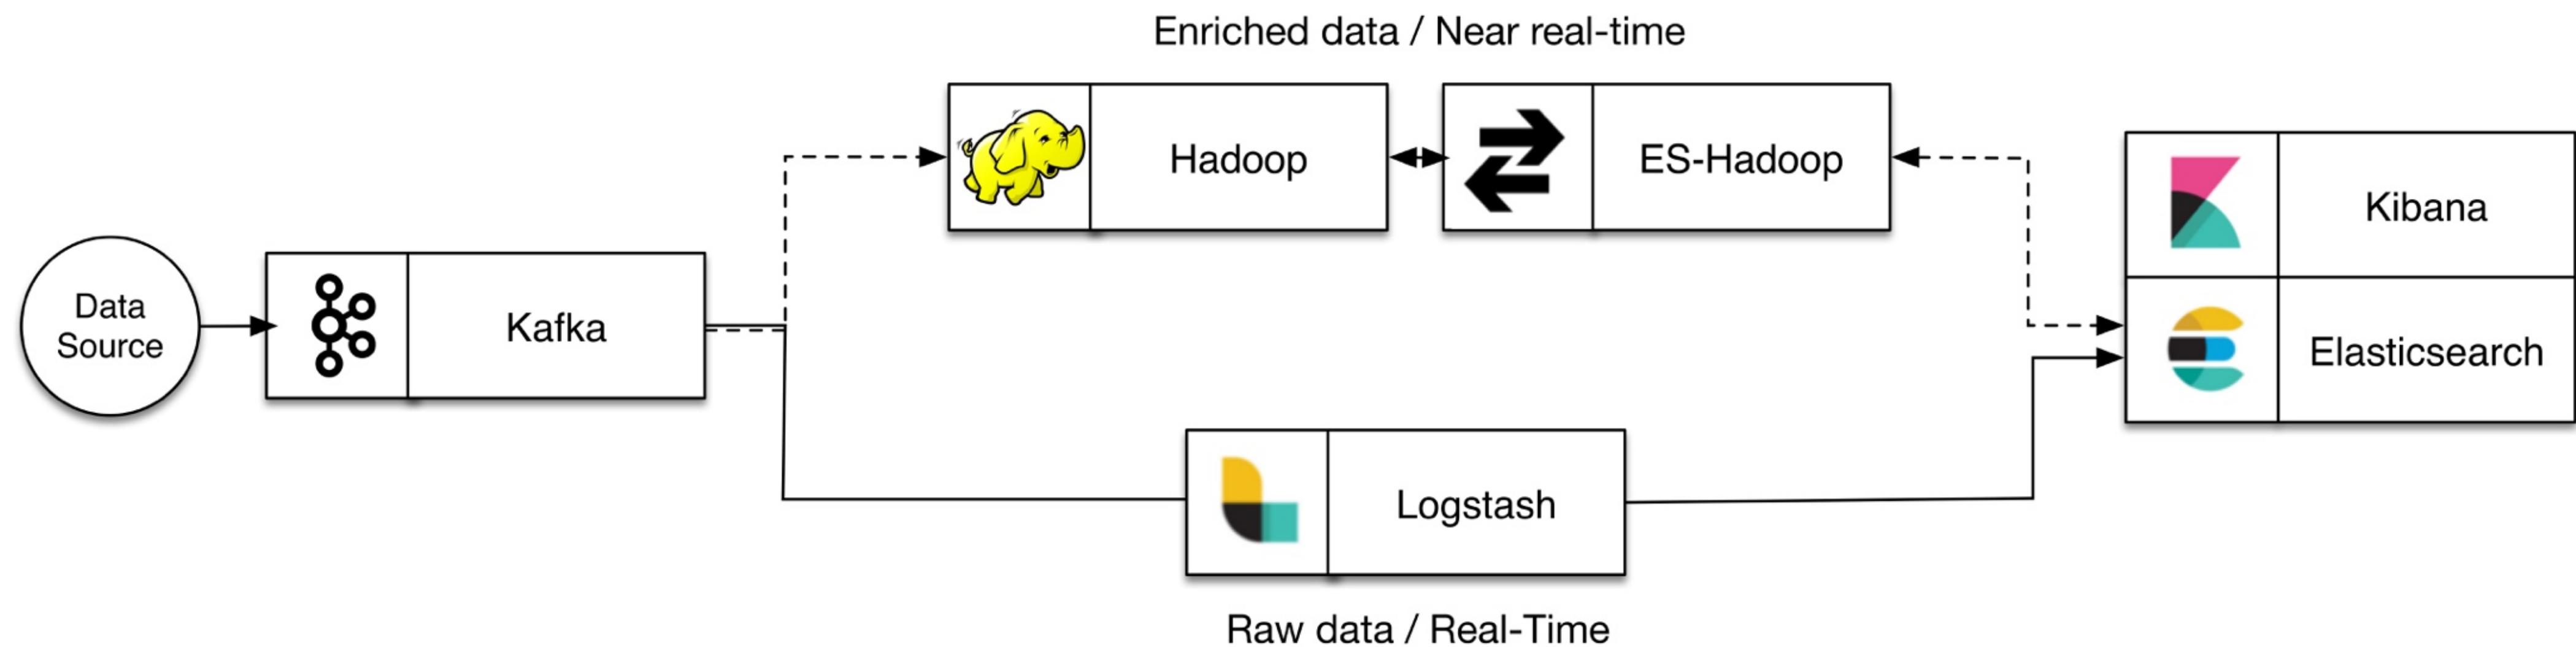
Bioinformatics Company



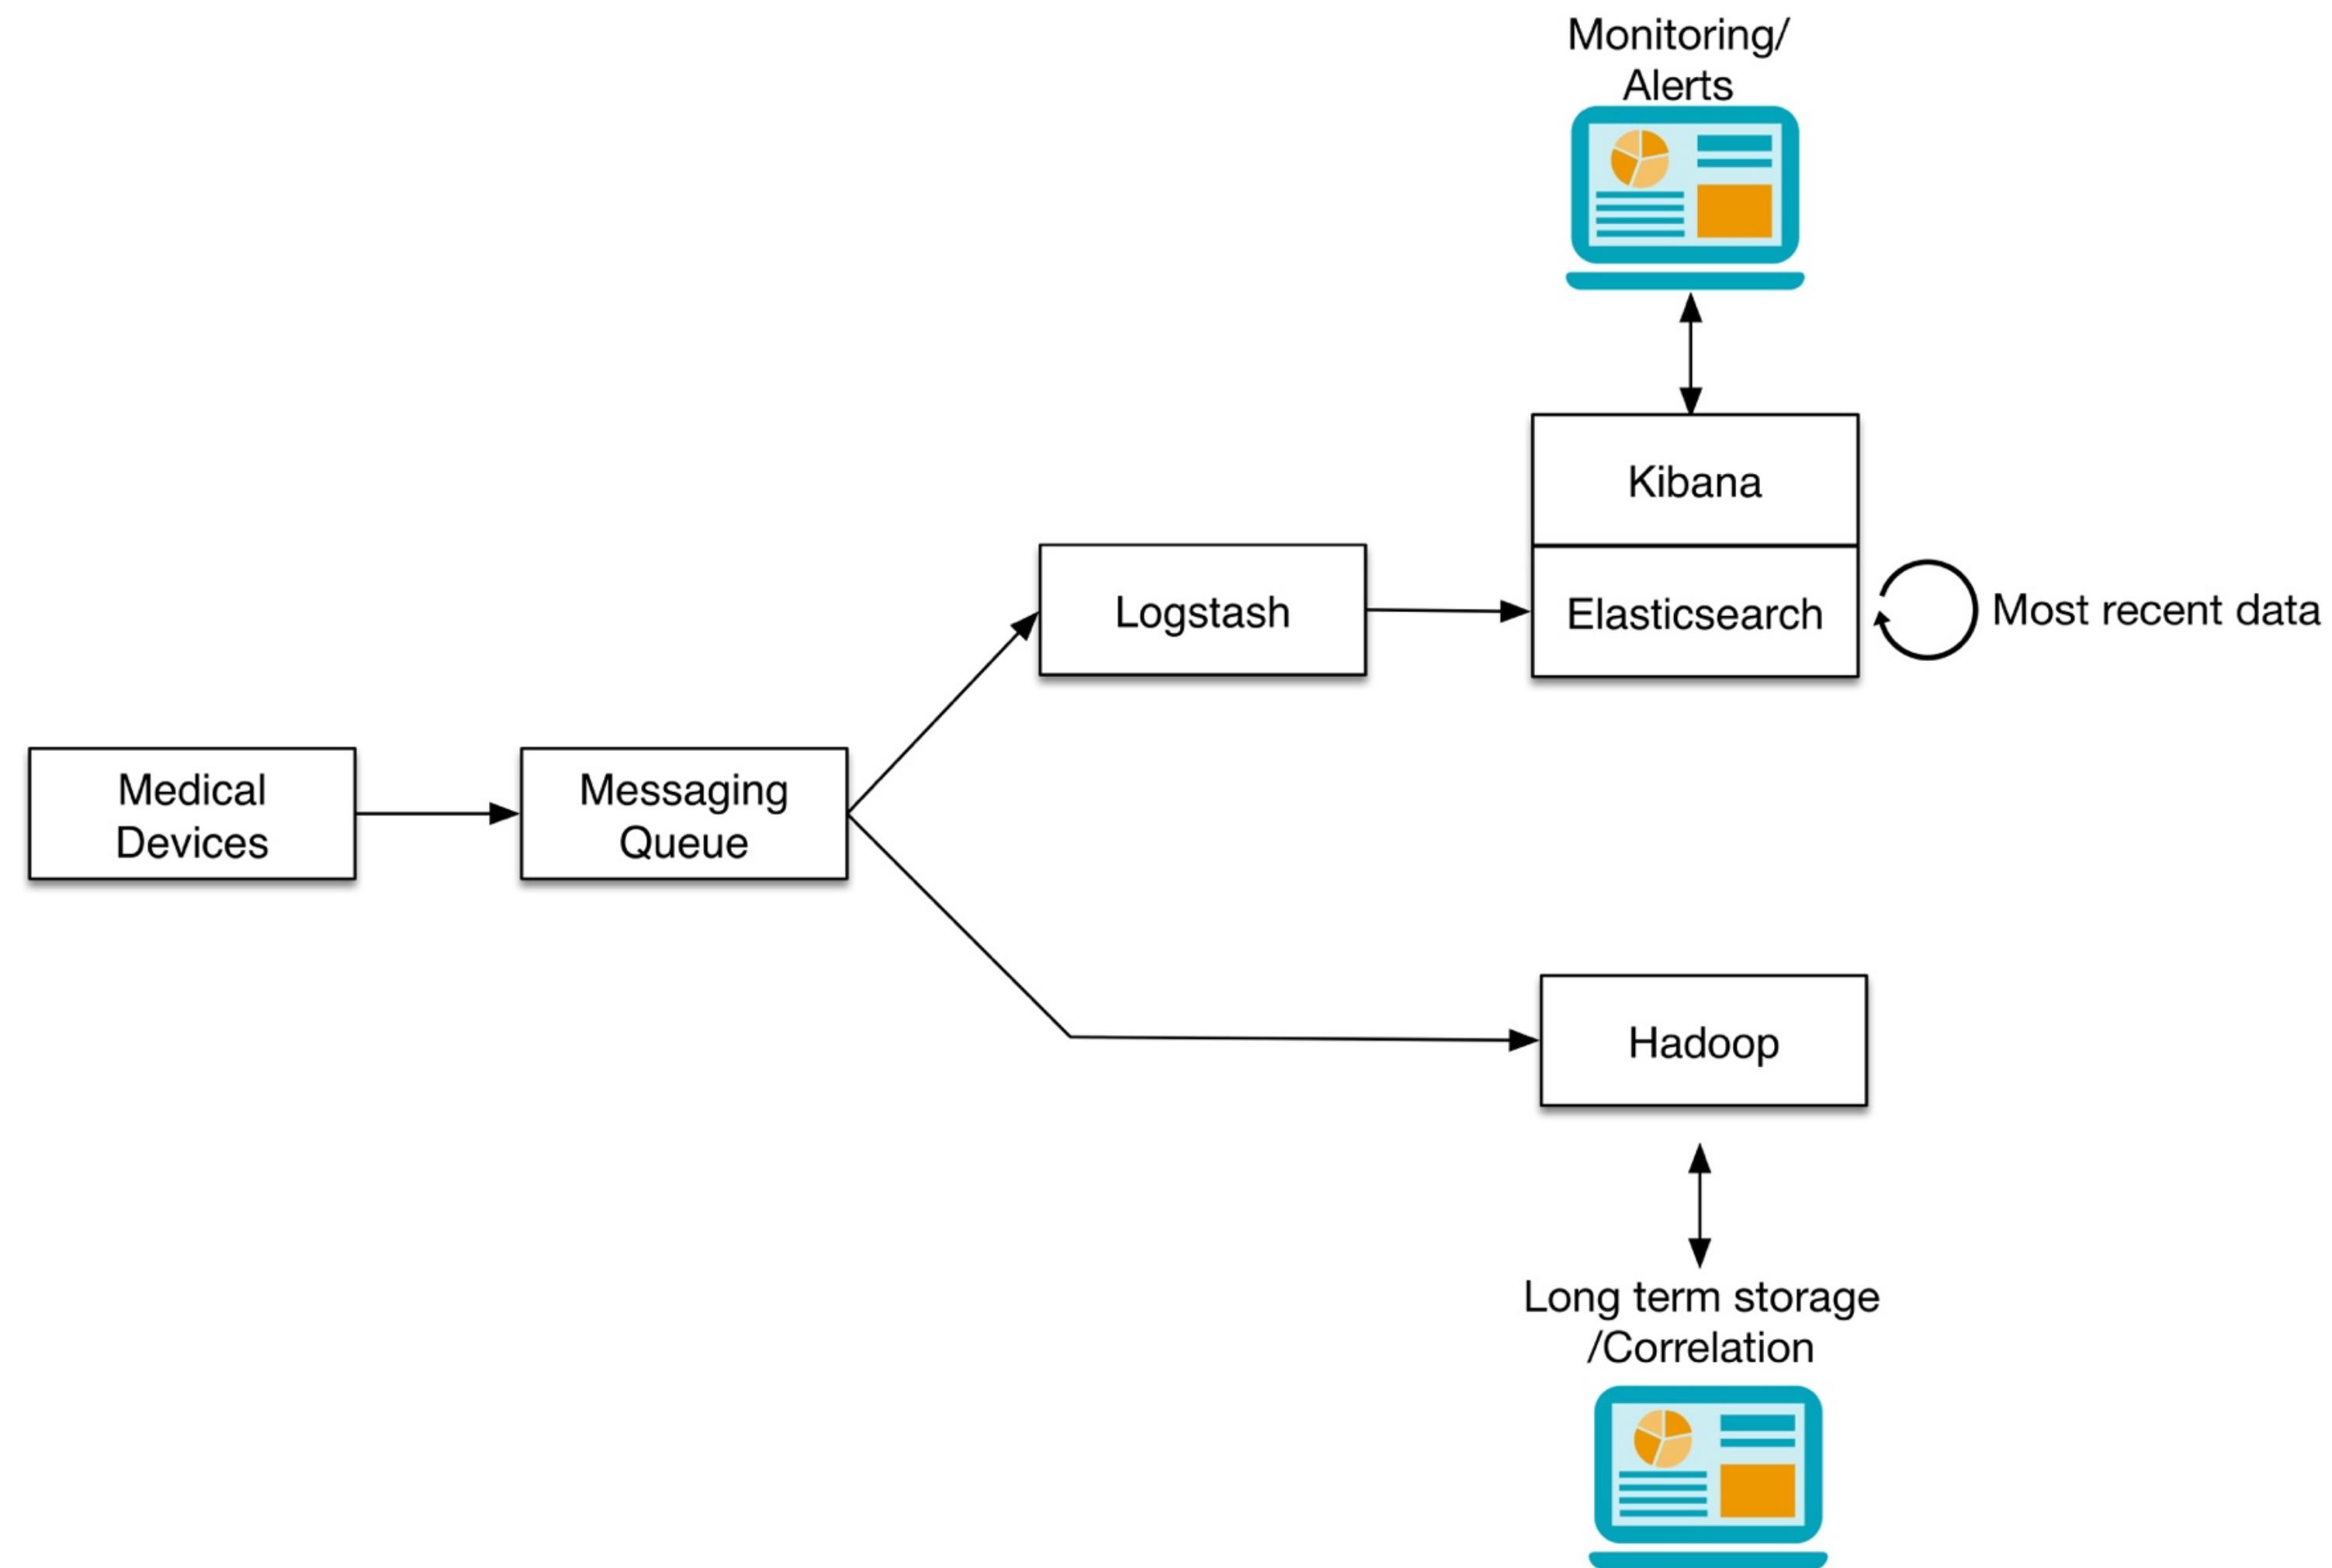
Use Case #2

Raw Real-time

Raw real-time Architecture



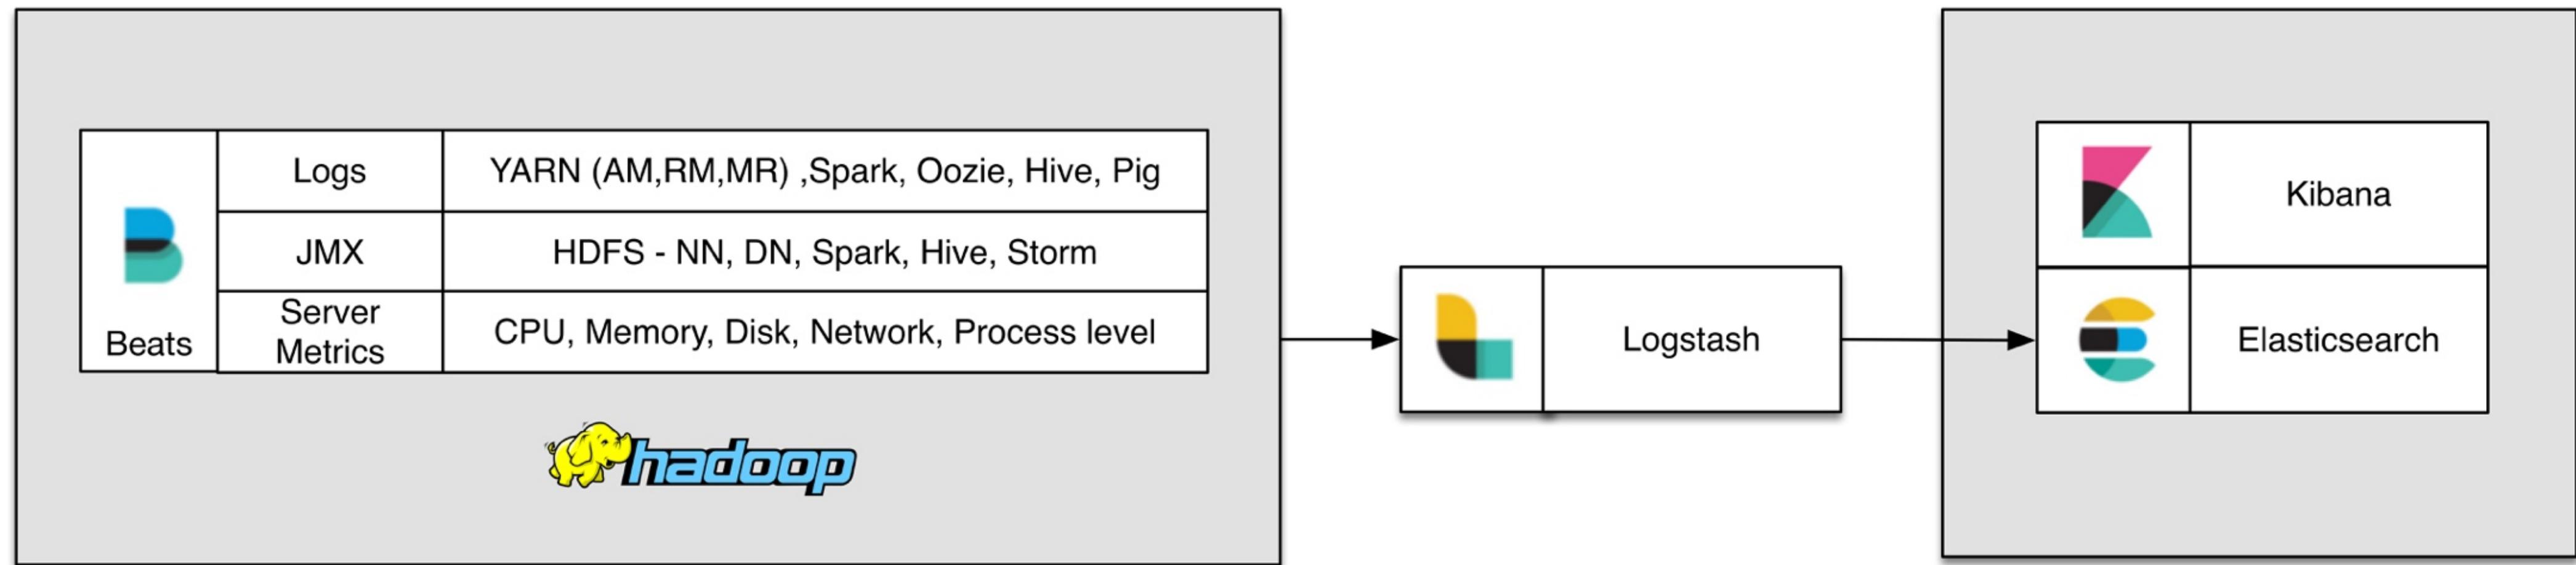
Medical Devices Company



Use Case #3

Monitoring Hadoop

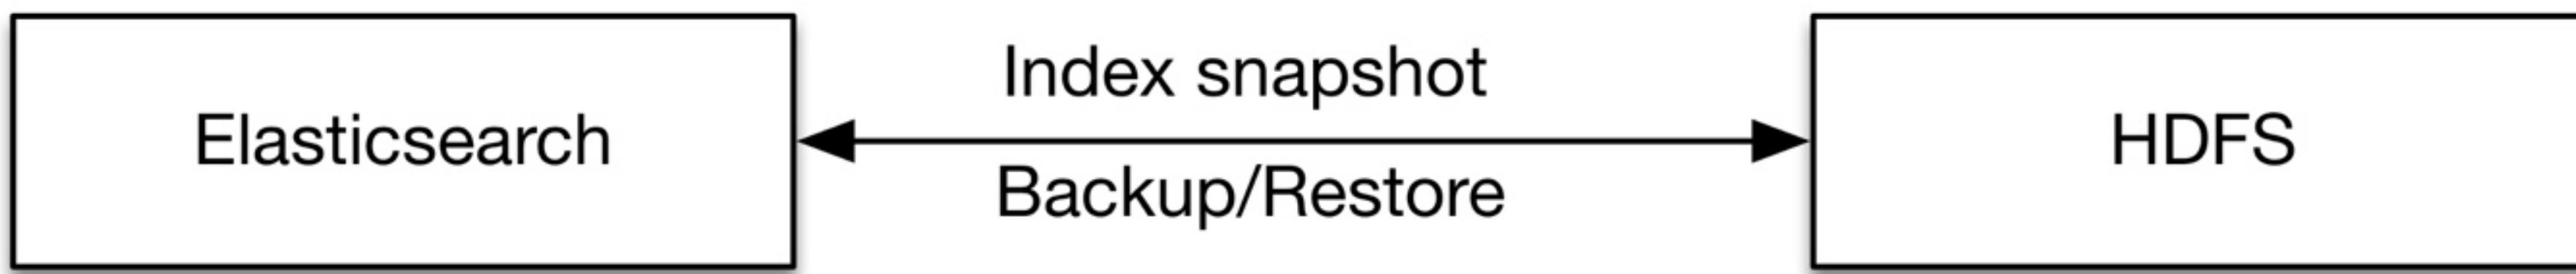
Monitoring Hadoop



Use Case #4

HDFS as a backup store

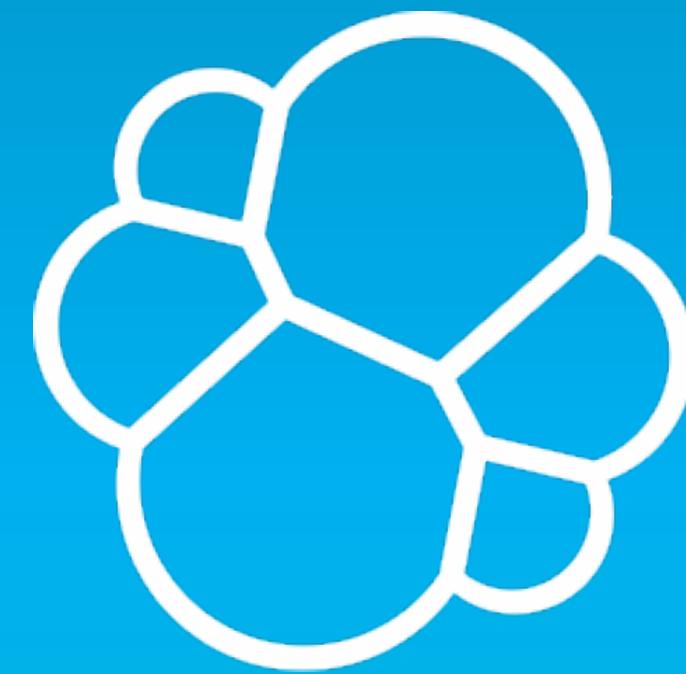
HDFS as a backup store



Thank You!

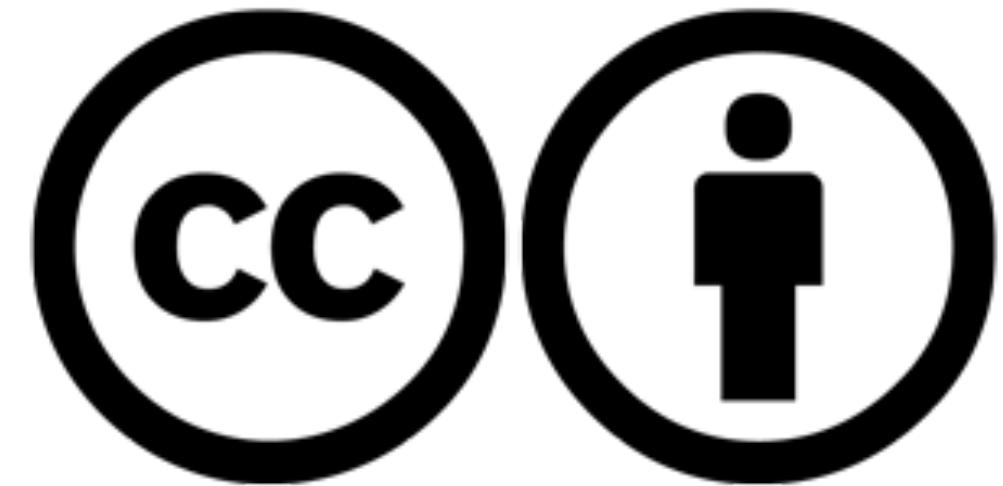
Visit us on our forums!

<https://discuss.elastic.co/c/elasticsearch-and-hadoop>



elastic

[www.elastic.c
o](http://www.elastic.co)



Except where otherwise noted, this work is licensed under
<http://creativecommons.org/licenses/by-nd/4.0/>

Creative Commons and the double C in a circle are registered trademarks of Creative Commons in the United States and other countries. Third party marks and brands are the property of their respective holders.